

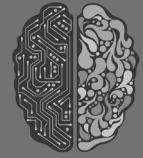
Introduction to Data Science and Artificial Intelligence

State-of-the-Art in DS&AI

Assoc Prof Bo AN

Research area: artificial intelligence,
computational game theory, optimization
www.ntu.edu.sg/home/boan
Email: boan@ntu.edu.sg
Office: N4-02b-55

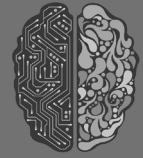




Lesson Outline

- A brief history of AI
- The state of the art





What is Artificial Intelligence?

AI is intelligence demonstrated by machines, in contrast to the **natural intelligence (NI)** displayed by humans and other animals.



Birth of Artificial Intelligence

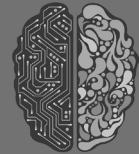
“We can only see a short distance ahead, but we can see plenty there that needs to be done.”

*Alan Turing
Father of Modern Computer Science*



Image source: <https://www.flickr.com/photos/23925401@N06/23022750663>





Founding Fathers of AI

1956 Dartmouth Conference



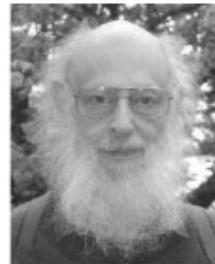
John McCarthy



Marvin Minsky



Claude Shannon



Ray Solomonoff



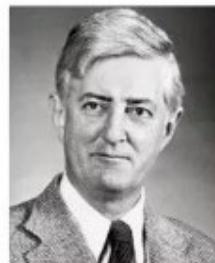
Alan Newell



Herbert Simon



Arthur Samuel



Oliver Selfridge



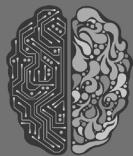
Nathaniel Rochester



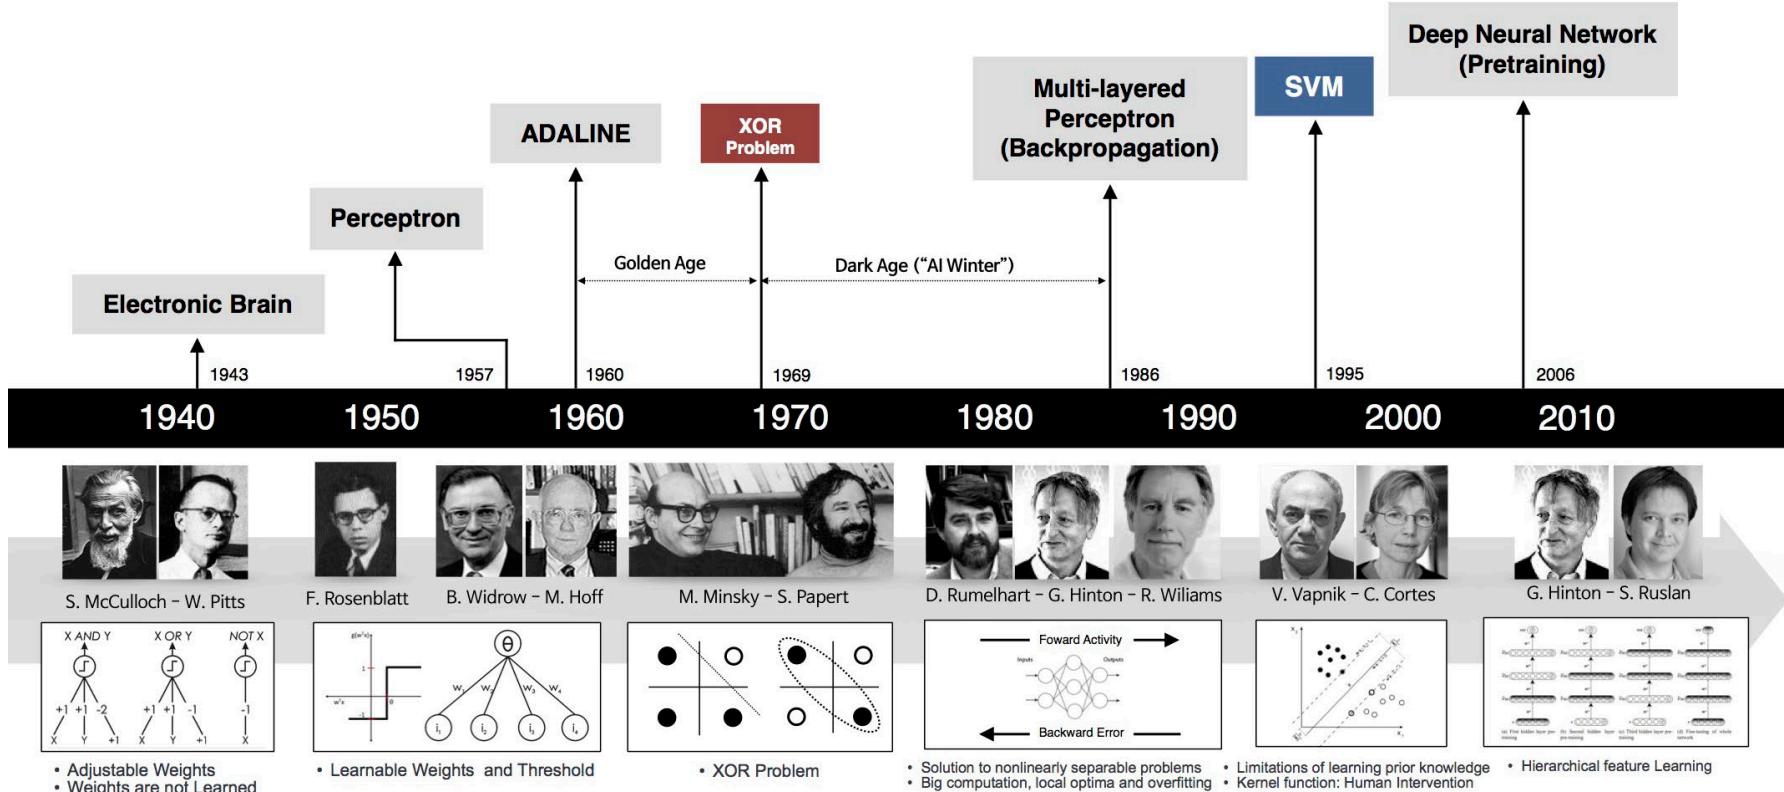
Trenchard More

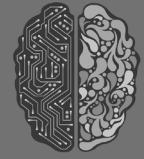
Image source: <https://www.scienceabc.com/innovation/what-is-artificial-intelligence.html>





Timeline of AI Development





Views of AI

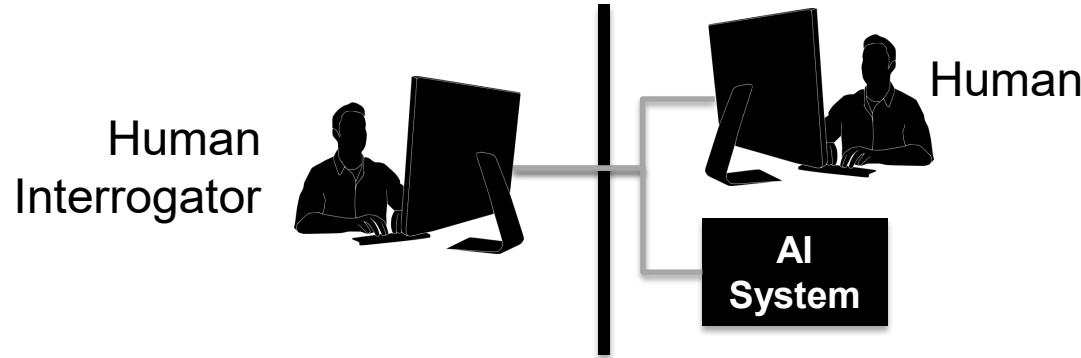
Views of AI fall into four categories:

1. Thinking Humanly
2. Acting Humanly
3. Thinking Rationally
4. Acting Rationally



Acting Humanly: Turing Test

- Turing (1950) "Computing machinery and intelligence"
- Operational test for intelligent behavior: the Imitation Game



- Suggested major components of AI: knowledge, reasoning, language understanding, learning

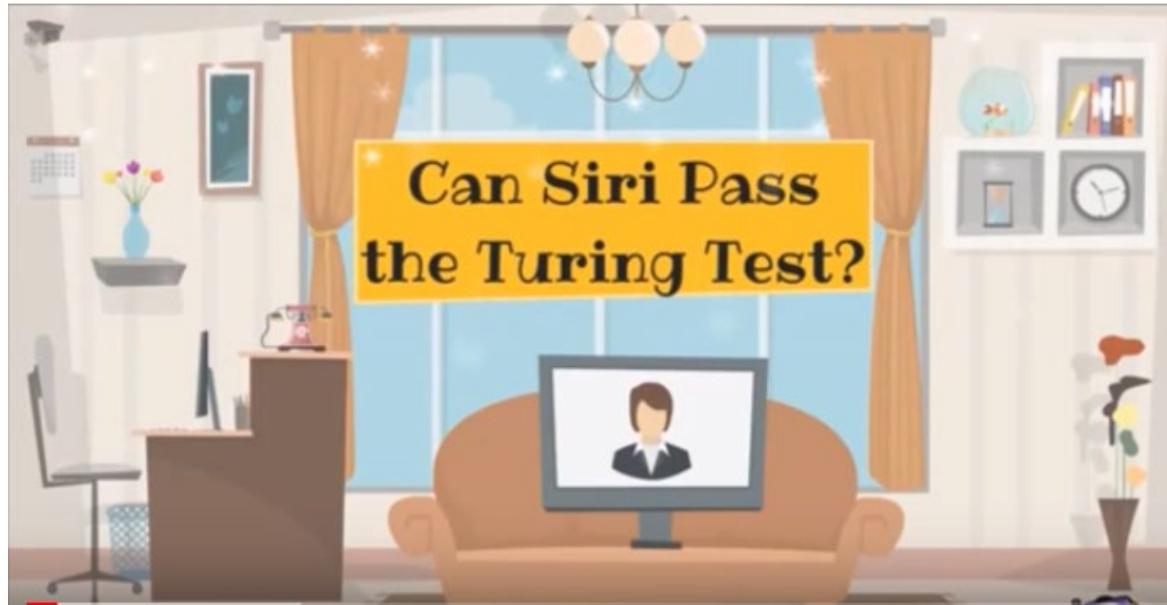
Image source: <https://svgsilh.com/image/2763393.html>





Can Siri Pass The Turing Test?

<https://www.youtube.com/watch?v=0qcHhBSbyVw>



<Video from previous slide will be
inserted in learning sequence for
students to watch>





Thinking Humanly: Cognitive Modeling

- 1960s "cognitive revolution": information-processing psychology
- Requires scientific theories of internal activities of the brain
- Both approaches (Cognitive Science and Cognitive Neuroscience) are now distinct from AI

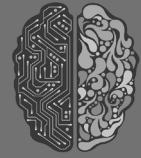




Thinking Rationally: "Laws of Thought"

- Aristotle: what are correct arguments/thought processes?
 - Several Greek schools developed various forms of logic: notation and rules of derivation for thoughts; may or may not have proceeded to the idea of mechanization
- Direct line through mathematics and philosophy to modern AI
- Problems:
 - Not all intelligent behavior is mediated by logical deliberation
 - What is the purpose of thinking? What thoughts should I have?



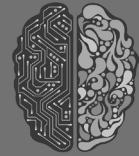


Acting Rationally: Rational Agents

- **Rational** behavior: doing the right thing
- The right thing: that which is expected to maximise goal achievement, given the available information
- This course is about designing rational agents
- For any given class of environments and tasks, we seek the agent(s) with the best performance

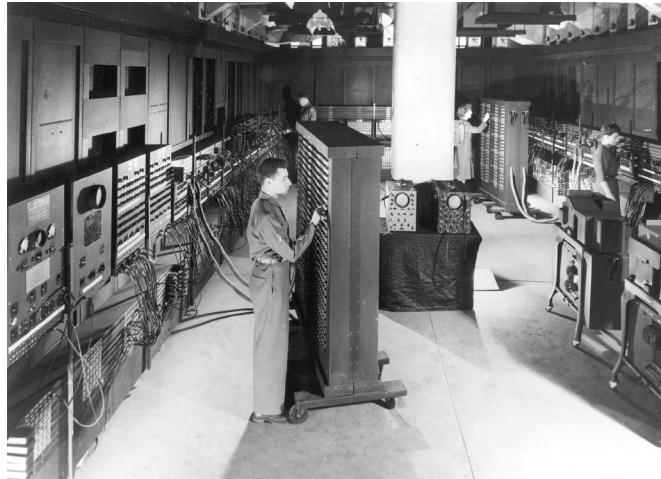






Computer Chess

Garry Kasparov VS Deep Blue



ENIAC 1946

VS



Deep Blue 1997

Image Sources:

<https://www.flickr.com/photos/gageskidmore/37137186253>

[https://en.wikipedia.org/wiki/Deep_Blue_\(chess_computer\)#/media/File:Deep_Blue.jpg](https://en.wikipedia.org/wiki/Deep_Blue_(chess_computer)#/media/File:Deep_Blue.jpg)





Robot Soccer

TED talk by
Peter Stone
from UT Austin

[http://www.youtube.com/
watch?v=FXhw0_iKwQ](http://www.youtube.com/watch?v=FXhw0_iKwQ)

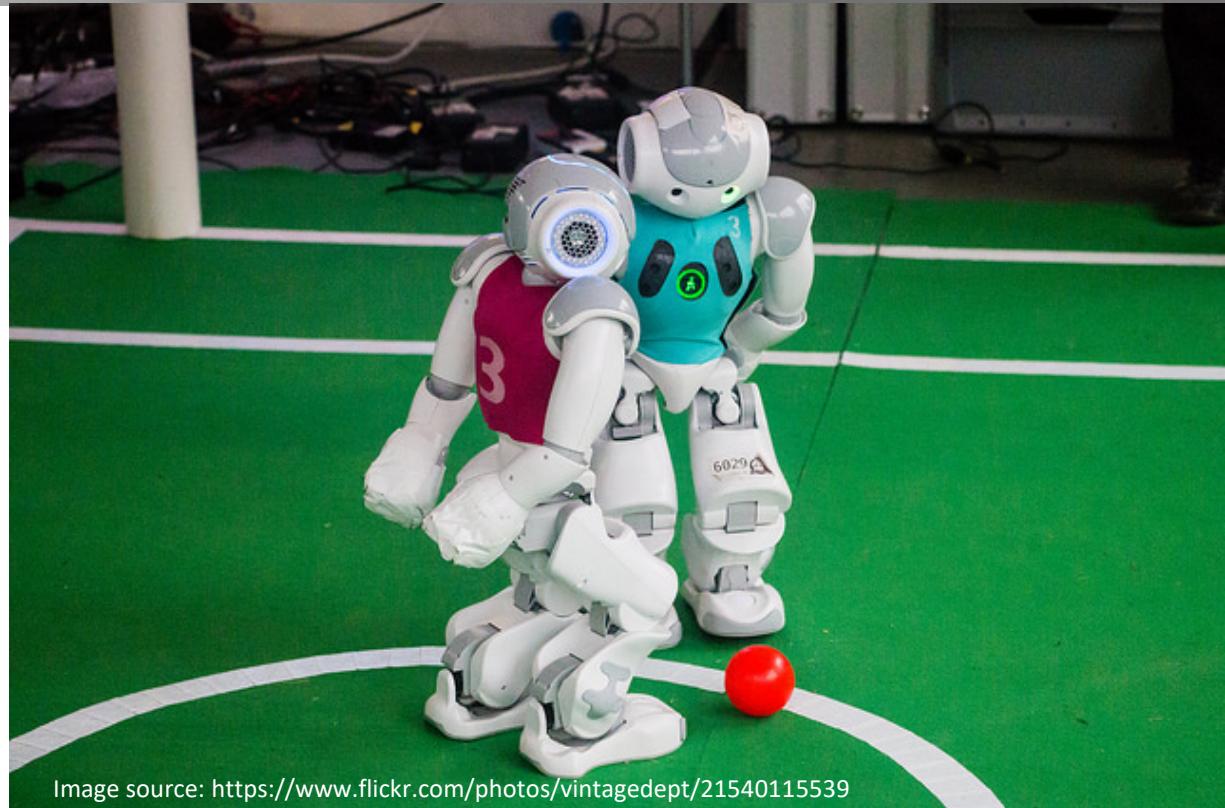


Image source: <https://www.flickr.com/photos/vintagedept/21540115539>



<Video from previous slide will be
inserted in learning sequence for
students to watch>





Game Show



Image source: https://commons.wikimedia.org/wiki/File:IBM_Watson_w_Jeopardy.jpg

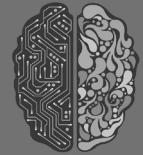
IBM's Watson
Destroys Humans in
Jeopardy

[https://www.youtube.com/watch?
v=P18EdAKuC1U](https://www.youtube.com/watch?v=P18EdAKuC1U)



<Video from previous slide will be
inserted in learning sequence for
students to watch>





AlphaGo vs World Champions

March 9 – 15, 2016 (Lee Sedol)

- Time limit: 2 hours
- Venue: Seoul, Four Seasons Hotel
- AlphaGo Wins (4:1)

May 23 – 27, 2017 (Ke Jie)

- Venue: Wuzhen, China
- AlphaGo Wins (3:0)

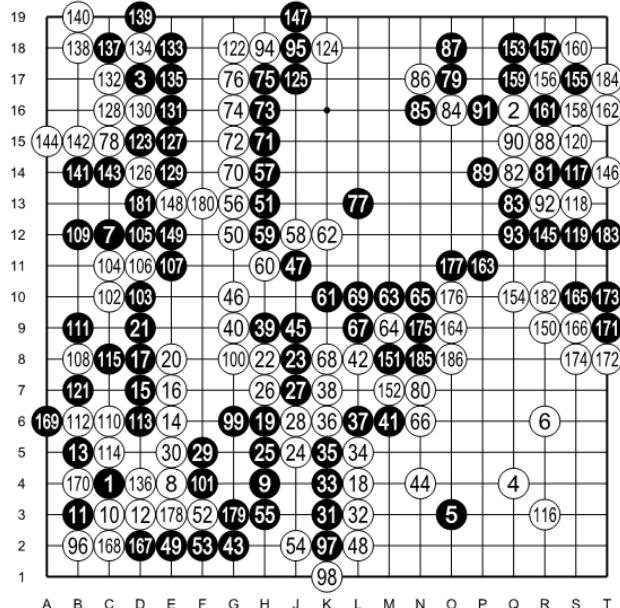


Image source: [https://commons.wikimedia.org/wiki/File:Lee_Sedol_\(B\)_vs_AlphaGo_\(W\)_-_Game_1_-_BW.jpg](https://commons.wikimedia.org/wiki/File:Lee_Sedol_(B)_vs_AlphaGo_(W)_-_Game_1_-_BW.jpg)





Libratus vs World Champions

The first AI to defeat top human poker players

January 11 to 31, 2017

- Venue: Pittsburgh
- 120,000 hands
- Has nothing to do with deep learning
- Algorithms for solving large scale games







Google Driverless Car

TED talk by Sebastian Thrun from Stanford

[https://www.youtube.com/
watch?v=bp9KBrH8H04](https://www.youtube.com/watch?v=bp9KBrH8H04)



Image source: https://commons.wikimedia.org/wiki/File:Google_driverless_car_at_intersection.gk.jpg



<Video from previous slide will be
inserted in learning sequence for
students to watch>





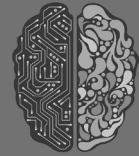
Google Robot Dog



<https://www.youtube.com/watch?v=4NzcB6TMzjw>

Image source: https://commons.wikimedia.org/wiki/File:Bio-inspired_Big_Dog_quadruped_robot_is_being_developed_as_a_mule_that_can_traverse_difficult_terrain.tiff





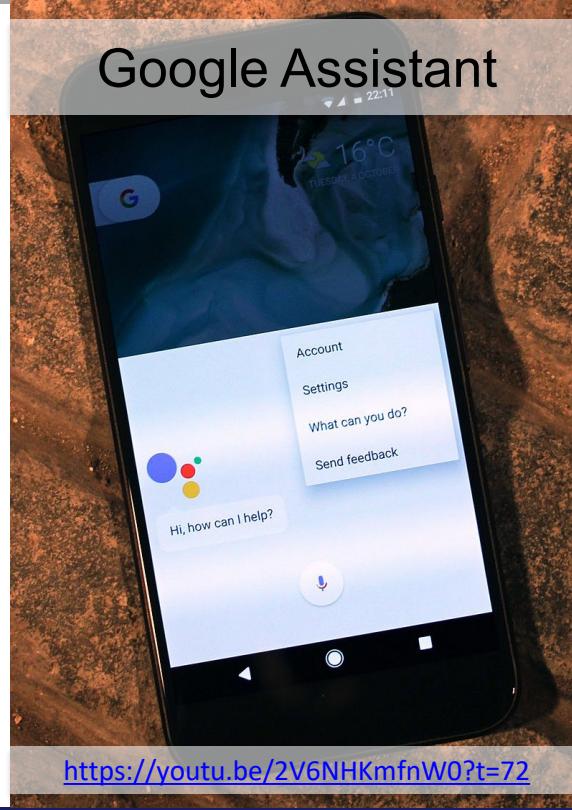
Natural Language Processing

Mi Box Voice Assistant



<https://youtu.be/S38ZJELrMmQ?t=16>

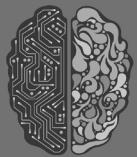
Google Assistant



<https://youtu.be/2V6NHKmfW0?t=72>

Alexa Amazon





Manufacturing

Tesla

[https://www.youtube.com/
watch?v=8_lfxPI5ObM](https://www.youtube.com/watch?v=8_lfxPI5ObM)



Image source: <https://www.flickr.com/photos/jurvetson/7408451314>



<Video from previous slide will be
inserted in learning sequence for
students to watch>





Artificial Intelligence for Finance



Source: World Economic Forum, 2015

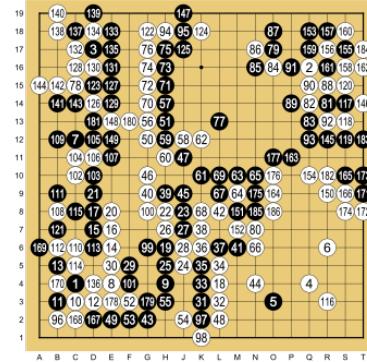


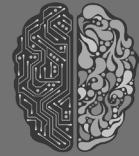


AI for Complex Interactions

Recent AI breakthrough

IMAGENET





What's Next?

For AI Complex Interactions

- Stochastic, open environment
- Multiple players
- Sequential decision, online
- Strategic (selfish) behaviour
- Distributed optimisation

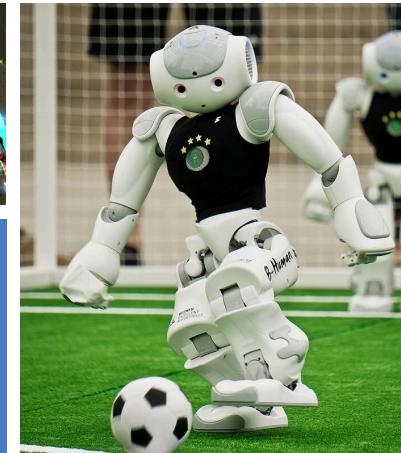
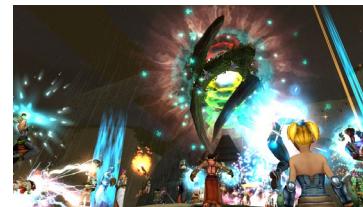


Image sources:

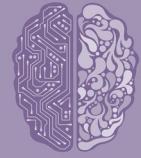
- https://en.wikipedia.org/wiki/RoboCup_Standard_Platform_League#/media/File:SPL_Team_B-Human,_RoboCup_2016.jpg
- https://en.wikipedia.org/wiki/Massively_multiplayer_online_role-playing_game#/media/File:Ryzom_screenshot_4.png
- <https://pixabay.com/en/network-iot-internet-of-things-782707/>
- <https://www.flickr.com/photos/thebetterday4u/37556363291>
- [https://en.wikipedia.org/wiki/Google_Assistant#/media/File:Android_Assistant_on_the_Google_Pixel_XL_smartphone_\(29526761674\).jpg](https://en.wikipedia.org/wiki/Google_Assistant#/media/File:Android_Assistant_on_the_Google_Pixel_XL_smartphone_(29526761674).jpg)
- <https://www.flickr.com/photos/stockcatalog/40095307924>
- <https://pixabay.com/en/woman-people-coffee-portrait-3083379/>
- <https://pixabay.com/en/microphone-audio-micro-recording-2104091/>
- <https://www.flickr.com/photos/gleonhard/34046646705>
- [https://commons.wikimedia.org/wiki/File:Lee_Sedol_\(B\)_vs_AlphaGo_\(W\)_-_Game_1.svg](https://commons.wikimedia.org/wiki/File:Lee_Sedol_(B)_vs_AlphaGo_(W)_-_Game_1.svg)

Introduction to Data Science and Artificial Intelligence **Intelligent Agents**

Assoc Prof Bo AN

Research area: artificial intelligence,
computational game theory, optimization
www.ntu.edu.sg/home/boan
Email: boan@ntu.edu.sg
Office: N4-02b-55





Lesson Outline

- What is an agent?
- How can one describe the task/problem for the agent?
- What are the properties of the task environment for the agent?
- What are the different basic kinds of agents for intelligent systems?
- Problem formulation

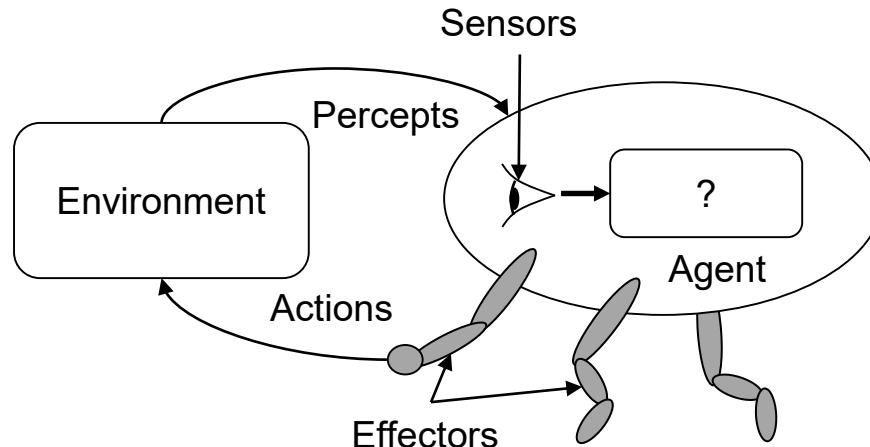


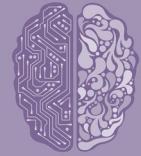


Agent

An **agent** is an entity that

- **Perceives** through sensors (e.g. eyes, ears, cameras, infrared range sensors)
- **Acts** through effectors (e.g. hands, legs, motors)





Rational Agents

- A rational agent is one that does the **right** thing
- **Rational action**: action that maximises the expected value of an objective **performance** measure given the percept sequence to date
- Rationality depends on
 - performance measure
 - everything that the agent has perceived so far
 - built-in knowledge about the environment
 - actions that can be performed





Example: Google X2: Driverless Taxi

- **Percepts:** video, speed, acceleration, engine status, GPS, radar, ...
- **Actions:** steer, accelerate, brake, horn, display, ...
- **Goals:** safety, reach destination, maximise profits, obey laws, passenger comfort, ...
- **Environment:** Singapore urban streets, highways, traffic, pedestrians, weather, customers, ...



Image source: https://en.wikipedia.org/wiki/Waymo#/media/File:Waymo_Chrysler_Pacifica_in_Los_Altos,_2017.jpg





Example: Medical Diagnosis System

- **Percepts:** symptoms, findings, patient's answers, ...
- **Actions:** questions, medical tests, treatments, ...
- **Goals:** healthy patient, faster recovery, minimise costs, ...
- **Environment:** Patient, hospital, clinic, ...

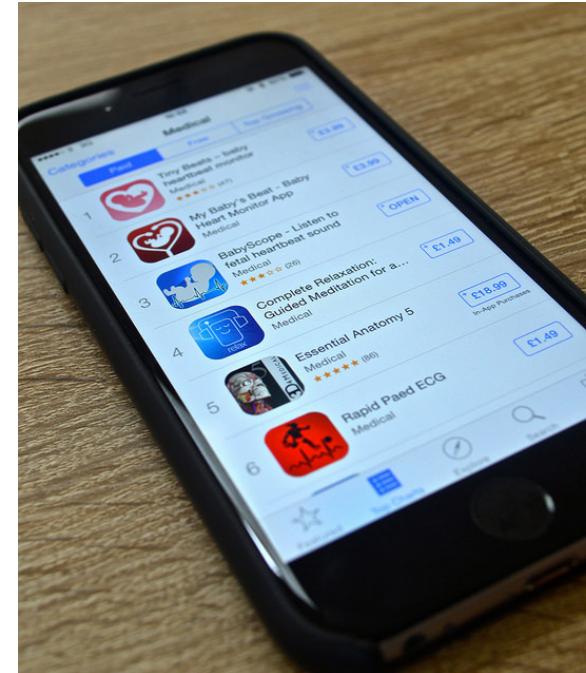
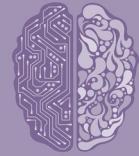


Image source: <https://www.flickr.com/photos/134647712@N07/20008817459>





Autonomous Agents

- Do **not** rely entirely on built-in knowledge about the environment (i.e. not entirely pre-programmed)
- Otherwise,
 - The agent will **only** operate successfully when the built-in knowledge are all correct
 - Adapt to the environments through experience

Example: Driverless Car

- Learn to drive in driving center
- Drive at NTU
- Drive on public roads
- Drive in highways
- Drive in City Hall



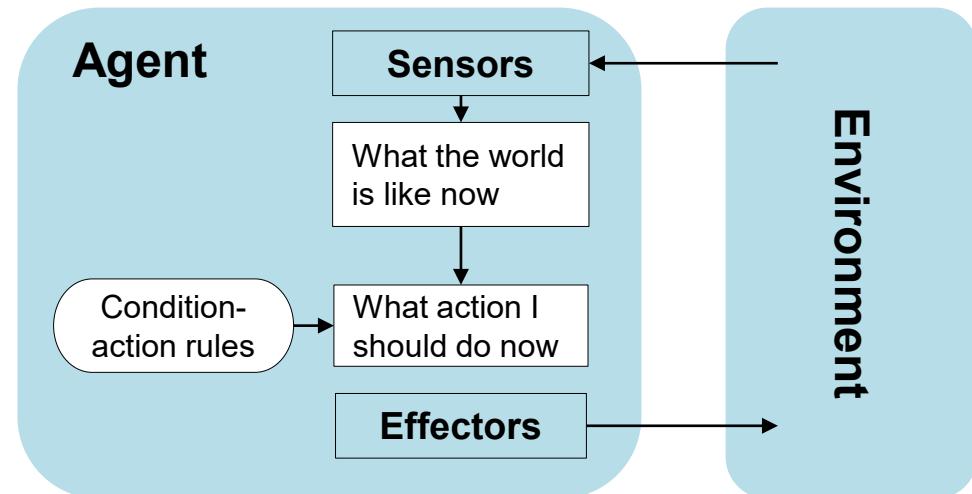


Simple Reflex Agents

Example

If car-in-front-is-braking then initiate-braking

1. Find the **rule** whose condition matches the current situation (as defined by the percept)
2. Perform the action associated with that rule



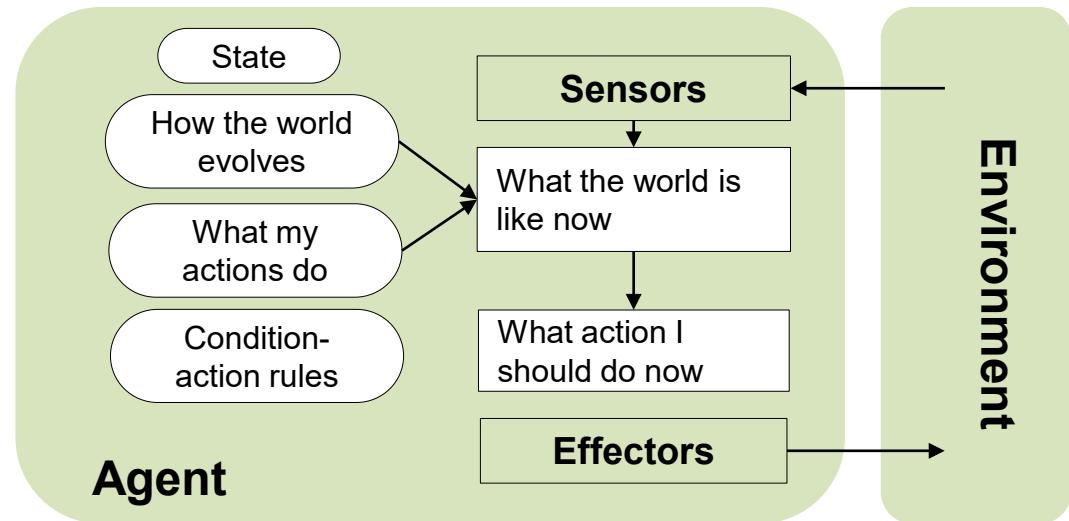


Reflex Agents with State

Example

If yesterday-at-NTU and no-traffic-jam-now then go-Orchard

1. Find the rule whose condition matches the current situation (as defined by the percept and the stored internal state)
2. Perform the action associated with that rule



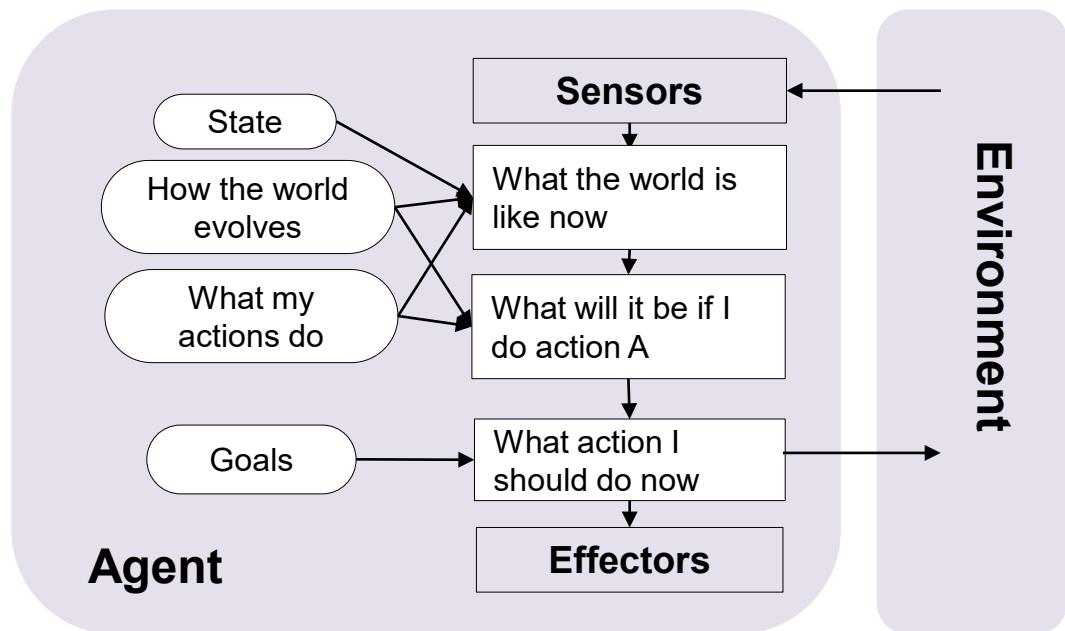


Goal-Based Agents

Needs some sort of **goal** information

Example: Driverless Taxi

- At a junction (known state), should I go left, right, or straight on?
- Reach Orchard (Destination)?



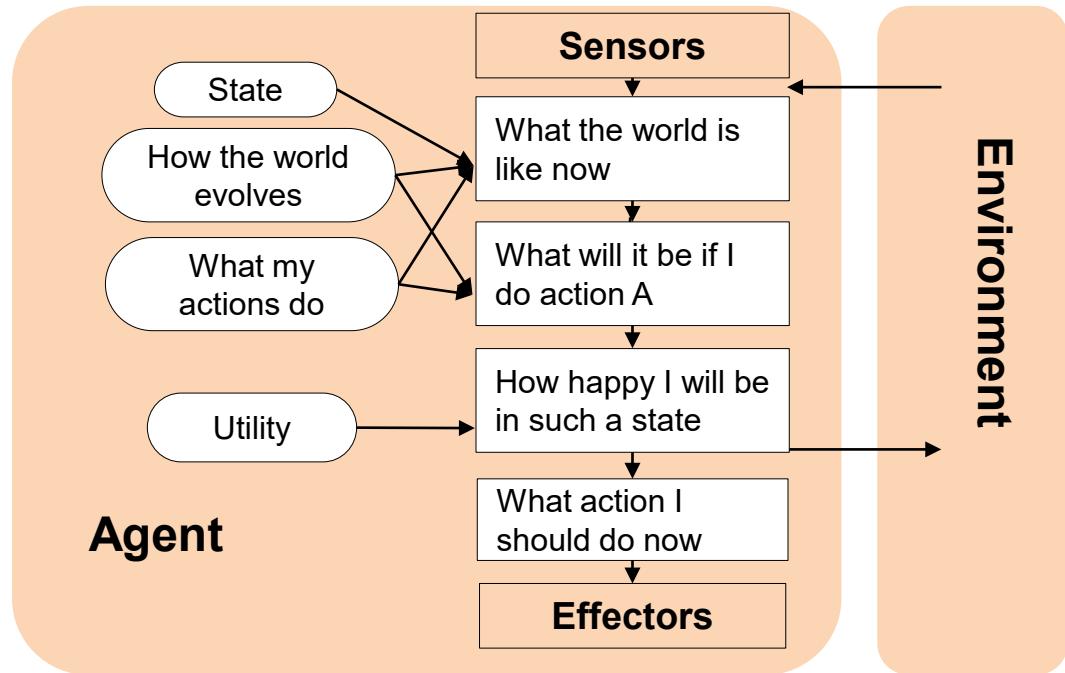


Utility-Based Agents

- There may be **many** action sequences that can achieve the same goal, which action sequence should it take?
- How happy will agent be if it attains a certain state? → Utility

Example: Driverless Taxi

- Go to Orchard (Destination) via PIE? AYE?
- Which one charges lower fare?



Question to think:

- Why is problem formulation (e.g., characterize percept, goals etc) important?







Types of Environment

Accessible (vs
inaccessible)

Agent's sensory apparatus gives it access to the complete state of the environment

Deterministic (vs
nondeterministic)

The next state of the environment is completely determined by the current state and the actions selected by the agent

Episodic (vs
Sequential)

Each episode is not affected by the previous taken actions

Static (vs
dynamic)

Environment does not change while an agent is deliberating

Discrete (vs
continuous)

A limited number of distinct percepts and actions

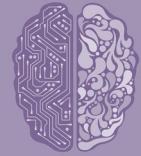




Example: Driverless Taxi

Accessible?	No. Some traffic information on road is missing
Deterministic?	No. Some cars in front may turn right suddenly
Episodic?	No. The current action is based on previous driving actions
Static?	No. When the taxi moves, Other cars are moving as well
Discrete?	No. Speed, Distance, Fuel consumption are in real domains

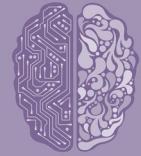




Example: Chess

Accessible?	Yes. All positions in chessboard can be observed
Deterministic?	Yes. The outcome of each movement can be determined
Episodic?	No. The action depends on previous movements
Static?	Yes. When there is no clock, when are you considering the next step, the opponent can't move; Semi . When there is a clock, and time is up, you will give up the movement
Discrete?	Yes. All positions and movements are in discrete domains

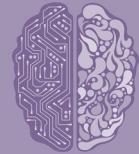




Example: Minesweeper

Accessible?	No. Mines are hidden
Deterministic?	No. Mines are randomly assigned in different positions
Episodic?	No. The action is based on previous outcomes
Static?	Yes. When are you considering the next step, no changes in environment
Discrete?	Yes. All positions and movements are in discrete domains





Slot machines - One-Armed Bandit



Accessible?	No
Deterministic?	No
Episodic?	Yes
Static?	Yes

Image source: https://commons.wikimedia.org/wiki/File:Las_Vegas_slot_machines.jpg

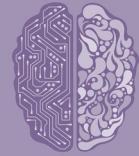


Question to think:

- Can you find another scenario whose environment is episodic?







Design of Problem-Solving Agent

Idea

- Systematically considers the **expected outcomes** of different possible sequences of actions that lead to states of known value
- Choose the best one
 - shortest journey from A to B?
 - most cost effective journey from A to B?





Design of Problem-Solving Agent

Steps

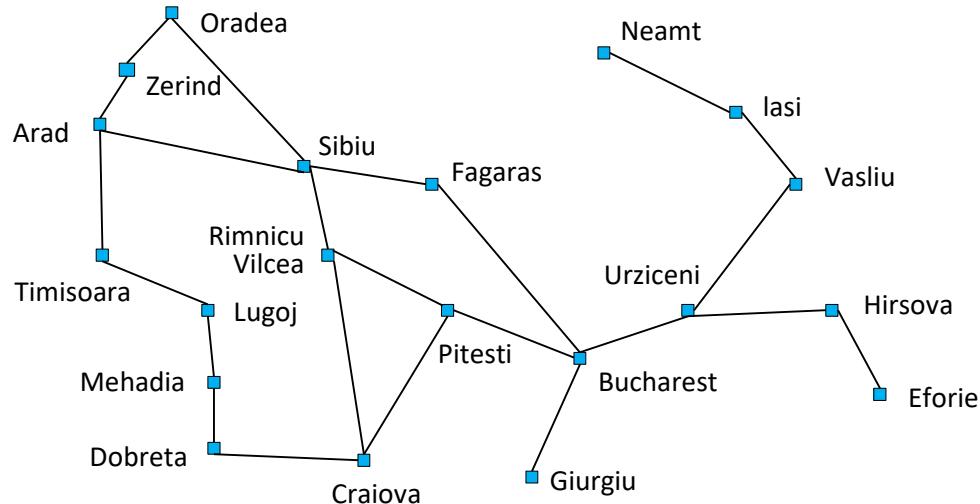
1. Goal formulation
2. Problem formulation
3. Search process
 - No knowledge → uninformed search
 - Knowledge → informed search
4. Action execution (follow the recommended route)





Goal-based Agent: Example

On holiday in Romania



- Currently in Arad (**Initial state**). Flight leaves tomorrow from Bucharest.
- **Goal:** be in Bucharest (other factors: cost, time, most scenic route, etc.)
- **State:** be in a city (defined by the map)
- **Action:** transition between states (highways defined by the map)



Example: Vacuum Cleaner Agent

- Robotic vacuum cleaners move autonomously
- Some can come back to a docking station to charge their batteries
- A few are able to empty their dust containers into the dock as well



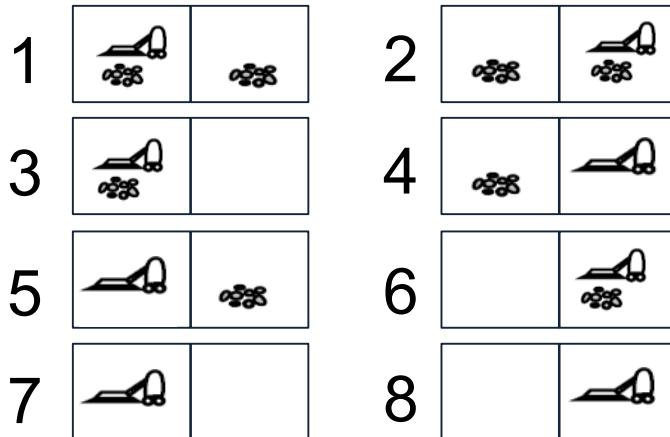
Image source: https://commons.wikimedia.org/wiki/File:iRobot_Roomba_780.jpg



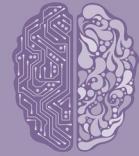


Example: A Simple Vacuum World

Two locations, each location may or may not contain dirt, and the agent may be in one location or the other.



- 8 possible world states
- Possible actions: **left**, **right**, and **suck**
- Goal: clean up all dirt → Two goal states, i.e. {7, 8}



Well-Defined Formulation

Definition of a problem	The information used by an agent to decide what to do
Specification	<ul style="list-style-type: none">• Initial state• Action set, i.e. available actions (successor functions)• State space, i.e. states reachable from the initial state<ul style="list-style-type: none">• Solution path: sequence of actions from one state to another• Goal test predicate<ul style="list-style-type: none">• Single state, enumerated list of states, abstract properties• Cost function<ul style="list-style-type: none">• Path cost $g(n)$, sum of all (action) step costs along the path
Solution	A path (a sequence of operators leading) from the Initial-State to a state that satisfies the Goal-Test





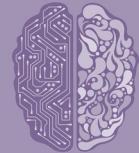
Measuring Problem-Solving Performance

Search Cost

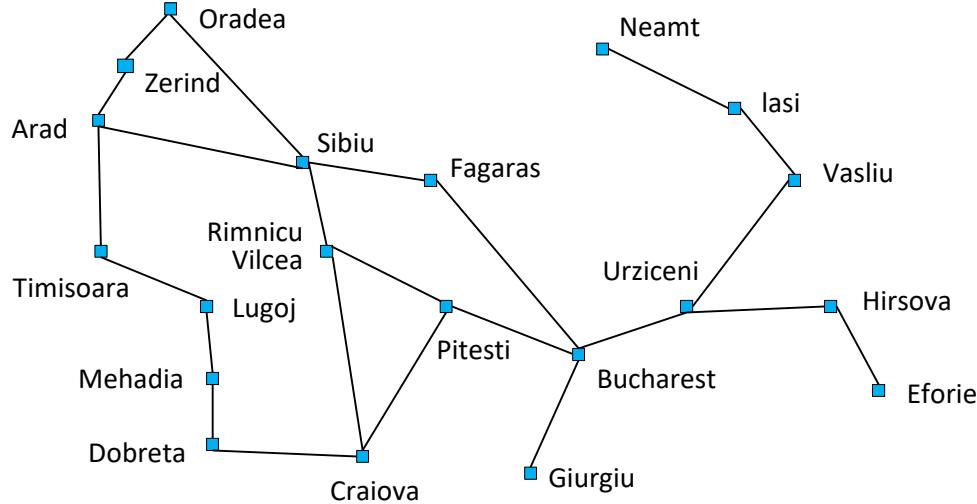
- What does it cost to find the solution?
 - e.g. How long (time)? How many resources used (memory)?

Total cost of problem-solving

- Search cost ("offline") + Execution cost ("online")
- Trade-offs often required
 - Search a very long time for the optimal solution, or
 - Search a shorter time for a "good enough" solution



Single-State Problem Example

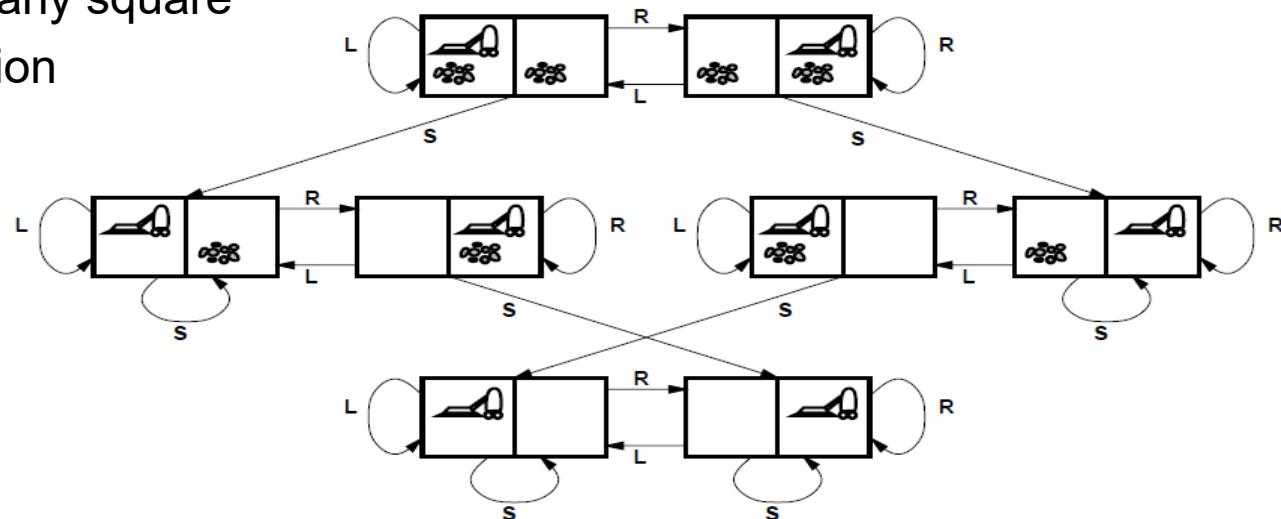


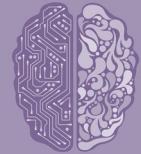
- **Initial state:** e.g., “at Arad”
- Set of possible **actions** and the corresponding next states
 - e.g., Arad → Zerind
- **Goal test:**
 - explicit (e.g., $x = \text{"at Bucharest"}$)
- **Path cost** function
 - e.g., sum of distances, number of operators executed solution: a sequence of operators leading from the initial state to a goal state

Example: Vacuum World (Single-state Version)



- **Initial state:** one of the eight states shown previously
- **Actions:** left, right, suck
- **Goal test:** no dirt in any square
- **Path cost:** 1 per action





Example: 8-puzzle

- **States:** integer locations of tiles
 - number of states = $9!$
- **Actions:** move blank left, right, up, down
- **Goal test:** = goal state (given)
- **Path cost:** 1 per move

Start state

5	4	
6	1	8
7	3	2

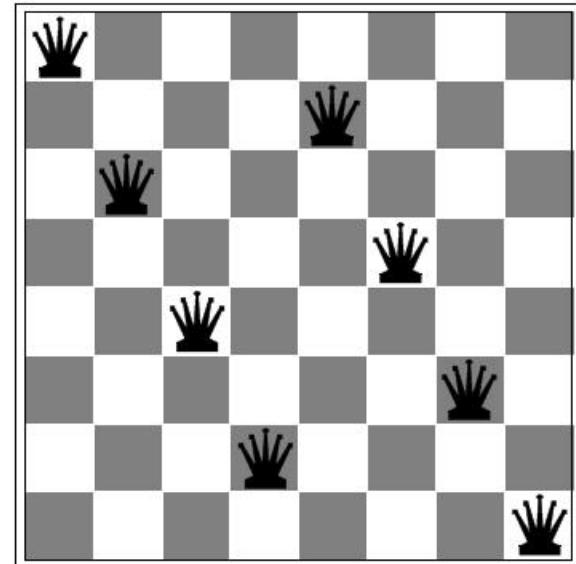
Goal state

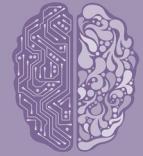
1	2	3
8		4
7	6	5



Example: 8-queens

- **States:** Any arrangement of 0 to 8 queens on the board
- **Actions:** Add a queen to any empty square
- **Goal test:** 8 queens are on the board, none attacked
- **Path cost:** Not necessary





Real-World Problems

Route finding problems:

- Routing in computer networks
- Robot navigation
- Automated travel advisory
- Airline travel planning

Touring problems:

- Traveling Salesperson problem
- “Shortest tour”: visit every city exactly once



Image source: https://commons.wikimedia.org/wiki/File:Font_Awesome_5_solid_route.svg



Introduction to Data Science and Artificial Intelligence

Solving Problems by Search: Uninformed Search and Informed Search

Assoc Prof Bo AN

Research area: artificial intelligence,
computational game theory, optimization
www.ntu.edu.sg/home/boan
Email: boan@ntu.edu.sg
Office: N4-02b-55





Lesson Outline

- Uninformed search strategies
- Informed search strategies
 - Greedy search
 - A * search





Review: Well-Defined Formulation

Definition of a problem	The information used by an agent to decide what to do
Specification	<ul style="list-style-type: none">Initial stateAction set, i.e. available actions (successor functions)State space, i.e. states reachable from the initial state<ul style="list-style-type: none">Solution path: sequence of actions from one state to anotherGoal test predicate<ul style="list-style-type: none">Single state, enumerated list of states, abstract propertiesCost function<ul style="list-style-type: none">Path cost $g(n)$, sum of all (action) step costs along the path
Solution	A path (a sequence of operators leading) from the Initial-State to a state that satisfies the Goal-Test





Search Algorithms

- Exploration of state space by generating successors of already-explored states
 - Frontier: candidate nodes for expansion
 - Explored set

1

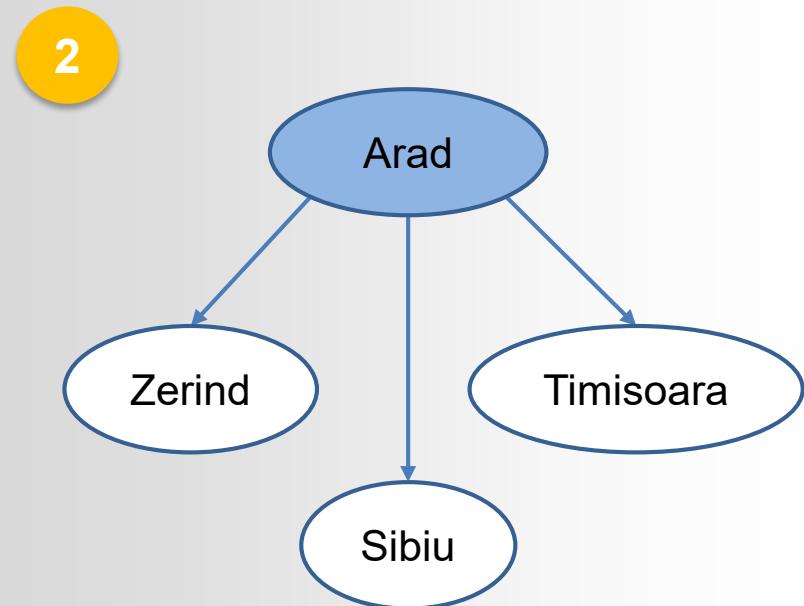
Arad





Search Algorithms

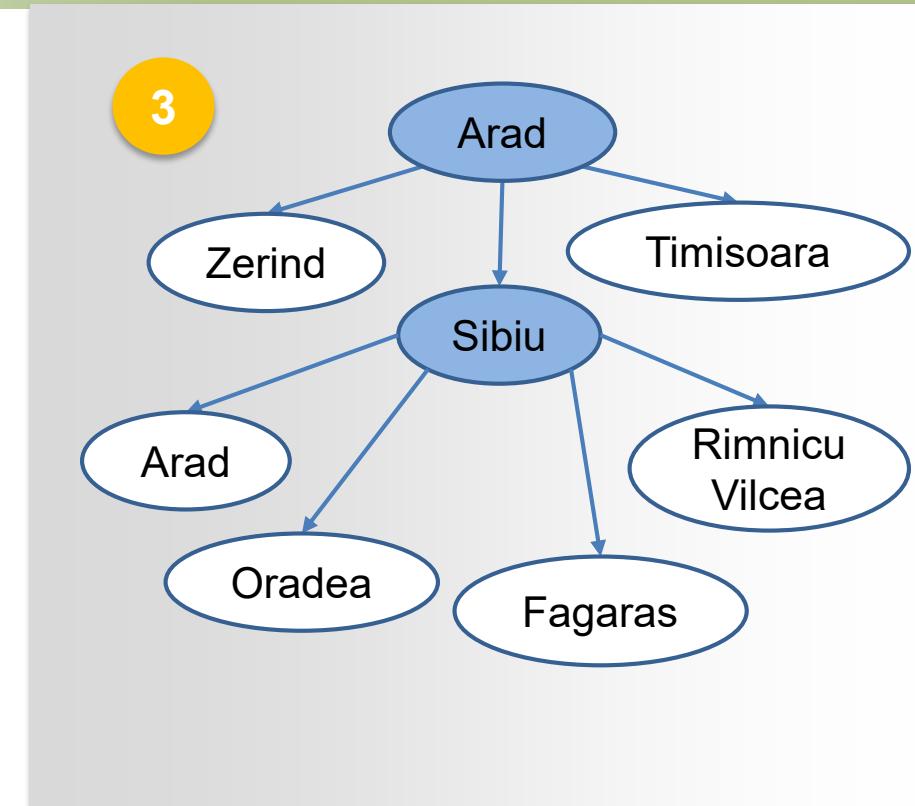
- Exploration of state space by generating successors of already-explored states
 - Frontier: candidate nodes for expansion
 - Explored set





Search Algorithms

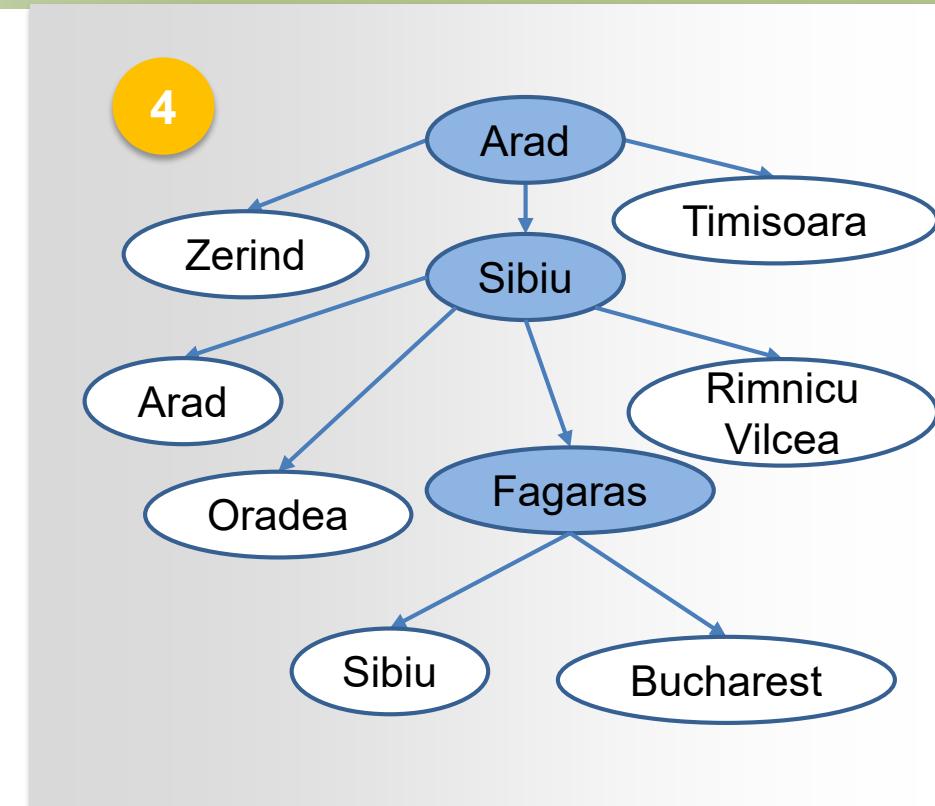
- Exploration of state space by generating successors of already-explored states
 - Frontier: candidate nodes for expansion
 - Explored set





Search Algorithms

- Exploration of state space by generating successors of already-explored states
 - Frontier: candidate nodes for expansion
 - Explored set





Search Strategies

- A **strategy** is defined by picking the order of node expansion.
- Strategies are evaluated along the following dimensions:

Completeness	Does it always find a solution if one exists?
Time Complexity	How long does it take to find a solution: the number of nodes generated
Space Complexity	Maximum number of nodes in memory
Optimality	Does it always find the best (least-cost) solution?



Search Strategies

- Branching factor
 - Maximum number of successors of any node
 - Or average branching factor



Uninformed vs Informed

Uninformed search strategies

- Use **only** the information available in the problem definition
 1. Breadth-first search
 2. Uniform-cost search
 3. Depth-first search
 4. Depth-limited search
 5. Iterative deepening search



Informed search strategies

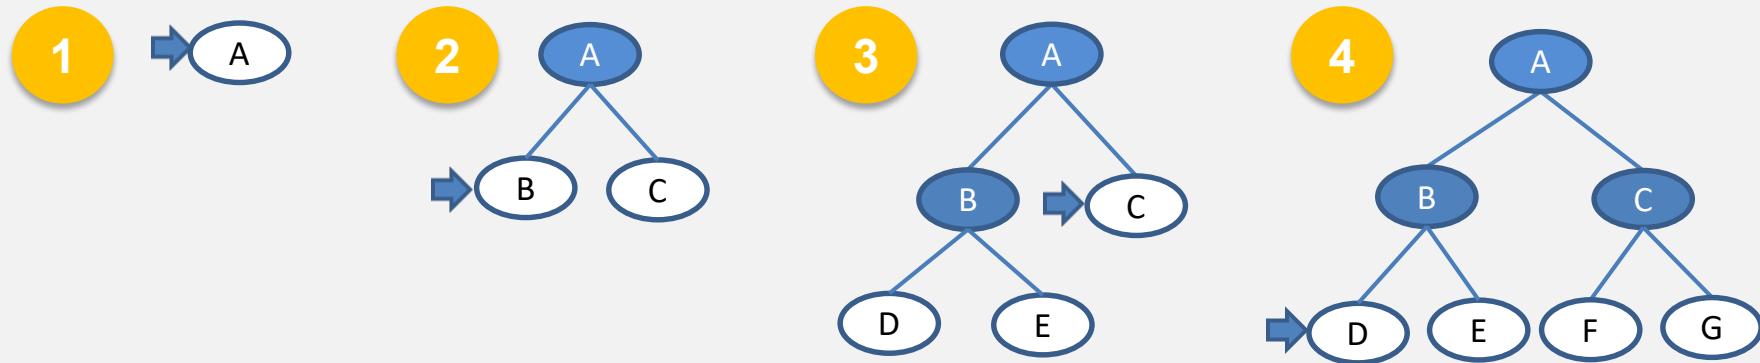
- Use **problem-specific knowledge** to guide the search
- Usually more efficient





Breadth-First Search

Expand **shallowest** unexpanded node which can be implemented by a First-In-First-Out (FIFO) queue



- Denote
- b: maximum branching factor of the search tree
 - d: depth of the least-cost solution
 - Complete: Yes
 - Optimal: Yes when **all** step costs equally





Complexity of BFS

- Hypothetical state-space, where every node can be expanded into b new nodes, solution of path-length d
- Time: $1 + b + b^2 + b^3 + \dots + b^d = O(b^d)$
- Space: (keeps every node in memory) $O(b^d)$ are equal

Depth	Nodes	Time		Memory	
0	1	1	millisecond	100	bytes
2	111	0.1	seconds	11	kilobytes
4	11111	11	seconds	1	kilobytes
6	10^6	18	minutes	111	megabyte
8	10^8	31	hours	11	gigabytes
10	10^{10}	128	days	1	terabyte
12	10^{12}	35	years	111	terabytes
14	10^{14}	3500	years	11111	terabytes

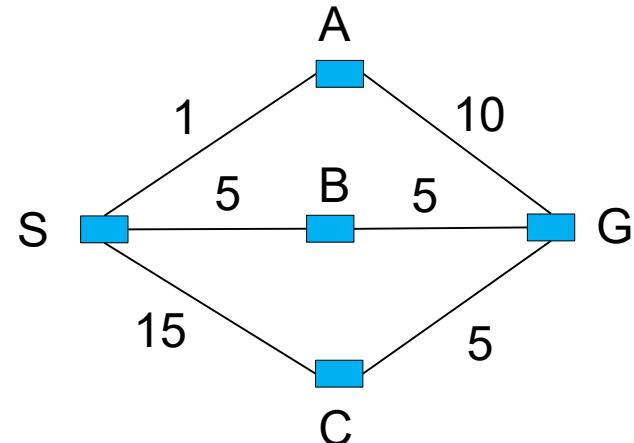




Uniform-Cost Search

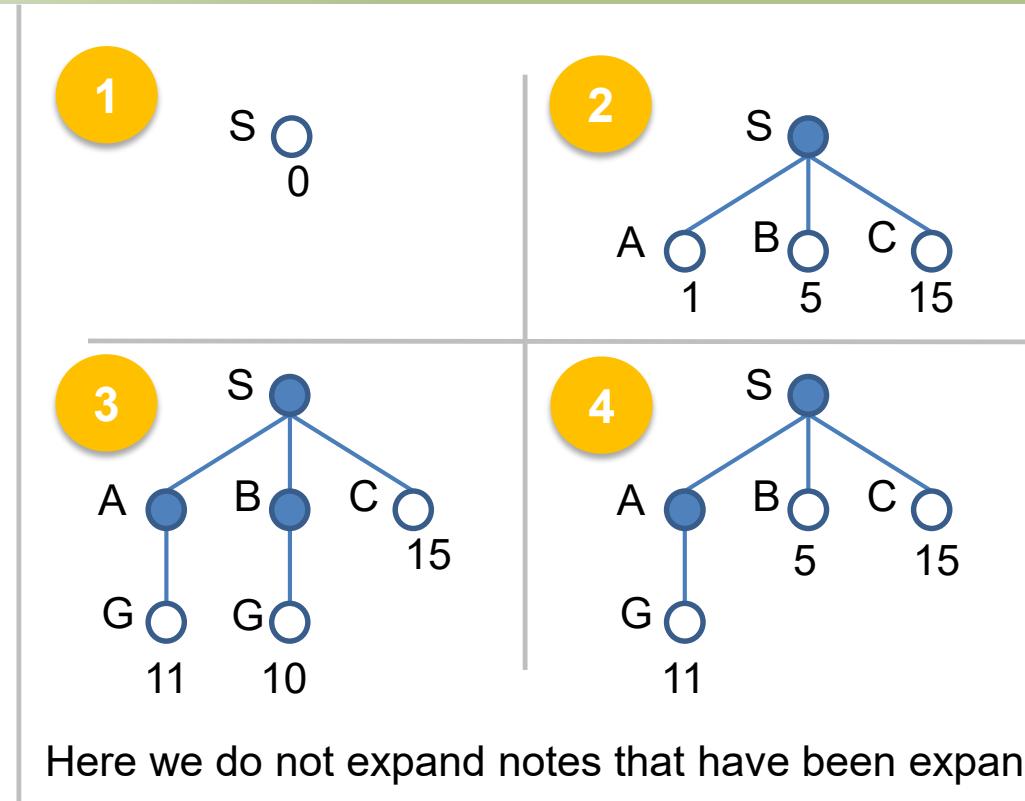
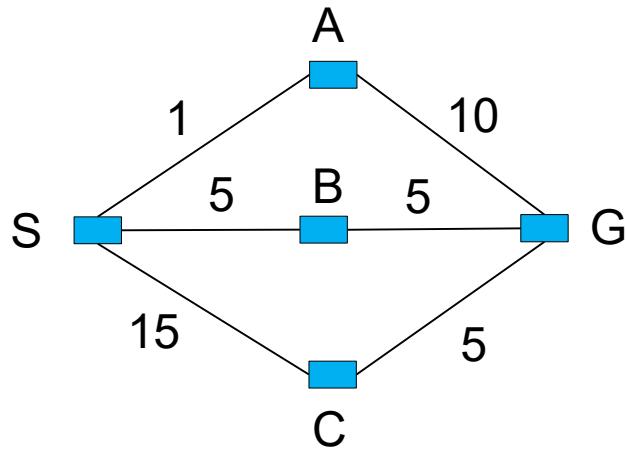
To consider **edge costs**, expand unexpanded node with the **least** path cost g

- Modification of breath-first search
- Instead of First-In-First-Out (FIFO) queue, using a priority queue with path cost $g(n)$ to order the elements
- BFS = UCS with $g(n) = \text{Depth}(n)$





Uniform-Cost Search





Uniform-Cost Search

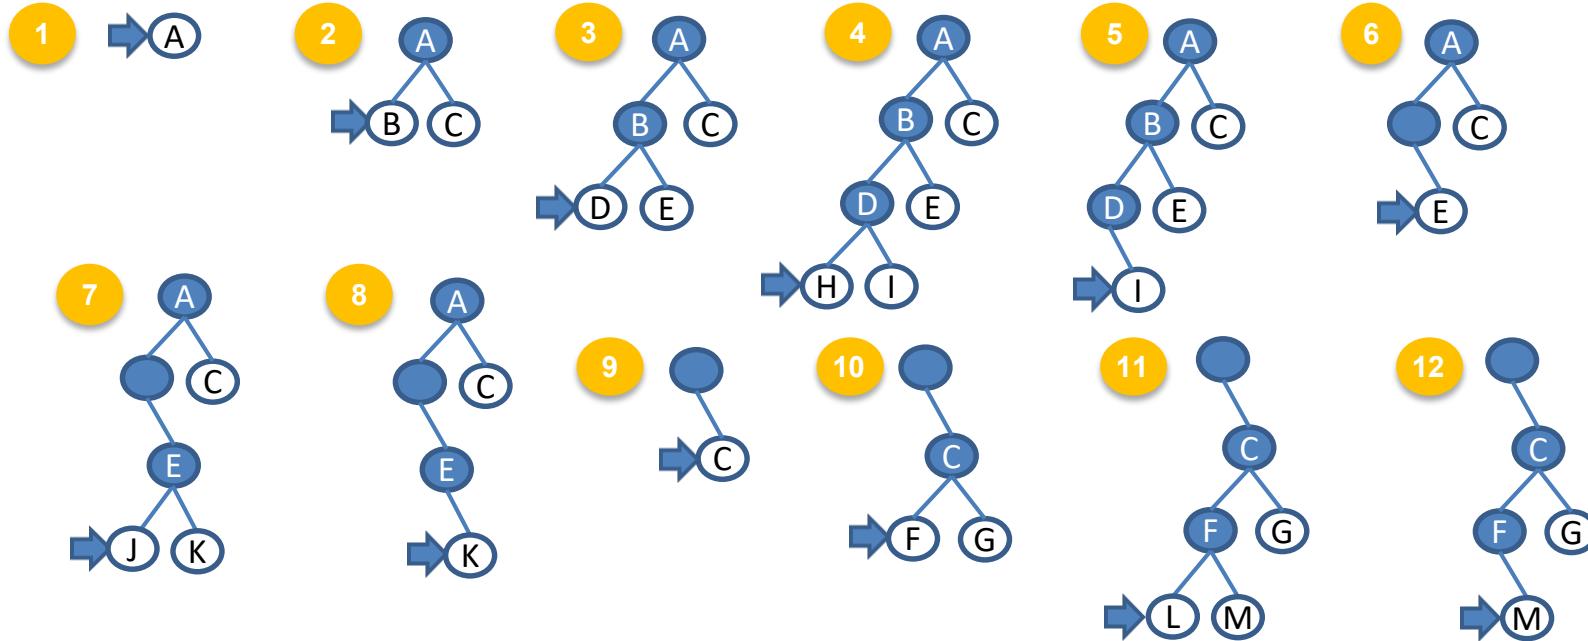
Complete	Yes
Time	# of nodes with path cost $g \leq$ cost of optimal solution (eqv. # of nodes pop out from the priority queue)
Space	# of nodes with path cost $g \leq$ cost of optimal solution
Optimal	Yes





Depth-First Search

Expand deepest unexpanded node which can be implemented by a Last-In-First-Out (LIFO) stack, Backtrack only when no more expansion





Depth-First Search

Denote

- m: maximum depth of the state space

Complete	<ul style="list-style-type: none">• infinite-depth spaces: No• finite-depth spaces with loops:<ul style="list-style-type: none">• with repeated-state checking: Yes• finite-depth spaces without loops: Yes
Time	$O(b^m)$ If solutions are dense, may be much faster than breadth-first
Space	$O(bm)$
Optimal	No





Depth-Limited Search

To avoid infinite searching, Depth-first search with a **cutoff** on the max depth / of a path

Complete	Yes, if $I \geq d$
Time	$O(b^I)$
Space	$O(bI)$
Optimal	No





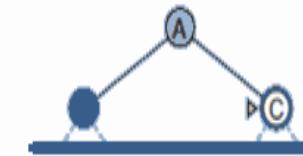
Iterative Deepening Search

Iteratively estimate the max depth / of DLS one-by-one

Limit = 0



Limit = 1

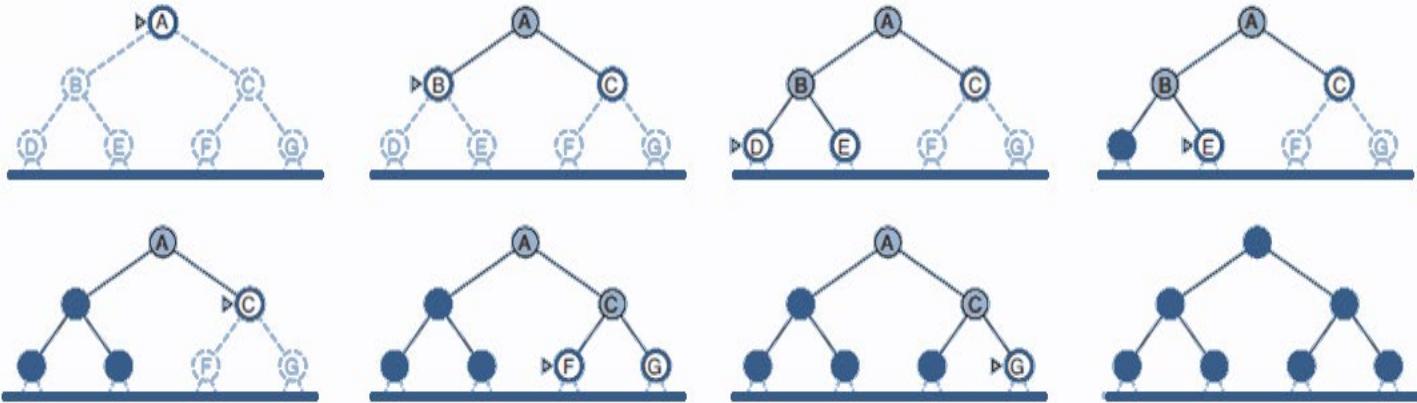




Iterative Deepening Search

Iteratively estimate the max depth / of DLS one-by-one

Limit = 2

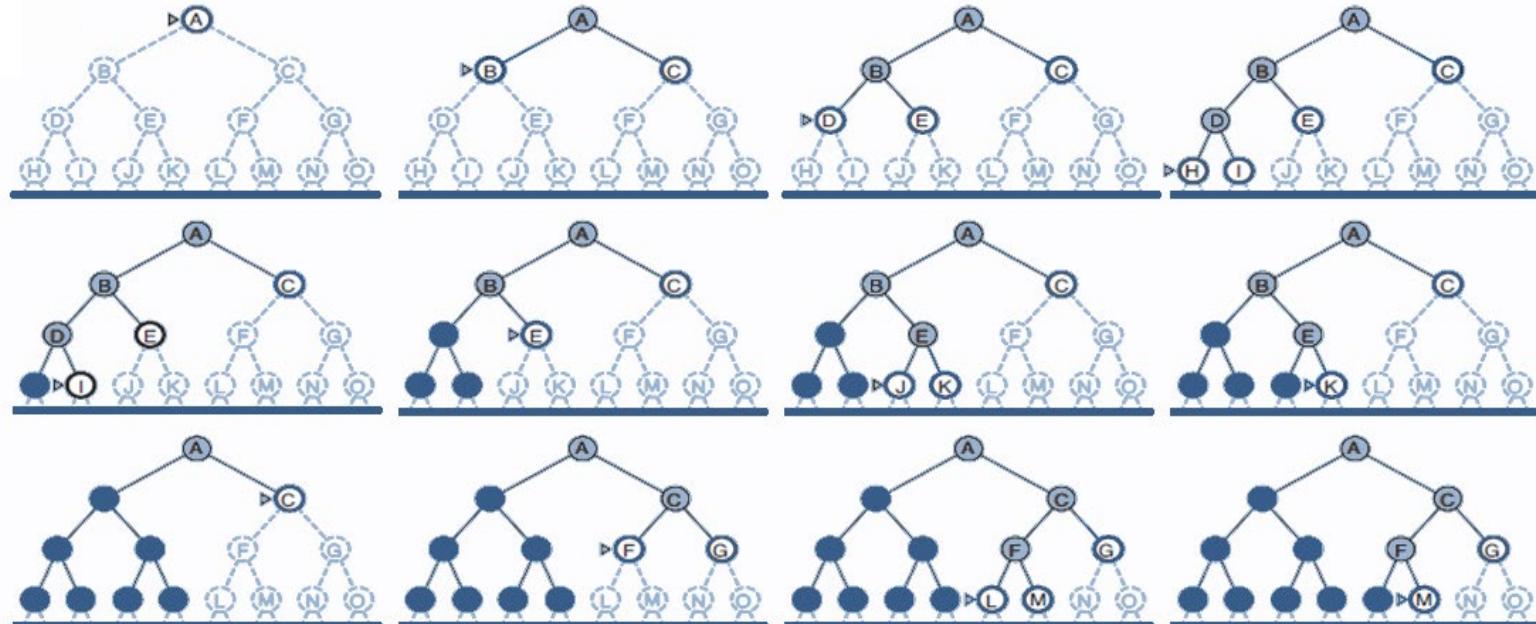




Iterative Deepening Search

Iteratively estimate the max depth / of DLS one-by-one

Limit = 3





Iterative Deepening Search...

```
Function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution sequence
  inputs: problem, a problem
```

```
    for depth 0 to  $\infty$  do
      if DEPTH-LIMITED-SEARCH(problem, depth) succeeds then return its result
    end
    return failure
```

Complete	Yes
Time	$O(b^d)$
Space	$O(bd)$
Optimal	Yes





Summary (we make assumptions for optimality)

Criterion	Breadth-first	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Time	b^d	b^d	b^m	b^l	b^d	$b^{d/2}$
Space	b^d	b^d	bm	bl	bd	$b^{d/2}$
Optimal	Yes	Yes	No	No	Yes	Yes
Complete	Yes	Yes	No	Yes, if $l \geq d$	Yes	Yes



Question to think:

- If a search strategy is optimal, is it also complete?







General Search

Uninformed search strategies

- Systematic generation of new states (→Goal Test)
- Inefficient (exponential space and time complexity)



Informed search strategies

- Use problem-specific knowledge
 - To decide the order of node expansion
- Best First Search: expand the most desirable unexpanded node
 - Use an evaluation function to estimate the “desirability” of each node





Evaluation function

- Path-cost function $g(n)$
 - Cost from initial state to current state (search-node) n
 - No information on the cost toward the goal
- Need to estimate cost to the closest goal
- “Heuristic” function $h(n)$
 - Estimated cost of the cheapest path from n to a goal state $h(n)$
 - Exact cost cannot be determined
 - depends only on the state at that node
 - $h(n)$ is not larger than the real cost (admissible)



Greedy Search

Expands the node that **appears** to be closest to goal

- Evaluation function $h(n)$: estimate of cost from n to **goal**
- Function Greedy-Search(problem) returns solution
 - Return Best-First-Search(problem, h) // $h(goal) = 0$

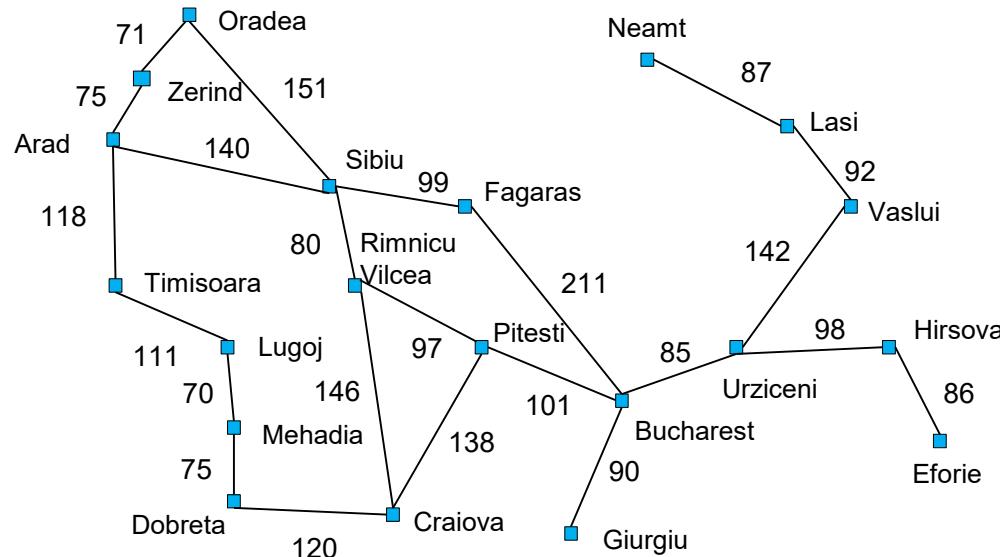
Question: How to estimation the cost from **n** to **goal**?

Answer: Recall that we want to use problem-specific knowledge



Example: Route-finding from Arad to Bucharest

$h(n)$ = straight-line distance from n to Bucharest



- Useful but potentially fallible (**heuristic**)
 - Heuristic functions are problem-specific

Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Efoire	161
Fagaras	176
Giurgiu	77
Hirsova	151
Lasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374



Example

a) The initial state



366

Straight-line distance to Bucharest

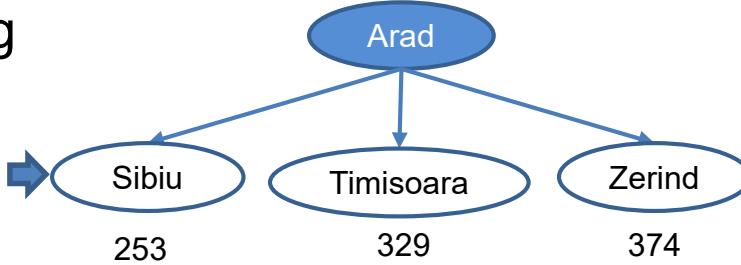
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Efoire	161
Fagaras	176
Giurgiu	77
Hirsova	151
Lasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374





Example

b) After
expanding
Arad



Straight-line distance to Bucharest

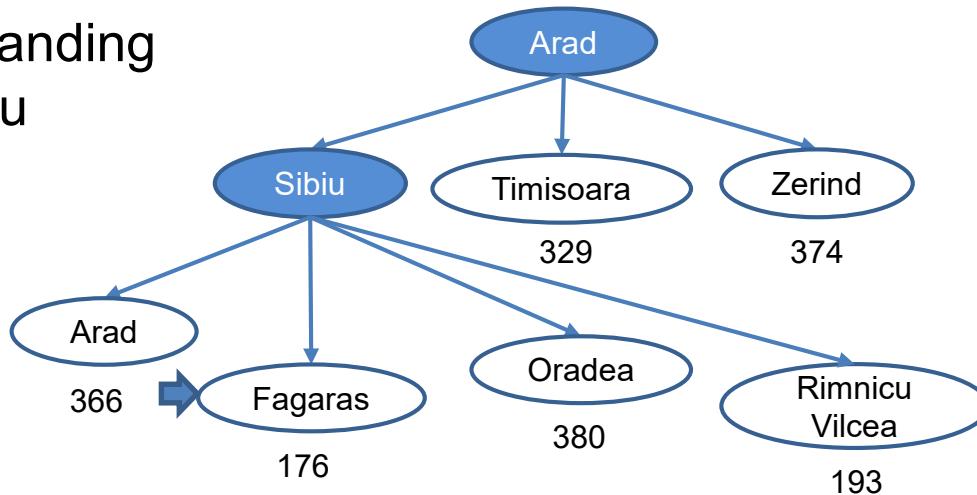
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Efoire	161
Fagaras	176
Giurgiu	77
Hirsova	151
Lasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374





Example

c) After
expanding
Sibiu



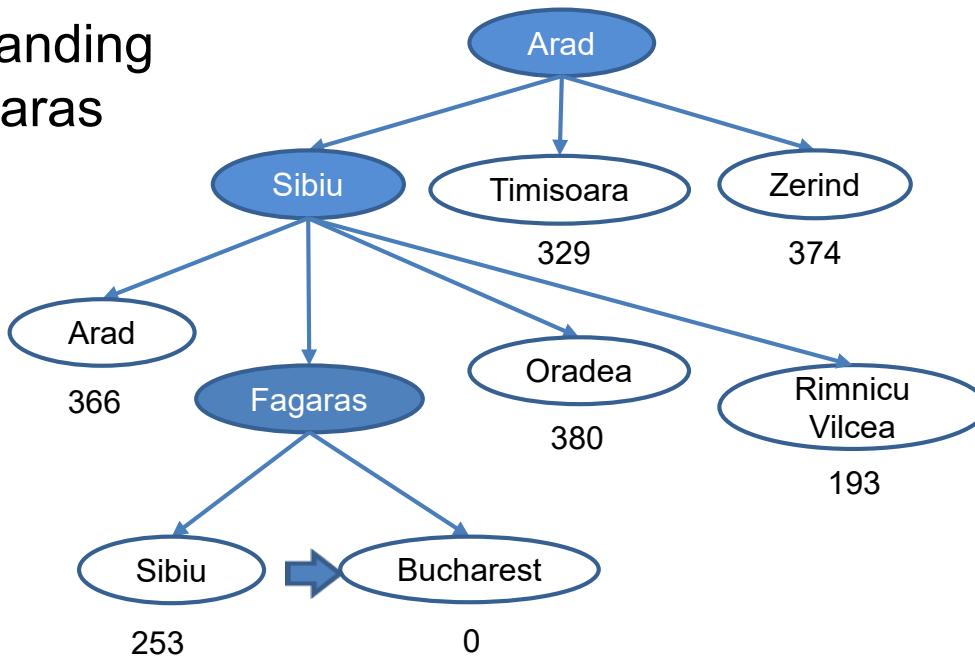
Straight-line distance to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Efoire	161
Fagaras	176
Giurgiu	77
Hirsova	151
Lasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374



Example

d) After
expanding
Fagaras



Straight-line distance to Bucharest

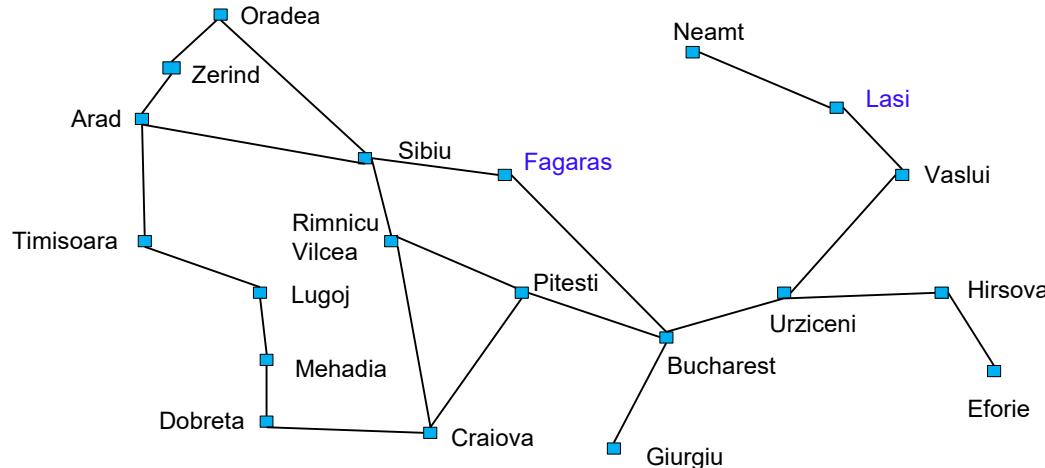
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Efoire	161
Fagaras	176
Giurgiu	77
Hirsova	151
Lasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374



Complete?

Question: Is this approach complete?

Example: Find a path from Lasi to Fagaras



Answer: No





Greedy Search...

- m: maximum depth of the search space

Complete	No
Time	$O(b^m)$
Space	$O(b^m)$ (keeps all nodes in memory)
Optimal	No



Question to think:

- Is it possible to combine functions $g(n)$ and $h(n)$ in one search strategy?







A * Search

- Uniform-cost search
 - $g(n)$: cost to reach n (**Past Experience**)
 - optimal and complete, but can be very inefficient
- Greedy search
 - $h(n)$: cost from n to goal (**Future Prediction**)
 - neither optimal nor complete, but cuts search space considerably





A * Search

Idea: Combine Greedy search with Uniform-Cost search

Evaluation function: $f(n) = g(n) + h(n)$

- $f(n)$: estimated total cost of path through n to goal (**Whole Life**)
- If $g = 0 \rightarrow$ greedy search; If $h = 0 \rightarrow$ uniform-cost search
- Function A* Search(problem) returns solution
 - Return **Best-First-Search(problem, $g + h$)**

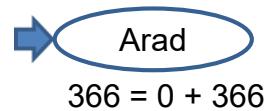




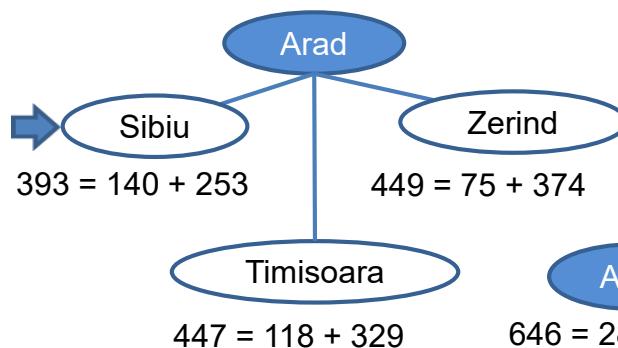
Example: Route-finding from Arad to Bucharest

Best-first-search with evaluation function $g + h$

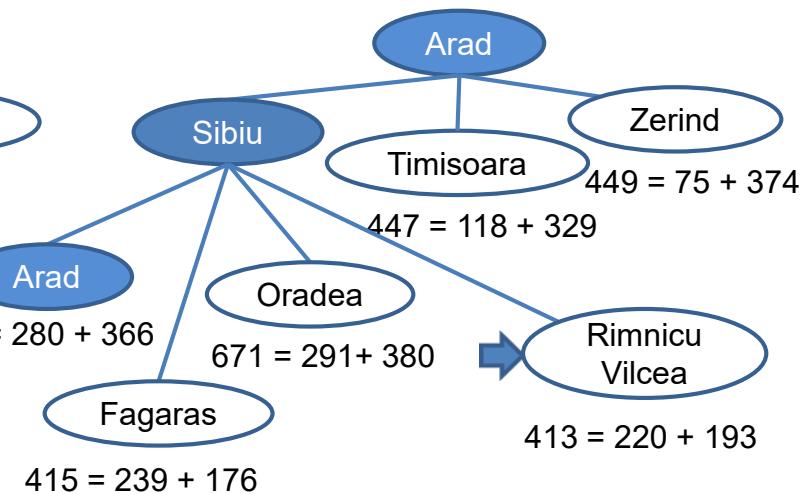
(a) The initial state



(b) After expanding Arad



(c) After expanding Sibiu

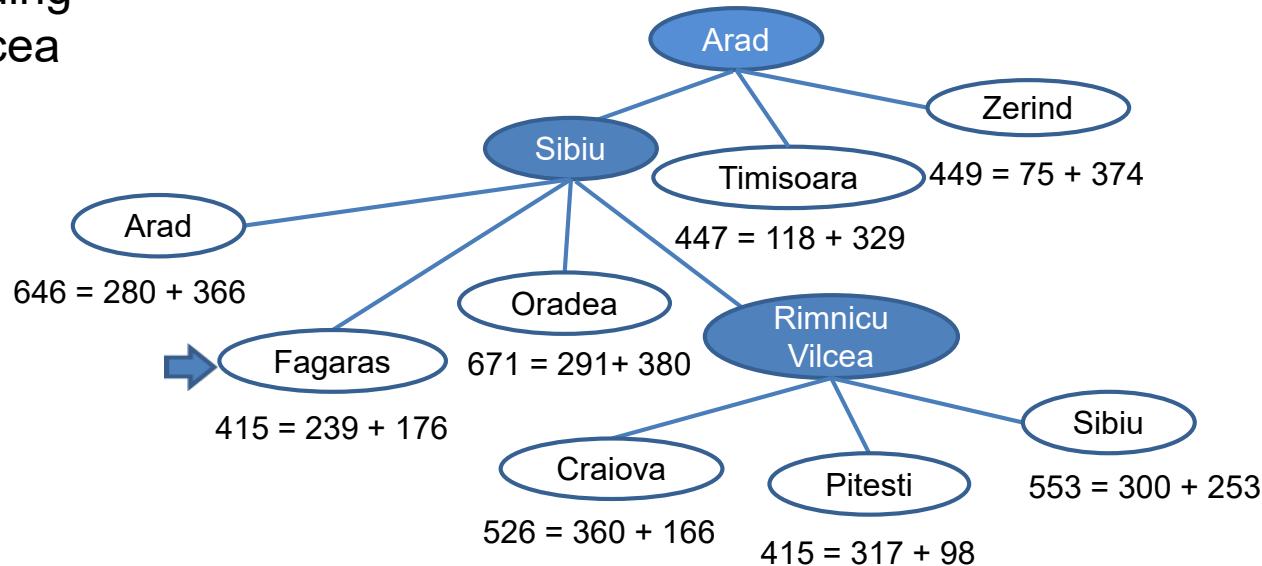




Example: Route-finding from Arad to Bucharest

Best-first-search with evaluation function $g + h$

(d) After expanding
Rimnicu Vilcea

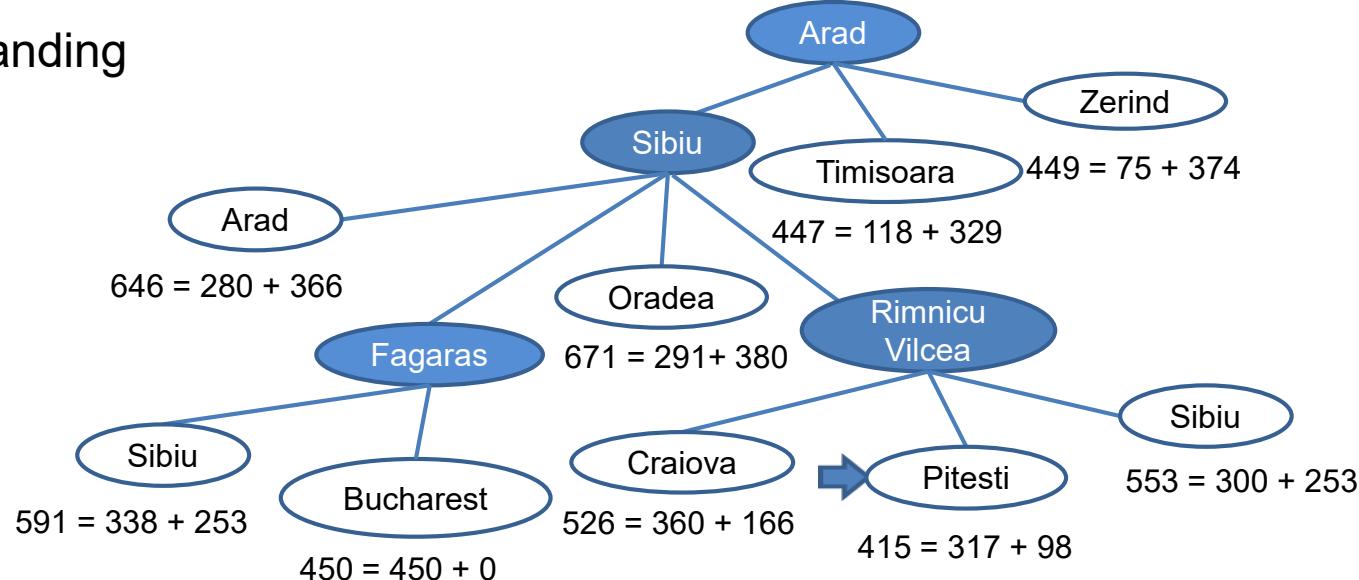




Example: Route-finding from Arad to Bucharest

Best-first-search with evaluation function $g + h$

(e) After expanding
Fagaras

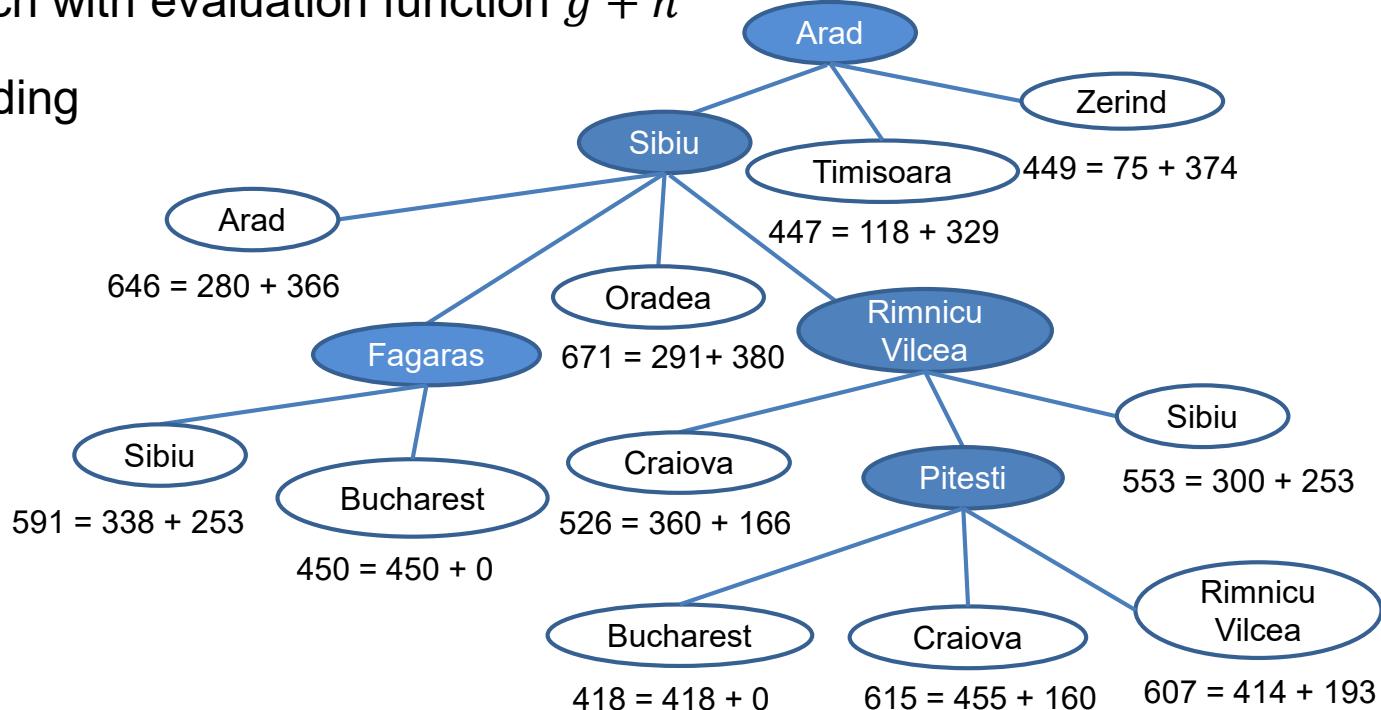




Example: Route-finding from Arad to Bucharest

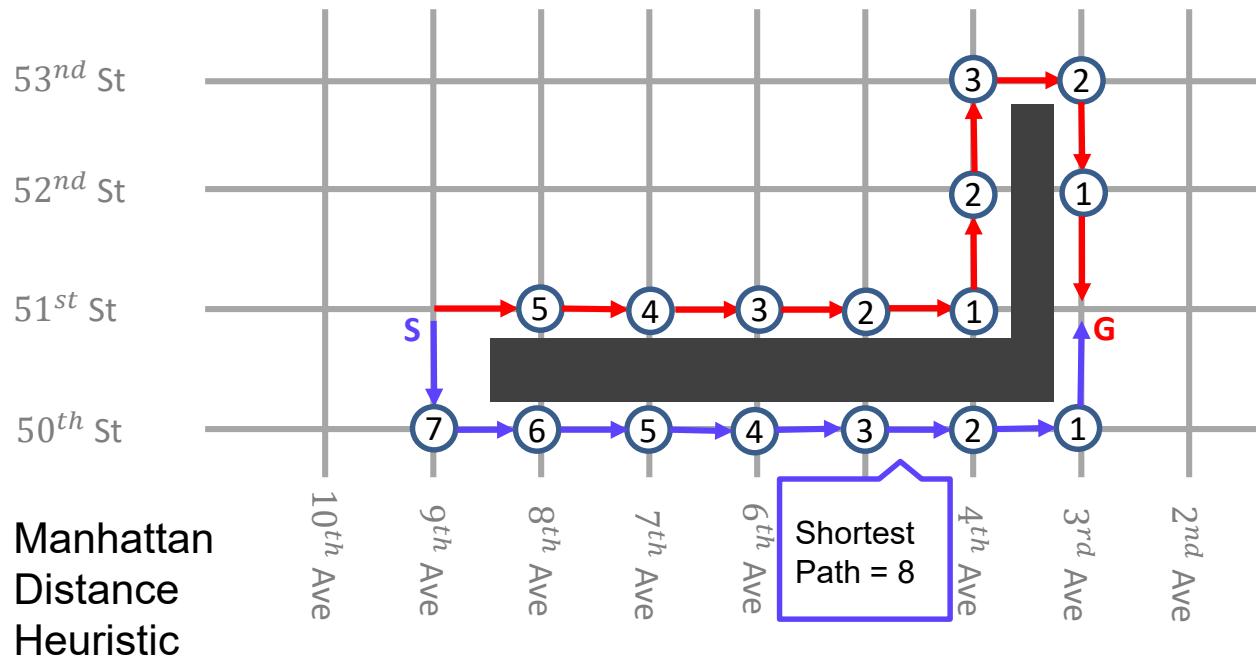
Best-first-search with evaluation function $g + h$

(f) After expanding
Pitesti





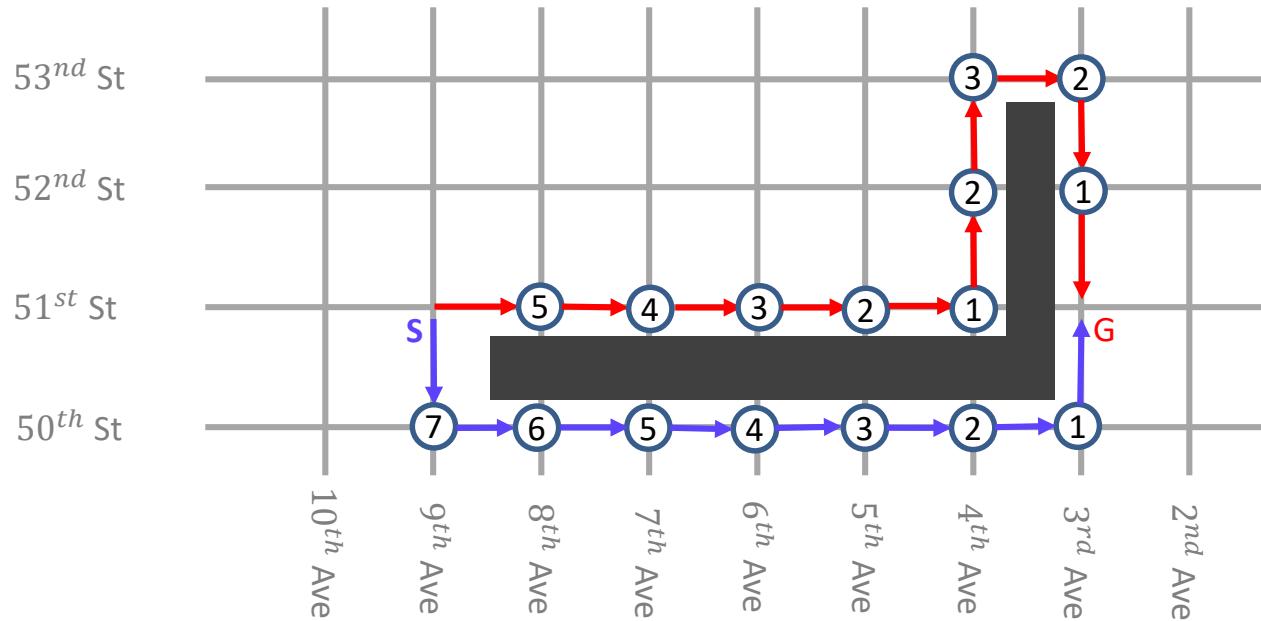
Example: Route-finding in Manhattan



Manhattan
Distance
Heuristic

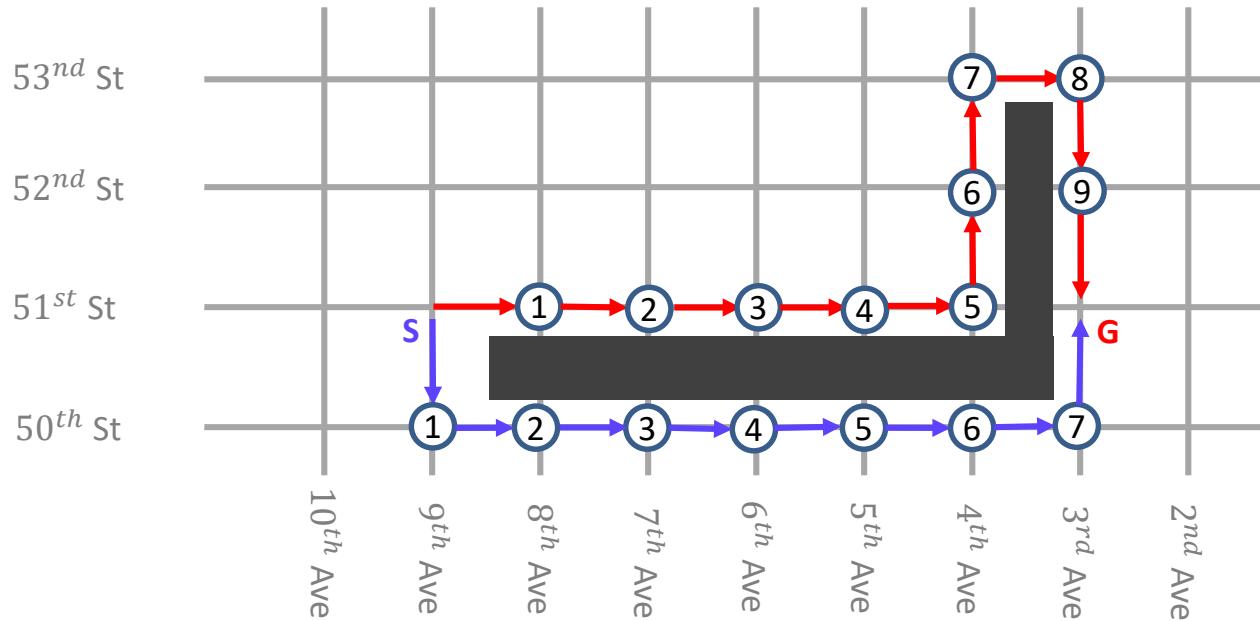


Example: Route-finding in Manhattan (Greedy)



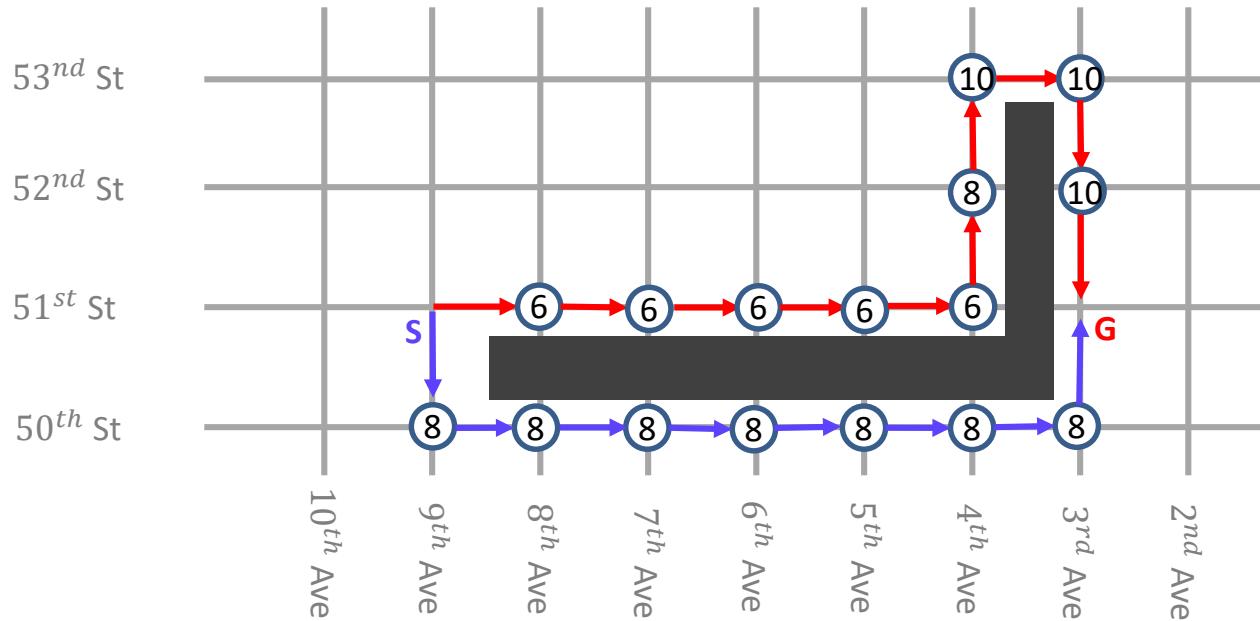


Example: Route-finding in Manhattan (UCS)



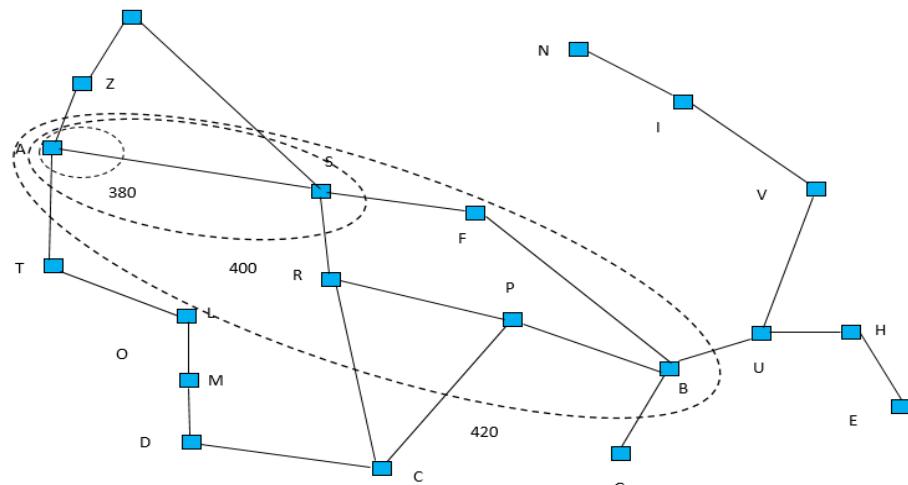


Example: Route-finding in Manhattan (A*)





Complexity of A*



Time	Exponential in length of solution
Space	(all generated nodes are kept in memory) Exponential in length of solution

With a good heuristic, significant savings are still possible compared to uninformed search methods

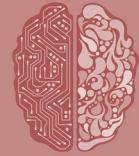
Introduction to Data Science and Artificial Intelligence

Constraint Satisfaction and Game Playing

Assoc Prof Bo AN

Research area: artificial intelligence,
computational game theory, optimization
www.ntu.edu.sg/home/boan
Email: boan@ntu.edu.sg
Office: N4-02b-55





Lesson Outline

- Constraint Satisfaction
 - Backtracking search
 - Forward checking and constraint propagation
- Game Playing
 - Games as search problems
 - Minimax search strategy
 - Evaluation functions

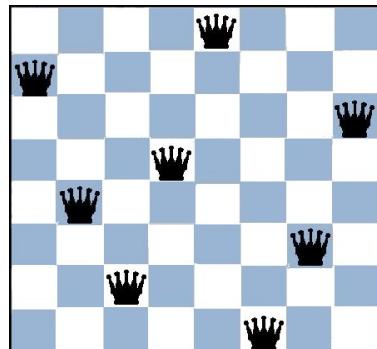




Constraint Satisfaction Problem (CSP)

Goal: discover some state that satisfies a given set of constraints

Example: 8-Queens Problem



Example: Cryptarithmetic Puzzle

$$\begin{array}{r} \text{S E N D} \\ + \text{M O R E} \\ \hline \text{M O N E Y} \end{array}$$



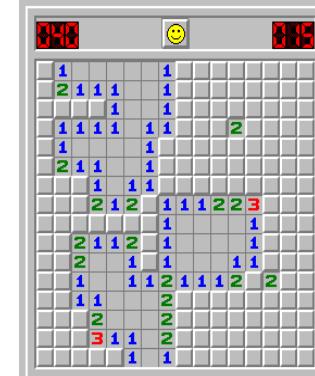
Constraint Satisfaction Problem (CSP)

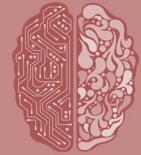
Goal: discover some state that satisfies a given set of constraints

Example: Sudoku

							1	3
			7					6
			5	9				
					9			
1		6						
					2			
7	4					5		
8				4				
			1					

Example: Minesweeper

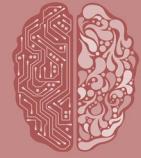




Examples: Real-world CSPs

- Assignment problems
 - e.g. who teaches what class
- Timetabling problems
 - e.g. which class is offered when and where?
- Hardware configuration
- Transportation scheduling
- Factory scheduling
- Floor-planning





CSP

State

- defined by **variables** V_i with **values** from domain D_i

Example: 8-queens

- Variables: locations of each of the eight queens
- Values: squares on the board

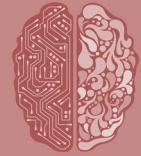
Goal test

- a set of **constraints** specifying allowable combinations of values for subsets of variables

Example: 8-queens

- Goal test: No two queens in the same row, column or diagonal



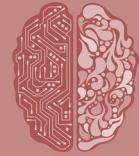


Example: Cryptarithmetic Puzzle

$$\begin{array}{r} \text{S E N D} \\ + \text{M O R E} \\ \hline \text{M O N E Y} \end{array}$$

- Variables: D, E, M, N, O, R, S, Y
- Domains: {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
- Constraints
 - $Y = D + E$ or $Y = D + E - 10$, etc.
 - $D \neq E$, $D \neq M$, $D \neq N$, etc.
 - $M \neq 0$, $S \neq 0$ (**unary** constraints: concern the value of a single variable)



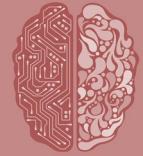


Example: Minesweeper

A	B	C
J	2	D
I	3	E
H	G	F

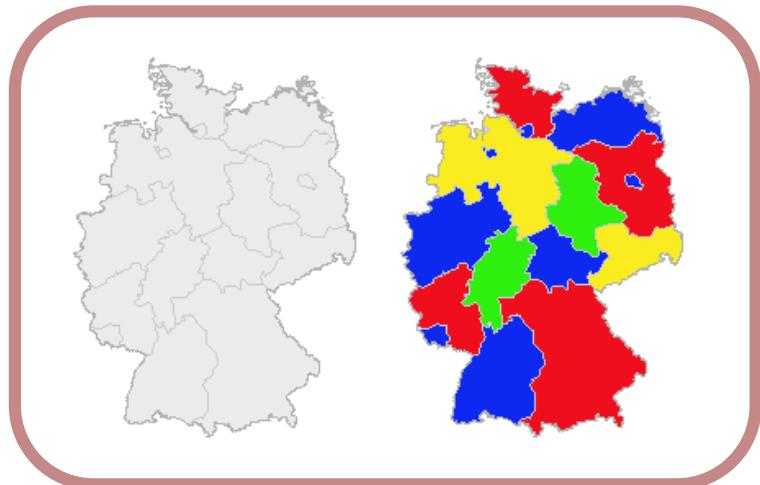
- Variables: The cells
- Domains: $\{0; 1\}$ representing {safe, mined}
- Constraints: Each cell has a number $m \in \{1, \dots, 8\}$ indicating the number of mines nearby, so m is equal to sum of value of neighbour cells



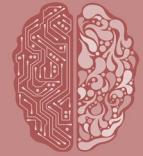


Example: Map Colouring

Colour a map so that no adjacent parts have the same colour



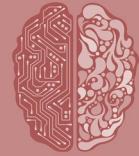
- Variables: Countries C_i
- Domains: {Red, Blue, Green}
- Constraints: $C_1 \neq C_2, C_1 \neq C_5$, etc.
 - binary constraints



Some Definitions

- A state of the problem is defined by an **assignment** of values to some or all of the variables.
- An assignment that does not violate any constraints is called a **consistent** or **legal** assignment.
- A **solution** to a CSP is an assignment with every variable given a value (**complete**) and the assignment satisfies all the constraints.

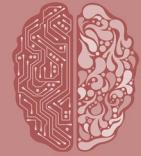




Applying Standard Search

- **States**: defined by the values assigned so far
- **Initial state**: all variables unassigned
- **Actions**: assign a value to an unassigned variable
- **Goal test**: all variables assigned, no constraints violated





Applying Standard Search

Question: How to represent constraints?

Answer: Explicitly (e.g., $D \neq E$)

Example

- Row the 1st queen occupies: $V_1 \in \{1, 2, 3, 4, 5, 6, 7, 8\}$
(similarly, for V_2)
- No-attack constraint for V_1 and V_2 :
 $\{ <1, 3>, <1, 4>, <1, 5>, \dots, <2, 4>, <2, 5>, \dots \}$

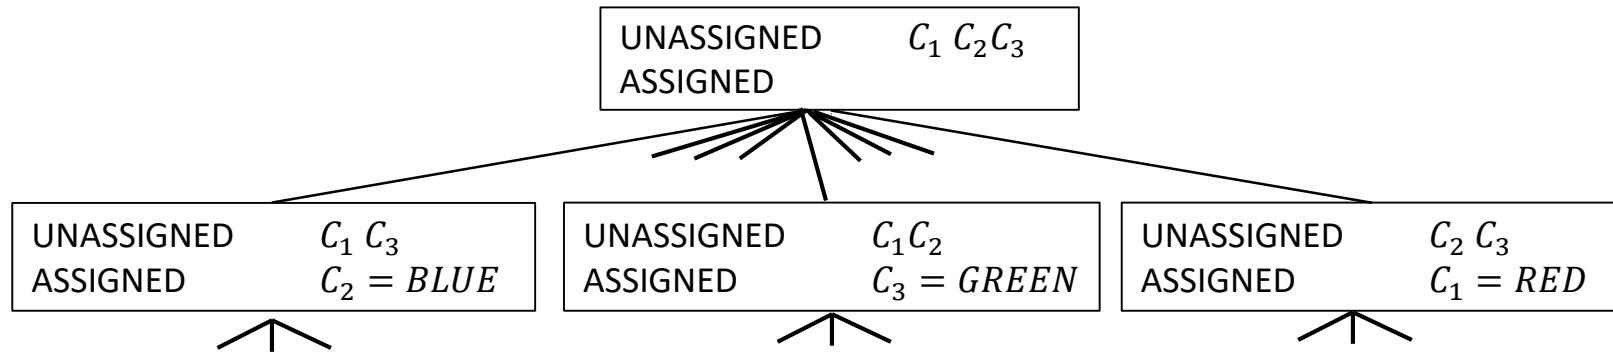
Implicitly: use a function to test for constraint satisfaction





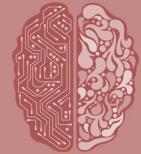
Applying Standard Search...

Example: map colouring



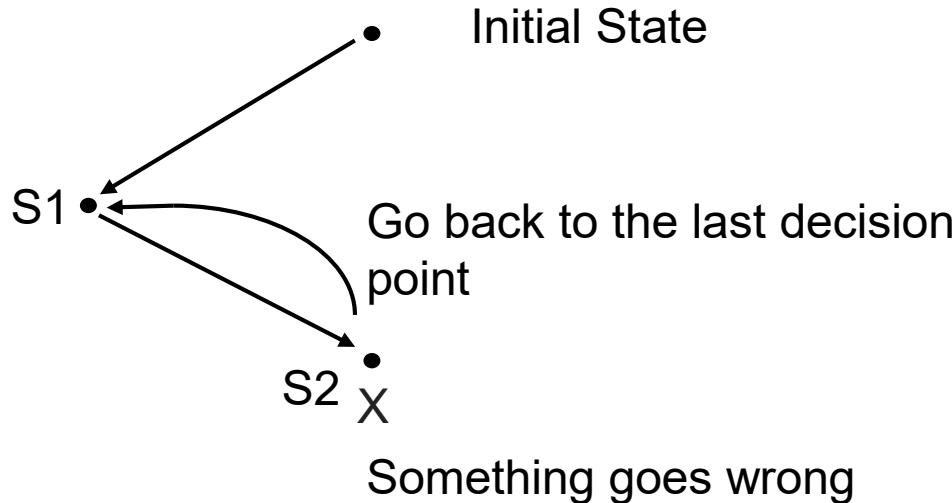
- Number of variables: n
- Max. depth of space: n
- Depth of solution state: n (all variables assigned)
- Search algorithm: depth-first search





Backtracking Search

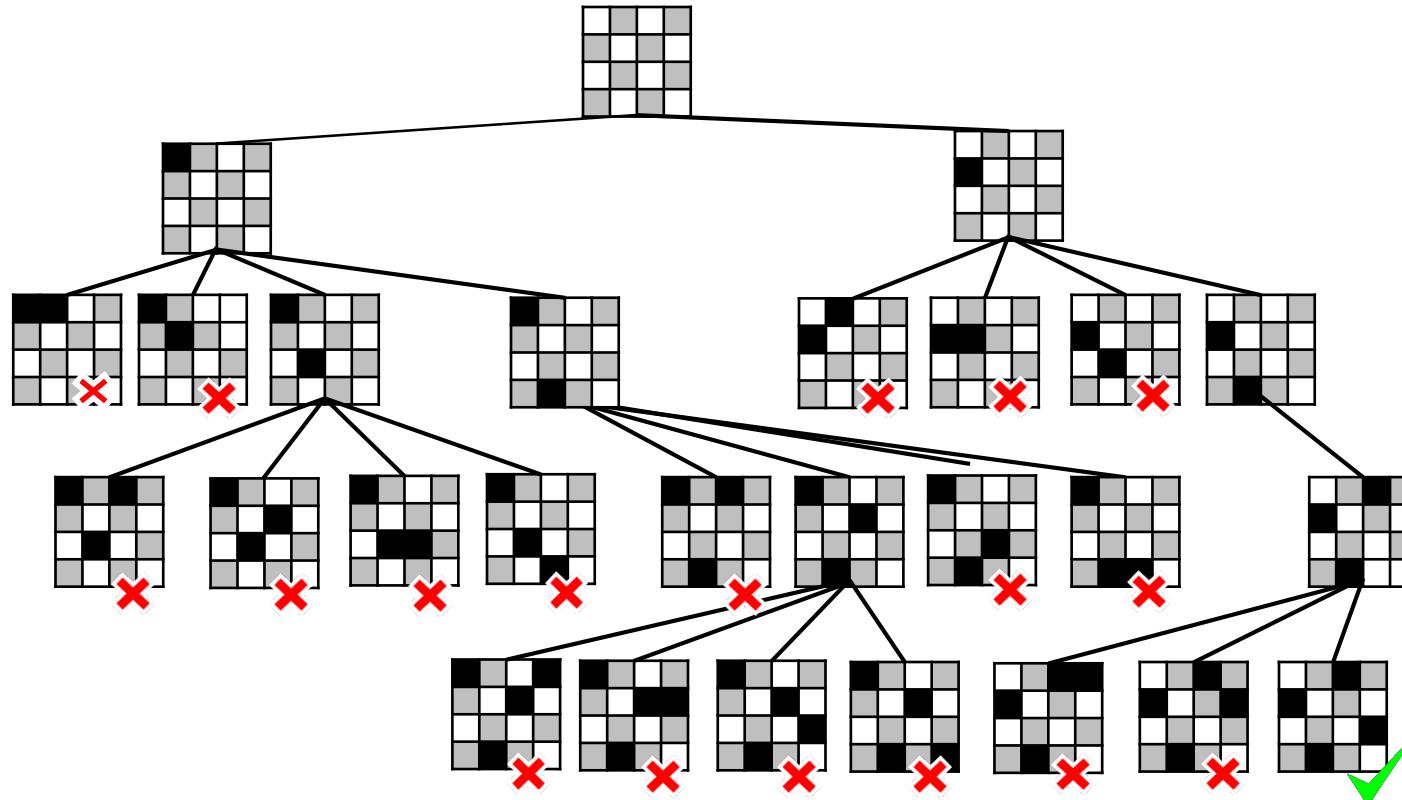
Backtracking search: Do not waste time searching when constraints have already been violated



- Before generating successors, check for constraint violations
- If yes, backtrack to try something else



Example (4-Queens)

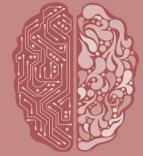


Question to think:

- Please think about whether there are ways for further improving search efficiency.







Heuristics for CSPs

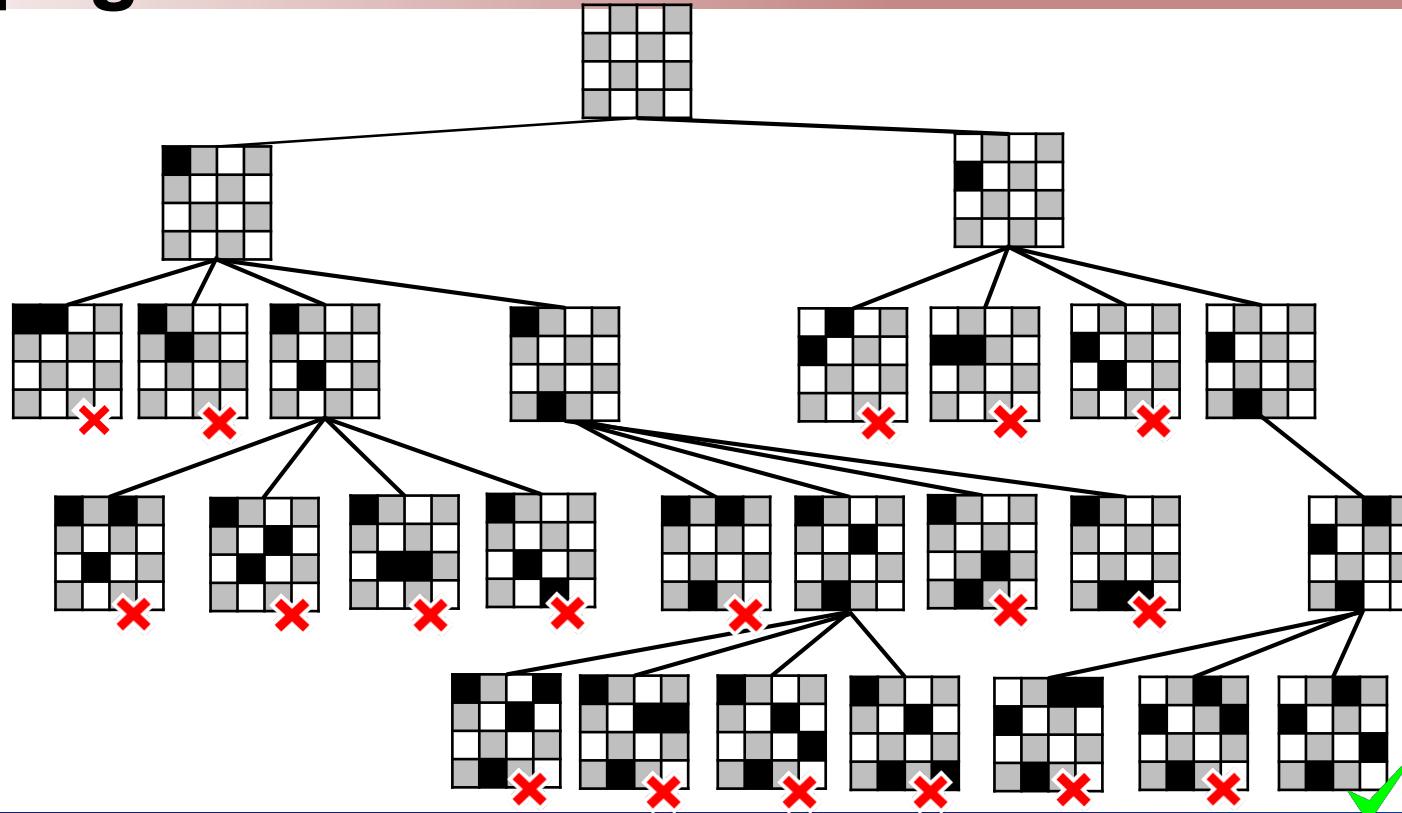
Plain backtracking is an uninformed algorithm!!

More intelligent search that takes into consideration

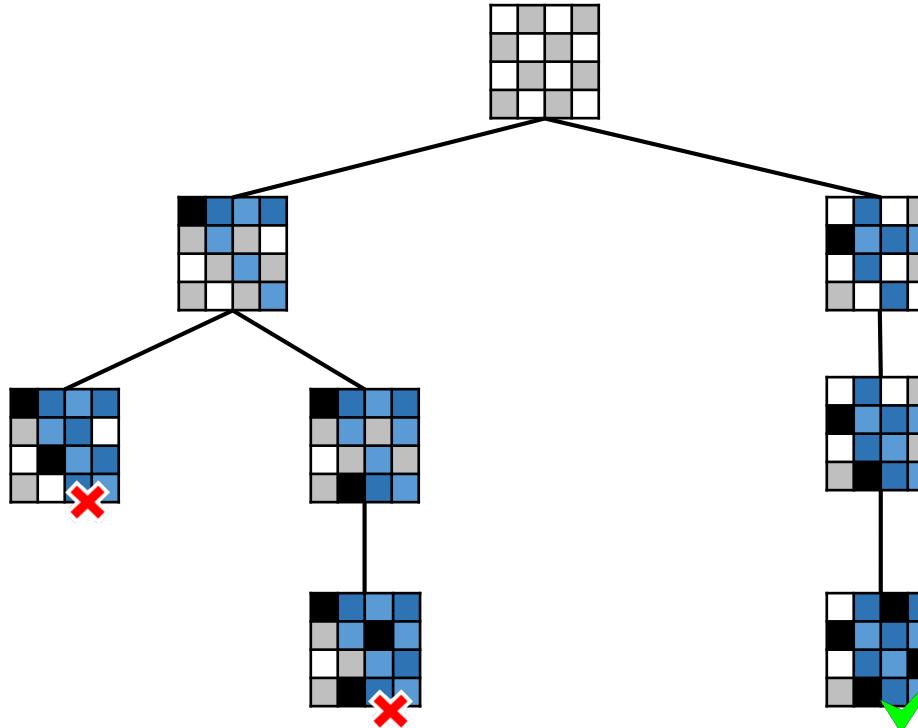
- Which variable to assign next
- What order of the values to try for each variable
- Implications of current variable assignments for the other unassigned variables
 - forward checking and **constraint propagation**

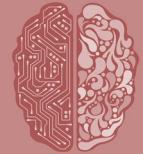
Constraint propagation: propagating the implications of a constraint on one variable onto other variables

Example (4-Queens) without Constraint Propagation

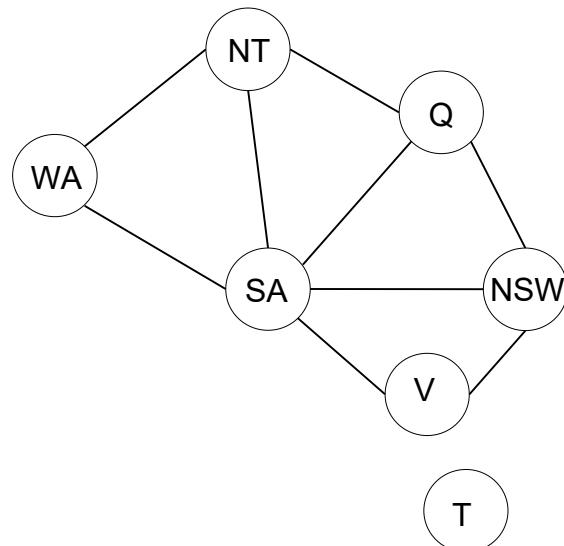


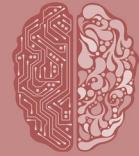
Search Tree of 4-Queens with Constraint Propagation





Example (Map Colouring)





Example (Map Colouring)...



WA

NT

Q

NSW

V

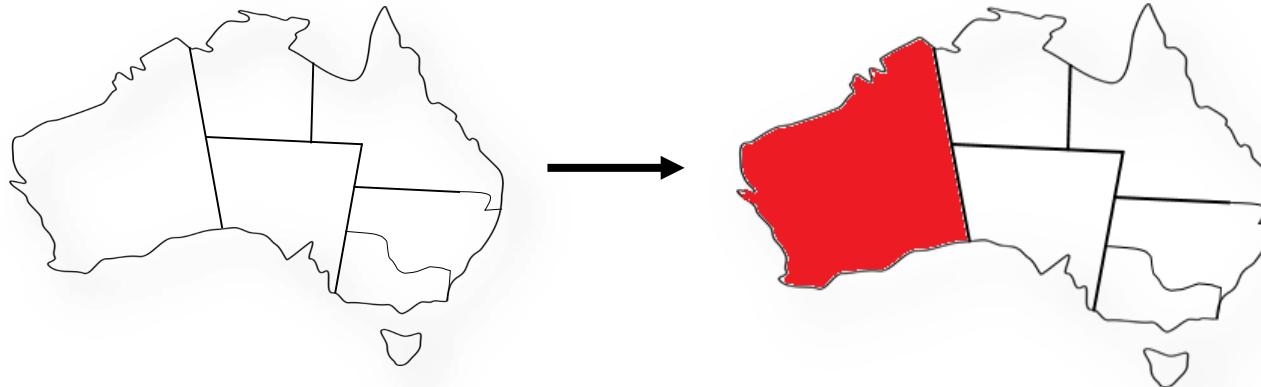
SA

T





Example (Map Colouring)...



WA

NT

Q

NSW

V

SA

T

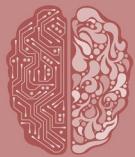




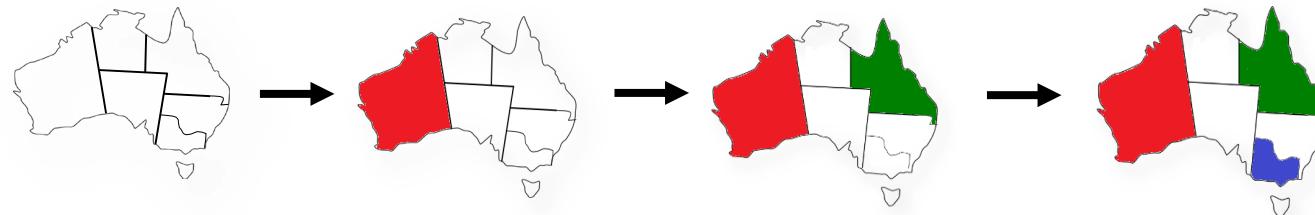
Example (Map Colouring)...



WA	NT	Q	NSW	V	SA	T



Example (Map Colouring)...



WA

NT

Q

NSW

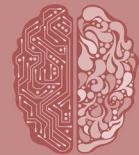
V

SA

T

■ Red ■ Green ■ Blue						
■ Red		■ Green ■ Blue	■ Red ■ Green ■ Blue	■ Red ■ Green ■ Blue	■ Green ■ Blue	■ Red ■ Green ■ Blue
■ Red		■ Blue	■ Green	■ Red ■ Blue	■ Blue	■ Red ■ Green ■ Blue
■ Red		■ Blue	■ Green		■ Blue	■ Red ■ Green ■ Blue

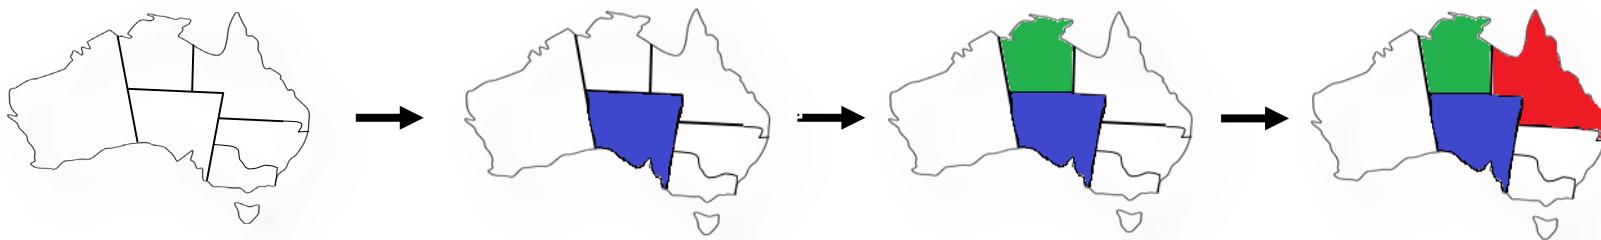




Most Constrained Variable

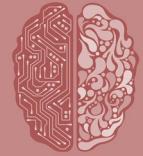
Or minimum remaining values (MRV) heuristic

Example: map colouring

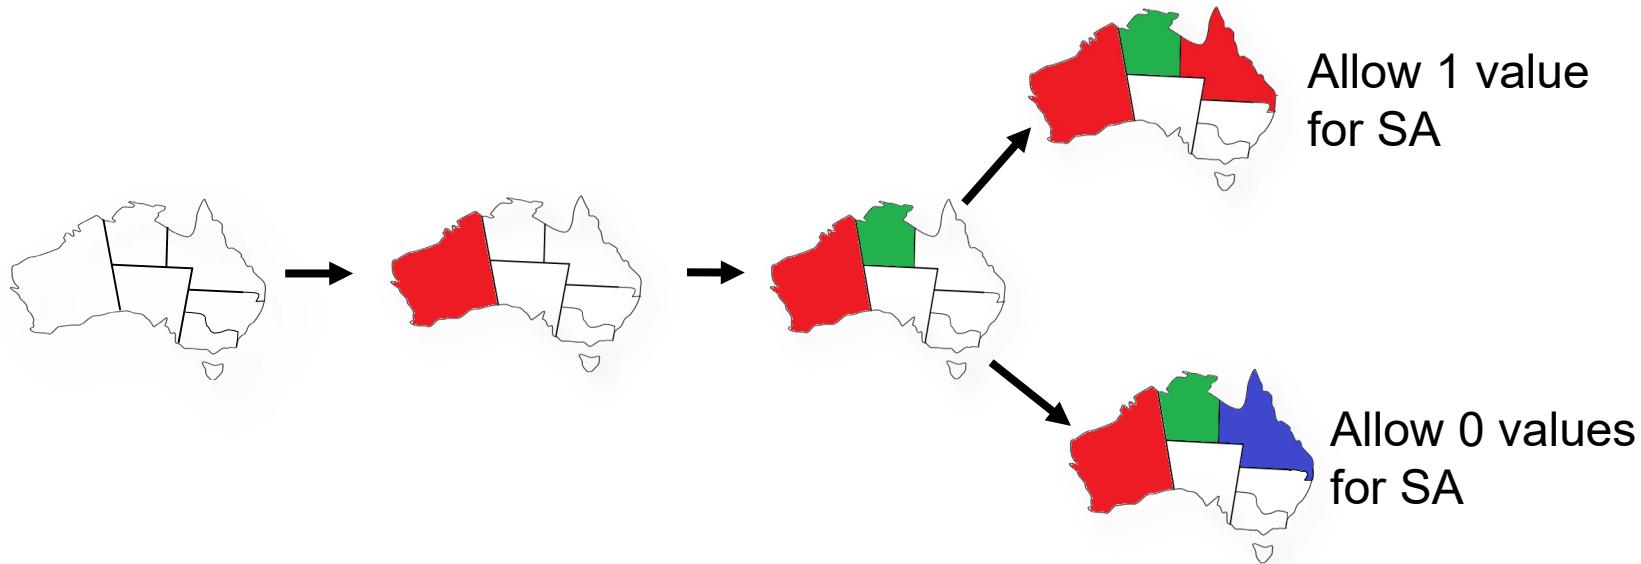


To reduce the branching factor on future choices by selecting the variable that is involved in the **largest number of constraints** on unassigned variables.



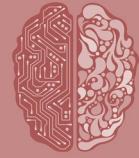


Least Constraining Value



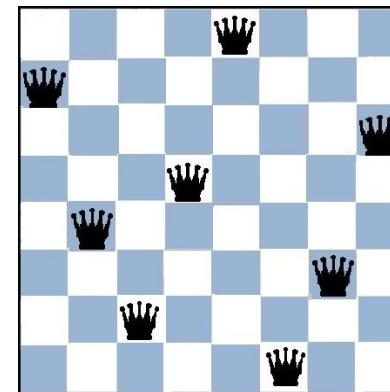
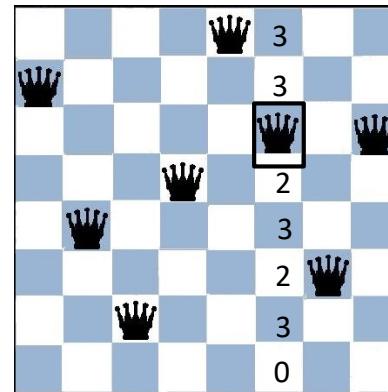
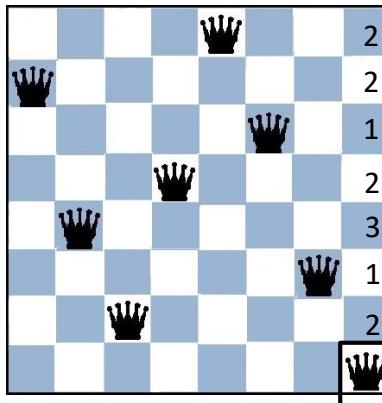
Choose the value that leaves maximum flexibility for subsequent variable assignments





Min-Conflicts Heuristic (8-queens)

- A local heuristic search method for solving CSPs
- Given an initial assignment, selects a variable in the scope of a violated constraint and assigns it to the value that minimises the number of violated constraints

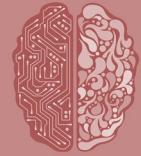


Question to think:

- Please think the intuitions for why the above ideas can improve search efficiency.







Games as Search Problems

Abstraction

- Ideal representation of real world problems
 - e.g. board games, chess, go, etc. as an abstraction of war games
 - Perfect information, i.e. fully observable
- Accurate formulation: state space representation

Uncertainty

- Account for the existence of **hostile** agents (players)
 - Other agents acting so as to diminish the agent's well-being
 - Uncertainty (about other agents' actions):
 - not due to the effect of non-deterministic actions
 - not due to randomness
- Contingency problem



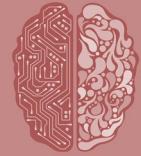


Games as Search Problems...

Complexity

- Games are abstract but not simple
 - e.g. chess: average branching factor = 35, game length > 50
→ complexity = 35^{50} (only 10^{40} for legal moves)
- Games are usually time limited
 - Complete search (for the optimal solution) not possible
→ uncertainty on actions desirability
 - Search efficiency is crucial





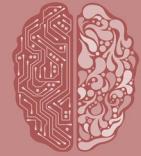
Types of Games

	Deterministic	Chance
Perfect information	Chess, Checkers, Go, Othello	Backgammon, Monopoly
Imperfect information		Bridge, Poker, Scrabble, Nuclear war

Perfect information

- each player has complete information about his opponent's position and about the choices available to him





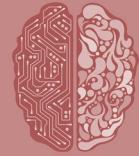
Game as a Search Problem

- Initial state: initial board configuration and indication of who makes the first move
- Operators: legal moves
- Terminal test: determines when the game is over
 - states where the game has ended: **terminal states**
- Utility function (payoff function): returns a numeric score to **quantify** the outcome of a game

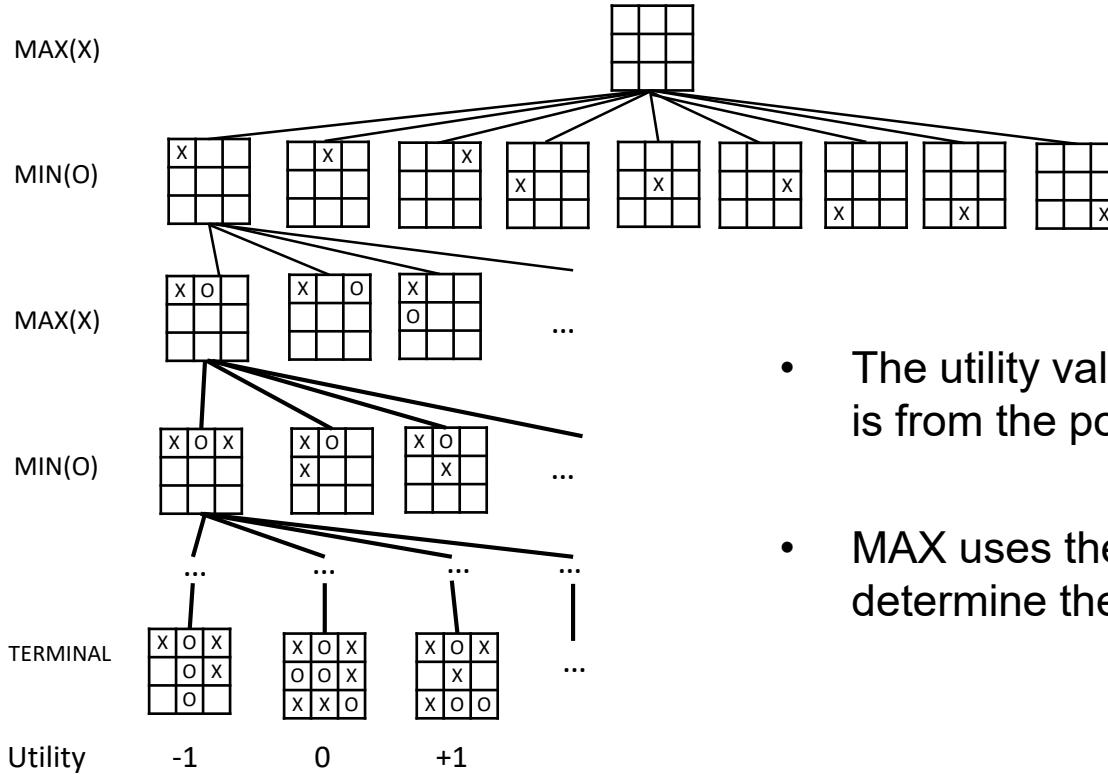
Example: Chess

Win (+1), loss(-1) or draw (0)

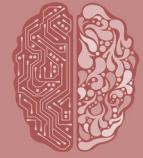




Game Tree for Tic-Tac-Toe

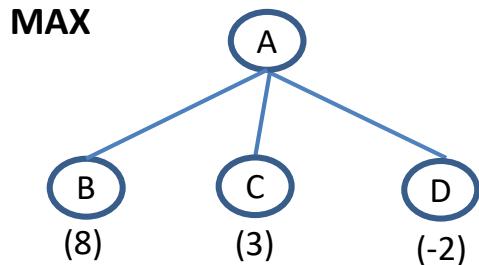


- The utility value of the terminal state is from the point of view of MAX
- MAX uses the search tree to determine the best move

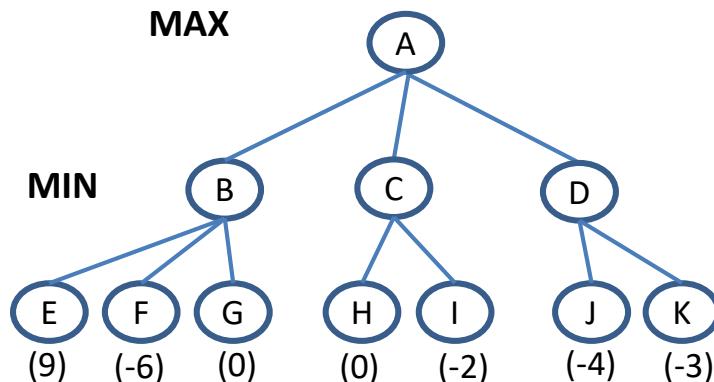


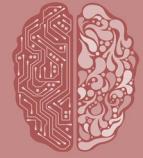
What Search Strategy?

One-play

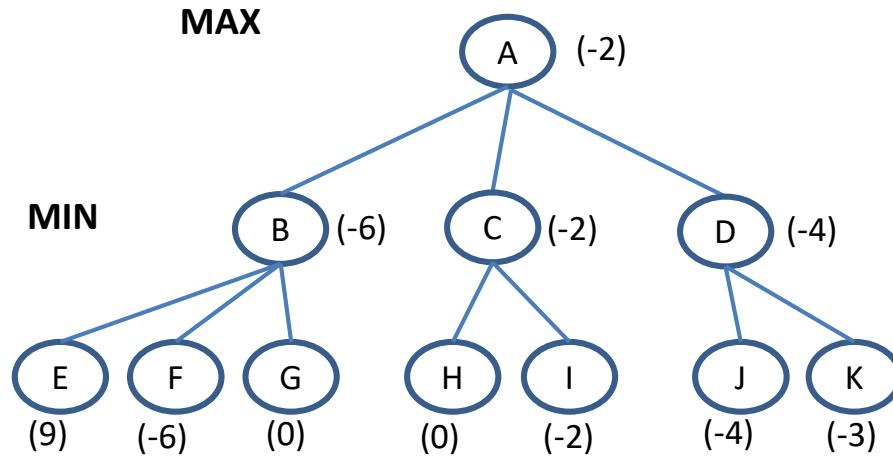


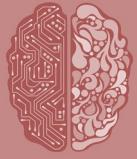
Two-play





What Search Strategy?...





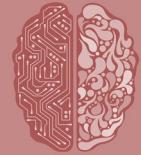
Minimax Search Strategy

Search strategy

- Find a sequence of moves that leads to a terminal state (goal)

Minimax search strategy

- Maximise one's own utility and minimise the opponent's
 - Assumption is that the opponent does the same

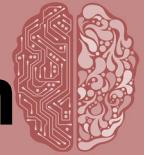


Minimax Search Strategy

3-step process

1. Generate the entire game tree down to terminal states
2. Calculate utility
 - a) Assess the utility of each terminal state
 - b) Determine the best utility of the parents of the terminal state
 - c) Repeat the process for their parents until the root is reached
3. Select the best move (i.e. the move with the highest utility value)





Perfect Decisions by Minimax Algorithm

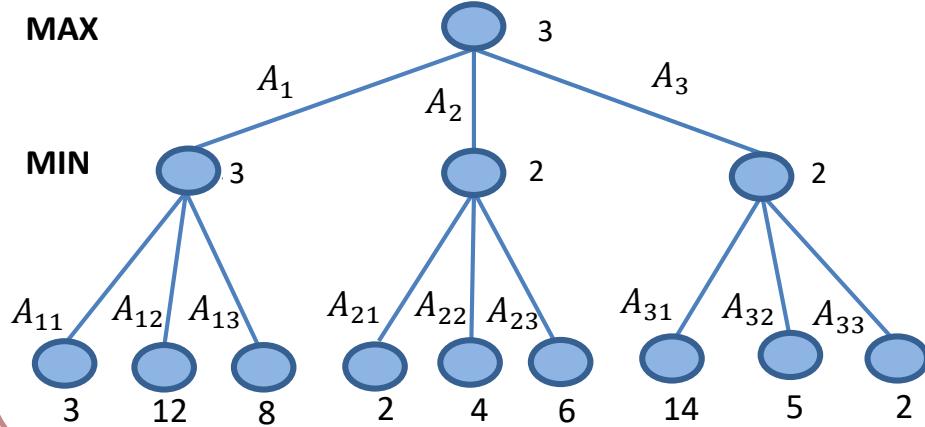
Perfect decisions: no time limit is imposed

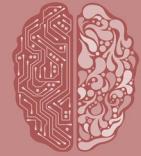
- generate the complete search tree

Two players: MAX and MIN

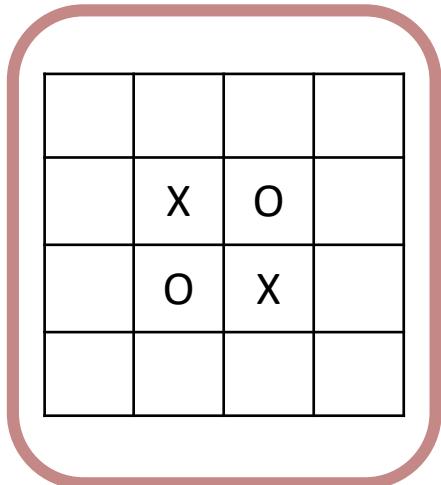
- Choose move with best achievable payoff against best play
- MAX tries to max the utility, assuming that MIN will try to min it

Example



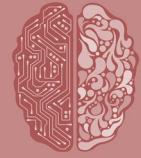


Othello 4



- A player can place a new piece in a position if there exists at least one straight (horizontal, vertical, or diagonal) occupied line between the new piece and another piece of the same kind, with one or more contiguous pieces from the opponent player between them
- After placing the new piece, the pieces from the opponent player will be captured and become the pieces from the same Player
- The player with the most pieces on the board wins





'X' plays first

X considers the game now

	X	O
O		X

O considers the game now

	X	O
X		X
X		

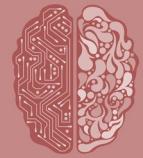
X considers the game now

O	O	O
X	X	
X		

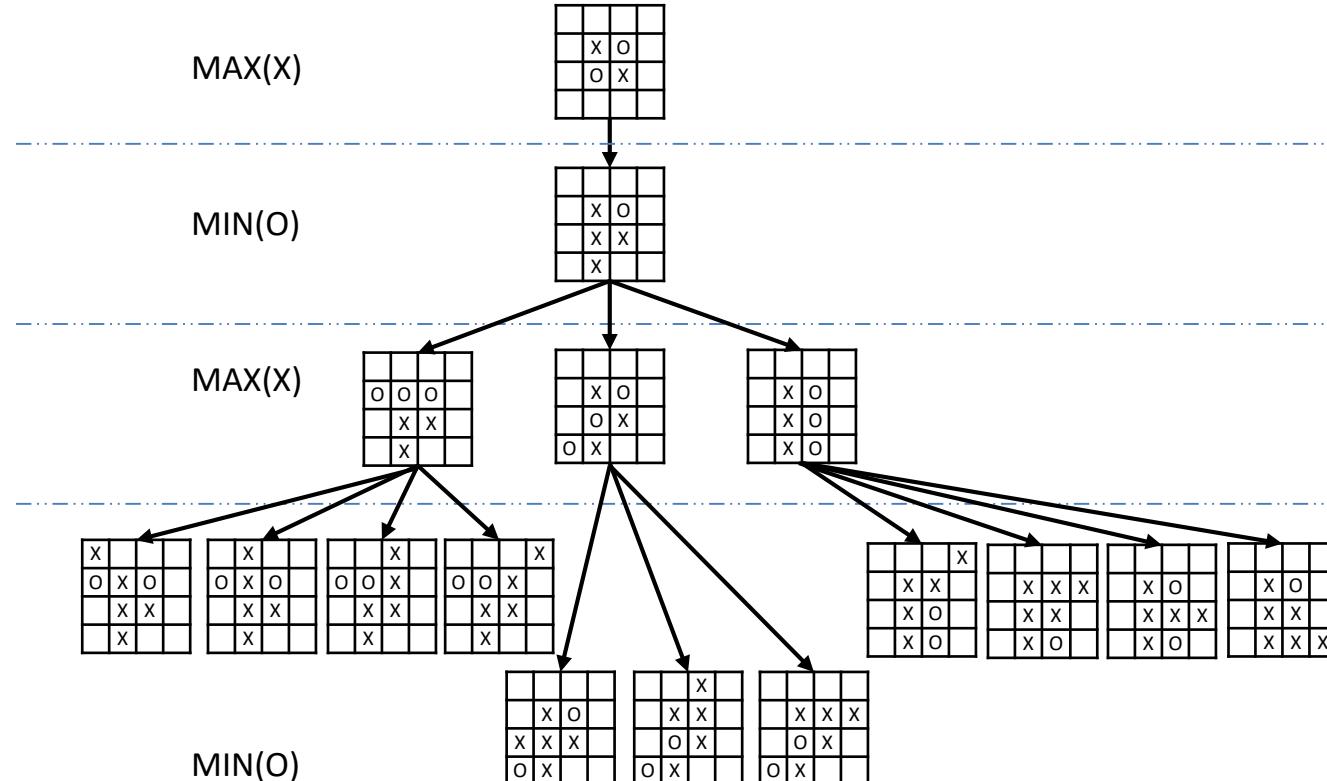
	X	O
O	X	
O	X	

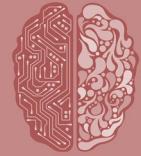
X	O	
X	O	
X	O	





Game Tree Othello 4





Imperfect Decisions

For chess, branching factor ≈ 35 , each player typically makes 50 moves \rightarrow for the complete game tree, need to examine 35^{100} positions

Time/space requirements \rightarrow complete game tree search is intractable \rightarrow impractical to make perfect decisions

Modifications to minimax algorithm

1. replace utility function by an estimated desirability of the position
 - Evaluation function
2. partial tree search
 - E.g., depth limit
 - Replace terminal test by a cut-off test





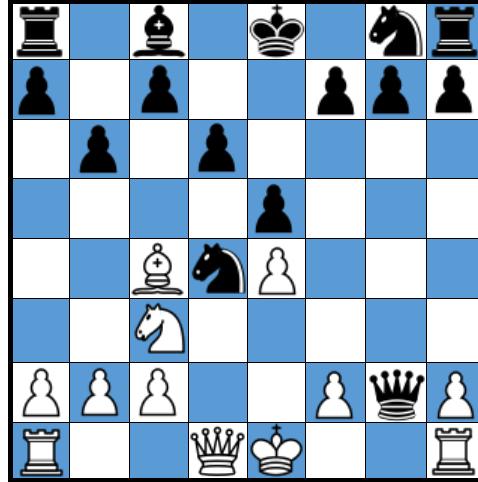
Evaluation Functions

Returns an **estimate** of the expected utility of the game from a given position



Black: to move

White: slightly better



White: to move

Black: winning

