

CE4062/CZ4062

Computer Security

Lecture 5: Operating System Security

Tianwei Zhang

Outline

- ▶ **Operating System Security Basis**
- ▶ **Security Protection Stages employed by OS**
- ▶ **UNIX Security Model**
- ▶ **Security Vulnerabilities in OS**

Outline

- ▶ **Operating System Security Basis**
- ▶ Security Protection Stages employed by OS
- ▶ UNIX Security Model
- ▶ Security Vulnerabilities in OS

OS Becomes More Complex

From single-user to multi-user

- ▶ DOS is truly single user
- ▶ MacOS, Linux, NT-based Windows are multi-user, but typically only one user in PCs.
- ▶ Cloud computing allows multiple users all over the world to run on the same OS, and they do not know each other.
- ▶ **Tradeoff: efficiency versus security**

From trusted apps to untrusted apps

- ▶ Simple real-time systems: only run one specific app
- ▶ Runs verified apps from trusted parties
- ▶ Modern PCs and smartphones: run apps from third-party developers
- ▶ **Tradeoff: functionality versus security**

Complex OS Brings More Challenges

Protecting a single computer with one user is easy

- ▶ Prevent everybody else from having access
- ▶ Encrypt all data with a key only one person knows


Sharing resources safely is hard

- ▶ Preventing some people from reading private data (e.g., medical record, financial data, employee information)
- ▶ Prevent some people from using too many resources (e.g., disk space, CPU core)
- ▶ Prevent some people from interfering with other programs (e.g., inserting keystrokes / modifying displays)

OS Responsibility

Functionalities

- ☐ Support multiple users concurrently
- ☐ Manage multiple apps concurrently
- ☐ Connect to the network
- ☐ Sharing data with different domains

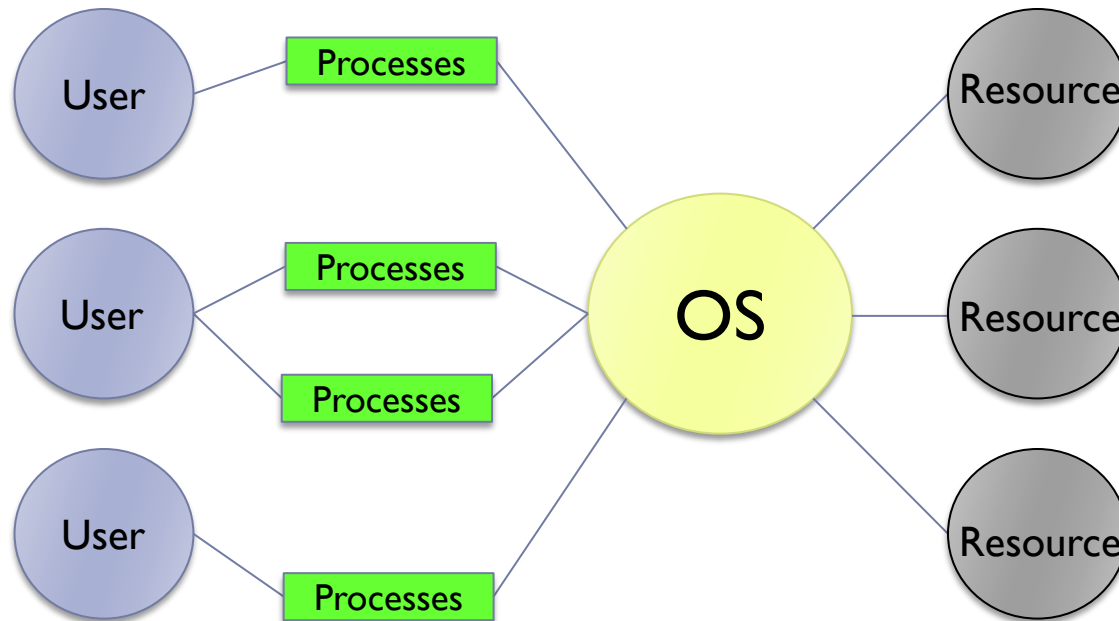


Security goals

- ☐ Protect users from each other
- ☐ Protect apps from each other
- ☐ Protect the system from the untrusted network
- ☐ Secure the data sharing

What's being protected? Resources

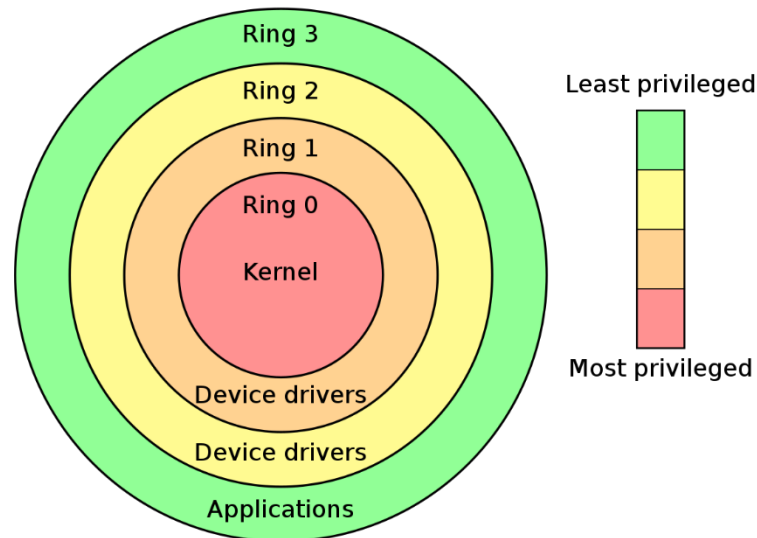
System is secure if **resources** are used and accessed as intended under all circumstances



Privileged Rings Inside OS

Operating modes

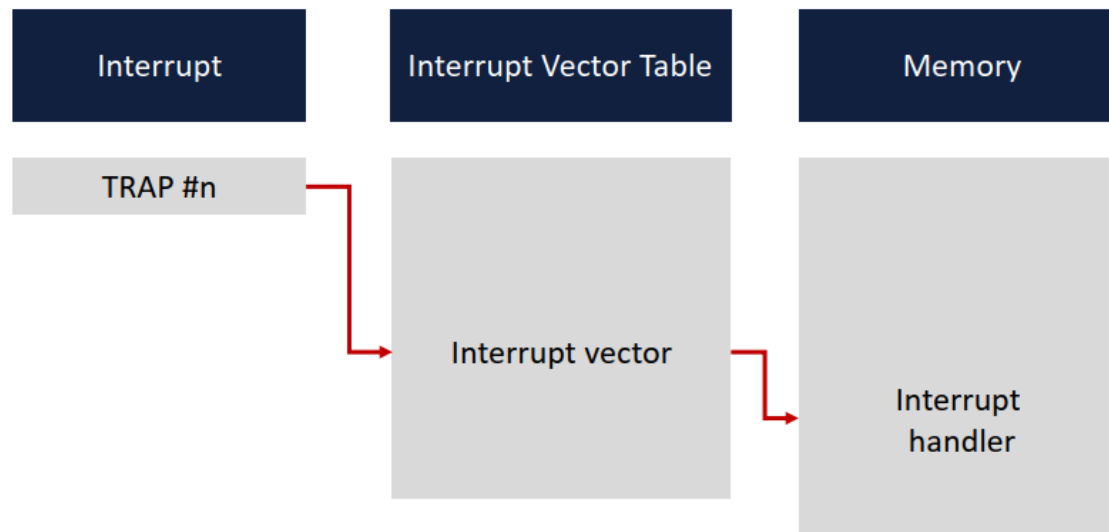
- ▶ Kernel mode has the highest privileges, running the critical functions and services
- ▶ Entities with the higher privilege levels cannot call the functions and access the objects in the lower privilege levels directly.
 - System call, interrupt, etc.
- ▶ Status flag allows system to work in different modes (context switching)



Interrupt

Context switch

- ▶ The user-level process generates a trap, which will incur an interrupt
- ▶ The CPU stores the process's states, and switches to the kernel mode by setting the status flag.
- ▶ The kernel handles the interrupt based on the trap address (interrupt vector) in an interrupt table.
- ▶ The CPU switches back to the user mode and restores the states.



Process Security

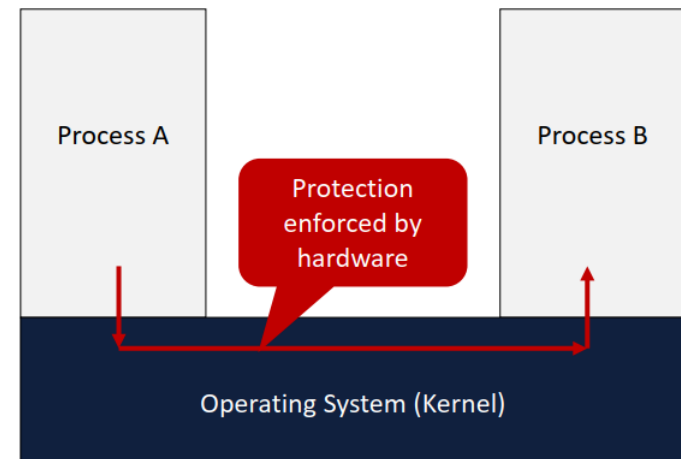
A program in execution, consisting of executable code, data, and the execution context, e.g., the contents of certain CPU registers.

Process isolation

- ▶ A process has its own address space
- ▶ Logical separation of processes as a basis for security: an **independent** process cannot affect or be affected by the execution of other processes

Process communication


- ▶ A **cooperating** process can affect or be affected by the execution of other processes
- ▶ Such processes have to communicate with each other to share data
 - Inter-Process Communication
- ▶ A context switch between processes can be an expensive and insecure operation



Memory Management

Memory layout

- ▶ Each process is allocated a segment of memory for storing data and computation results
- ▶ Memory scope is restricted by **Base** and **Limit**
- ▶ Divided into memory pages of equal lengths.
- ▶ A process is not allowed to access memory pages not belonging to it
 - The OS will check if each memory access is allowed.
 - A page fault will be generated if a memory access is illegal.



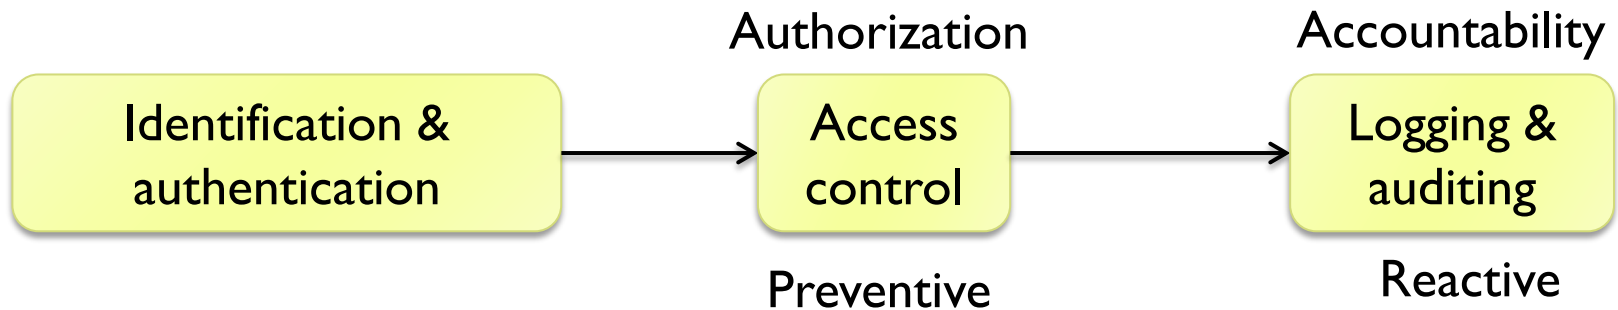
Outline


- ▶ Operating System Security Basis
- ▶ **Security Protection Stages employed by OS**
- ▶ UNIX Security Model
- ▶ Security Vulnerabilities in OS

Security Protection from OS

OS is responsible for protecting the apps running on it

- ▶ OS controls what users/processes can do






How does a computer know if I am a correct user?

- ▶ **What you know?** password, PIN, public/private keys...
- ▶ **What you have?** smartcard, hardware tokens...
- ▶ **Who you are?** biometrics, face recognition, voice recognition...

How does the system conduct authentication?

- ▶ Compare the input credential with the stored one
- ▶ Allow entry when the credential matches



Principle of least privilege

- ▶ Users should only have access to the resources needed to perform the desired tasks.
- ▶ Too much privilege allows a malicious user to conduct unintended activities

Privilege separation

- ▶ Spit a system into different components, and each component is assigned with the least privilege for its tasks
- ▶ Limiting the privilege can prevent any attacker from taking over the entire system.

Security Policy

Specify (Subject, Object, Operation) triples

Subject

- ▶ Acting system principals.
- ▶ User, program, process...

Object

- ▶ Protected resources.
- ▶ File, memory, devices...

Operation

- ▶ How subject operates on objects
- ▶ Read, write, execute, delete...

Access Control Policy

Who sets policy?

- ▶ Users, with some system restrictions

Access Control Matrix

- ▶ Each column represents an object
- ▶ Each row represents a subject
- ▶ The entry shows the allowed verbs.

| | /etc | /homes | /usr |
|-------|---------------|---------------|---------------|
| Alice | Read | Read | Read Write |
| Bob | Read Write | Read Write | Read Write |
| Carl | None | None | Read |

How is policy enforced?

- ▶ OS exposes API to apps, with privileged operations
- ▶ Checks access control matrix when API functions are called

Update Access Control Matrix

Access Control Changes

- ▶ Grant capabilities: the owner of the object can grant rights to other users.
- ▶ Revoke capabilities: subjects can revoke the rights from others

Six Commands to Alter the Access Matrix

- ▶ **create subject s** : creates a new subject s .
- ▶ **create object o** : creates a new object o .
- ▶ **enter r into Ms,o** : adds right r to cell Ms,o .
- ▶ **delete r from Ms,o** : deletes right r from cell Ms,o .
- ▶ **destroy subject s** : deletes subject s . The column and row for s in M are also deleted.
- ▶ **destroy object o** : deletes object o . The column for o in M is also deleted.

More Representations

Access Control List (ACLs)

- ▶ For one object, which subject has accesses to it? (check the column in the Access Control Matrix)

Capability

- ▶ For one subject, which objects it has capability to access? (check rows in the Access Control Matrix)

| | /etc | /homes | /usr |
|-------|------------|------------|------------|
| Alice | Read | Read | Read Write |
| Bob | Read Write | Read Write | Read Write |
| Carl | None | None | Read |

Most systems use both

- ▶ ACLs for opening an object, e.g. `fopen()`
- ▶ Capabilities for performing operations, e.g. `read()`


Data Sharing

Problem: multiple users want to access the same file or data

- ▶ Give each user the corresponding permissions.
- ▶ When a new user joins, the permissions have to be granted again.
- ▶ When permissions are changed, need to alter each user.

Solution: group

- ▶ Set permissions for the group instead of the user
- ▶ A user joining the group will have the corresponding permissions.
- ▶ A user quitting the group will loss the corresponding permissions.
- ▶ Easier to manage and update.



Audit trail

- ▶ Recording all protection-orientated activities, important to understanding what happened, why, and catching things that shouldn't

| | |
|-------------------------|--|
| /usr/adm/lastlog | Records the last time a user has logged in; displayed with finger |
| /var/adm/utmp | Records accounting information used by the who command. |
| /var/adm/wtmp | Records every time a user logs in or logs out; displayed with the last command. |
| /var/adm/acct | Records all executed commands; displayed with lastcomm |
| /var/log/ | In modern Linux systems, log files are located in there |

Outline

- ▶ Operating System Security Basis
- ▶ Security Protection Stages employed by OS
- ▶ **UNIX Security Model**
- ▶ Security Vulnerabilities in OS

UNIX

A family of multitasking, multiuser computer operating systems

- ▶ Developed by Dennis Ritchie & Ken Thompson at AT&T Bell Labs in 1969
- ▶ Originally for small multi-user computers in a friendly network environment;
- ▶ Later scaled up to commercial; improved gradually.

Several flavors of Unix: Unix-like OS

- ▶ Vendor versions differ in the way some security controls are managed & enforced.
- ▶ Solars, FreeBSD, macOX, ...

Security was not a primary design goal.

- ▶ Dominant goals were modularity, portability and efficiency.
- ▶ Now it provides sufficient security mechanisms that have to be properly configured and administered.

Users

A system can have many accounts

- ▶ Service accounts: running background processes
- ▶ User accounts: tied to each person

User identifiers (UIDs)

- ▶ A 16-bit number (size of UID values varies for different systems)
- ▶ 0 reserved for root, 1-99 for other predefined accounts, 100-999 for system accounts/groups. User accounts start from 1000.
 - e.g., root (0); bin (1); daemon (2); mail (8); news (9); diego (261)

Superuser

A special privileged principal

- ▶ **UID 0** and usually the username **root**
- ▶ All security checks are turned off for superuser
 - All access control mechanisms turned off
 - Can become an arbitrary user
 - Can change system clock
- ▶ Some restrictions remain but can be overcome:
 - Cannot write to a read-only file system but can remount it as writeable.
 - Cannot decrypt passwords but can reset them.

Responsibility

- ▶ Used by the operating system for essential tasks like login, recording the audit log, or access to I/O devices.
- ▶ A major weakness of Unix; an attacker achieving superuser status effectively takes over the entire system

Protecting Superuser

- ▶ Privilege separation; create users like **uucp** or **daemon** to deal with networking; if a special users is compromised, not all is lost.
- ▶ root should not be used as a personal account.

User Information

User account stored in `/etc/passwd`:

- ▶ Username: : used when user logs in, 1–32 characters long
- ▶ Password: 'x' indicates that encrypted password is stored in `/etc/shadow`
- ▶ UID: the user ID
- ▶ GID: the primary group ID
- ▶ Name: a comment field
- ▶ Homedir: the path the user will be in when they log in
- ▶ Shell: the absolute path of a command or shell

```
root:x:0:0:root:/root:/bin/bash
daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin
bin:x:2:2:bin:/bin:/usr/sbin/nologin
sys:x:3:3:sys:/dev:/usr/sbin/nologin
sync:x:4:65534:sync:/bin:/bin/sync
games:x:5:60:games:/usr/games:/usr/sbin/nologin
man:x:6:12:man:/var/cache/man:/usr/sbin/nologin
lp:x:7:7:lp:/var/spool/lpd:/usr/sbin/nologin
mail:x:8:8:mail:/var/mail:/usr/sbin/nologin
news:x:9:9:news:/var/spool/news:/usr/sbin/nologin
landscape:x:111:117::/var/lib/landscape:/usr/sbin/nologin
usbmux:x:112:46:usbmux daemon,,,:/var/lib/usbmux:/usr/sbin/nologin
tianweiz:x:1000:1000:Tianwei Zhang:/home/tianweiz:/bin/bash
```

Groups

Users belong to one or more groups.

- ▶ Group info stored in `/etc/group`
 - group name
 - password
 - GID
 - list of users
- ▶ Each user belongs to a primary group. The ID can be found in `/etc/passwd`.

Group is convenient for access control decisions.

- ▶ It is easier for users to share the same access control permission or resources.
- ▶ e.g., put all users allowed to access email in a group called mail

```
root:x:0:
daemon:x:1:
bin:x:2:
sys:x:3:
adm:x:4:syslog,tianweiz
tty:x:5:
disk:x:6:
lp:x:7:
mail:x:8:
news:x:9:
uucp:x:10:
man:x:12:
proxy:x:13:
kmem:x:15:
dialout:x:20:
fax:x:21:
voice:x:22:
cdrom:x:24:tianweiz
floppy:x:25:
tape:x:26:
sudo:x:27:tianweiz
audio:x:29:
dip:x:30:tianweiz
www-data:x:33:
backup:x:34:
operator:x:37:
list:x:38:
```

Processes

A process has a process ID (PID)

- ▶ New processes generated with **exec** or **fork**.

A process has two types of User ID (UID) and Group ID (GID)

- ▶ **Real UID/GID**: typically the UID/GID of the user that logs into the system
- ▶ **Effective UID/GID**: inherited from the parent process or from the file being executed. This determines the permissions for this process

| Process | UID | | GID | |
|------------|------|-----------|--------|-----------|
| | Real | Effective | Real | Effective |
| /bin/login | root | root | system | system |

User **dieter** logs on; the login process verifies the password and changes its UID and GID:

| | | | | |
|------------|--------|--------|-------|-------|
| /bin/login | dieter | dieter | staff | staff |
|------------|--------|--------|-------|-------|

The login process executes the user's login shell:

| | | | | |
|-----------|--------|--------|-------|-------|
| /bin/bash | dieter | dieter | staff | staff |
|-----------|--------|--------|-------|-------|

From the shell, the user executes a command, e.g. **ls**

| | | | | |
|---------|--------|--------|-------|-------|
| /bin/ls | dieter | dieter | staff | staff |
|---------|--------|--------|-------|-------|

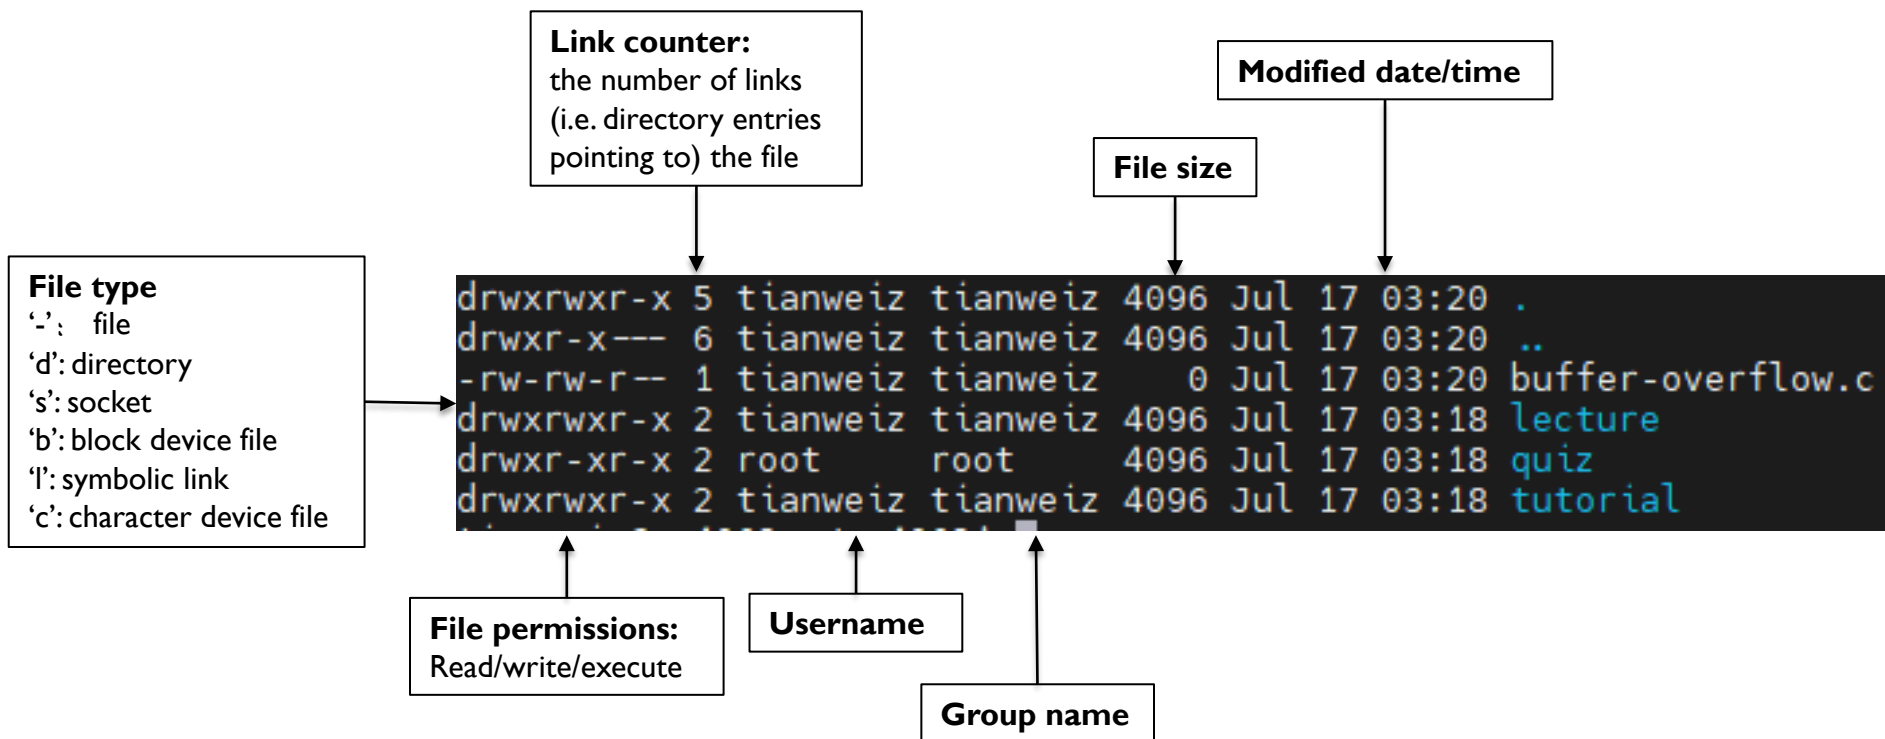
The User executes command **su** to start a new shell as root:

| | | | | |
|-----------|--------|------|-------|--------|
| /bin/bash | dieter | root | staff | system |
|-----------|--------|------|-------|--------|

File, Directory and Devices

Files, directories, memory devices, I/O devices are uniformly treated as **resources**

- ▶ These resources are the objects of access control.
- ▶ Each resource has a single user owner and group owner



File Permission

Three permissions with three subjects

- ▶ Read, Write, Execute
- ▶ Owner, Group, Other
- ▶ Examples:
 - `rw-r--r--`: read and write access for the owner, read access for group and other.
 - `rwX-----`: read, write, and execute access for the owner, no rights to group and other.

Octal Representation

- ▶ `rw-r--r--`: `110 100 100`: `644`
- ▶ `rwXrwxrwx`: `111 111 111`: `777`

Adjust permission:

- ▶ Users can change the permissions:
 - `chmod 754 filename`
 - `chmod u+wx,g+rx,g-w,o+r,o-wx filename`
- ▶ root can change the ownerships:
 - `chown user:group filename`

Controlled Invocation

Superuser privilege is required to execute certain OS functions

- ▶ Example: password changing
 - User passwords are usually stored in the file `/etc/shadow`
 - This file is owned by the root superuser. So a regular user does not have R/W access to it.
 - When a user wants to change his password with the program `passwd`, this program needs to give him additional permissions to write to `/etc/shadow`

SUID: a special flag for a program

- ▶ If SUID is enabled, then user who executes this program will inherit the permissions of that program's owner.
- ▶ `passwd` has such SUID enabled. When a user executes this program to change his password, he gets additional permissions to write the new password to `/etc/shadow`

```
root@cx4062:~# ls -al /usr/bin/passwd
-rwsr-xr-x 1 root root 59976 Mar 14 08:59 /usr/bin/passwd
```

The execute permission of the owner is given as **s** instead of **x**

Security of Controlled Invocation

Many other SUID programs with the owner of root

- ▶ `/bin/passwd`: change password
- ▶ `/bin/login`: login program
- ▶ `/bin/at`: batch job submission
- ▶ `/bin/su`: change UID program

Potential dangers

- ▶ As the user has the program owner's privileges when running a SUID program, the program should only do what the owner intended
- ▶ By tricking a SUID program owned by root to do unintended things, an attacker can act as the root

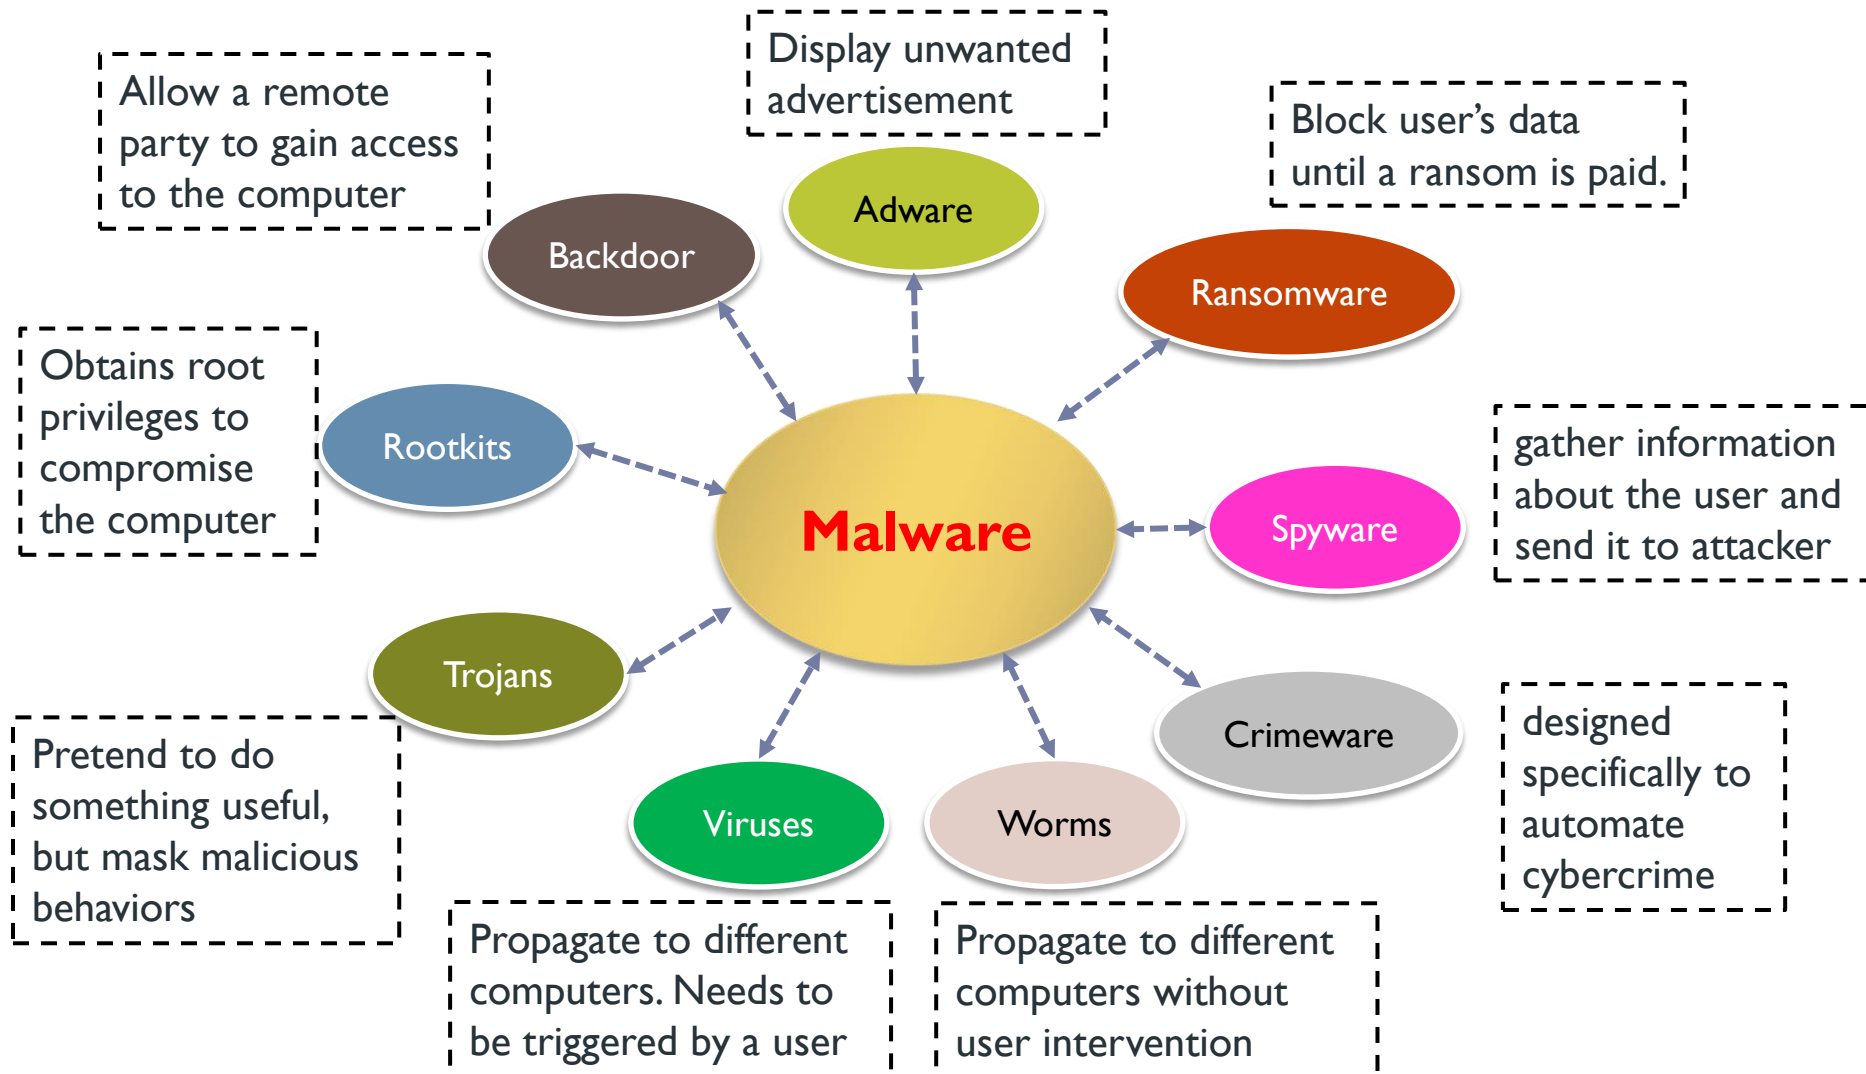
Security consideration

- ▶ All user input (including command line arguments and environment variables) must be processed with extreme care
- ▶ Programs should have SUID status only if it is really necessary.
- ▶ The integrity of SUID programs must be monitored.

Outline

- ▶ Operating System Security Basis
- ▶ Security Protection Stages employed by OS
- ▶ UNIX Security Model
- ▶ **Security Vulnerabilities in OS**

Different Kinds of Malware



Rootkit

Malware that obtains root privileges to compromise the computer

- ▶ Root user does not go through any security checks, and can perform any actions to the system
 - Insert and execute arbitrary malicious code in the system's code path
 - Hide its existence, e.g., malicious process, files, network sockets, from being detected.

How can the attacker gain the root privileges?


- ▶ Vulnerabilities in the software stack: buffer overflow, format string...

There are some common techniques for rootkits to compromise the systems.

System-call Table

A system call is an interface that allows a user-level process to request functions or services from the kernel level.


- ▶ Examples:
 - Process control
 - File management
 - Device management



How to issue a system call?

- ▶ System call table: a table of pointers in the kernel region, to different system call functions.
- ▶ A user process passes the index of the system call and corresponding parameters with the following API:

`syscall(SYS_call, arg1, arg2, ...);`



Highjack System-call Table

Rootkit change pointers of certain entries in the system-call table.

- ▶ Other processes calling these system calls will execute the attacker's code


An example

- ▶ **syscall_open** is used to display the running process (**ps** command)
- ▶ Rootkit redirects this system call to **new_syscall_open**
 - When the object to be opened matches the malicious name, return NULL to hide it
 - Otherwise, call normal **old_syscall_open**

```
1 struct file sysmap = open("System.map-version");
2 long *syscall_addr = read_syscall_table(sysmap);

4 old_syscall_open = syscall_addr[__NR_open];
5 syscall_addr[__NR_open] = new_syscall_open();

7 malicious_object_name = {"xingyi", "bind_shell",
8                           "reverse_shell"...};
9 int new_syscall_open(char *object_name) {
10     if strstr(object_name, malicious_object_name)
11         return NULL;
12     return old_syscall_open(object_name)
13 }
```



Compromise System Call Functions

In addition to change the pointer, rootkit can directly change the system call function.


- ▶ **syscall_open** is used to display the running process (**ps** command)
- ▶ Replace the first 7 bytes of **syscall_open** as jump to **malicious_open**.
 - This faked system call will issue malicious function, restore the original system call and then call the correct one.

```
1 struct file sysmap = open("System.map-version");
2 long *syscall_addr = read_syscall_table(sysmap);
3 syscall_open = syscall_addr[__NR_open];

5 char old_syscall_code[7];
6 memcpy(old_syscall_code, syscall_open, 7);

8 char pt[4];
9 memcpy(pt, (long)malicious_open, 4);
10 char new_syscall_code[7] =
11 {"\xbd", pt[0], pt[1], pt[2], pt[3], // movl %pt, %ebp
12 "\xff", "\xe5"};                // jmp %ebp
13 memcpy(syscall_open, new_syscall_code, 7);

15 int malicious_open(char *object_name) {
16     malicious_function();
17     memcpy(syscall_open, old_syscall_code, 7);
18     return syscall_open(object_name);
19 }
```



Highjack Interrupt Descriptor Table

An interrupt is a signal from the hardware or software to notify the processor that something needs to be handled immediately.

- ▶ After receiving the signal, the processor issues the interrupt handler
- ▶ Interrupt Descriptor Table (IDT): a table of pointers to different interrupt handler functions

Rootkit can alter the pointer in the IDT to make the processor execute wrong functions.

```
1  unsigned char idtr[6];
2  unsigned long idt_addr;
3  __asm__ volatile ("sidt %0" : "=m" (idtr));
4  idt_addr = *((unsigned long *)&idtr[2]);

6  old_idt_handler = idt_addr[handler_id];
7  idt_addr[handler_id] = new_idt_handler;

9  void new_idt_handler(args){
10     malicious_function();
11     return old_idt_handler(args);
12 }
```

