

SC4000/CZ4041/CE4041: Machine Learning

Lesson 5: Decision Tree

Kelly KE

School of Computer Science and Engineering,
NTU, Singapore

Acknowledgements: slides are adapted from the lecture notes of the books “Introduction to Machine Learning” (Chap. 9) and “Introduction to Data Mining” (Chap. 4).

An Illustrative Example

- Consider the problem of predicting whether a loan applicant will repay his/her loan obligation (no cheat) or become delinquent (cheat).

Predefined categories

Example



An Illustrative Example (cont.)

- **Training set:** constructed by examining the records of previous borrowers.

<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Bayesian classifiers: $P(y|x)$

→ $f(x) = y$

Decision tree



Motivation of Decision Trees

- Suppose a new applicant submits a loan application. How do we decide whether to approve or reject the application?

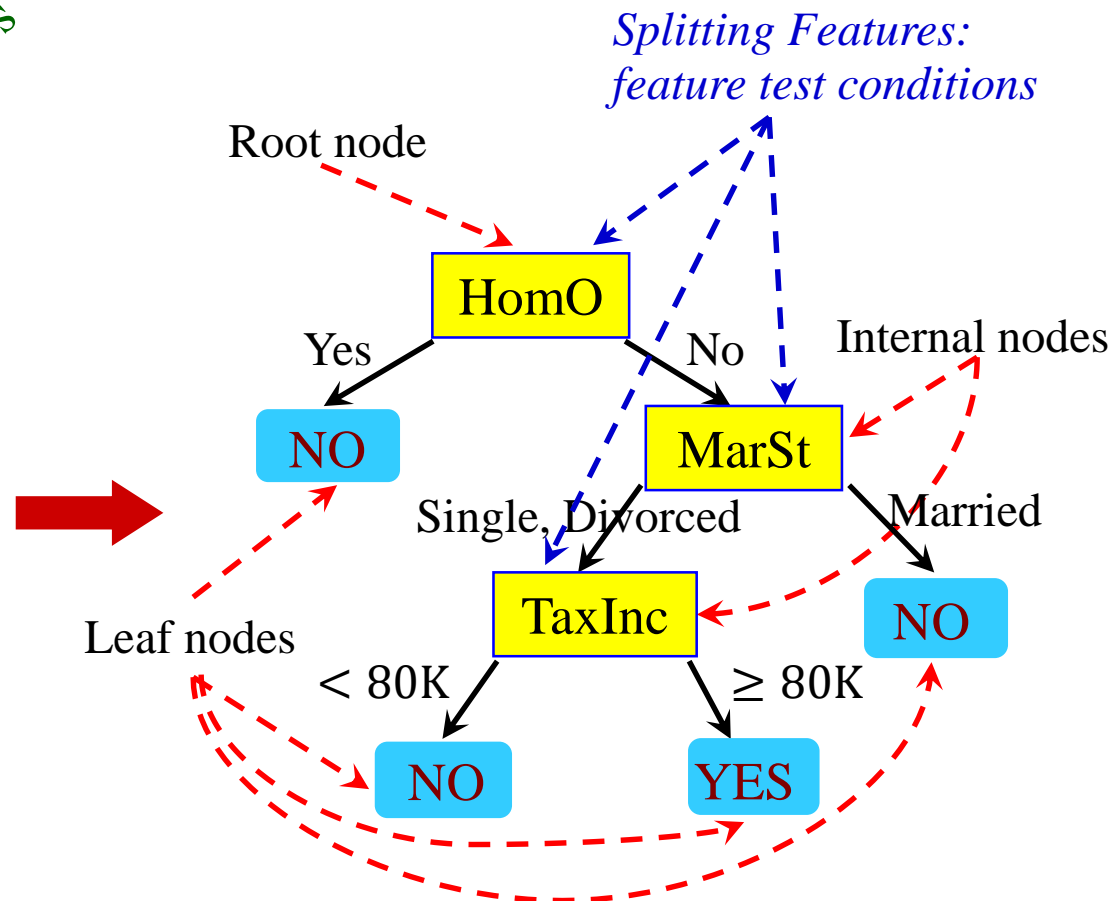
Home Owner	Marital Status	Taxable Income	Cheat
No	Married	80K	?

- To pose a series of questions about the profile of the applicant:
 - Whether the applicant is a home owner? If yes, then he/she may repay the loan obligation with a high probability.
 - After that, we may ask a follow-up question: what is the applicant's marital status? ...

Example of a Decision Tree

<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

binary discrete continuous class



<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Induction
(induce a tree)

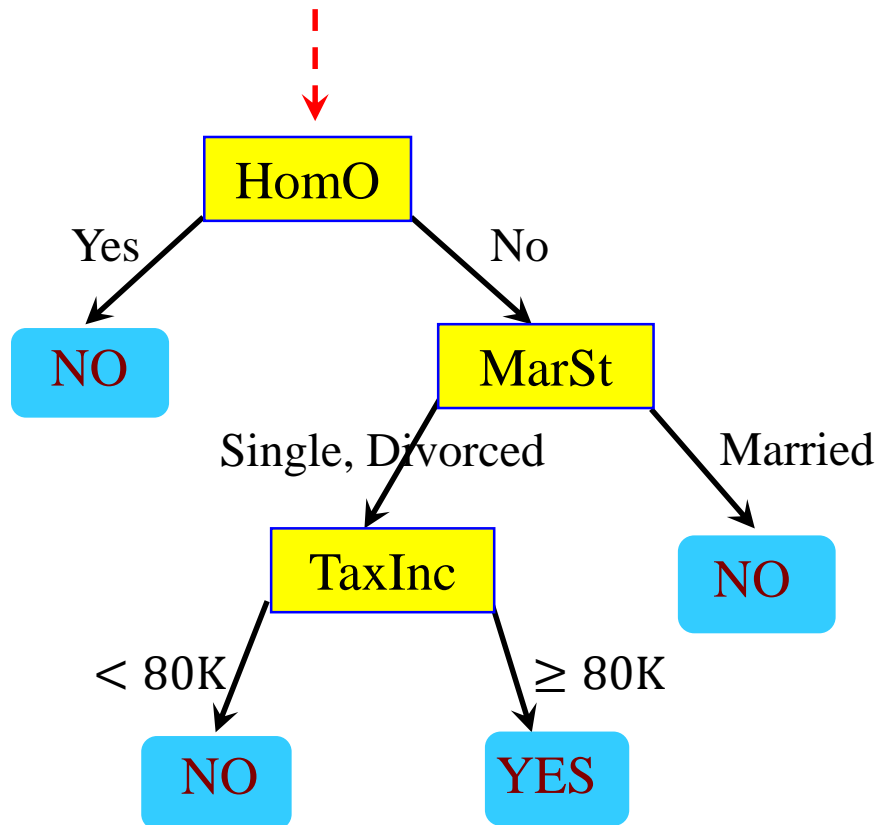
Decision
Tree

<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
11	No	Single	55K	?
12	Yes	Divorce	80K	?
13	Yes	Married	110K	?
14	No	Single	95K	?
15	No	Married	67K	?

Deduction
(apply the tree)

Apply Decision Tree to Test Data

Start from the root of tree.



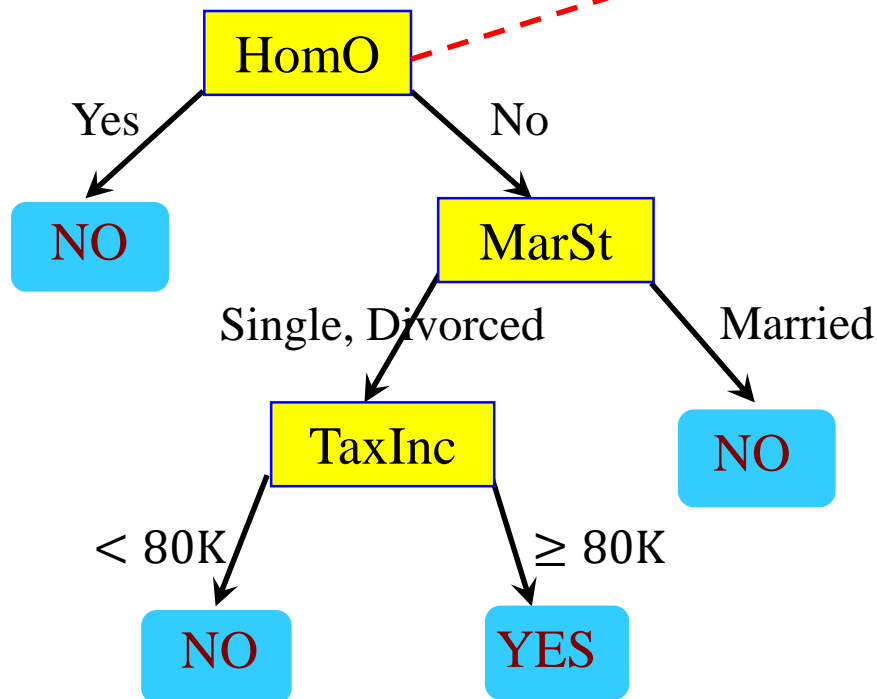
Test Data

Home Owner	Marital Status	Taxable Income	Cheat
No	Married	80K	?

Apply Model to Test Data

Test Data

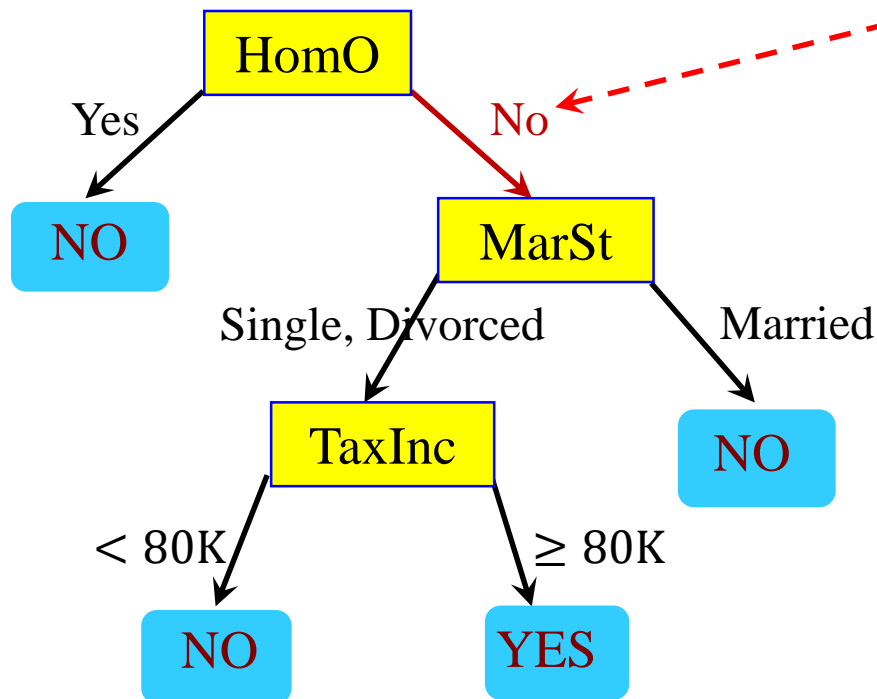
Home Owner	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data

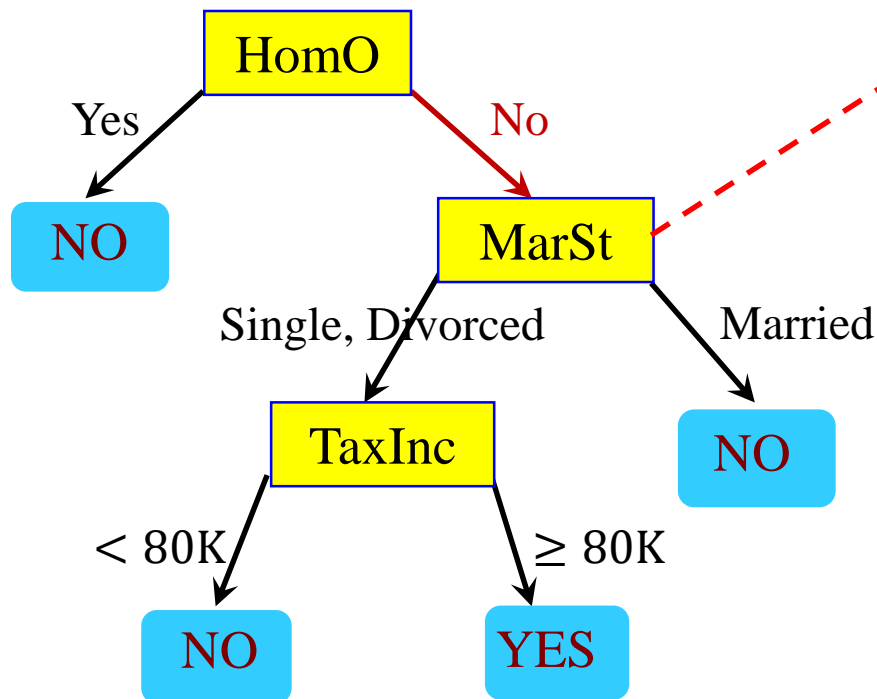
Home Owner	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data

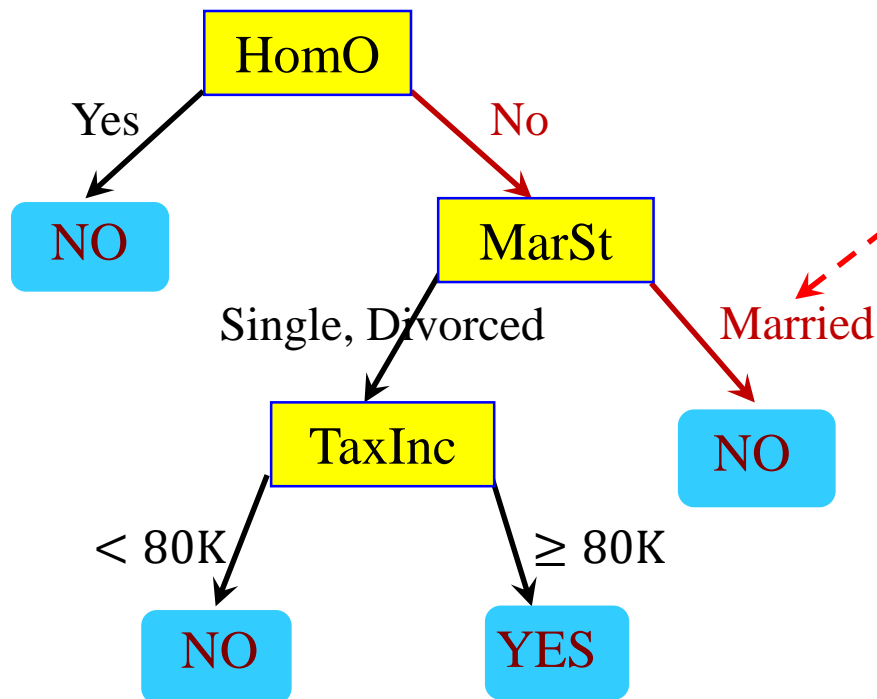
Home Owner	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data

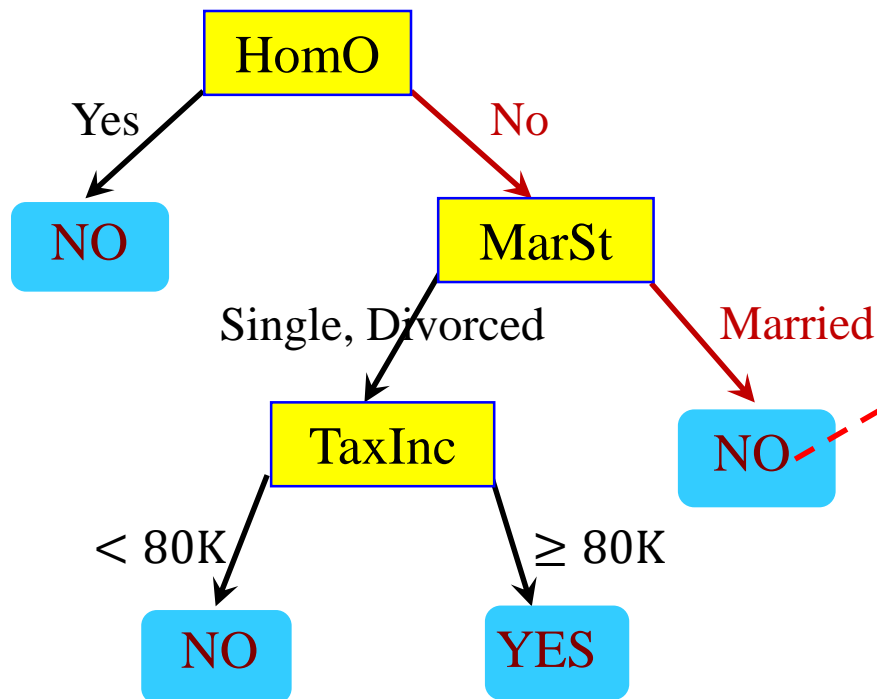
Home Owner	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data

Home Owner	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Assign "No" to Cheat

<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Tree induction algorithm

Induction
(induce a tree)

Decision
Tree

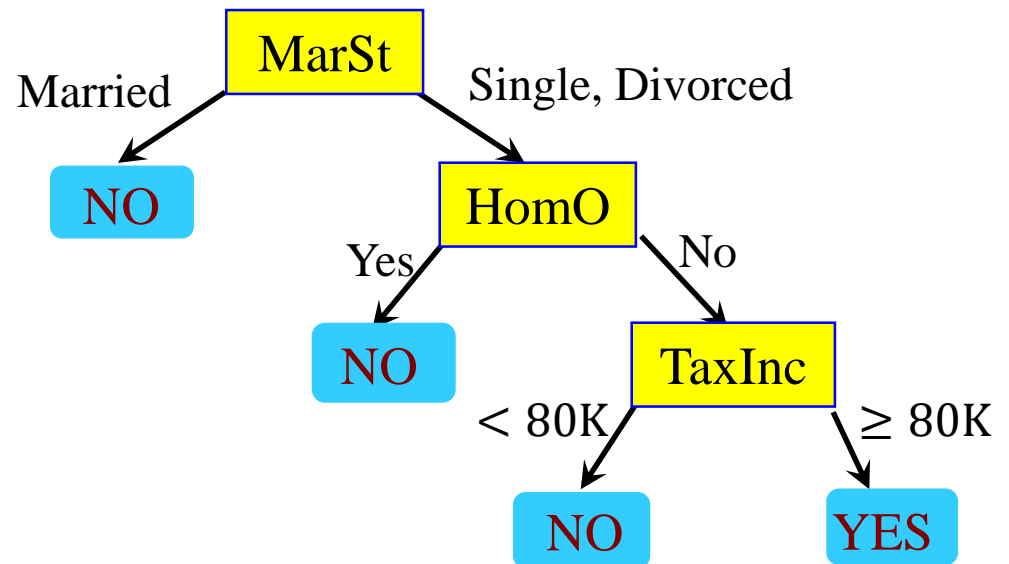
Deduction
(apply the tree)

<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
11	No	Single	55K	?
12	Yes	Divorce	80K	?
13	Yes	Married	110K	?
14	No	Single	95K	?
15	No	Married	67K	?

Another Example of Decision Tree

<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

binary
discrete
continuous
class

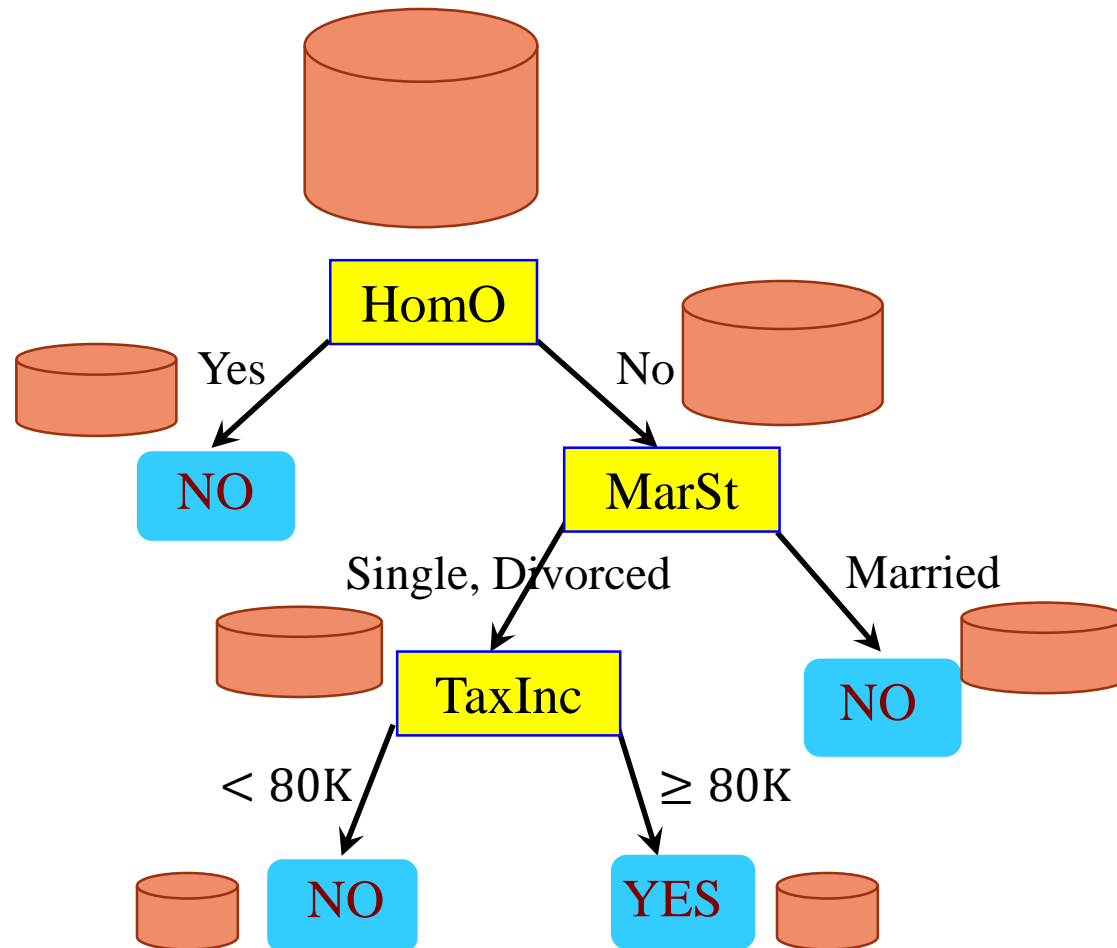


There could be more than one tree that fits the same data!

Decision Tree Induction Algorithms

- CART, ID3, C4.5, etc.
- High-level basic idea:
 - Given a new application, the goal of asking a series of questions (checking properties of the profile) one by one is to find similar profiles (applicants) in the past until it is confident to make a decision based on the labels (whether cheat or no cheat) of the similar applicants.
 - A tree can be learned by splitting training data into subsets based on outcomes of a feature test.
 - This process is recursively applied on each derived subset until the subset at a node has all the same class or there is no improvement for prediction.

Illustration



High-level Algorithm

<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Initial tree

Don't
Cheat

Default class: **Don't Cheat**
(most of the borrowers
successfully repaid their loans)



Home Owner test
condition is selected

HomO

High-level Algorithm (cont.)

<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Initial tree

Don't
Cheat

Default class: **Don't Cheat**
(most of the borrowers
successfully repaid their loans)

Home Owner test
condition is selected

HomO


Yes

No




High-level Algorithm (cont.)

<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
4	Yes	Married	120K	No
7	Yes	Divorced	220K	No

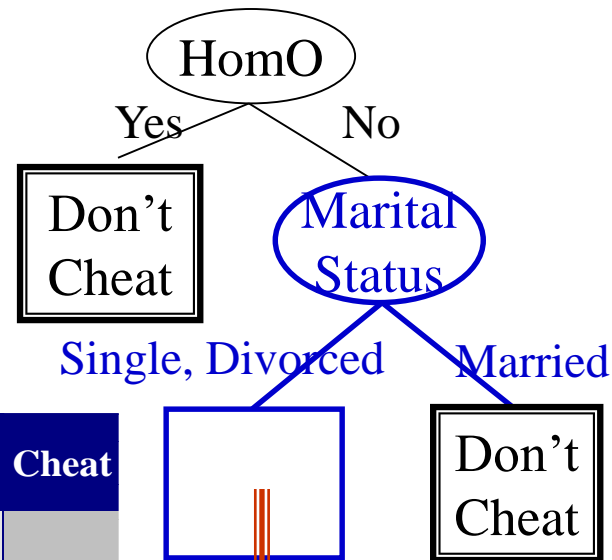


Marital Status
test condition
is selected




<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
2	No	Married	100K	No
3	No	Single	70K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

High-level Algorithm (cont.)



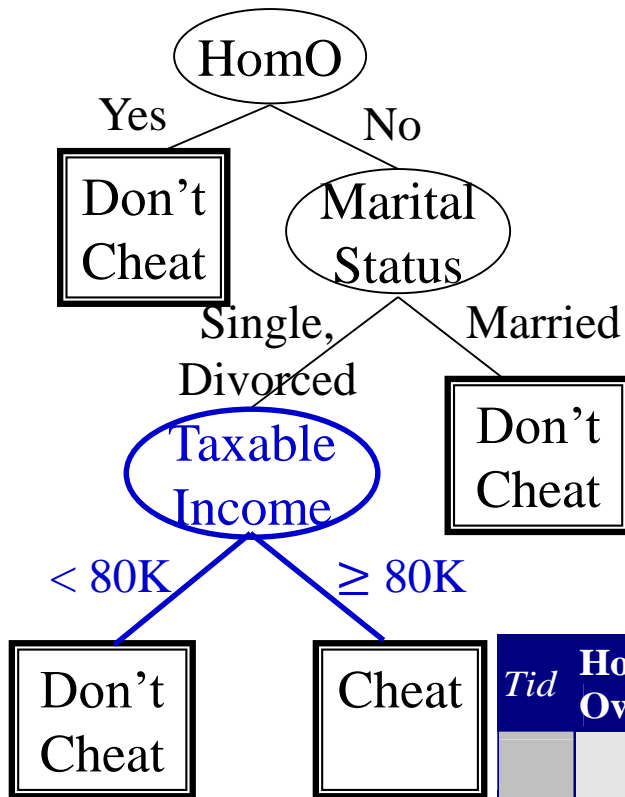
<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
3	No	Single	70K	No
5	No	Divorced	95K	Yes
8	No	Single	85K	Yes
10	No	Single	90K	Yes

Taxable Income test condition is selected



<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
2	No	Married	100K	No
6	No	Married	60K	No
9	No	Married	75K	No

High-level Algorithm (cont.)

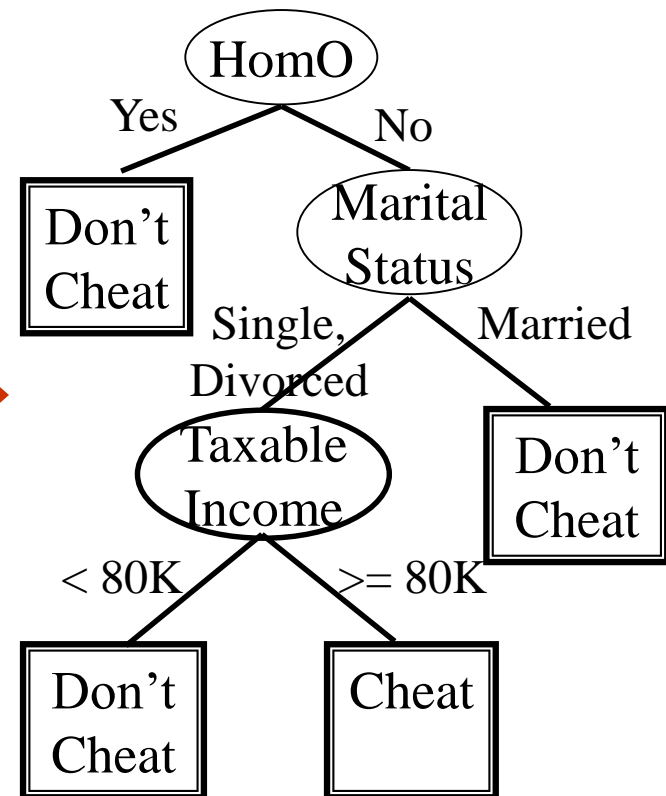


<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
3	No	Single	70K	No

<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
5	No	Divorced	95K	Yes
8	No	Single	85K	Yes
10	No	Single	90K	Yes

High-level Algorithm (cont.)

<i>Tid</i>	Home Owner	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes



High-level Algorithm: Summary

- Let D_t be the set of training records that reach a node t
- General Procedure:
 - If D_t contains records that belong to the same class y_t , then t is a leaf node labeled as y_t (e.g., Yes or No)
 - Otherwise if D_t is an empty set, then t is a leaf node labeled by the default class, y_d (e.g., No)
 - Otherwise if D_t contains records that belong to more than one class, then a feature is selected to conduct condition test to split the data into smaller subsets.
 - A child node is created for each outcome of the test condition and the records in D_t are distributed to the children based on the outcomes
 - Recursively apply the procedure to each subset

Tree Induction

- Greedy strategy
 - Split the records based on a feature test that optimizes certain criterion
- Issues
 - Determine how to split the records
 - How to specify the feature test condition?
 - How to determine the best split?
 - Determine when to stop splitting

Tree Induction


- Greedy strategy
 - Split the records based on a feature test that optimizes certain criterion
- Issues
 - Determine how to split the records
 - [How to specify the feature test condition?](#)
 - How to determine the best split?
 - Determine when to stop splitting

How to Specify Test Condition?

- Depends on feature types
 - Discrete
 - Continuous
- Depends on number of ways to split
 - Binary split
 - Multi-way split


Splitting Based on Binary Features

- Generate two potential outcomes

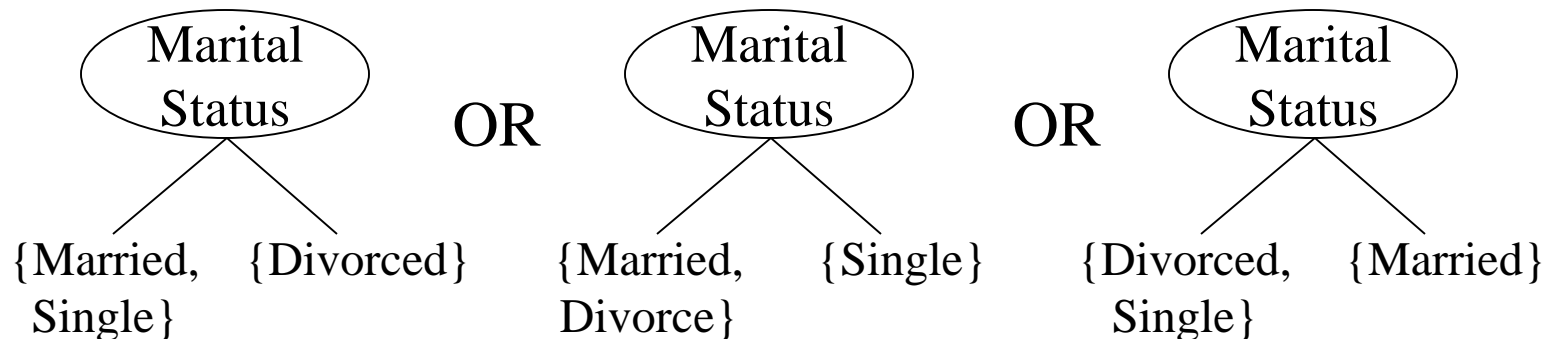


Splitting Based on Discrete Features (more than two distinct values)


- Multi-way split: Use as many partitions as distinct values



- Binary split: Divide values into two subsets. Need to find optimal partitioning




Splitting Based on Continuous Features



(a) Binary split

$(x_j < v)$ or $(x_j \geq v)$



(a) Multi-way split

Discretization

- Consider all possible splits and find the best cut
- Can be very computationally intensive

Tree Induction

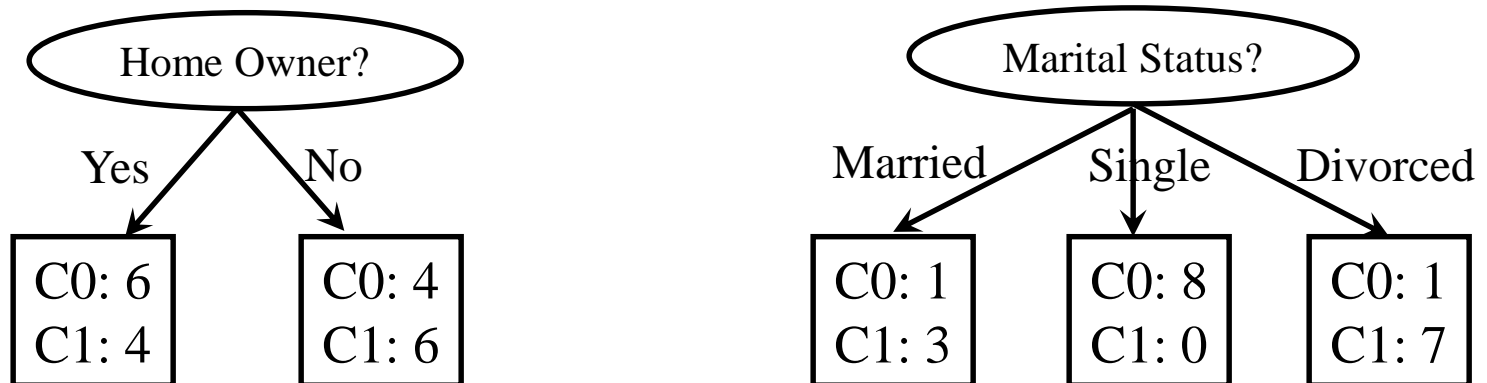
- Greedy strategy
 - Split the records based on a feature test that optimizes certain criterion
- Issues
 - Determine how to split the records
 - How to specify the attribute test condition?
 - [How to determine the best split?](#)
 - Determine when to stop splitting

How to Determine the Best Split

Before Splitting:

10 records of class 0
10 records of class 1

After Splitting:



Which condition test is the best?

How to Determine the Best Split (cont.)

- An intuitive idea:
 - Nodes with homogeneous class distribution are preferred
- Need a measure of node impurity

C0: 5
C1: 5

Non-homogeneous,
High degree of impurity

C0: 9
C1: 1

Homogeneous,
Low degree of impurity


- A split criterion is defined in terms of the difference in degree of node impurity before and after splitting

Measure of Impurity: Entropy

- Entropy at a given node t : The probability of class c at the node t

$$\text{Entropy}(t) = - \sum_c \boxed{P(y = c; t)} \log_2 P(y = c; t)$$


- Entropy for a given node t , with binary classes:



Entropy Properties

- Entropy at a given node t :

$$\text{Entropy}(t) = - \sum_c P(y = c; t) \log_2 P(y = c; t)$$

- Maximum: $\log_2 C$  Total number of all possible values of y , i.e., #classes
when records are equally distributed among all classes
- Minimum: 0
when all records belong to one class

Examples of Computing Entropy

$$\text{Entropy}(t) = - \sum_c P(y = c; t) \log_2 P(y = c; t)$$

Y₁	0
Y₂	6

$$P(Y_1) = \frac{0}{6} = 0 \quad P(Y_2) = \frac{6}{6} = 1$$
$$\text{Entropy} = -0 \log_2(0) - 1 \log_2(1) = -0 - 0 = 0$$

Y₁	1
Y₂	5

$$P(Y_1) = \frac{1}{6} \quad P(Y_2) = \frac{5}{6}$$
$$\text{Entropy} = -\left(\frac{1}{6}\right) \log_2 \left(\frac{1}{6}\right) - \left(\frac{5}{6}\right) \log_2 \left(\frac{5}{6}\right) = 0.65$$

Y₁	2
Y₂	4

$$P(Y_1) = \frac{2}{6} \quad P(Y_2) = \frac{4}{6}$$
$$\text{Entropy} = -\left(\frac{2}{6}\right) \log_2 \left(\frac{2}{6}\right) - \left(\frac{4}{6}\right) \log_2 \left(\frac{4}{6}\right) = 0.92$$

Information Gain: Motivation

- Recall: a split criterion should be defined in terms of the difference in degree of node impurity before and after splitting
- Information Gain:

$$\Delta_{\text{info}} = \text{Entropy}(\text{parent node}) - \text{Entropy}(\text{children nodes})$$

- Choose a feature whose condition test maximizes the gain (minimizes the weighted average impurity measures of the child nodes)


A motivating example of Information Gain

Node P

-1	5
+1	5

Entropy before splitting

$$\text{Entropy } E_p = 1$$



#instances in node A_1

#instances in node A_2

$$E_A = \frac{\#A_1}{\#P} E_{A_1} + \frac{\#A_2}{\#P} E_{A_2}$$

#instances in node P

$$= \frac{3}{10} 0.92 + \frac{7}{10} 0.99$$

$$= 0.97$$

Use the **Information Gain**: $E_p - E_A$ to measure the difference of impurity before and after splitting on A

$$E_p - E_A = 0.03$$

Goal: find a feature with maximum information gain to conduct condition test

A motivating example
of Information Gain

Node P

-1	5
+1	5

Entropy before splitting

Entropy $E_p = 1$

A

Yes

No

Node A_1

Node A_2

-1	2
+1	1

-1	3
+1	4

Entropy

Entropy

$$E_{A_1} = 0.92$$

$$E_{A_2} = 0.99$$

Entropy after splitting on A

$$E_A = \frac{3}{10} \times 0.92 + \frac{7}{10} \times 0.99 = 0.97$$

B

Yes

No

Node B_1

Node B_2

-1	1
+1	5

-1	4
+1	0

Entropy

Entropy

$$E_{B_1} = 0.65$$

$$E_{B_2} = 0$$

Entropy after splitting on B

$$E_B = \frac{6}{10} \times 0.65 + \frac{4}{10} \times 0 = 0.39$$

Information Gain: $E_p - E_A = 0.03$

<

Information Gain: $E_p - E_B = 0.61$



Information Gain: Definition

- Suppose a parent node t is split into P partitions (children)
- Information Gain:

$$\Delta_{\text{info}} = \text{Entropy}(t) - \sum_{j=1}^P \frac{n_j}{n} \text{Entropy}(j)$$

Number of examples at node t (points to n)


Number of examples at child j (points to n_j)

- To choose a feature whose condition test maximizes the gain (minimizes the weighted average impurity measures of the children nodes)


Information Gain: Practice

- Suppose features A and B are binary
- Based on Information Gain, which feature should be selected to split data?

	Parent
Y_1	6
Y_2	6



	N1	N2
Y_1	5	1
Y_2	2	4



	N3	N4
Y_1	6	0
Y_2	3	3

Entropy(Parent)

$$= -\left(\frac{6}{12}\right)\log_2\left(\frac{6}{12}\right) - \left(\frac{6}{12}\right)\log_2\left(\frac{6}{12}\right) = 1$$

	Parent
Y_1	6
Y_2	6

$$\text{Entropy(N1)} = -\left(\frac{5}{7}\right)\log_2\left(\frac{5}{7}\right) - \left(\frac{2}{7}\right)\log_2\left(\frac{2}{7}\right) = 0.8631$$

$$\text{Entropy(N2)} = -\left(\frac{1}{5}\right)\log_2\left(\frac{1}{5}\right) - \left(\frac{4}{5}\right)\log_2\left(\frac{4}{5}\right) = 0.7219$$

$$\text{Entropy}(\text{Split}_A) = \left(\frac{7}{12}\right) \times 0.8631 + \left(\frac{5}{12}\right) \times 0.7219 = 0.8043$$

Split on A

	N1	N2
Y_1	5	1
Y_2	2	4

$$\text{Entropy(N3)} = -\left(\frac{3}{9}\right)\log_2\left(\frac{3}{9}\right) - \left(\frac{6}{9}\right)\log_2\left(\frac{6}{9}\right) = 0.9183$$

$$\text{Entropy(N4)} = -\left(\frac{0}{3}\right)\log_2\left(\frac{0}{3}\right) - \left(\frac{3}{3}\right)\log_2\left(\frac{3}{3}\right) = 0$$

$$\text{Entropy}(\text{Split}_B) = \left(\frac{9}{12}\right) \times 0.9183 + \left(\frac{3}{12}\right) \times 0 = 0.6887$$

Split on B

	N3	N4
Y_1	6	0
Y_2	3	3

$$\begin{aligned} &\text{Entropy}(\text{Parent}) \\ &= -\left(\frac{6}{12}\right)\log_2\left(\frac{6}{12}\right) - \left(\frac{6}{12}\right)\log_2\left(\frac{6}{12}\right) = 1 \end{aligned}$$

	Parent
Y_1	6
Y_2	6

Split on A

$$\text{Entropy}(\text{Split}_A) = \left(\frac{7}{12}\right) \times 0.8631 + \left(\frac{5}{12}\right) \times 0.7219 = 0.8043$$

	N1	N2
Y_1	5	1
Y_2	2	4

Split on B

$$\text{Entropy}(\text{Split}_B) = \left(\frac{9}{12}\right) \times 0.9183 + \left(\frac{3}{12}\right) \times 0 = 0.6887$$

	N3	N4
Y_1	6	0
Y_2	3	3

$$\Delta_{\text{info}}(A) = 1 - 0.8043 = 0.1957$$

$$\Delta_{\text{info}}(B) = 1 - 0.6887 = 0.3113$$

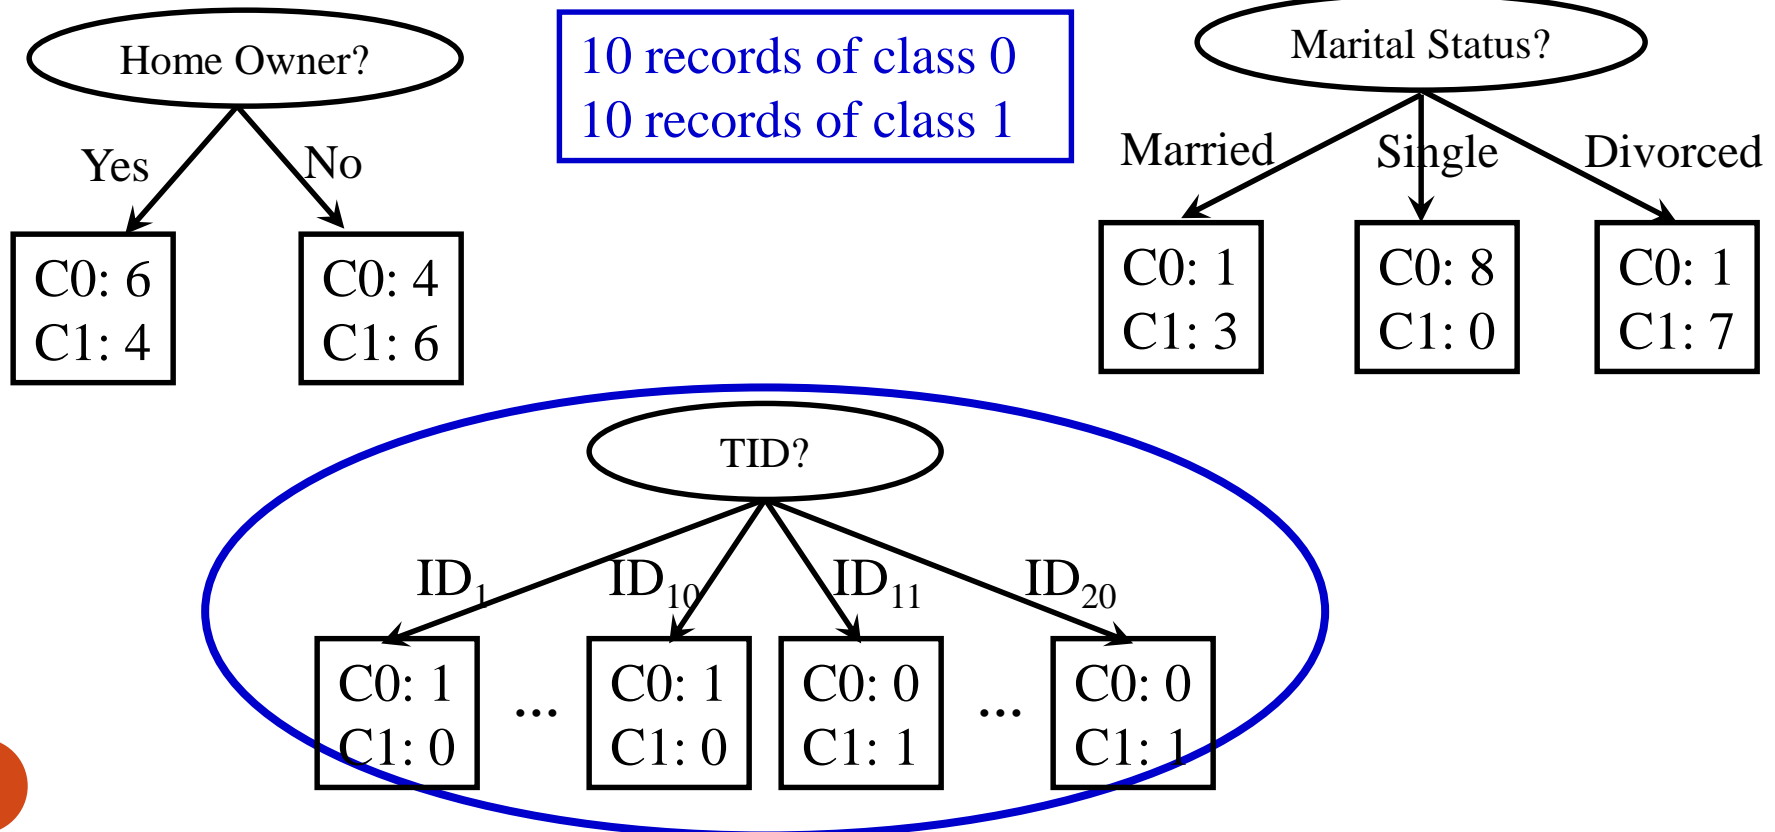
Choose B to conduct
condition test to split data

Information Gain: Limitation

- Disadvantage: tends to prefer splits that result in large number of partitions, each being small but pure

Before Splitting:

10 records of class 0
10 records of class 1



Splitting Based on Entropy (cont.)

- Gain Ratio:

$$\Delta_{\text{InfoR}} = \frac{\Delta_{\text{info}}}{\text{SplitINFO}}$$

Penalty on large number
of small partitions

Maximum: $\log_2 P$

where

$$\text{SplitINFO} = - \sum_{i=1}^P \frac{n_i}{n} \log_2 \left(\frac{n_i}{n} \right)$$

- Parent node t of n instances is split into P partitions (children), n_i is the number of instances in partition i
- Higher entropy partitioning (large number of small partitions) is penalized!

Refer to Question 1 in tutorial for practice

Tree Induction

- Determine how to split the data
 - How to specify the feature test condition?
 - How to determine the best split?
- Determine when to stop splitting

Stopping Criteria for Tree Induction

- Stop expanding a node when all the data instances belong to the same class
 - Ideal case but not always possible
- Stop expanding a node when all the data instances have similar feature values
- Early termination
 - Useful to avoid the overfitting issue (will be introduced next week)

Decision Tree Classifiers: Summary

- Easy to interpret
- Efficient in both training and testing
- Effective for datasets with a lot of categorical features
- Used as a base classifier in many ensemble learning approaches (will be introduced in the 2nd half)

Implementation --- scikit-learn

- API: `sklearn.tree.DecisionTreeClassifier`

<https://scikit-learn.org/stable/modules/classes.html#module-sklearn.tree>

sklearn.tree: Decision Trees

The `sklearn.tree` module includes decision tree-based models for classification and regression.

User guide: See the [Decision Trees](#) section for further details.

`tree.DecisionTreeClassifier(*
[, criterion, ...])` A decision tree classifier.

`tree.DecisionTreeRegressor(*
[, criterion, ...])` A decision tree regressor.

`tree.ExtraTreeClassifier(*
[, criterion, ...])` An extremely randomized tree classifier.

`tree.ExtraTreeRegressor(*
[, criterion, ...])` An extremely randomized tree regressor.

`tree.export_graphviz(decision_tree[, ...])` Export a decision tree in DOT format.

`tree.export_text(decision_tree, *[, ...])` Build a text report showing the rules of a decision tree.

Plotting

`tree.plot_tree(decision_tree, *
[, ...])` Plot a decision tree.

An Example

```
>>> from sklearn import tree
```

```
>>> ...
```

Data preprocessing

```
>>> dtC = tree.DecisionTreeClassifier()
```

```
>>> dtC.fit(X, y)
```

```
>>> pred= dtC.predict(X_t)
```

Thank you!