

SC4000/CZ4041/CE4041: Machine Learning

Lesson 6b: *K*-NN Classifiers

Kelly KE

School of Computer Science and Engineering,
NTU, Singapore

Acknowledgements: some figures are adapted from the lecture notes of the book
“Introduction to Data Mining” (Chap. 5)

Training data

		F1	F2	F3	F4	F5	F6	...	
y_1	+1	1	1	0	0	1	0	...	x_1
y_2	-1	0	0	1	1	0	1	...	x_2
...									
y_N	-1	0	1	0	0	1	1	...	x_N

Some classification algorithm

Predicted label

$$f: x \rightarrow y$$
$$y^* = f(x^*)$$

Test data


	F1	F2	F3	F4	F5	F6	...	
?	1	1	0	0	1	0	...	x^*

Inductive Learning

Training data

		F1	F2	F3	F4	F5	F6	...	
y_1	+1	1	1	0	0	1	0	...	x_1
y_2	-1	0	0	1	1	0	1	...	x_2
...									
y_N	-1	0	1	0	0	1	1	...	x_N

Predicted label



Test data

	F1	F2	F3	F4	F5	F6	...	
?	1	1	0	0	1	0	...	x^*


Lazy Learning

Instance Based Classifiers

- K -Nearest Neighbors classifier
 - Uses K “closest” points (nearest neighbors) for performing classification


Nearest-Neighbor Classifiers

Unknown instance




- Requires three things
 - The set of stored labeled instances
 - Distance metric to compute distance between instances
 - The value of K , the number of nearest neighbors to retrieve
- To classify an unknown instance:
 - Compute distance to all the training instances
 - Identify K nearest neighbors
 - Use class labels of nearest neighbors to determine the class label of the unknown instance (e.g., by taking majority vote)


Definition of Nearest Neighbors



(a) 1-nearest
neighbor



(b) 2-nearest
neighbor



(c) 3-nearest
neighbor

K -nearest neighbors of an instance x are data points that have the K smallest distance to x

Distance Metric

- Compute distance between two data points in a d -dimensional space:

\mathbf{x}_i			
x_{i1}	x_{i2}	...	x_{id}

\mathbf{x}_j			
x_{j1}	x_{j2}	...	x_{jd}

- Euclidean distance

Inner product between \mathbf{x}_i and \mathbf{x}_j :

$$\mathbf{x}_i \cdot \mathbf{x}_j = \sum_{k=1}^d x_{ik} x_{jk}$$


$$\begin{aligned} d(\mathbf{x}_i, \mathbf{x}_j) &= \sqrt{\sum_{k=1}^d (x_{ik} - x_{jk})^2} = \sqrt{(\mathbf{x}_i - \mathbf{x}_j) \cdot (\mathbf{x}_i - \mathbf{x}_j)} \\ &= \|(\mathbf{x}_i - \mathbf{x}_j)\|_2 \end{aligned}$$

L_2 norm distance

$$L_2 \text{ norm of } \mathbf{x}_i: \|\mathbf{x}_i\|_2 = \sqrt{\sum_{k=1}^d x_{ik}^2}$$

Value of K

- Choosing the value of K :
 - If K is too small, sensitive to noise points
 - If K is too large, neighborhood may include points from other classes



Determine Class Label

- Determine the class from nearest neighbor list
 - Take the majority vote of class labels among the K -nearest neighbors
- Given test data \mathbf{x}^* , majority voting:

$$y^* = \arg \max_c \sum_{(\mathbf{x}_i, y_i) \in \mathcal{N}_{\mathbf{x}^*}} I(c = y_i)$$

Indicator function that returns 1 if its input is true, otherwise 0

Nearest neighbors of the test instance \mathbf{x}^*

- Every neighbor has the same impact on the classification
- This indeed makes the algorithm sensitive to the choice of K

Revised Voting Scheme


- Solution: distance-weighted voting
 - Weight the influence of each nearest neighbor \mathbf{x}_i according to its distance to the test instance \mathbf{x}^* :

$$w_i = \frac{1}{d(\mathbf{x}^*, \mathbf{x}_i)^2}$$

$$y^* = \arg \max_c \sum_{(\mathbf{x}_i, y_i) \in \mathcal{N}_{\mathbf{x}^*}} w_i \times I(c = y_i)$$

Example

Consider a binary classification problem, and a 5-NN classifier



Training record	Class label	Distance to test instance
1	+	3
2	+	3.5
3	+	4
4	-	1.5
5	-	2

- Majority voting:

$$+: 3 > -: 2$$

- Distance-weighted voting:



Tutorial

Other Issues

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{\sum_{k=1}^d (x_{ik} - x_{jk})^2}$$

- Scaling issues
 - Feature may need to be scaled to prevent distance from being dominated by some features
 - Example:
 - height of a person may vary from 1.5m to 1.8m
 - weight of a person may vary from 40kg to 200kg
 - income of a person may vary from \$10K to \$1M
 - Solution: normalization on features of different scales

Normalization

- Min-max normalization: to $[min_{new}, max_{new}]$
 - Example: To normalize income ranging from \$12,000 to \$98,000 to $[0.0, 1.0]$, what is the value for \$73,600 after normalization?

$$v_{new} = \frac{v_{old} - min_{old}}{max_{old} - min_{old}} (max_{new} - min_{new}) + min_{new}$$

$$73,600 \quad \longrightarrow \quad \frac{73600 - 12000}{98000 - 12000} (1.0 - 0) + 0 = 0.716$$

Normalization (cont.)

- Standardization (z-score normalization) (μ : mean, σ : standard deviation):

$$v_{new} = \frac{v_{old} - \mu_{old}}{\sigma_{old}} \quad \longrightarrow \quad \mu_{new} = 0, \text{ and } \sigma_{new} = 1$$

- Example: Let $\mu = 54,000$, $\sigma = 16,000$. What is the value for \$73,600 after standardization?

$$\frac{73600 - 54000}{16000} = 1.225$$

Summary of NN Classifier

- The K -NN classifier is a lazy learner
 - It does not build models explicitly.
 - “Training” is very efficient.
 - Classifying unknown test instances is relatively expensive.

Implementation

```
>>> from sklearn.neighbors import KNeighborsClassifier
```

...

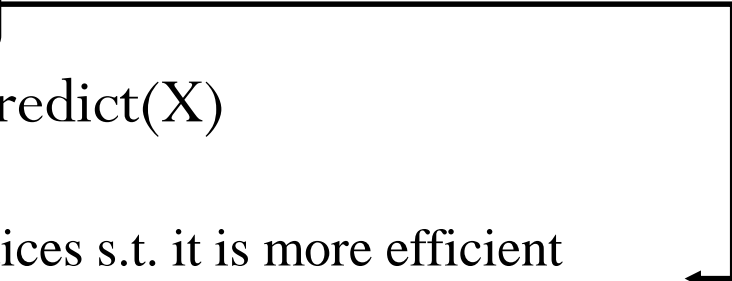
set number of neighbors

```
>>> knnC = KNeighborsClassifier(n_neighbors=3)
```

```
>>> knnC.fit(X, y)
```

```
>>> pred = knnC.predict(X)
```

Build indices s.t. it is more efficient
when making predictions on test data



Thank you!