

CZ4041/SC4000: Machine Learning

Lesson 8b: Linear Regression

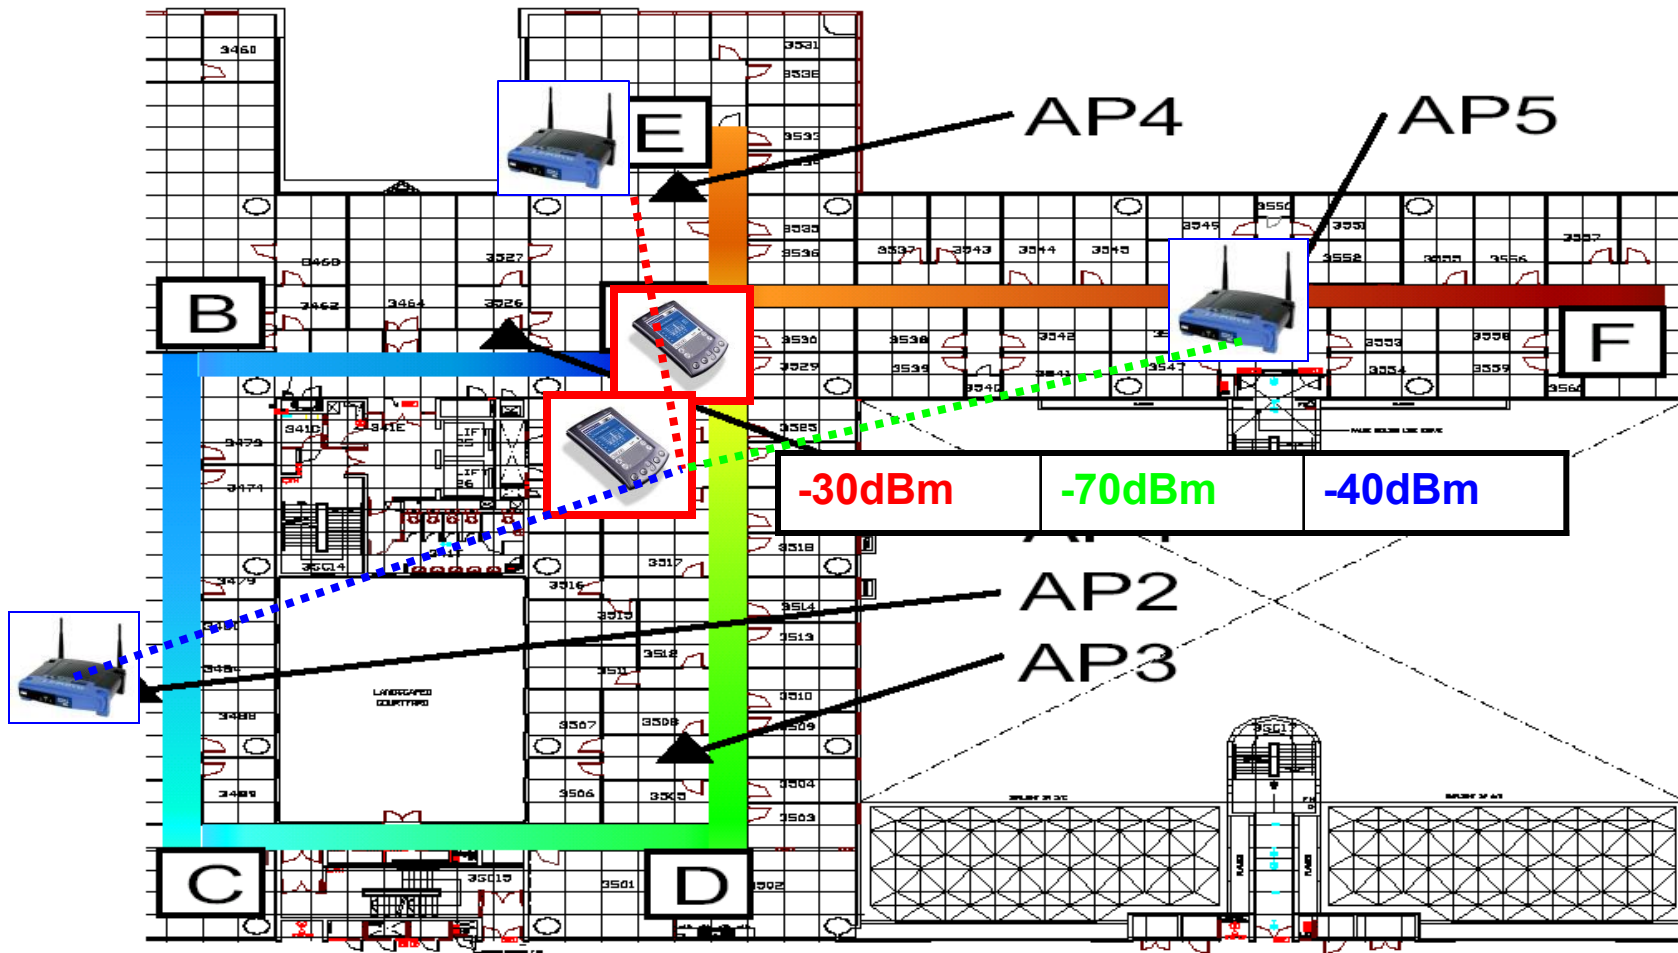
LI Boyang, Albert

School of Computer Science and Engineering,
NTU, Singapore


Acknowledgements: Slides are modified from the version prepared by Dr. Sinno Pan.

Regression: Example

Indoor WiFi localization



Linear Fitting




Linear Regression Model

- A special case, an instance is represented by one input feature
- To learn a linear function $f(x)$ in terms of w (drop bias term b for simplicity) from $\{x_i, y_i\}, i = 1, \dots, N$

$$f(x) = w \cdot x$$

- Such that the difference (i.e., error) between the predicted values $f(x_i)$'s and the ground-truth values y_i 's is as small as possible

Linear Regression Model (cont.)



Linear Regression Model (cont.)

- Suppose sum-of-squares error is used

$$\mathcal{L}(w) = \frac{1}{2} \sum_{i=1}^N (f(x_i) - y_i)^2 = \frac{1}{2} \sum_{i=1}^N (w \times x_i - y_i)^2$$


- Learn the linear model in terms of w by minimizing the error

$$w^* = \arg \min_w \mathcal{L}(w)$$

Linear Regression Model (cont.)

- To solve the unconstrained minimization problem, we can set the derivative of $E(w)$ w.r.t. w to zero


$$\frac{\partial \mathcal{L}(w)}{\partial w} = \frac{\partial \left(\frac{1}{2} \sum_{i=1}^N (w \times x_i - y_i)^2 \right)}{\partial w} = 0$$




$$\sum_{i=1}^N (w \times x_i - y_i) \times x_i = 0$$

Linear Regression Model (cont.)

$$\sum_{i=1}^N (w \times x_i - y_i) \times x_i = 0$$



$$w \sum_{i=1}^N x_i^2 - \sum_{i=1}^N y_i \times x_i = 0$$



$$w = \frac{\sum_{i=1}^N y_i \times x_i}{\sum_{i=1}^N x_i^2}$$

Linear Regression Model (cont.)

- To learn a linear function $f(\mathbf{x})$ in more general form in terms of \mathbf{w} and b ,

$$f(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} + b$$

- By defining $w_0 = b$, and $X_0 = 1$, \mathbf{w} and \mathbf{x} are of $d + 1$ dimensions

$$f(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x}$$

Linear Regression Model (cont.)

- Suppose sum-of-squares error is used

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N (f(\mathbf{x}_i) - y_i)^2 = \frac{1}{2} \sum_{i=1}^N (\mathbf{w} \cdot \mathbf{x}_i - y_i)^2$$

- Learn the linear model in terms of \mathbf{w} by minimizing the error

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

$\|\mathbf{w}\|_2^2 = \mathbf{w} \cdot \mathbf{w}$

Positive tradeoff parameter

A regularization term to control the complexity of the model

Regularized Linear Regression

- To solve the unconstrained minimization problem, we can set the derivative of $\mathcal{L}(\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$ w.r.t. \mathbf{w} to zero

$$\frac{\partial \left(\mathcal{L}(\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|_2^2 \right)}{\partial \mathbf{w}} = \mathbf{0}$$

- We can obtain a closed-form solution for \mathbf{w} by the above equations

Closed-Form Solution

- Denote by $\mathbf{X} = (\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N)^T$

$$\mathbf{X} = \begin{pmatrix} \boxed{x_{10}} & \cdots & x_{N0} \\ \vdots & \ddots & \vdots \\ x_{1d} & \cdots & x_{Nd} \end{pmatrix}^T = \begin{pmatrix} \boxed{x_{10} \cdots x_{1d}} \\ \vdots & \ddots & \vdots \\ x_{N0} & \cdots & x_{Nd} \end{pmatrix}$$

- And by $\mathbf{y} = \begin{pmatrix} y_1 \\ \vdots \\ y_N \end{pmatrix}$

- The closed-form solution for \mathbf{w} :

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

How to induce
this closed-form
solution?

Tutorial



Implementation Example

```
>>> from sklearn.linear_model import Ridge
```

```
...
```

Referred to as λ on
our lecture notes

```
>>> rlr = Ridge(alpha=0.1)
```

```
>>> rlr.fit(X, y)
```

```
>>> pred_train_rlr = rlr.predict(X)
```

Evaluation


- Root Mean Square Error (RMSE)

$$\sqrt{\frac{1}{N} \sum_{i=1}^N (f(\mathbf{x}_i) - y_i)^2}$$

- Mean Absolute Error (MAE)

$$\frac{1}{N} \sum_{i=1}^N |f(\mathbf{x}_i) - y_i|$$

Nonlinear Fitting



Kernel trick

Thank you!