

# CZ4041/SC4000: Machine Learning

## Lesson 9: Ensemble Learning

Li Boyang, Albert

School of Computer Science and Engineering,  
NTU, Singapore

Acknowledgements: some figures are adapted from the lecture notes of the books “Introduction to Data Mining” (Chap. 5). Slides are modified from the version prepared by Dr. Sinno Pan.

# Ensemble Methods

- Objective:
  - To improve model performance in terms of accuracy by aggregating the predictions of multiple models
- How to do it?
  - Construct a set of base models from the training data
  - Make predictions by combining the predicted results made by each base model

*“Two heads are better than one”*

# Stories of Success

- Data mining competitions on Kaggle
  - Winning teams employ ensembles of classifiers






**Feature Engineering**  $\Longrightarrow$  **Ensemble Learning**

# Why Ensemble Work?

- Suppose there are 3 base binary classifiers
  - Each classifier has error rate,  $\varepsilon = 0.35$  or accuracy  $\text{acc} = 0.65$
  - Given a test instance, if we choose any one of these classifiers to make prediction, the probability that the classifier makes a wrong prediction is 35%

Base classifiers:



A test instance:



# Why Ensemble Work (cont.)

- Consider to combine the 3 base classifiers to make a prediction on a test instance using a majority vote
- The ensemble makes a wrong prediction only if more than 1 (i.e. 2 or 3) of the base classifiers predict incorrectly

# Why Ensemble Work (cont.)

| $x$             | $f_1$ | $f_2$ | $f_3$ | $f_M$ |   |
|-----------------|-------|-------|-------|-------|---|
| Truth label: -1 | +1    | +1    | +1    | +1    | ✗ |
|                 | +1    | +1    | -1    | +1    | ✗ |
|                 | +1    | -1    | +1    | +1    | ✗ |
|                 | -1    | +1    | +1    | +1    | ✗ |
|                 | +1    | -1    | -1    | -1    | ✓ |
|                 | -1    | +1    | -1    | -1    | ✓ |
|                 | -1    | -1    | +1    | -1    | ✓ |
|                 | -1    | -1    | -1    | -1    | ✓ |

The combined model makes a wrong prediction if at least two of the three base classifiers make a wrong prediction at the same time

# Why Ensemble Work (cont.)

- Assuming the three base classifiers are independent
  - The decision of one classifier does not tell us anything about another classifier.
- Each classifier has an error rate of  $\epsilon = 35\%$
- What is the overall error rate of the ensemble?



# Why Ensemble Work (cont.)

| $x$                               | $f_1$  | $f_2$  | $f_3$  | $f_M$ |   |
|-----------------------------------|--------|--------|--------|-------|---|
| Truth label: -1                   | +1 35% | +1 35% | +1 35% | +1    | ✗ |
| 3 lines with the same probability | +1 35% | +1 35% | -1 65% | +1    | ✗ |
|                                   | +1     | -1     | +1     | +1    | ✗ |
|                                   | -1     | +1     | +1     | +1    | ✗ |
|                                   | +1     | -1     | -1     | -1    | ✓ |
|                                   | -1     | +1     | -1     | -1    | ✓ |
|                                   | -1     | -1     | +1     | -1    | ✓ |
|                                   | -1     | -1     | -1     | -1    | ✓ |

The combined model makes a wrong prediction if at least two of the three base classifiers make a wrong prediction at the same time



# Why Ensemble Work (cont.)

- Therefore, probability that the ensemble classifier makes a wrong prediction is:

$$\sum_{i=2}^3 \binom{3}{i} \varepsilon^i (1 - \varepsilon)^{3-i} = 3 \times 0.35^2 \times 0.65 + 1 \times 0.35^3 \times 1 = 0.2817$$

- Case 1: when there are two exact classifiers make wrong predictions, the probability is — Two classifiers make wrong predictions

All possible combination  $\longrightarrow \left( \binom{3}{2} \varepsilon^2 (1 - \varepsilon)^{3-2} \right)$  ← The rest one makes correct prediction

- Case 2: when all the three classifiers make wrong predictions, the probability is  $\binom{3}{3} \varepsilon^3 (1 - \varepsilon)^{3-3}$

- That is the accuracy of the ensemble classifier is 71.83%

$$\varepsilon_{f_i} = 35\% \longrightarrow \varepsilon_M = 28.17\%$$

# Why Ensemble Work (cont.)

- Suppose there are 25 independent base classifiers
  - Therefore, probability that the ensemble classifier makes a wrong prediction is:

$$\sum_{i=13}^{25} \binom{25}{i} \varepsilon^i (1 - \varepsilon)^{25-i} = 0.06$$


- That is the accuracy of the ensemble classifier is 94%

$$\varepsilon_{f_i} = 35\%$$

$$\varepsilon_M = 6\%$$



# Error Rate of Base Classifiers



The base classifiers are identical (perfectly correlated)

The base classifiers are independent

**Observation:** the ensemble classifier performs worse than the base classifiers when the base classifier error rate is larger than 0.5

Error rate of an ensemble of 25 binary classifiers for different base classifier error rates

# Correlation between Base Classifiers

- Sometimes the base classifiers are not completely independent.
- An extreme case is when they are perfectly correlated. That is, they always have the same predictions.

| $f_1$ | $f_2$ |
|-------|-------|
| +1    | +1    |
| -1    | -1    |
| -1    | -1    |
| +1    | +1    |
| +1    | +1    |

Perfectly positively correlated



| $f_1$ | $f_2$ |
|-------|-------|
| +1    | -1    |
| -1    | +1    |
| -1    | +1    |
| +1    | -1    |
| +1    | -1    |

Perfectly negatively correlated

# Review: Random Variables

- Consider two random variables  $X$  and  $Y$ , which may take on values from  $\{0, 1\}$ .
  - Random variables are variables whose values are randomly assigned (from some unspecified random experiment).
- Imagine  $X$  and  $Y$  are the results of two coin tosses.
  - Head for the first toss  $\Rightarrow X = 1$
  - Tail for the first toss  $\Rightarrow X = 0$
  - Head for the second toss  $\Rightarrow Y = 1$
  - Tail for the second toss  $\Rightarrow Y = 0$
- The probabilities are denoted as  $P(X = 0)$ ,  $P(X = 1)$ ,  $P(Y = 0)$ ,  $P(Y = 1)$

# Review: Joint Probability

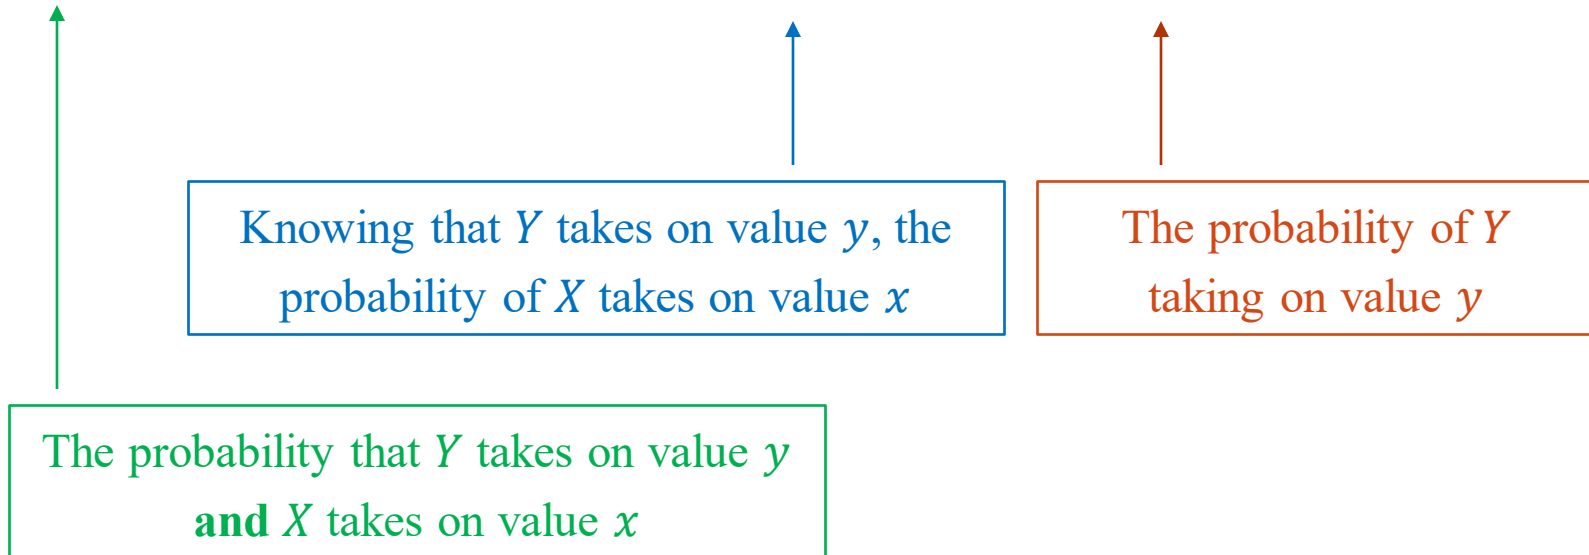
- $P(X = 0, Y = 0)$  is the probability of both toss getting tail.
- What about  $P(X = 1, Y = 0)$ ?
- What about  $P(Y = 0, Y = 1)$ ?

# Review: Conditional Probability

- $P(X = 0|Y = 0)$  is the probability of the first toss being a tail, knowing the second toss is a tail.
- What about  $P(Y = 0|X = 1)$ ?
- What about  $P(X = 0|X = 0)$ ?

# Review: Conditional Probability

- Conditional and joint probabilities are related as
- $P(X = x, Y = y) = P(Y = y|X = x)P(X = x)$
- $P(X = x, Y = y) = P(X = x|Y = y)P(Y = y)$






# Review: Independence

- If  $X$  and  $Y$  are independent, then we have
- $P(X = x|Y = y) = P(X = x)$ 
  - $Y$  provides no knowledge about  $X$
  - E.g., no matter the outcome of the first coin toss, the second toss is still 50/50
  - OR, no matter the second toss, the first toss is still 50/50
- Equivalently,  $P(X = x, Y = y) = P(X = x)P(Y = y)$
- Why is this equivalent?


$$\begin{aligned} P(X = x, Y = y) &= P(X = x|Y = y)P(Y = y) \\ &= P(X = x)P(Y = y) \end{aligned}$$

# Review: Correlation between RVs



One variable tells us something about the other variable.

# Difficult Input Causes Correlation



- These are supposedly digits from 0 to 9.
- Most models probably make incorrect predictions on these data points.
- Having multiple models would not help too much here.

# Necessary Conditions

- Two necessary conditions for an ensemble classifier to perform better than a single classifier:
  1. The base classifiers are not perfectly correlated with each other.
    - In practice, the base classifiers are usually somewhat correlated.
  2. The base classifiers should do better than a classifier that performs random guessing (e.g., for binary classification, accuracy should be better than 0.5)



Tutorial

# Ensemble Methods

- How to generate a set of base classifiers?
  - By manipulating the training set: multiple training sets are created by resampling the original data according to some sampling distribution. A classifier is then trained from each training set, such as [Bagging, Boosting](#)
  - By manipulating the input features: a subset of input features is chosen to form each training set. A classifier is then built from each training set, such as [Random forest](#)
  - By manipulating the learning algorithm(s): applying the algorithm several times on the same training data using different parameters or applying different algorithms
- How to combine the base classifiers for predictions?


# General Procedure

1. Let  $D$  denote the original training data,  $k$  denote the number of base classifiers, and  $T$  be the test dataset.
2. **for**  $i = 1$  to  $k$  **do**
3.     Train a base classifier  $f_i$  from  $D$
4. **end for**
5. **for** each test instance  $\mathbf{x} \in T$  **do**
6.     Generate  $f_1(\mathbf{x})$ ,  $f_2(\mathbf{x})$ , ..., and  $f_k(\mathbf{x})$
7.     Calculate  $f_M(\mathbf{x}) = \text{Merge}(f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_k(\mathbf{x}))$
8. **end for**

For example, majority voting  
(can be other schemes)

# Bagging


- Known as bootstrap aggregating (bootstrapping), to repeatedly sample with replacement according to a uniform probability distribution



- Build classifier on each bootstrap sample, which is of the same size of the original data
- Use majority voting to determine the class label of ensemble classifier

# Bagging (cont.)

Index of an instance





# Bagging (cont.)

Index of an instance



|           |   |   |   |   |   |   |   |   |   |    |
|-----------|---|---|---|---|---|---|---|---|---|----|
| Test Data | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
|-----------|---|---|---|---|---|---|---|---|---|----|

|                   |   |   |   |   |   |   |   |   |   |   |
|-------------------|---|---|---|---|---|---|---|---|---|---|
| Round 1 ( $f_1$ ) | + | + | + | + | - | - | - | + | - | - |
|-------------------|---|---|---|---|---|---|---|---|---|---|

|                   |   |   |   |   |   |   |   |   |   |   |
|-------------------|---|---|---|---|---|---|---|---|---|---|
| Round 2 ( $f_2$ ) | - | + | - | + | - | - | + | - | + | + |
|-------------------|---|---|---|---|---|---|---|---|---|---|

|                   |   |   |   |   |   |   |   |   |   |   |
|-------------------|---|---|---|---|---|---|---|---|---|---|
| Round 3 ( $f_3$ ) | + | - | - | + | + | - | - | + | + | - |
|-------------------|---|---|---|---|---|---|---|---|---|---|

|                    |   |   |   |   |   |   |   |   |   |   |
|--------------------|---|---|---|---|---|---|---|---|---|---|
| Ensemble ( $f_M$ ) | + | + | - | + | - | - | - | + | + | - |
|--------------------|---|---|---|---|---|---|---|---|---|---|

**Majority Voting**

# Bagging (cont.)

- Suppose a training set  $D$  contains  $N$  examples
- A training instance has a probability of  $1 - \frac{1}{N}$  of *not* being selected
- Its probability of ending up *not* in a training set  $D_i$  is  $\left(1 - \frac{1}{N}\right)^N \approx \frac{1}{e} = 0.368$
- A bootstrap sample  $D_i$  contains approximately 63.2% of the original training data

# Implementation Example

```
>>> from sklearn.ensemble import BaggingClassifier
```

```
>>> from sklearn.svm import SVC
```

Classification algorithm used in bagging

...

Specify how many  
base classifiers

```
>>> bagC = BaggingClassifier(base_estimator=SVC(), n_estimators=10)
```

```
>>> bagC.fit(X, y)
```

```
>>> pred = bagC.predict(X)
```

Specify a base classification  
algorithm (usually decision  
tree is used)

# Boosting

- Principles:
  - Boost a set of weak learners to a strong learner
  - Make instances currently misclassified more important
- Generally,
  - To adaptively change the distribution of training data so that the base classifiers will focus more on previously misclassified records

# Boosting (cont.)

- Specifically,
  - Initially, all  $N$  instances are assigned equal weights
  - Unlike bagging, weights may change at the end of each boosting round
  - In each boosting round, after the weights are assigned to the training instances,
    - Draw a bootstrap sample from the original data by using the weights as a sampling distribution to build a model

# Boosting: Procedure

1. Initially, all instances are assigned equal weights  $\frac{1}{N}$ , so that they are equally likely to be chosen for training. A sample is drawn uniformly to obtain a new training set.
2. A classifier is induced from the training set, and used to classify all the examples in the original training set
3. The weights of the training instances are updated at the end of each boosting round
  - Instances that are wrongly classified will have their weights increased
  - Instances that are classified correctly will have their weights decreased
4. Repeat Steps 2 and 3 until the stopping condition is met
5. Finally, the ensemble is obtained by aggregating the base classifiers obtained from each boosting round

# Boosting: Example

Initially, all the instances are assigned the same weights.

$\frac{1}{10}$   $\frac{1}{10}$   $\frac{1}{10}$   $\frac{1}{10}$   $\frac{1}{10}$   $\frac{1}{10}$   $\frac{1}{10}$   $\frac{1}{10}$   $\frac{1}{10}$   $\frac{1}{10}$

|                      |   |   |   |   |   |   |   |   |   |    |
|----------------------|---|---|---|---|---|---|---|---|---|----|
| <b>Original Data</b> | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
|----------------------|---|---|---|---|---|---|---|---|---|----|

$D$

+1 +1 +1 +1 -1 -1 -1 +1 +1 -1

Uniformly randomly sample using bootstrapping (sampling with replacement)

|                |   |   |   |   |   |   |   |    |   |   |
|----------------|---|---|---|---|---|---|---|----|---|---|
| <b>Round 1</b> | 7 | 3 | 2 | 8 | 7 | 9 | 4 | 10 | 6 | 3 |
|----------------|---|---|---|---|---|---|---|----|---|---|

$D_1$

Not selected in Round 1

A classifier built from the data

$f_1$

Perform the classifier on all original instances

|                      |   |   |   |   |   |   |   |   |   |    |
|----------------------|---|---|---|---|---|---|---|---|---|----|
| <b>Original Data</b> | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
|----------------------|---|---|---|---|---|---|---|---|---|----|

**True Label:** +1 +1 +1 +1 -1 -1 -1 +1 +1 -1

**Prediction:** +1 -1 +1 -1 -1 -1 -1 +1 +1 -1

Misclassified

Weights increased  
(relatively) for instances  
not selected in Round 1

Weights increased for  
misclassified instances

Weights decreased for  
correctly classified instances

|               |                |               |                |               |                |                |                |                |                |                |
|---------------|----------------|---------------|----------------|---------------|----------------|----------------|----------------|----------------|----------------|----------------|
| Original Data | $\frac{1}{10}$ | $\frac{1}{4}$ | $\frac{1}{20}$ | $\frac{1}{4}$ | $\frac{1}{10}$ | $\frac{1}{20}$ | $\frac{1}{20}$ | $\frac{1}{20}$ | $\frac{1}{20}$ | $\frac{1}{20}$ |
|               | 1              | 2             | 3              | 4             | 5              | 6              | 7              | 8              | 9              | 10             |
|               | +1             | +1            | +1             | +1            | -1             | -1             | -1             | +1             | +1             | -1             |

$D$

Randomly sample based on the weights of each instance

|         |   |   |   |   |   |   |   |   |   |   |
|---------|---|---|---|---|---|---|---|---|---|---|
| Round 2 | 5 | 4 | 9 | 4 | 2 | 5 | 1 | 7 | 4 | 2 |
|---------|---|---|---|---|---|---|---|---|---|---|

$D_2$

A classifier built from the data


$f_2$

Perform the classifier  
on all original instances

|               |    |    |    |    |    |    |    |    |    |    |
|---------------|----|----|----|----|----|----|----|----|----|----|
| Original Data | 1  | 2  | 3  | 4  | 5  | 6  | 7  | 8  | 9  | 10 |
| True Label:   | +1 | +1 | +1 | +1 | -1 | -1 | -1 | +1 | +1 | -1 |
| Prediction:   | +1 | +1 | +1 | -1 | -1 | -1 | -1 | +1 | +1 | -1 |



Randomly sample based on the updated weights of each instance



As the boosting rounds proceed, examples that are the hardest to classify tend to become even more prevalent, e.g., instance 4

|                    |                           |            |              |      |              |      |            |              |            |            |
|--------------------|---------------------------|------------|--------------|------|--------------|------|------------|--------------|------------|------------|
|                    | Index of an instance<br>↓ |            |              |      |              |      |            |              |            |            |
| Test Data          | 1                         | 2          | 3            | 4    | 5            | 6    | 7          | 8            | 9          | 10         |
| Round 1 ( $f_1$ )  | +                         | +          | +            | +    | -            | -    | -          | +            | -          | -          |
| Round 2 ( $f_2$ )  | -                         | +          | -            | +    | -            | -    | +          | -            | +          | +          |
| Round 3 ( $f_3$ )  | +                         | -          | -            | +    | +            | -    | -          | +            | +          | -          |
| Ensemble ( $f_M$ ) | +4                        | +2.8       | +1           | +5.8 | +3           | -5.8 | +1.8       | +4           | +4.8       | +1.8       |
|                    | v.s.<br>-1.8              | v.s.<br>-3 | v.s.<br>-4.8 |      | v.s.<br>-2.8 |      | v.s.<br>-4 | v.s.<br>-1.8 | v.s.<br>-1 | v.s.<br>-4 |
| Prediction         | +                         | -          | -            | +    | +            | -    | -          | +            | +          | -          |

Weights for each classifier  
↓

1  
1.8  
3

## Weighted Voting

Note: here the values of the classifiers weights as well as the updated instance weights in each boosting round are just examples. Different boosting algorithms have different ways to compute these weights.

# Implementation Example

```
>>> from sklearn.ensemble import AdaBoostClassifier
```

```
>>> from sklearn.svm import SVC
```

AdaBoost, a classic  
boosting algorithm

...

Specify how many  
base classifiers

```
>>> adaC = AdaBoostClassifier (base_estimator=SVC(), n_estimators=10 )
```

```
>>> adaC.fit(X, y)
```

Specify a base  
classification algorithm

```
>>> pred= adaC.predict(X)
```

# Random Forests

- A class of ensemble methods specifically designed for decision tree classifiers
- Random Forests grow many trees
- Each tree is generated based on a random subset of features
- Final result on classifying a new instance – voting
  - Forest chooses the classification result having the most votes (over all the trees in the forest)

# Random Forests (cont.)

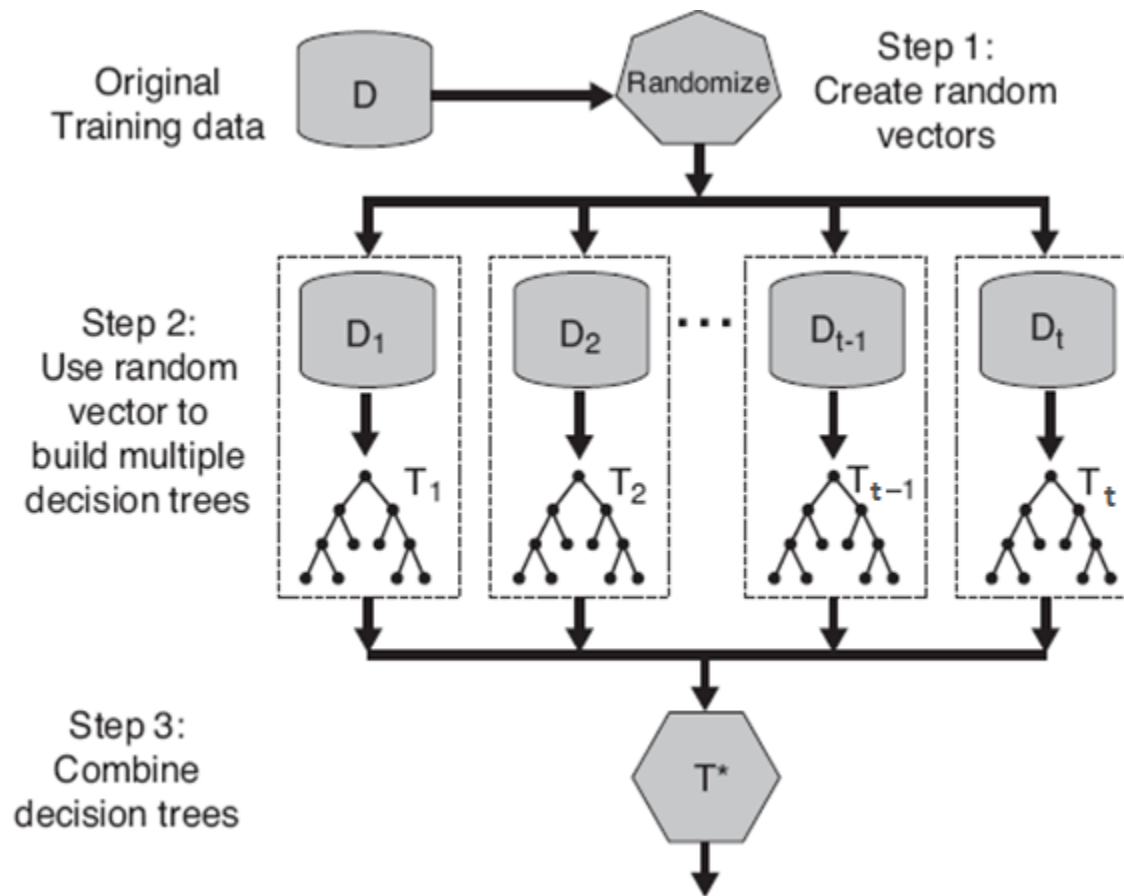



Illustration of random forests

# Random Forests: Algorithm

- Choose  $T$ : number of trees to grow
- Choose  $m' < m$  ( $m$  is the number of total features): number of features used to calculate the best split at each node (typically 20%)
- For each tree
  - Choose a training set via bootstrapping
  - For each node, randomly choose  $m'$  features and calculate the best split
  - Fully grown and not pruned
- Use majority vote among all the trees


# Example

Original training dataset




Full set of 8 input features


|       |       |       |       |       |       |       |       |
|-------|-------|-------|-------|-------|-------|-------|-------|
| $F_1$ | $F_2$ | $F_3$ | $F_4$ | $F_5$ | $F_6$ | $F_7$ | $F_8$ |
|-------|-------|-------|-------|-------|-------|-------|-------|



Decision Tree 1



Decision Tree 2



Decision Tree 3

# Random Forests: Discussions

- Bagging + random features
- Improve Accuracy
  - Incorporate more diversity
- Improve Efficiency
  - Searching among subsets of features is much faster than searching among the complete set



# Implementation Example

```
>>> from sklearn.ensemble import RandomForestClassifier
```

...

```
>>> rfC = RandomForestClassifier(n_estimators = 20, max_depth = 3)
```

Set parameter  
for the base tree

```
>>> rfC.fit(X, y)
```

Specify how many  
base classifiers

```
>>> pred= rfC.predict(X)
```

# Combination Methods

- Voting
  - Majority voting
  - Weighted voting
- Average
  - Simple average
  - Weighted average
- Combining by learning

# Voting

- Majority voting:
  - Takes the class label that receives the largest number of votes as the final winner
- Weighted voting:
  - A generalized version of majority voting by introducing weights for each classifier

# Average

- Simple average:

$$f_M(\mathbf{x}) = \frac{1}{T} \sum_{i=1}^T f_i(\mathbf{x})$$

- Weighted average:

$$f_M(\mathbf{x}) = \sum_{i=1}^T w_i f_i(\mathbf{x})$$

$$\text{where } w_i \geq 0, \text{ and } \sum_{i=1}^T w_i = 1$$

- Prediction:

$$\hat{y}(\mathbf{x}) = \begin{cases} 1, & f_M(\mathbf{x}) \geq 0 \\ -1, & \text{otherwise} \end{cases}$$

# Combining by Learning

- Stacking:
  - A general procedure where a learner is trained to combine the individual learners
  - Individual learners: first-level learners
  - Combiner: second-level learner, or meta-learner

# Combining by Learning (cont.)

- Suppose given a binary classification problem,

$m$  input features

$N$  training instances

|       | $X_1$ | $X_2$ | ... | $X_m$ | $Y$ |
|-------|-------|-------|-----|-------|-----|
| $x_1$ | 3     | 1     | ... | 0.3   | 1   |
| $x_2$ | 4     | 1     | ... | 0.6   | -1  |
| ...   | ...   | ...   | ... | ...   | ... |
| $x_N$ | -10   | 0     | ... | 0.2   | 1   |

$k$  base classifiers are trained

$f_1 \quad f_2 \quad \dots \quad f_k$

Predicted value of classifier  $f_1$  on the instance  $x_1$

$k$  base classifiers

Labels

$N$  instances

|       | $f_1$ | $f_2$ | ... | $f_k$ | $Y$ |
|-------|-------|-------|-----|-------|-----|
| $x_1$ | 0.6   | 1     | ... | 0.3   | 1   |
| $x_2$ | 0.3   | 1     | ... | 0.6   | -1  |
| ...   | ...   | ...   | ... | ...   | ... |
| $x_N$ | 0.1   | -1    | ... | 0.2   | 1   |

A meta classifier is learned:  $f_M: \mathbb{R}^k \rightarrow \{-1, 1\}$

Test instance

|       | $X_1$ | $X_2$ | ... | $X_m$ |
|-------|-------|-------|-----|-------|
| $x^*$ | 0.8   | -1    | ... | 0.4   |

Apply  $f_1, \dots, f_k$

|           | $f_1$ | $f_2$ | ... | $f_k$ |
|-----------|-------|-------|-----|-------|
| $x^{*'} $ | 0.3   | 0.6   | ... | 1     |

Prediction  $f_M(x^{*'})$

# Combining by Learning: Summary

- Represent each training instances using classifier-generated outputs

$$\mathbf{x}'_i = (f_1(\mathbf{x}_i), f_2(\mathbf{x}_i), \dots, f_k(\mathbf{x}_i))$$

- Learn a meta classifier  $f_M$  from  $\{\mathbf{x}'_i, y_i\}, i = 1, \dots, N$ .
- For a test instance  $\mathbf{x}^*$ 
  - Represent it by  $\mathbf{x}^{*'} = (f_1(\mathbf{x}^*), f_2(\mathbf{x}^*), \dots, f_k(\mathbf{x}^*))$
  - Use  $f_M$  to make a prediction  $f_M(\mathbf{x}^{*'})$



**Thank you!**