


# Autoencoders

Xingang Pan

潘新钢

<https://xingangpan.github.io/>

<https://twitter.com/XingangP>



# Outline

- Supervised vs unsupervised learning
- Autoencoders
  - Denoising autoencoders
  - Undercomplete and overcomplete autoencoders
  - Sparse autoencoders
  - Other encoders

# Supervised vs Unsupervised Learning

## Supervised Learning

**Data:**  $(x, y)$

$x$  is data,  $y$  is label

**Goal:** Learn a *function* to map  $x \rightarrow y$

**Examples:** Classification, regression,  
object detection, semantic  
segmentation, image captioning, etc.

Classification



Cat

# Supervised vs Unsupervised Learning

## Supervised Learning


**Data:**  $(x, y)$

$x$  is data,  $y$  is label

**Goal:** Learn a *function* to map  $x \rightarrow y$

**Examples:** Classification, regression,  
object detection, semantic  
segmentation, image captioning, etc.

## Object Detection



**DOG, DOG, CAT**

# Supervised vs Unsupervised Learning

## Supervised Learning


**Data:**  $(x, y)$

$x$  is data,  $y$  is label

**Goal:** Learn a *function* to map  $x \rightarrow y$

**Examples:** Classification, regression,  
object detection, semantic  
segmentation, image captioning, etc.

## Semantic Segmentation



GRASS, CAT, TREE, SKY

# Supervised vs Unsupervised Learning

## Supervised Learning

**Data:**  $(x, y)$

$x$  is data,  $y$  is label

**Goal:** Learn a *function* to map  $x \rightarrow y$

**Examples:** Classification, regression, object detection, semantic segmentation, image captioning, etc.

Image captioning



*A cat sitting on a suitcase on the floor*

# Supervised vs Unsupervised Learning

## Supervised Learning

**Data:**  $(x, y)$

$x$  is data,  $y$  is label

**Goal:** Learn a *function* to map  $x \rightarrow y$

**Examples:** Classification, regression,  
object detection, semantic  
segmentation, image captioning, etc.

## Unsupervised Learning

**Data:**  $x$


Just data, no labels!

**Goal:** Learn some underlying  
hidden *structure* of the data

**Examples:** Clustering,  
dimensionality reduction, feature  
learning, density estimation, etc.

# Supervised vs Unsupervised Learning

## Clustering (e.g. K-Means)



## Unsupervised Learning

**Data:** x


Just data, no labels!

**Goal:** Learn some underlying hidden *structure* of the data

**Examples:** Clustering, dimensionality reduction, feature learning, density estimation, etc.

# Supervised vs Unsupervised Learning

Dimensionality Reduction  
(e.g. Principal Components Analysis)



## Unsupervised Learning

**Data:**  $x$



Just data, no labels!

**Goal:** Learn some underlying hidden *structure* of the data

**Examples:** Clustering, dimensionality reduction, feature learning, density estimation, etc.

# Supervised vs Unsupervised Learning

## Density Estimation



## Unsupervised Learning

**Data:**  $x$



Just data, no labels!

**Goal:** Learn some underlying hidden *structure* of the data

**Examples:** Clustering, dimensionality reduction, feature learning, density estimation, etc.

# Supervised vs Unsupervised Learning

## Feature Learning (e.g. autoencoders)



## Unsupervised Learning

**Data:**  $x$

Just data, no labels!

**Goal:** Learn some underlying hidden *structure* of the data



**Examples:** Clustering, dimensionality reduction, feature learning, density estimation, etc.

# Autoencoders

Unsupervised method for learning feature vectors from raw data  $x$ , without any labels

Features should extract useful information (maybe object identities, properties, scene type, etc) that we can use for downstream tasks

**Originally:** Linear + nonlinearity (sigmoid)  
**Later:** Deep, fully-connected  
**Later:** ReLU CNN





Input Data

# Autoencoders

**Problem:** How can we learn this feature transform from raw data?

Features should extract useful information (maybe object identities, properties, scene type, etc) that we can use for downstream tasks  
But we can't observe features!

**Originally:** Linear + nonlinearity (sigmoid)  
**Later:** Deep, fully-connected  
**Later:** ReLU CNN




Input Data

# Autoencoders


- **Problem:** How can we learn this feature transform from raw data?

Idea: Use the features to reconstruct the input data with a **decoder**

“Autoencoding” = encoding itself



# Autoencoders




An autoencoder is a neural network that is trained to attempt to copy its input to its output. Its hidden layer describes a *code* that represents the input.

The network consists of two parts: an **encoder** and a **decoder**.

# Autoencoders

Given an input  $x$ , the hidden-layer performs the encoding function  $\mathbf{h} = \emptyset(x)$  and the decoder  $\varphi$  produces the reconstruction  $y = \varphi(\mathbf{h})$ .




If the autoencoder succeeds:

$$\mathbf{y} = \varphi(\mathbf{h}) = \varphi(\emptyset(x)) = \mathbf{x}$$

In order to be useful, autoencoders are designed to be **unable to copy exactly** and enable to copy only inputs that resembles the training data.

Since the model is forced to prioritize which aspects of the input should be copied, the hidden-layer often **learns useful properties of the data**.

# Autoencoders



Reverse mapping from the hidden layer to the output can be optionally constrained to be the same as the input to hidden-layer mapping. That is, if encoder weight matrix is  $\mathbf{W}$ , the decoder weigh matrix is  $\mathbf{W}^T$ .

Hidden layer and output layer activation can be then written as

$$\begin{aligned}\mathbf{h} &= f(\mathbf{W}^T \mathbf{x} + \mathbf{b}) \\ \mathbf{y} &= f(\mathbf{W}\mathbf{h} + \mathbf{c})\end{aligned}$$

$f$  is usually a sigmoid.

# Training autoencoders

The **cost function of reconstruction** can be measured by many ways, depending on the appropriate distributional assumptions on the inputs.

Learning of autoencoders is **unsupervised** as no specific targets are given.

Given a training set  $\{\mathbf{x}_p\}_{p=1}^P$ .

The mean-square-error cost is usually used if the data is assumed to be continuous and Gaussian distributed:

$$J_{mse} = \frac{1}{P} \sum_{p=1}^P \|\mathbf{y}_p - \mathbf{x}_p\|^2$$

where  $\mathbf{y}_p$  is the output for input  $\mathbf{x}_p$  and  $\|\cdot\|$  denotes the magnitude of the vector.

# Training autoencoders

If the inputs are interpreted as **bit vectors** or **vectors of bit probabilities**, **cross-entropy of the reconstruction** can be used:

$$J_{\text{cross-entropy}} = - \sum_{p=1}^P (x_p \log y_p + (1 - x_p) \log(1 - y_p))$$

Learning are done by using gradient descent:

$$\mathbf{W} \leftarrow \mathbf{W} - \alpha \nabla_{\mathbf{W}} J$$

$$\mathbf{b} \leftarrow \mathbf{b} - \alpha \nabla_{\mathbf{b}} J$$

$$\mathbf{c} \leftarrow \mathbf{c} - \alpha \nabla_{\mathbf{c}} J$$

# Example 1

Given input patterns:

$(1, 0, 1, 0, 0)$ ,  $(0, 0, 1, 1, 0)$ ,  $(1, 1, 0, 1, 1)$  and  $(0, 1, 1, 1, 0)$

Design an autoencoder with 3 hidden units.

Show the representations of inputs at the hidden layer upon convergence.

$$\mathbf{H} = f(\mathbf{XW} + \mathbf{B})$$

$$\mathbf{Y} = f(\mathbf{HW}^T + \mathbf{C})$$

$$\mathbf{O} = \mathbf{1}(\mathbf{Y} > 0.5)$$


Gradient descent on entropy:

$$J = - \sum_{p=1}^4 (x_p \log y_p + (1 - x_p) \log(1 - y_p))$$


$$bit\ errors = \sum_{p=1}^4 \sum_{k=1}^5 1(x_{pk} \neq o_{pk})$$

eg10.1.ipynb

# Example 1



# Example 1



# Example 1

```
▶ # Display weights and biases
print(f'W:\n {autoencoder.W.data}\n')
print(f'b:\n {autoencoder.b.data}\n')
print(f'b_prime:\n {autoencoder.b_prime.data}\n')

→ W:
    tensor([[-3.6867, -0.3808,  2.6520],
           [-0.2545, -3.6158, -1.8657],
           [ 1.7148,  3.0138,  0.9254],
           [ 1.9011, -0.9880, -3.5277],
           [-2.4830, -2.5980, -1.1112]])


b:
    tensor([ 0.1496,  0.5881, -0.1604])

b_prime:
    tensor([ 1.2448,  2.3414, -0.7427,  2.1173,  0.9872])
```

# Denoising Autoencoders (DAE)

A *denoising autoencoder* (DAE) receives corrupted data points as inputs.

It is trained to predict the original uncorrupted data points as its output.



The idea of DAE is that in order to force the hidden layer to **discover more robust features** and prevent it from simply learning the identity. We train the autoencoder to reconstruct the input from a corrupted version of it.

In other words, DAE attempts to encode the input (preserve the information about input) and attempts to **undo the effect of corruption process** applied to the input of the autoencoder.

# DAE

For training DAE:

- First, input data is corrupted to mimic the noise in the images:  $x \rightarrow \tilde{x}$

$\tilde{x}$  is the corrupted version of data. The corruption process simulates the distribution of data

- The network is trained to produce uncorrupted data:  $y \rightarrow x$

# DAE: Corrupting inputs

To obtain corrupted version of input data, each input  $x_i$  of input data is added with **additive or multiplicative noise**.

**Additive noise:**

$$\tilde{x}_i = x_i + \varepsilon$$

where noise  $\varepsilon$  is Gaussian distributed:  $\varepsilon \sim N(0, \sigma^2)$

And  $\sigma$  is the standard deviation that determines the S/N ratio. Usually used for continuous data.

**Multiplicative noise:**


$$\tilde{x}_i = \varepsilon x_i$$

where noise  $\varepsilon$  could be Binomially distributed:  $\varepsilon \sim Binomial(p)$

And  $p$  is the probability of ones and  $1 - p$  is the probability of zeros (noise). Usually, used for binary data.

# Example 2: Denoising autoencoders

The MNIST database of gray level images of handwritten digits:  
<http://yann.lecun.com/exdb/mnist/>



Each image is 28x28 size.

Intensities are in the range [0, 255] and normalized to [0, 1].

Training set: 60,000, Test set: 10,000

To build a DAE with 500 hidden units.


# Example 2a: DAE (multiplicative noise)

Original data	Corrupted data
7 2 1 0 4 1 4	7 2 1 0 4 1 4
9 5 9 0 6 9 0	9 5 9 0 6 9 0
1 5 9 7 3 4 9	1 5 9 7 3 4 9
6 6 5 4 0 7 4	6 6 5 4 0 7 4
0 1 3 1 3 4 7	0 1 3 1 3 4 7
2 7 1 2 1 1 7	2 7 1 2 1 1 7
4 2 3 5 1 2 4	4 2 3 5 1 2 4

DAE attempts to predict the input data from corrupted input data with multiplicative noise with binomial distribution.  
Corruption level  $1 - p = 10\%$

eg10.2a.ipynb


# Example 2a: DAE (multiplicative noise)



500 hidden units  
 $\alpha = 0.1$

# Example 2a: DAE (multiplicative noise)

Weights learned by the hidden-layer



# Example 2a: DAE (multiplicative noise)

Sample of reconstructed test images

Input data							Reconstructed (noise-filtered) data						
7	2	1	0	4	1	4	7	2	1	0	4	1	4
9	5	9	0	6	9	0	9	5	9	0	6	9	0
1	5	9	7	3	4	9	1	5	9	7	3	4	9
6	6	5	4	0	7	4	6	6	5	4	0	7	4
0	1	3	1	3	4	7	0	1	3	1	3	4	7
2	7	1	2	1	1	7	2	7	1	2	1	1	7
4	2	3	5	1	2	4	4	2	3	5	1	2	4

# Example 2b: DAE (additive noise)

Original data

7	2	1	0	4	1	4
9	5	9	0	6	9	0
1	5	9	7	3	4	9
6	6	5	4	0	7	4
0	1	3	1	3	4	7
2	7	1	2	1	1	7
4	2	3	5	1	2	4

Corrupted data

7	2	1	0	4	1	4
9	5	9	0	6	9	0
1	5	9	7	3	4	9
6	6	5	4	0	7	4
0	1	3	1	3	4	7
2	7	1	2	1	1	7
4	2	3	5	1	2	4




Additive noise with gaussian distribution.

Data is scaled between [0, 1], noise s.d = 0.1

eg10.2b.ipynb


## Example 2b: DAE (additive noise)



500 hidden units  
 $\alpha = 0.1$

# Example 2b: DAE (additive noise)

Features (weights) learned by the hidden-layer



## Example 2b: DAE (additive noise)

Sample of reconstructed images


Input data

7	2	1	0	4	1	4
9	5	9	0	6	9	0
1	5	9	7	3	4	9
6	4	5	4	0	7	4
0	1	3	1	3	4	7
2	7	1	2	1	1	7
4	2	3	5	1	2	4

Reconstructed (noise-filtered) data

7	2	1	0	4	1	4
9	5	9	0	6	9	0
1	5	9	7	3	4	9
6	4	5	4	0	7	4
0	1	3	1	3	4	7
2	7	1	2	1	1	7
4	2	3	5	1	2	4

# Autoencoders




Input dimension  $n$  and hidden dimension  $M$ :

If  $M < n$ , *undercomplete* autoencoders

If  $M > n$ , *overcomplete* autoencoders

# Undercomplete autoencoders




In undercomplete autoencoders, the hidden-layer has a **lower dimension** than the input layer.

By learning to approximate an  $n$ -dimensional inputs with  $M$  ( $< n$ ) number of hidden units, we obtain a **lower dimensional representation** of the input signals. The network reconstructs the input signals from the reduced-dimensional hidden representation.

Learning an undercomplete representation **forces the autoencoder to capture the most salient features**.

By limiting the number of hidden neurons, **hidden structures** of input data can be inferred from autoencoders. For example, correlations among input variables, learning principal components of data, etc.

# Overcomplete autoencoders




In overcomplete autoencoders, the hidden-layer has a **higher dimension** than the dimension of the input.

In order to learn useful information from overcomplete autoencoders, it is **necessary use some constraints** on its characteristics.

Even when the hidden dimensions are large, one can still explore interesting structures of inputs by introducing other constraints such as '**sparsity**' of input data.


# Regularizing autoencoders



Both undercomplete and overcomplete autoencoders fail to learn anything useful if the encoder and decoder are given too much capacity. Some form of constraints are needed in order to make them useful.

Deep autoencoders are only possible when some constraints are imposed on the cost function.

# Regularizing autoencoders



**Regularized autoencoders** incorporate a penalty to the cost function to learn interesting features from the input.

Regularized autoencoders provide the ability to train any autoencoder architecture successfully by choosing suitable code dimension and the capacity of the encoder and decoder.

With regularized autoencoders, one can use **larger model capacity** (for example, deeper autoencoders) and large code size.

# Regularizing autoencoders

Regularized autoencoders add an appropriate penalty function  $\Omega$  to the cost function:

$$J_1 = J + \beta \Omega(h)$$

where  $\beta$  is the penalty or regularization parameter. The penalty is usually **imposed on the hidden activations**.

A regularized autoencoder can be nonlinear and overcomplete but still learn something useful about the data distribution, even if the model capacity is great enough to learn trivial identity function.

The regularized loss function encourages the model to have other properties besides the ability to copy its input to its output.

# Sparse Autoencoders (SAE)

A **sparse autoencoder** (SAE) is simply an autoencoder whose training criterion involves the **sparsity penalty**  $\Omega_{\text{sparsity}}$  at the hidden layer:

$$J_1 = J + \beta \Omega_{\text{sparsity}}(h)$$

The sparsity penalty term makes the features (weights) learnt by the hidden-layer to be **sparse**.

With the sparsity constraint, one would constraint the neurons at the hidden layers to be **inactive for most of the time**.

We say that the neuron is “active” when its output is close to 1 and the neuron is “inactive” when its output is close to 0.

# Sparsity constraint

For a set  $\{x_p\}_{p=1}^P$  of input patterns, the **average activation**  $\rho_j$  of neuron  $j$  at the hidden-layer is given by:

$$\rho_j = \frac{1}{P} \sum_{p=1}^P h_{pj} = \frac{1}{P} \sum_{p=1}^P f(\mathbf{x}_p^T \mathbf{w}_j + b_j)$$

where  $h_{pj}$  is the activation of hidden neuron  $j$  for  $p$  th pattern, and  $\mathbf{w}_j$  and  $b_j$  are the weights and biases of the hidden neuron  $j$ .

We would like to enforce the constraint:  $\rho_j = \rho$  such that the **sparsity parameter**  $\rho$  is set to a **small value** close to zero (say 0.05).

That is, most of the time, hidden neuron activations are maintained at  $\rho$  on average. By choosing a smaller value for  $\rho$ , the neurons are **activated selectively to patterns and thereby learn sparse features**.

# Sparsity constraint

To achieve sparse activations at the hidden-layer, the **Kullback-Leibler (KL) divergence** is used as the sparsity constraint:

$$D(\mathbf{h}) = \sum_{j=1}^M \rho \log \frac{\rho}{\rho_j} + (1 - \rho) \log \frac{1 - \rho}{1 - \rho_j}$$

where  $M$  is the number of hidden neurons and  $\rho$  is the sparsity parameter.

KL divergence measures the **deviation of the distribution**  $\{\rho_j\}$  of activations at the hidden-layer from the uniform distribution of  $\rho$ .

The KL divergence is **minimum** when  $\rho_j = \rho$  for all  $j$ .

That is, when the average activations are uniform and equal to very low value  $\rho$ .

# Sparse Autoencoder (SAE)

The cost function for the sparse autoencoder (SAE) is given by

$$J_1 = J + \beta D(\mathbf{h})$$

where  $D(h)$  is the KL divergence of hidden-layer activations:

$$D(\mathbf{h}) = \sum_{j=1}^M \rho \log \frac{\rho}{\rho_j} + (1 - \rho) \log \frac{1 - \rho}{1 - \rho_j}$$

For gradient descent,

$$\nabla_{\mathbf{W}} J_1 = \nabla_{\mathbf{W}} J + \beta \nabla_{\mathbf{W}} D(\mathbf{h})$$

By chain rule:

$$\frac{\partial D(\mathbf{h})}{\partial \mathbf{w}_j} = \frac{\partial D(\mathbf{h})}{\partial \rho_j} \frac{\partial \rho_j}{\partial \mathbf{w}_j} \quad (\text{A})$$

where

$$D(\mathbf{h}) = \sum_{j=1}^M \rho \log \frac{\rho}{\rho_j} + (1 - \rho) \log \frac{1 - \rho}{1 - \rho_j}$$

and

$$\begin{aligned} \frac{\partial D(\mathbf{h})}{\partial \rho_j} &= \rho \frac{1}{\rho} \left( -\frac{\rho}{\rho_j^2} \right) + (1 - \rho) \frac{1}{\left( \frac{1 - \rho}{1 - \rho_j} \right)} \left( \frac{-(1 - \rho)}{\left( \frac{1 - \rho}{1 - \rho_j} \right)^2} \right) \\ &= -\frac{\rho}{\rho_j} + \frac{1 - \rho}{1 - \rho_j} \end{aligned} \quad (\text{B})$$

# SAE

Note:

$$\rho_j = \frac{1}{P} \sum_{p=1}^P f(u_{pj})$$

where  $u_{pj} = \mathbf{x}_p^T \mathbf{w}_j + b_j$  is the synaptic input of the  $j$  th neuron due to  $p$  th pattern.

$$\frac{\partial \rho_j}{\partial \mathbf{w}_j} = \frac{1}{P} \sum_p f'((u_{pj})) \frac{\partial (u_{pj})}{\partial \mathbf{w}_j} = \frac{1}{P} \sum_p f'((u_{pj})) \mathbf{x}_p \quad (c)$$

Substituting (B) and (C) in (A):

$$\frac{\partial D(\mathbf{h})}{\partial \mathbf{w}_j} = \frac{1}{P} \left( \frac{\rho}{\rho_j} + \frac{1 - \rho}{1 - \rho_j} \right) \sum_p f'((u_{pj})) \mathbf{x}_p$$

Substituting in (A) , we can find:

$$\nabla_{\mathbf{W}} D(\mathbf{h}) = (\nabla_{w_1} D(\mathbf{h}) \quad \nabla_{w_2} D(\mathbf{h}) \quad \cdots \quad \nabla_{w_M} D(\mathbf{h}))$$

The gradient of the constrained cost function:

$$\nabla_{\mathbf{W}} J_1 = \nabla_{\mathbf{W}} J + \beta \nabla_{\mathbf{W}} D(\mathbf{h})$$

Similarly,  $\nabla_{\mathbf{b}} J_1$  and  $\nabla_{\mathbf{c}} J_1$  can be derived.

Learning can be done as

$$\begin{aligned}\mathbf{W} &\leftarrow \mathbf{W} - \alpha \nabla_{\mathbf{W}} J_1 \\ \mathbf{b} &\leftarrow \mathbf{b} - \alpha \nabla_{\mathbf{b}} J_1 \\ \mathbf{c} &\leftarrow \mathbf{c} - \alpha \nabla_{\mathbf{c}} J_1\end{aligned}$$

# Example 3

Design the following autoencoder to process MNIST images:

1. Undercomplete autoencoder with 100 hidden units
2. Overcomplete autoencoder with 900 hidden units
3. Sparse autoencoder with 900 hidden units. Use sparsity parameter = 0.5

The MNIST database of gray level images of handwritten digits:


<http://yann.lecun.com/exdb/mnist/>

Each image is 28x28 size.


Intensities are in the range [0, 255].

Training set: 60,000

Test set: 10,000




# Example 3a: Undercomplete autoencoder




eg10.3a.ipynb

# Example 3a: Undercomplete autoencoder




# Example 3a: Undercomplete autoencoder

Test inputs




Reconstructed inputs




# Example 3a: Undercomplete autoencoder

Hidden layer activations of 49 test patterns




10x10 neurons

# Example 3b: Overcomplete autoencoder



# Example 3b: Overcomplete autoencoder

Learned weights



900 units

# Example 3b: Overcomplete autoencoder

Test inputs


7	2	1	0	4	1	4
9	5	9	0	6	9	0
1	5	9	7	3	4	9
6	4	5	4	0	7	4
0	1	3	1	3	4	7
2	7	1	2	1	1	7
4	2	3	5	1	2	4

Reconstructed inputs

7	2	1	0	4	1	4
9	5	9	0	6	9	0
1	5	9	7	3	4	9
6	4	5	4	0	7	4
0	1	3	1	3	4	7
2	7	1	2	1	1	7
4	2	3	5	1	2	4


# Example 3b: Overcomplete autoencoder

Hidden layer activations




30x30 neurons

# Example 3c: Sparse autoencoder



# Example 3c: Sparse autoencoder

Learned weights



# Example 3c: Sparse autoencoder

Test inputs


7	2	1	0	4	1	4
9	5	9	0	6	9	0
1	5	9	7	3	4	9
6	6	5	4	0	7	4
0	1	3	1	3	4	7
2	7	1	2	1	1	7
4	2	3	5	1	2	4

Reconstructed inputs

7	2	1	0	4	1	4
9	5	9	0	6	9	0
1	5	9	7	3	4	9
6	6	5	4	0	7	4
0	1	3	1	3	4	7
2	7	1	2	1	1	7
4	2	3	5	1	2	4

# Example 3c: Sparse autoencoder

Hidden layer activations



30x30 neurons


# Deep stacked autoencoders

Deep autoencoders can be built by **stacking autoencoders** one after the other.


Training of deep autoencoders is done in a step-by-step fashion **one layer at a time**:

- After training the first level of denoising autoencoder, the resulting hidden representation is used to train a second level of the denoising encoder.
- The second level hidden representation can be used to train the third level of the encoders.
- This process is repeated and deep stacked autoencoder can be realized.


# Deep stacked autoencoders



# Deep stacked autoencoders



Training the third hidden layer.  
Reconstructs second hidden  
layer output




Autoencoder with  
three hidden layers

# Semi-Supervised Classification


- Many images, but few ground truth labels

start unsupervised  
train autoencoder on many images




L2 Loss function:  
 $\|x - \hat{x}\|^2$

supervised fine-tuning  
train classification network on labeled images





# Context Encoders

[ Pathak et al., 2016 ]



# Context Encoders

[ Pathak et al., 2016 ]




(a) Central region

(b) Random block

(c) Random region


# Context Encoders

[ Pathak et al., 2016 ]



# Context Encoders

[ Pathak et al., 2016 ]



# Context Encoders

[ Pathak et al., 2016 ]

