

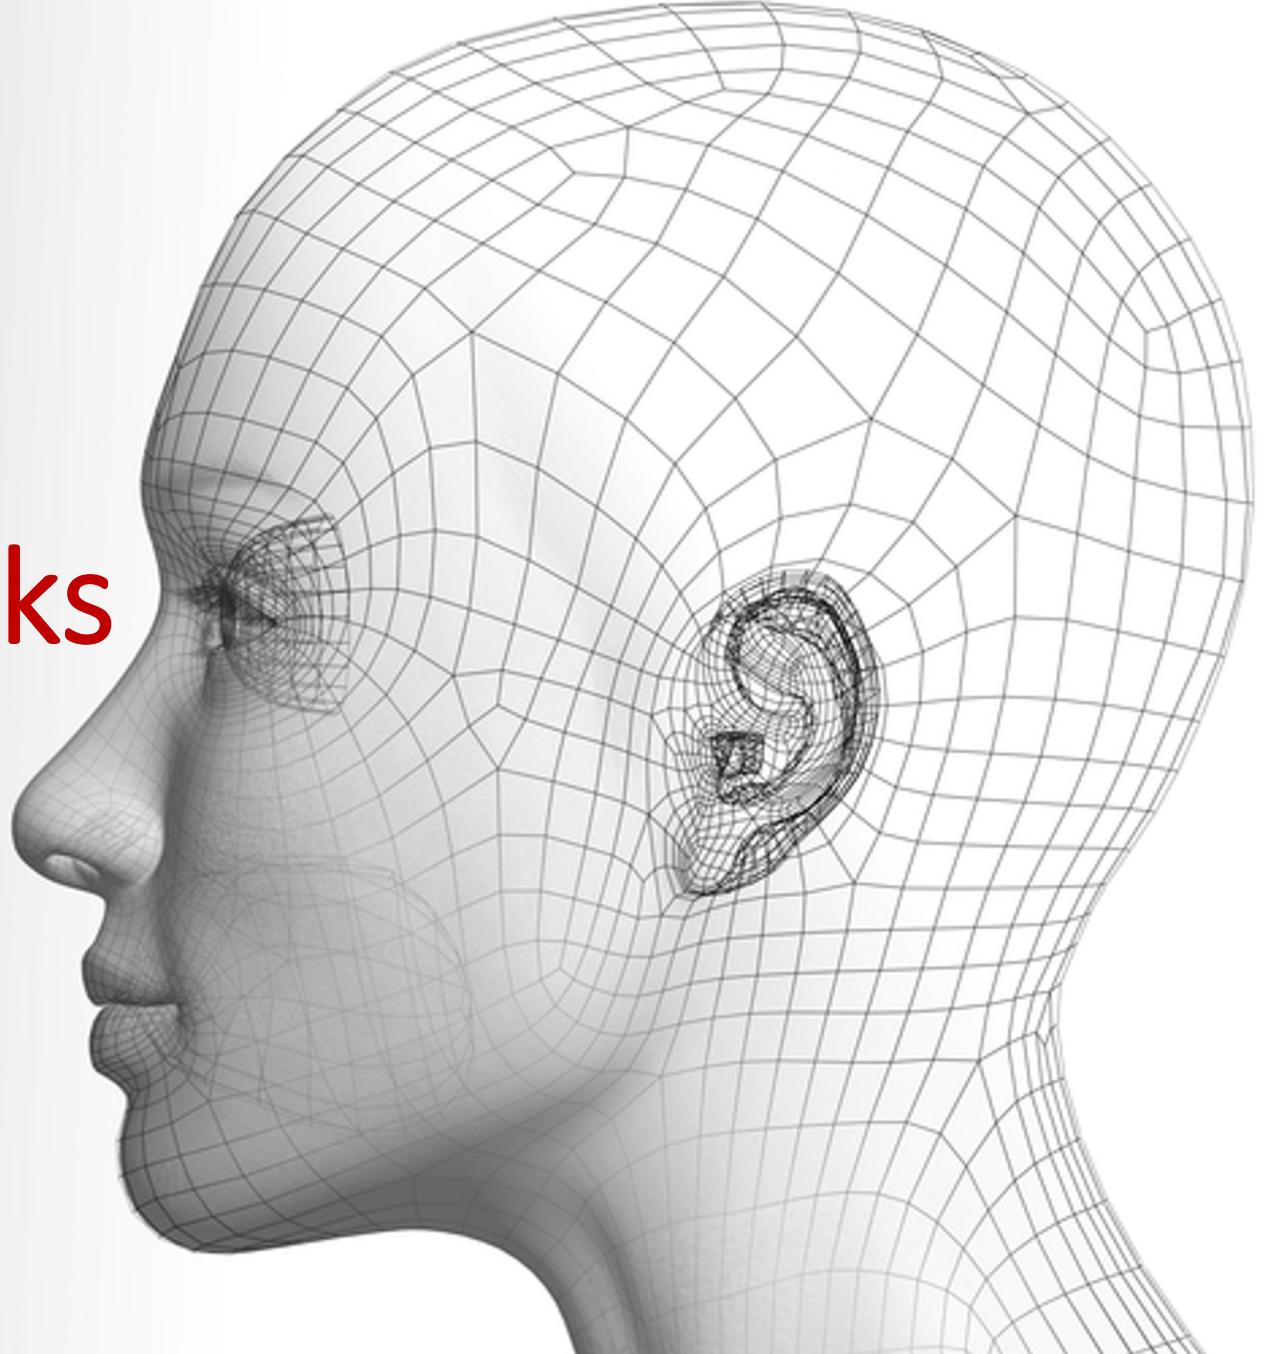
# Generative Adversarial Networks

Xingang Pan

潘新钢

<https://xingangpan.github.io/>

<https://twitter.com/XingangP>



# Outline

- GAN Basics
- GAN Training
- DCGAN
- Mode Collapse

# Supervised learning

**Data:**  $(x, y)$

$x$  is data,  $y$  is label

**Goal:** Learn a *function* to map  $x \rightarrow y$

**Examples:**

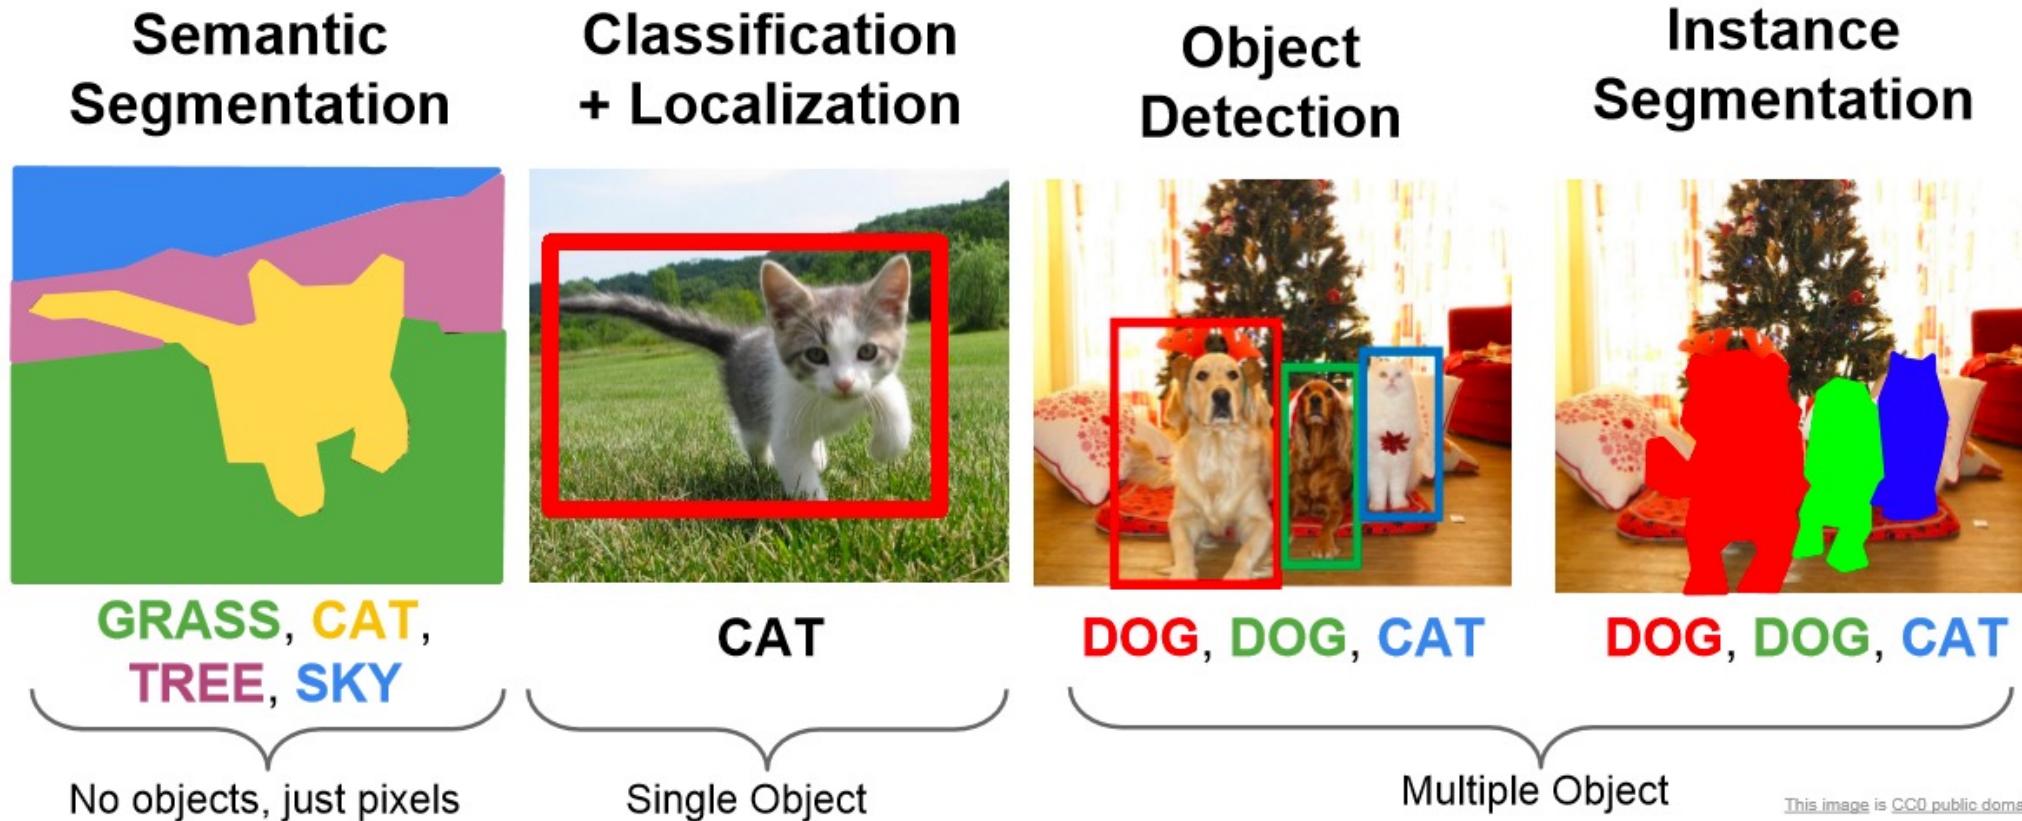
- Classification,
- regression
- object detection
- semantic segmentation,
- image captioning, etc.



→ **Cat**

Classification

# Applications of supervised learning

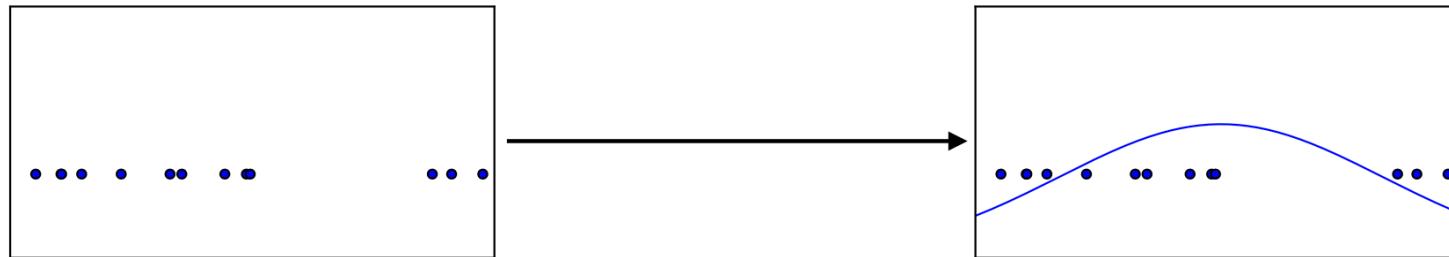


# Why generative models?

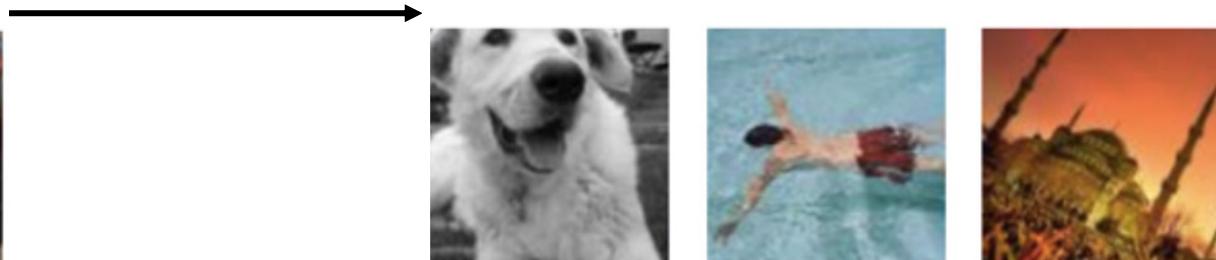
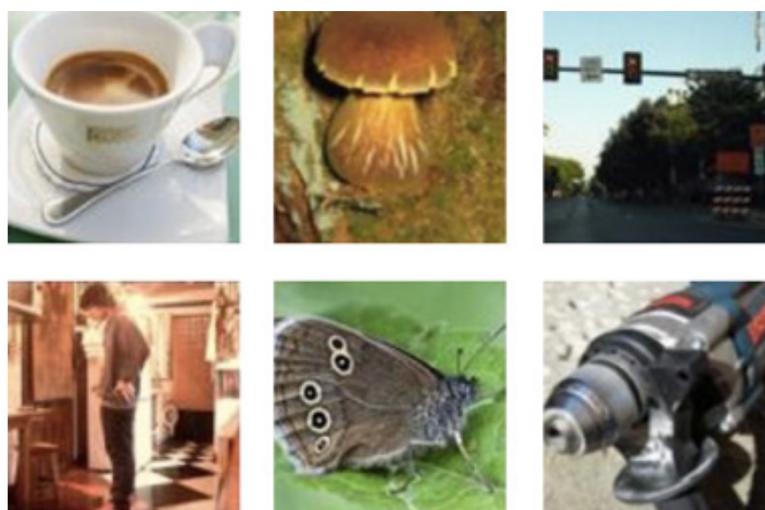
- We've only seen **discriminative models** so far
  - Given an data  $x$ , predict a label  $y$
  - Estimates  $p(y|x)$
- Discriminative models have several key limitations
  - Can't model  $p(x)$ , i.e., the probability of seeing a certain data
  - Thus, can't sample from  $p(x)$ , i.e., **can't generate new data**
- **Generative models** (in general) cope with all of above
  - Can model  $p(x)$
  - Can generate new data or images

# Generative modeling

Density estimation – a core problem in unsupervised learning



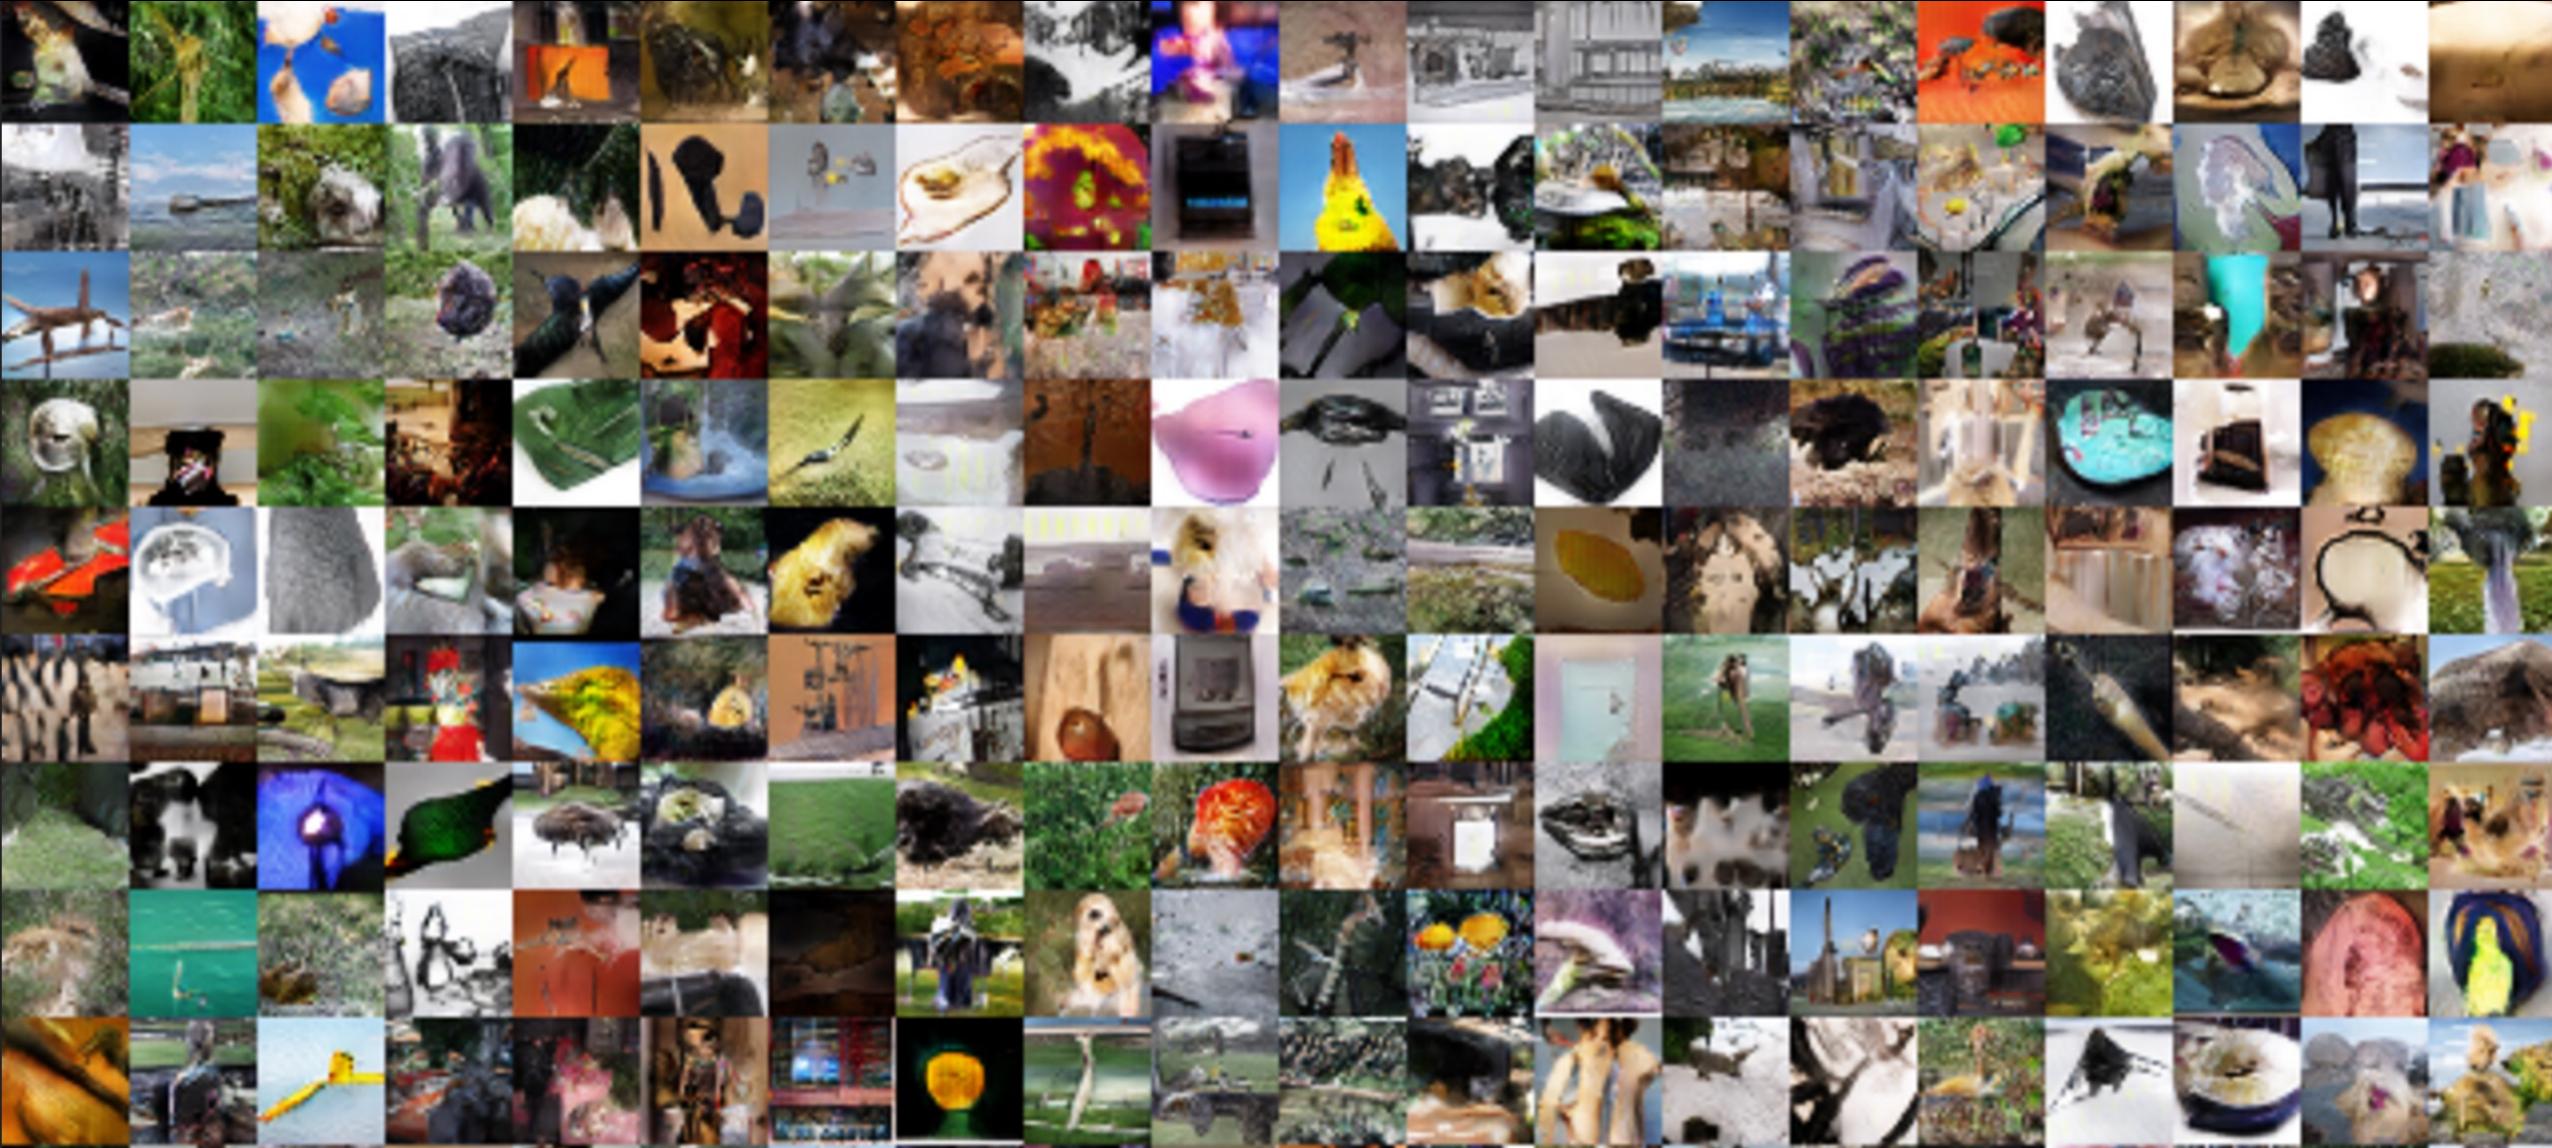
Sample generation



Training samples –  $p_{data}(x)$

Model samples –  $p_{model}(x)$   
(ideally  $p_{model}(x) = p_{data}(x)$ )

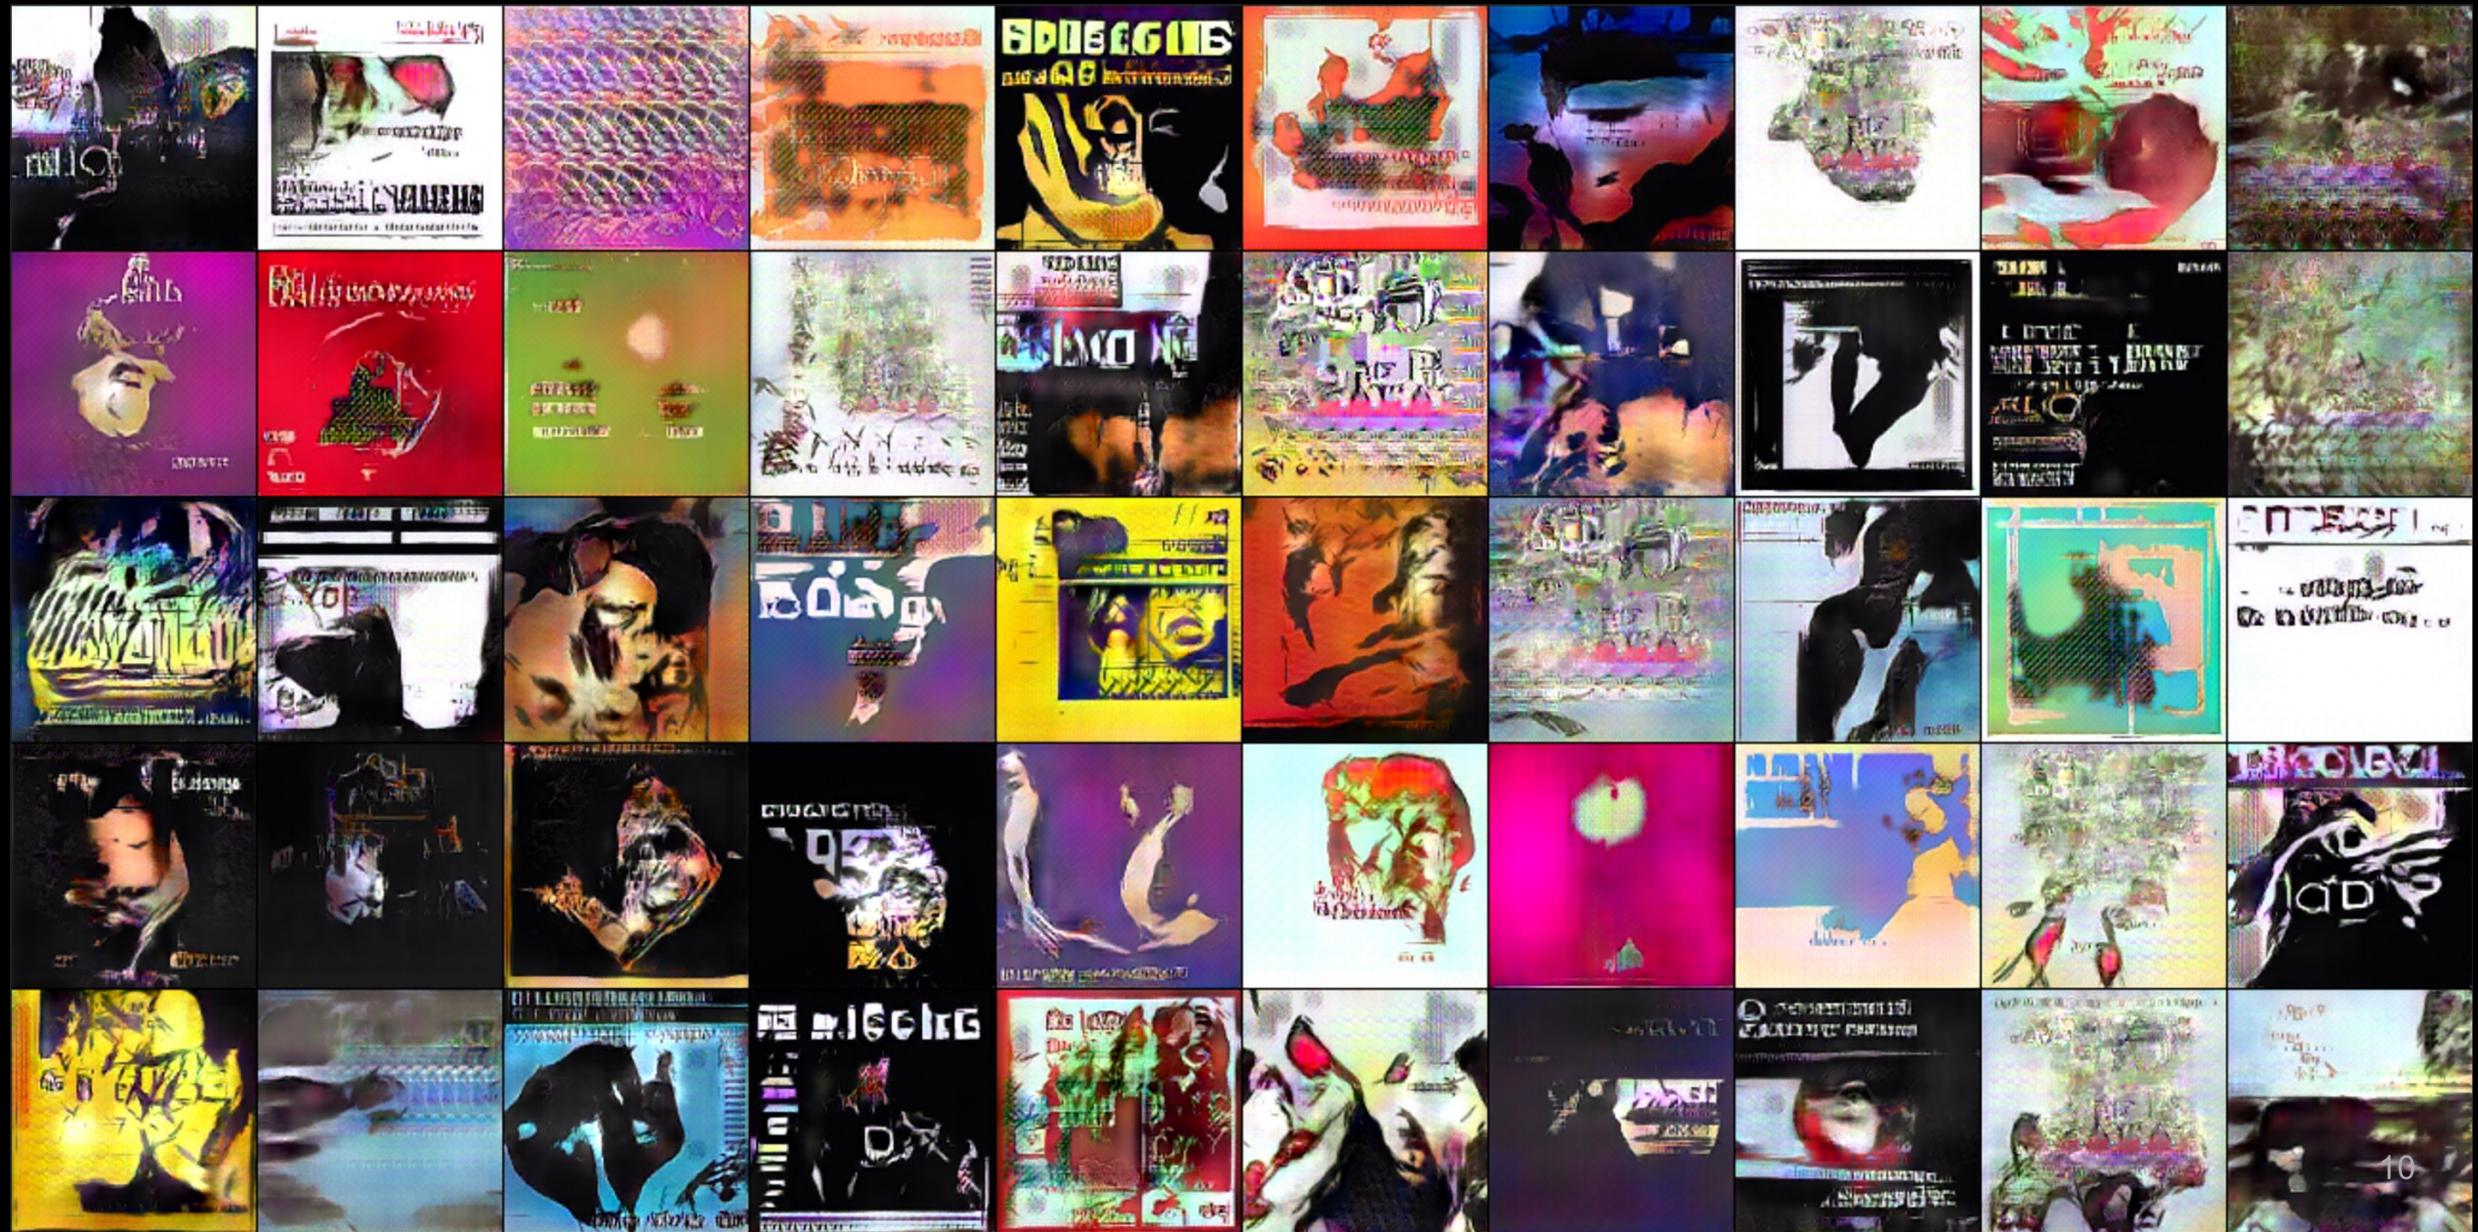
# Image generated - ImageNet



# Image generated - Bedrooms



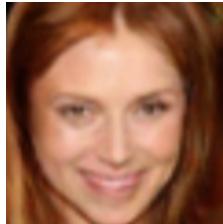
# Image generated - Albums



# The GAN Revolution 2014-2017



Conditional GAN



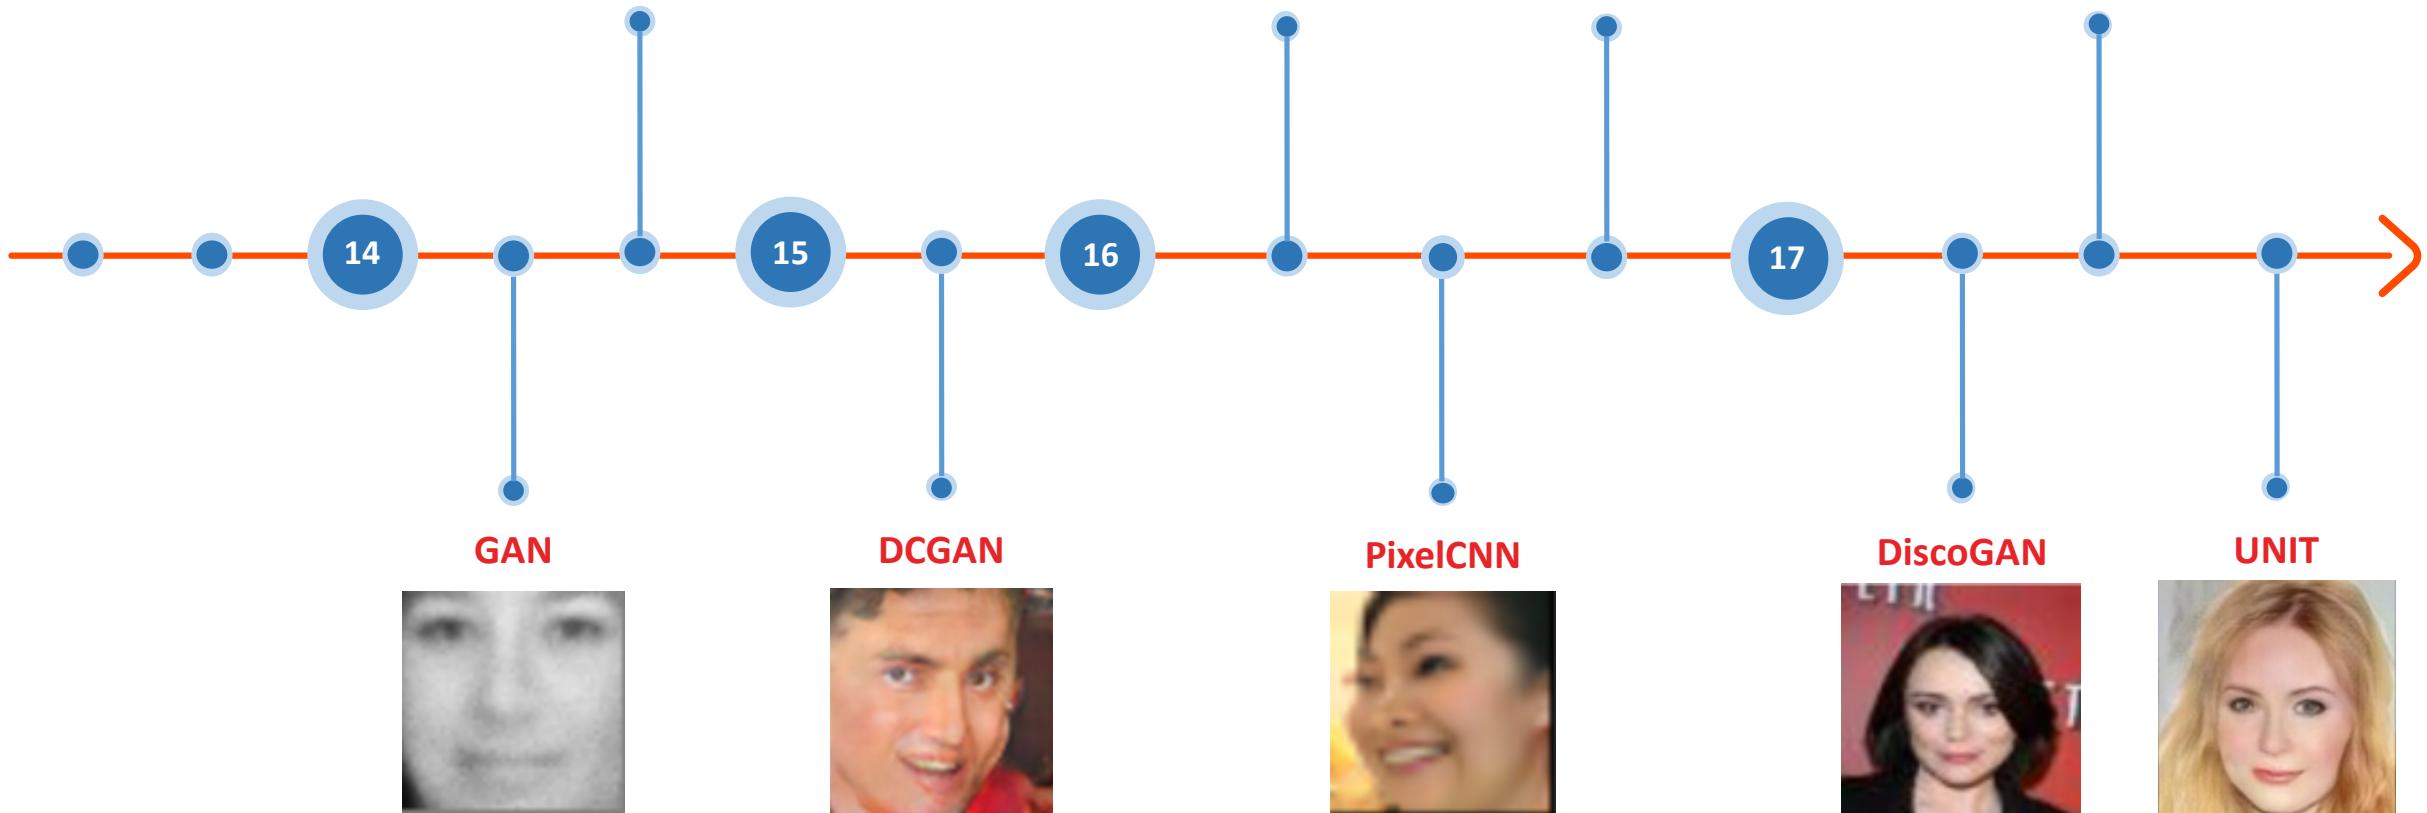
VAE/GAN



CoGAN



CycleGAN



# The GAN Revolution 2018 - Present



StarGAN



ELEGANT



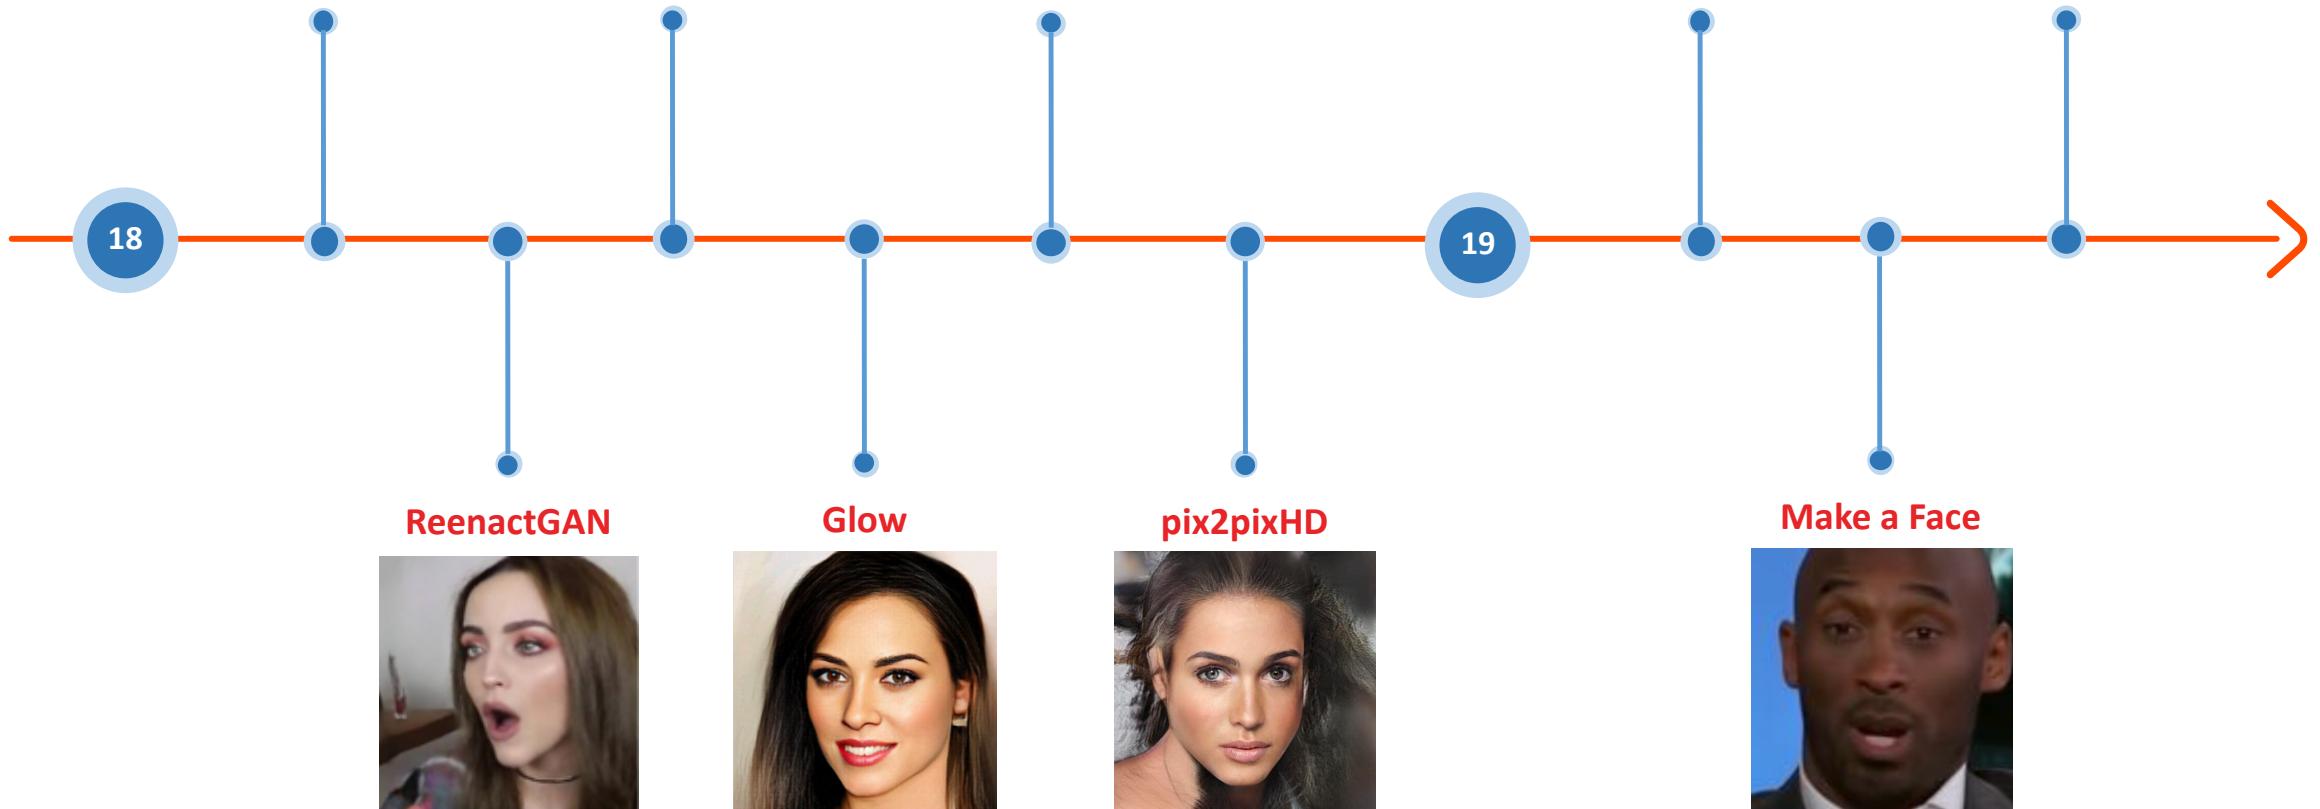
PGGAN



BeautyGlow



StyleGAN



# Results of StyleGAN



Credit: Tero Karras et al., A Style-Based Generator Architecture for Generative Adversarial Networks, CVPR 2019

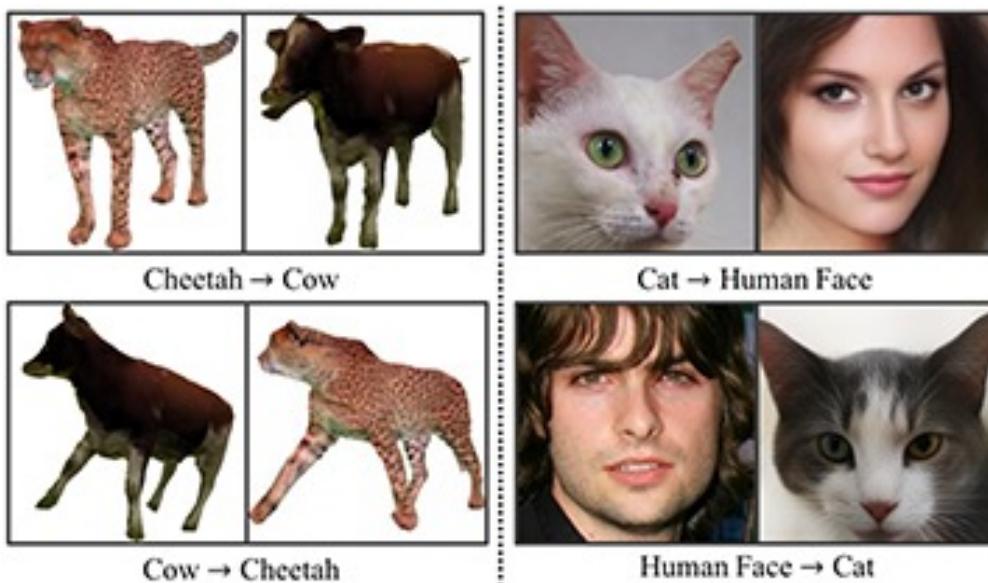
# Almost Everything can be Forged!



**CycleGAN**  
[Zhu et al., ICCV 2017]

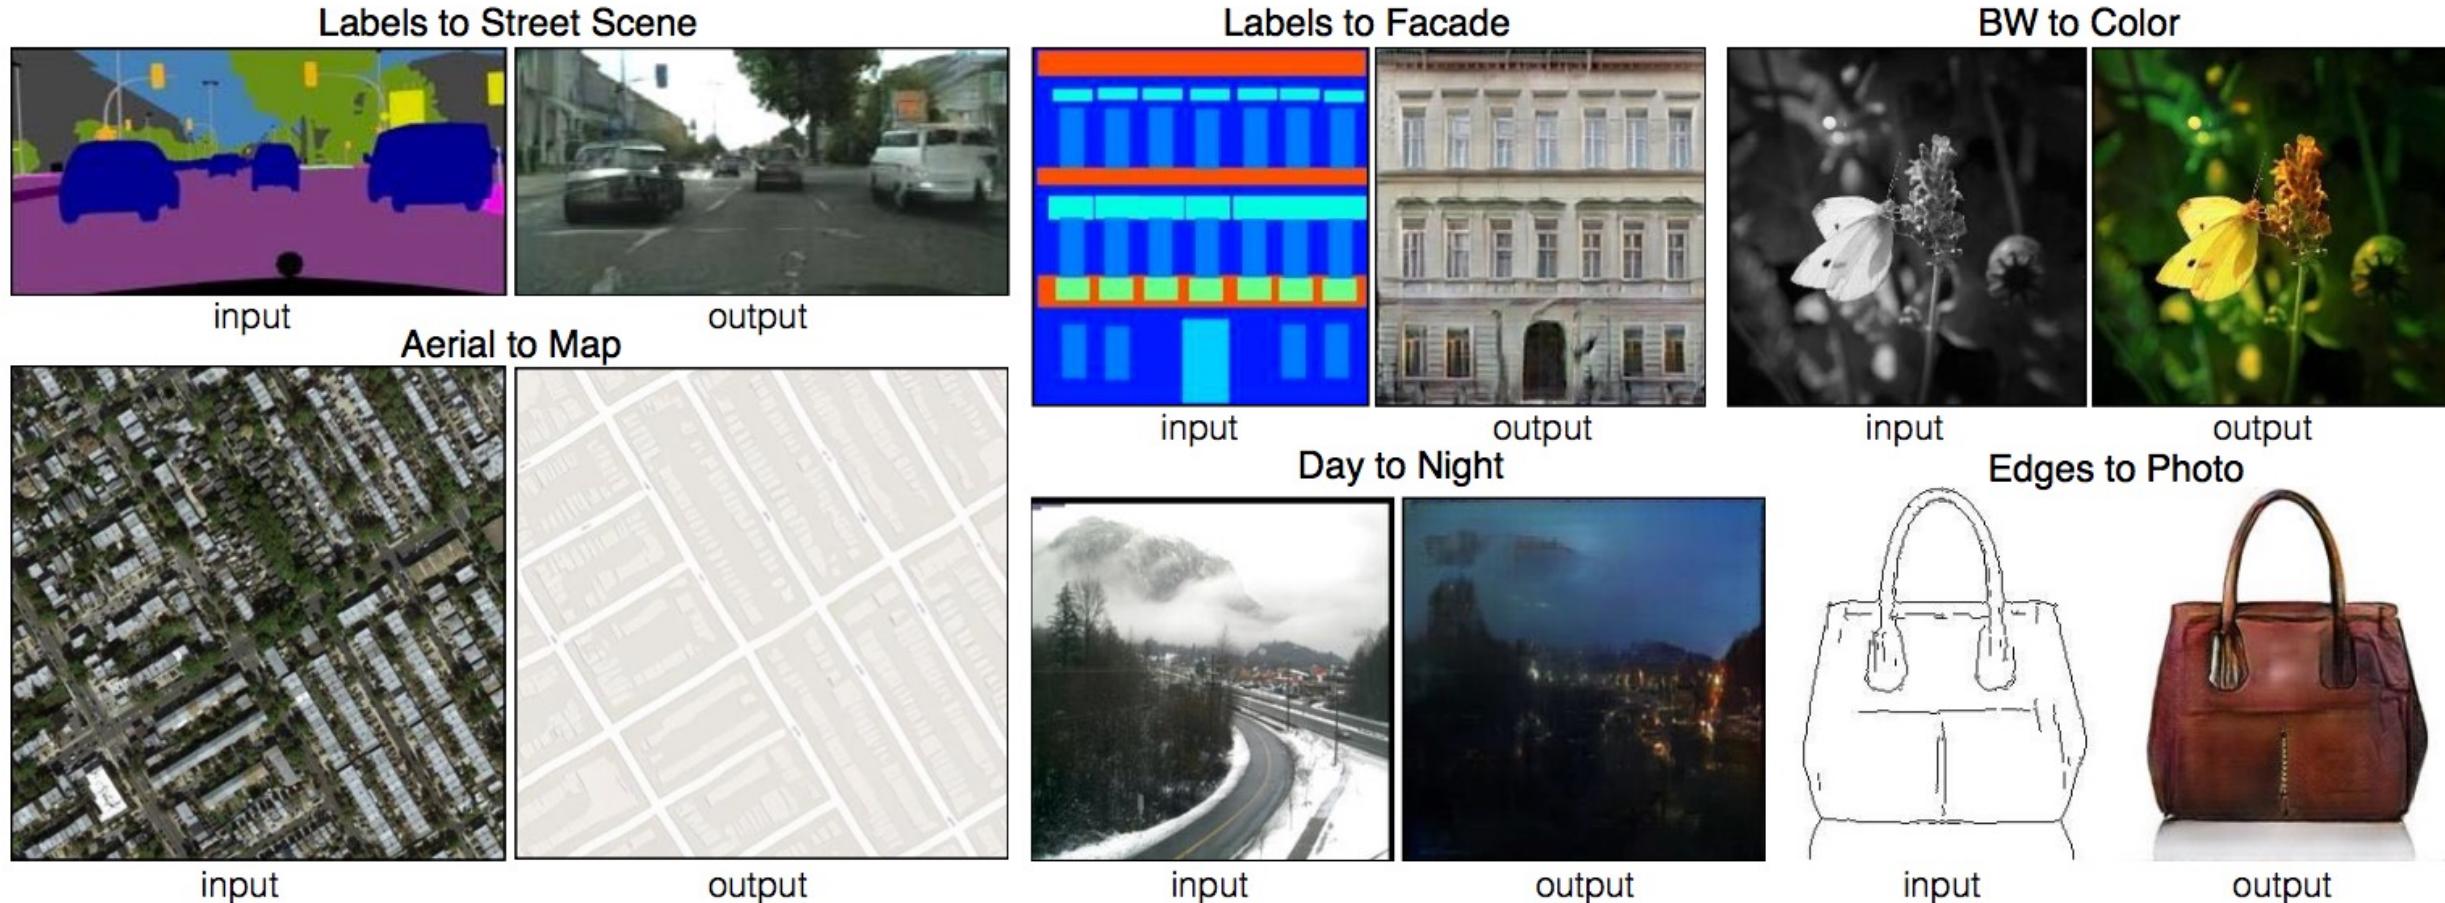


**StarGAN**  
[Choi et al., CVPR 2018]



**TransGaGa**  
[Wu et al., CVPR 2019]

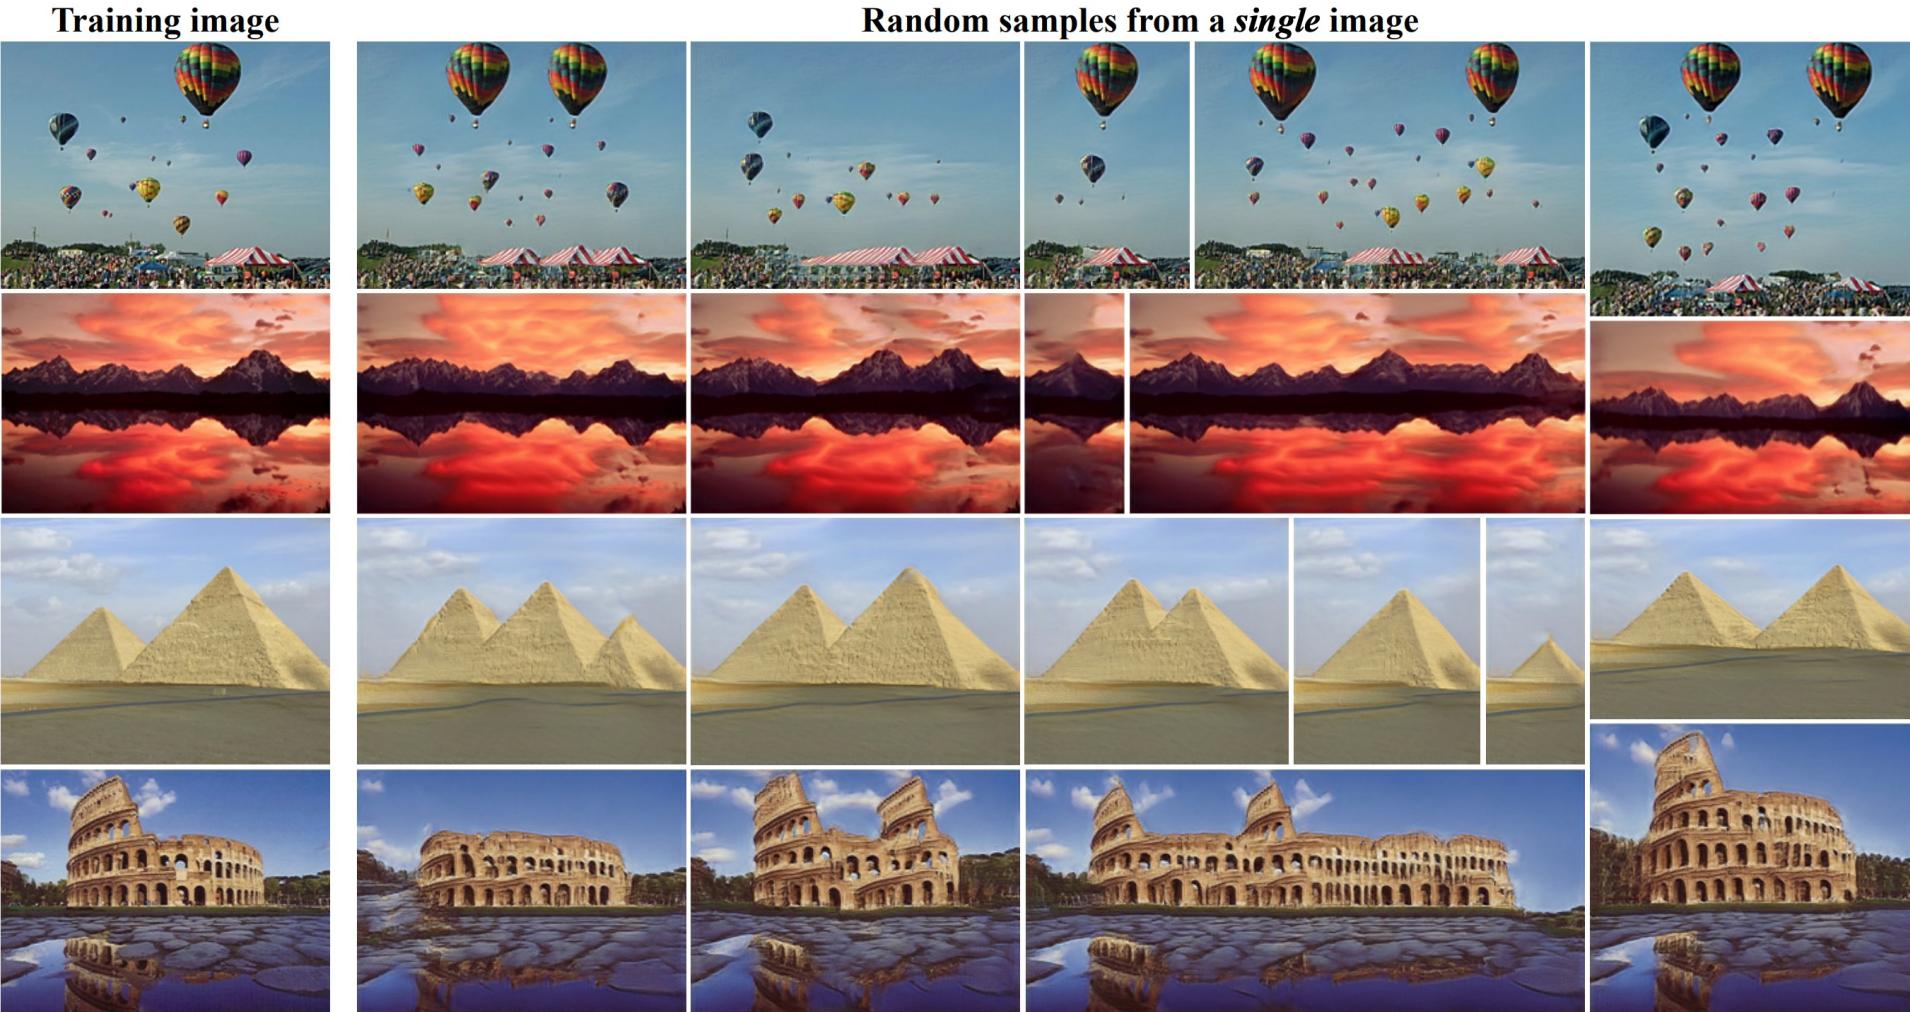
# Image-to-Image Translation



Isola et al. "Image-to-Image Translation with Conditional Adversarial Networks", arXiv:1611.07004

<https://phillipi.github.io/pix2pix/>

# SinGAN: Learning a Generative Model from a Single Natural Image



# Audio-driven Emotional Face Manipulation



Angry



Contempt



Fear



Input Video



Reference Style



Our Toonification Result

Edit upper length (StyleSpace)



Edit bottom length (StyleSpace)



Edit upper length (InterFaceGAN)



Edit bottom length (InterFaceGAN)

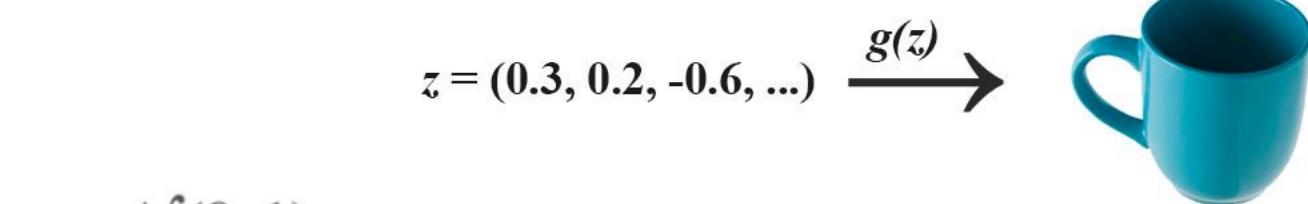


# Outline

- GAN Basics
- GAN Training
- DCGAN
- Mode Collapse

# GAN Basics

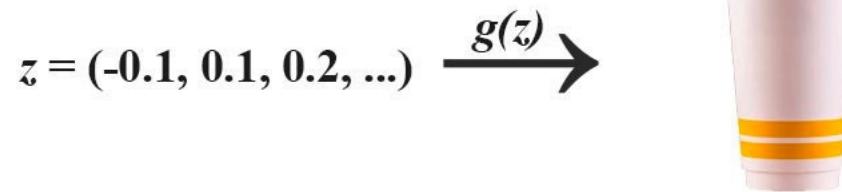
# Generative Adversarial Networks (GAN)



$z \sim \mathcal{N}(0, 1)$

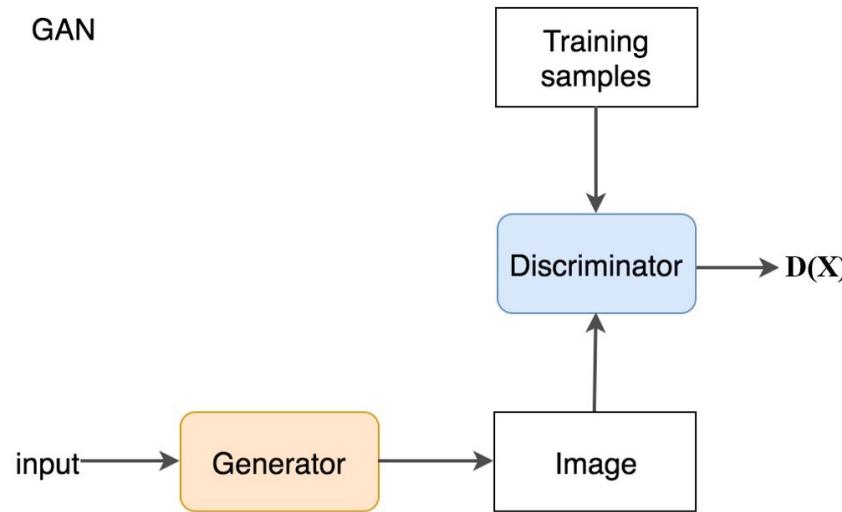
or

$z \sim U(-1, 1)$



GAN samples noise  $z$  using normal or uniform distribution and utilizes a deep network generator  $G$  to create an image  $x$  ( $x=G(z)$ )

# Generative Adversarial Networks (GAN)

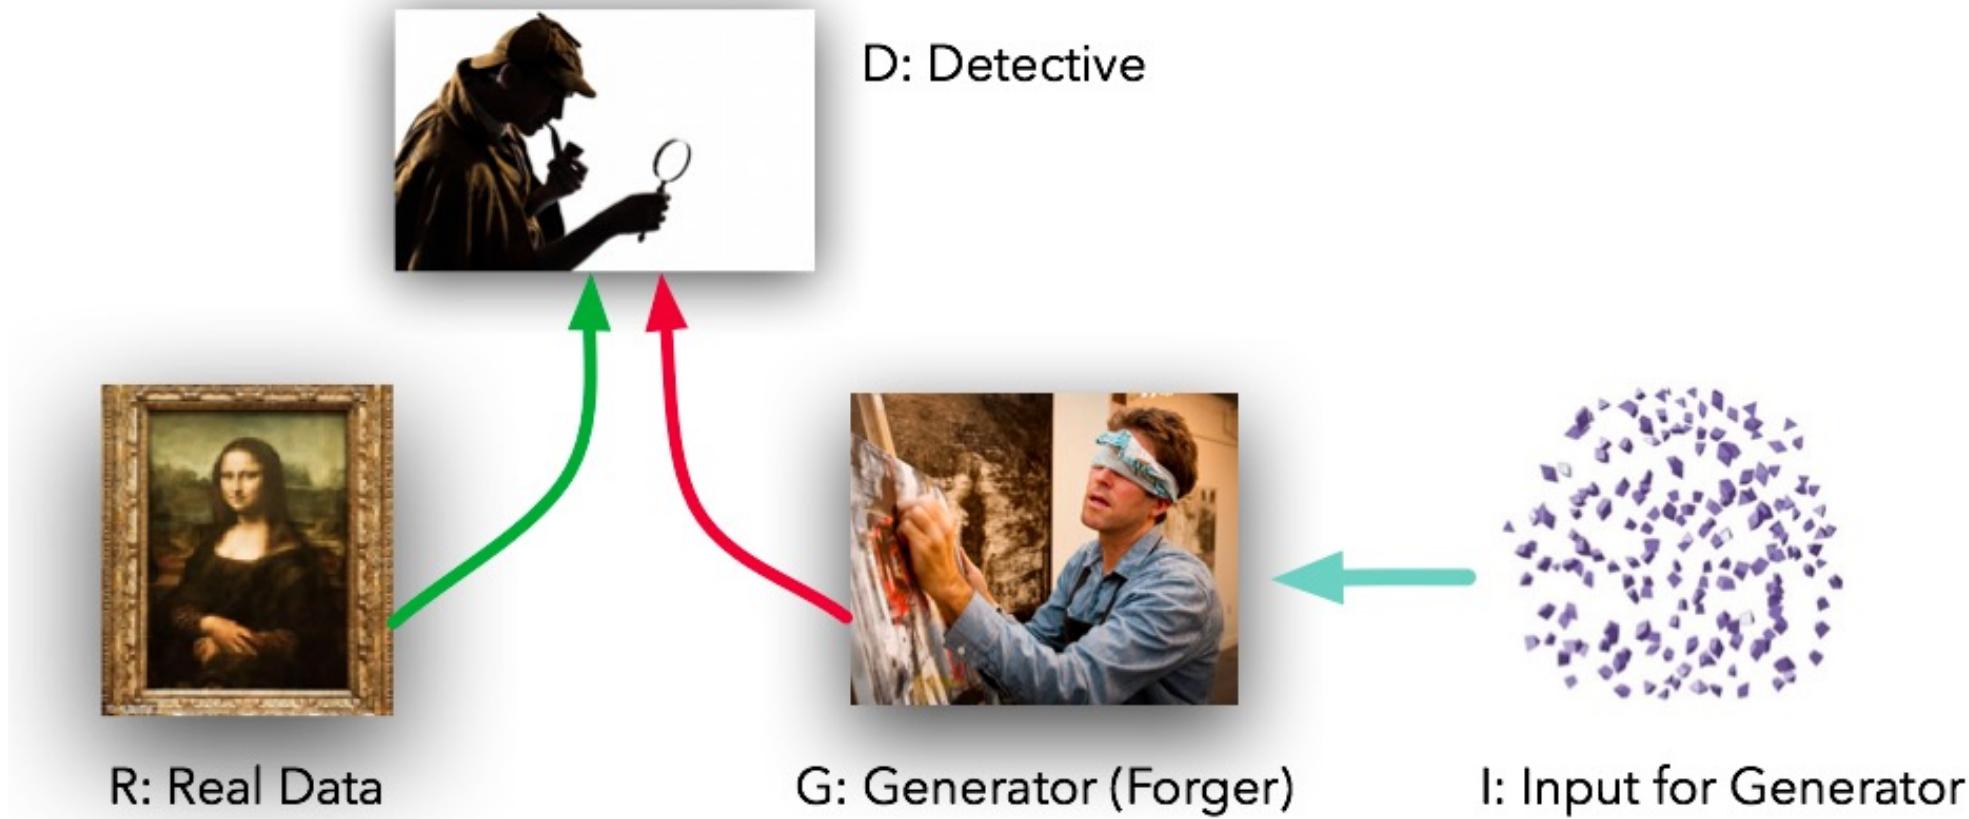


In GAN, we add a discriminator to distinguish whether the discriminator input is real or generated. It outputs a value  $D(x)$  to estimate the chance that the input is real.

**$D(x) = 1$  suggests  $x$  is real**

**$D(x) = 0$  suggests  $x$  is fake**

# Generative Adversarial Networks (GAN)



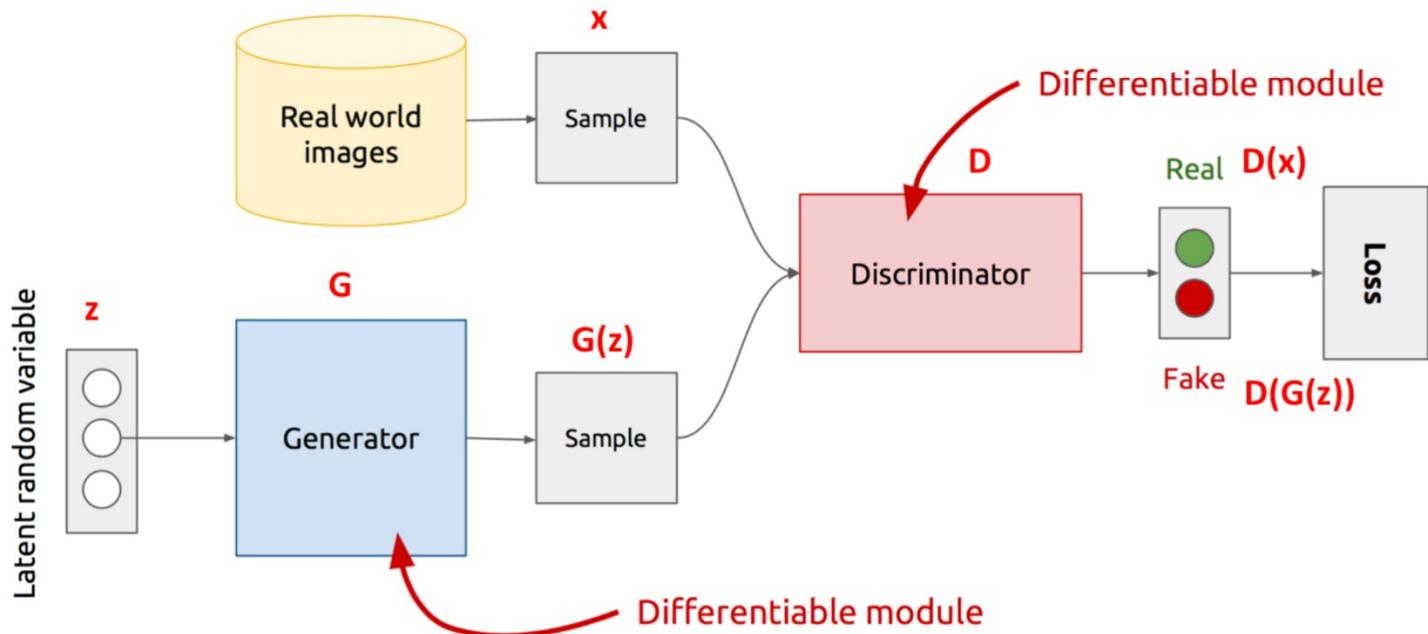
# Generative adversarial networks (GAN)

- **Generative model  $G$ :**

- Captures data distribution
- Fool  $D(G(z))$
- Generate an image  $G(z)$  such that  $D(G(z))$  is wrong (i.e.  $D(G(z)) = 1$ )

- **Discriminative model  $D$ :**

- Distinguishes between real and fake samples
- $D(x) = 1$  when  $x$  is a real image, and otherwise

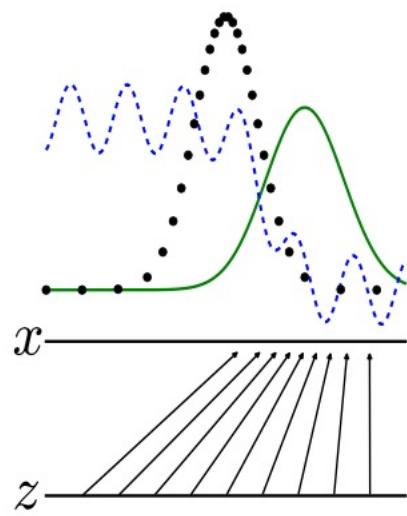


$z$  is some random noise (Gaussian/Uniform).

$z$  can be thought as the latent representation of the data.

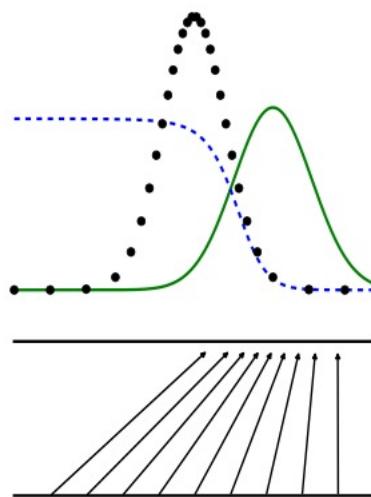
# Generative adversarial networks (GAN)

----- data generating distribution  
——— generative distribution  
- - - discriminative distribution



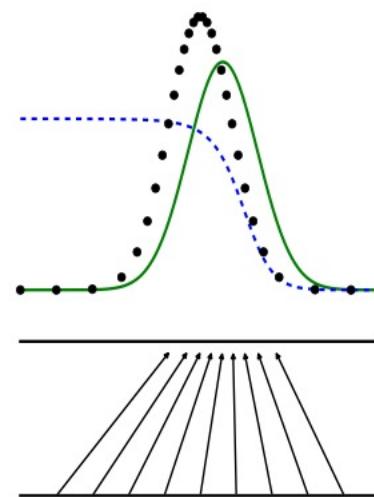
(a)

An adversarial pair near convergence



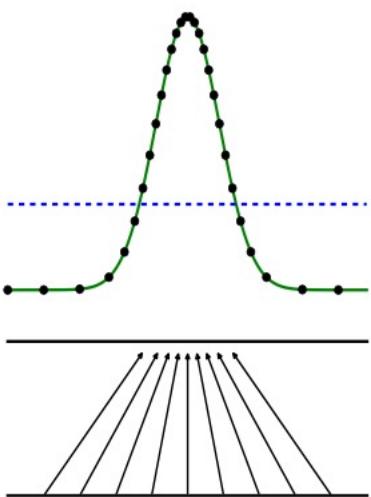
(b)

$D$  is trained to discriminate samples from data



(c)

After an update to  $G$ , gradient of  $D$  has guided  $G(z)$  to flow to regions that are more likely to be classified as data

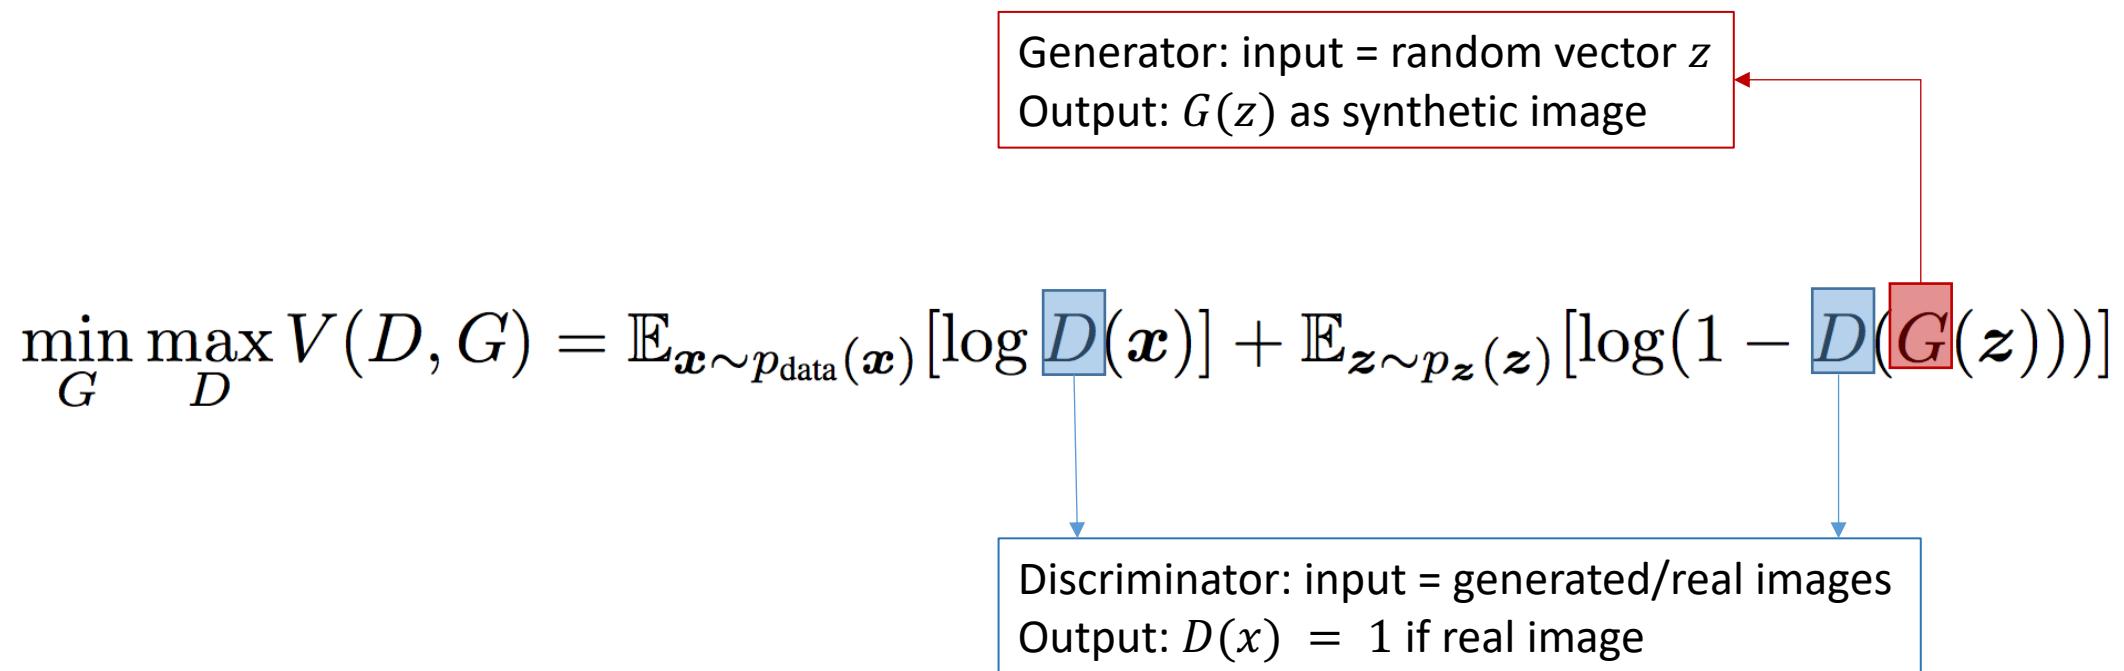


(d)

After several steps of training,  $G$  and  $D$  will reach a point at which both cannot improve because  $p_g = p_{data}$ . The discriminator is unable to differentiate between the two distributions, i.e.  $D(x) = 0.5$ .

# Generative adversarial networks (GAN)

- $D$  and  $G$  play the following two-player minimax game with value function  $V(D, G)$



# Outline

- GAN Basics
- GAN Training
- DCGAN
- Mode Collapse

# GAN Training

# Training procedure

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
- Update the discriminator by ascending its stochastic gradient:

maximize

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Update the generator by descending its stochastic gradient:

minimize

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))).$$

**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**  
**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

Average over  $m$  samples

Uniform noise vector (random numbers)

maximize

Real images

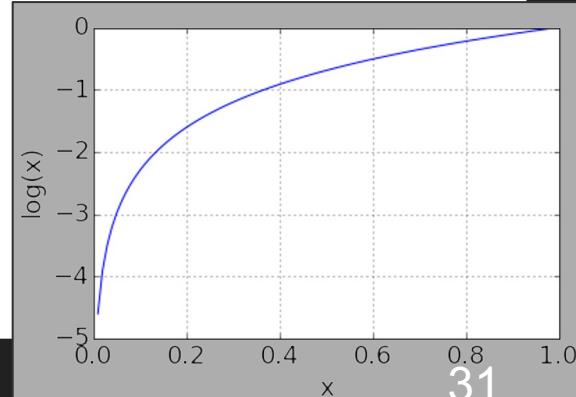
**Generator:**  
 input=*random numbers*,  
 output=*synthetic image*

**Discriminator:** (input=*generated/real image*,  
 output=*prediction of real image*)

**end for**

Let's do an example

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

Average over  $m$  samples

$$D(x) = 0.8 \\ \log(0.8) = -0.2$$

Imagine for a real image  $D(x)$  scores 0.8 that it is a real image (good)

Uniform noise vector (random numbers)

maximize

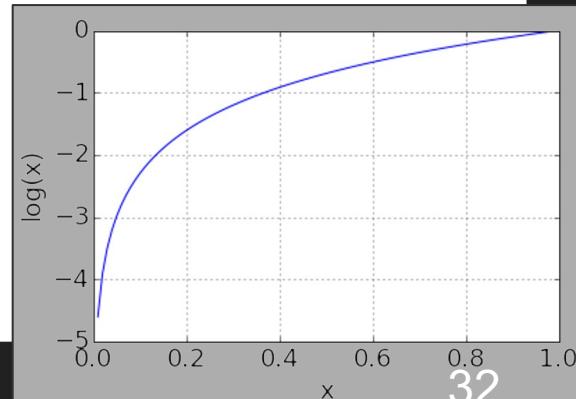
Real images

**Generator:**

input=random numbers,  
output=synthetic image

**end for**

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**  
**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**Discriminator:** (input=generated/real image, output=prediction of real image)

**end for**

Average over  $m$  samples

$$D(x) = 0.8 \\ \log(0.8) = -0.2$$

Then for a **generated** image,  $D(x)$  scores 0.2 that it is a real image (good)

**end for**

$$D(G(z)) = 0.2 \\ \log(1-0.2) = \log(0.8) = -0.2$$

Uniform noise vector (random numbers)

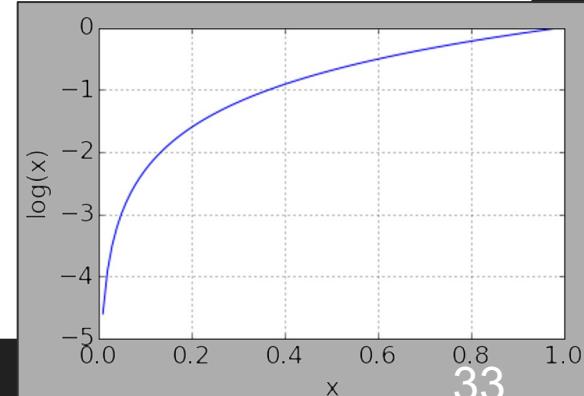
maximize

Real images

**Generator:**

input=random numbers, output=synthetic image

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

Average over  $m$  samples

**Discriminator:** (input=generated/real image, output=prediction of real image)

**end for**

$$D(x) = 0.8$$

$$\log(0.8) = -0.2$$

$$D(G(z)) = 0.2$$

$$\log(1-0.2) = \log(0.8) = -0.2$$

$$-0.2 + -0.2 = -0.4$$

Uniform noise vector (random numbers)

maximize

Real images

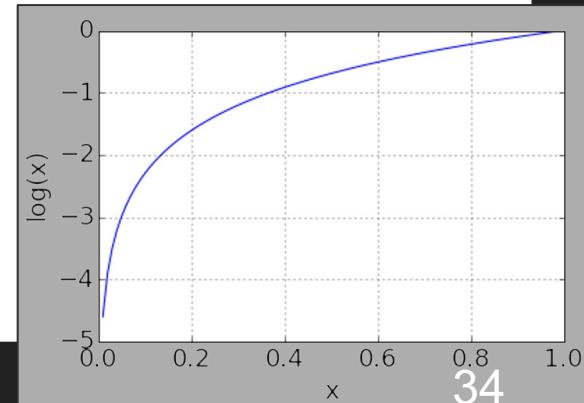
**Generator:**

input=random numbers,  
output=synthetic image

We add them together and this gives us a fairly high (-0.4) loss (note we ascend so we want to maximize this).

Also note that we are adding two negative numbers so 0 is the upper bound

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**  
**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

Average over  $m$  samples

Discriminator: (input=generated/real image, output=prediction of real image)

**end for**

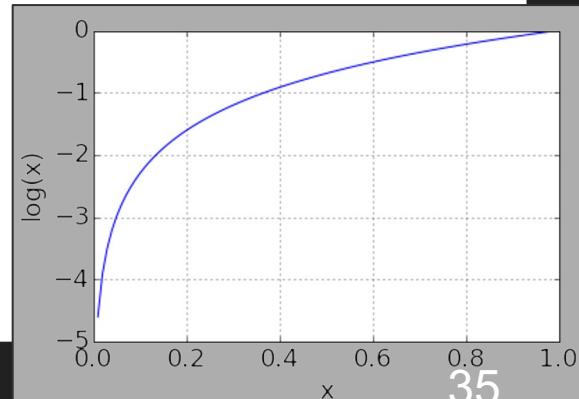
$$D(x) = 0.8 \\ \log(0.8) = -0.2$$

$$D(G(z)) = 0.2 \\ \log(1-0.2) = \log(0.8) = -0.2 \\ -0.2 + -0.2 = -0.4$$

**end for**

Let's do another example

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

Average over  $m$  samples

Discriminator: (input=generated/real image, output=prediction of real image)

$$D(x) = 0.8$$

$$\log(0.8) = -0.2$$

$$D(G(z)) = 0.2$$

$$\log(1-0.2) = \log(0.8) = -0.2$$

**end for**

$$-0.2 + -0.2 = -0.4$$

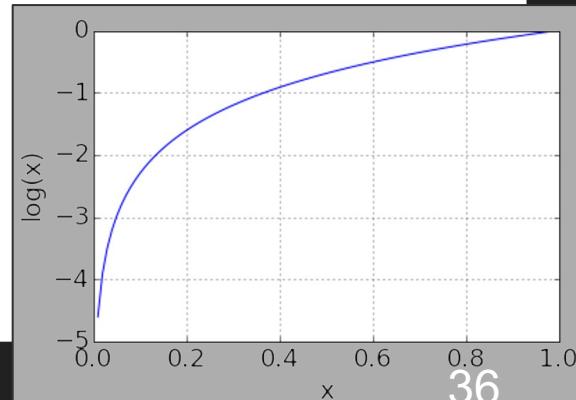
$$D(x) = 0.2$$

$$\log(0.2) = -1.6$$

$D(x)$  scores 0.2 that a real image is a real image (bad)

**Generator:**  
 input=random numbers,  
 output=synthetic image

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

Average over  $m$  samples

$$D(x) = 0.8$$

$$\log(0.8) = -0.2$$

$D(x)$  scores 0.8 that the generated image is a real image (bad)

$$\text{end for } -0.2 + -0.2 = -0.4$$

maximize

Uniform noise vector (random numbers)

Real images

**Generator:**  
 input=*random numbers*,  
 output=*synthetic image*

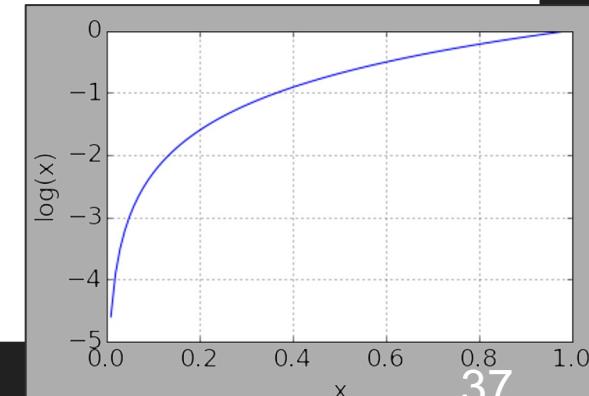
**Discriminator:** (input=*generated/real image*,  
 output=*prediction of real image*)

$$D(x) = 0.2$$

$$\log(0.2) = -1.6$$

$D(G(z)) = 0.8$   
 $\log(1-0.8) = \log(0.2) = -1.6$

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**  
**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

Average over  $m$  samples

$$D(x) = 0.8$$

$$\log(0.8) = -0.2$$

$$D(G(z)) = 0.2$$

$$\log(1-0.2) = \log(0.8) = -0.2$$

$$-0.2 + -0.2 = -0.4$$

**end for**

maximize

Uniform noise vector (random numbers)

Real images

**Generator:**  
 input=*random numbers*,  
 output=*synthetic image*

**Discriminator:** (input=*generated/real image*,  
 output=*prediction of real image*)

$$D(x) = 0.2$$

$$\log(0.2) = -1.6$$

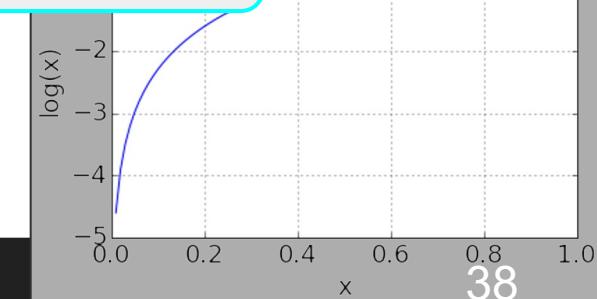
$$D(G(z)) = 0.8$$

$$\log(1-0.8) = \log(0.2) = -1.6$$

$$-1.6 + -1.6 = -3.2$$

Note showing  $\ln(x)$  not  $\log(x)$

These "bad" predictions combined gives a loss of -3.2



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**  
**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

Average over  $m$  samples

Uniform noise vector (random numbers)

maximize

Real images

**Generator:**  
 input=random numbers,  
 output=synthetic image

**Discriminator:** (input=generated/real image,  
 output=prediction of real image)

$$D(x) = 0.8$$

$$\log(0.8) = -0.2$$

$$D(G(z)) = 0.2$$

$$\log(1-0.2) = \log(0.8) = -0.2$$

**end for**

$$-0.2 + -0.2 = -0.4$$

$$D(x) = 0.2$$

$$\log(0.2) = -1.6$$

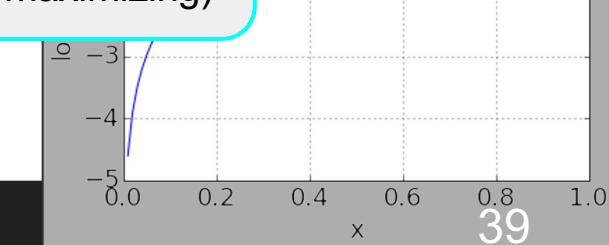
$$D(G(z)) = 0.8$$

$$\log(1-0.8) = \log(0.2) = -1.6$$

$$-1.6 + -1.6 = -3.2$$

Note showing  $\ln(x)$  not  $\log(x)$

Compare the loss to when we had "good" predictions (remember we are maximizing)



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**  
**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

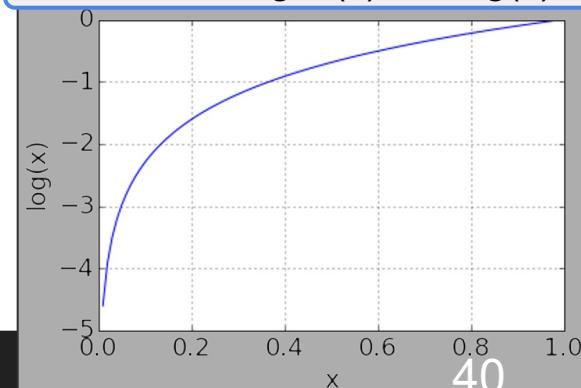
**Discriminator:** (input=generated/real image, output=prediction of real image)

**end for**

Average over  $m$  samples

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

Note showing  $\ln(x)$  not  $\log(x)$



**end for**

$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**Discriminator:** (input=generated/real image, output=prediction of real image)

**end for**

Average over  $m$  samples

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

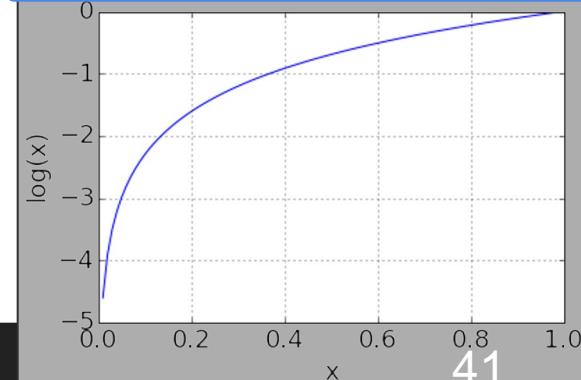
**end for**

Uniform noise vector (random numbers)

maximize

Real images

**Generator:**  
input=random numbers,  
output=synthetic image



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**  
**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**Discriminator:** (input=generated/real image, output=prediction of real image)

**end for**

Average over  $m$  samples

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$D(G(z)) = 0.2$$

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

$D(G(z))$  scores 0.2 that a generated image is a real image = bad :( We didn't fool D(.)

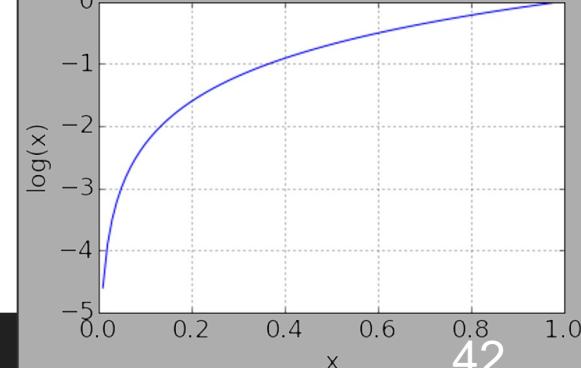
**end for**

Uniform noise vector (random numbers)

maximize

Real images

**Generator:**  
 input=random numbers,  
 output=synthetic image



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**Discriminator:** (input=generated/real image, output=prediction of real image)

**end for**

Average over  $m$  samples

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$D(G(z)) = 0.2$$

$$\log(1-0.2) = \log(0.8) = -0.2$$

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

Gets assigned a loss of -0.2

**end for**

Uniform noise vector (random numbers)

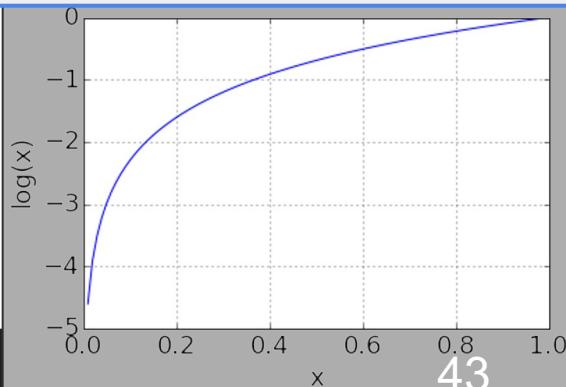
maximize

Real images

**Generator:**

input=random numbers,  
output=synthetic image

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**Discriminator:** (input=generated/real image, output=prediction of real image)

**end for**

Average over  $m$  samples

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by **descending** its stochastic gradient:

$$D(G(z)) = 0.2$$

$$\log(1-0.2) = \log(0.8) = -0.2$$

**minimize**

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

Notice here we want to minimize this loss function (not maximize like before)

**end for**

Uniform noise vector (random numbers)

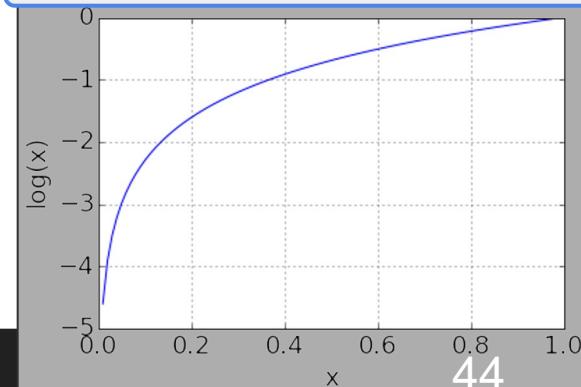
maximize

Real images

**Generator:**

input=random numbers,  
output=synthetic image

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**Discriminator:** (input=generated/real image, output=prediction of real image)

**end for**

Average over  $m$  samples

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by **descending** its stochastic gradient:

$$D(G(z)) = 0.2$$

**minimize**

$$\log(1-0.2) = \log(0.8) = -0.2$$

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

One more example ...

Uniform noise vector (random numbers)

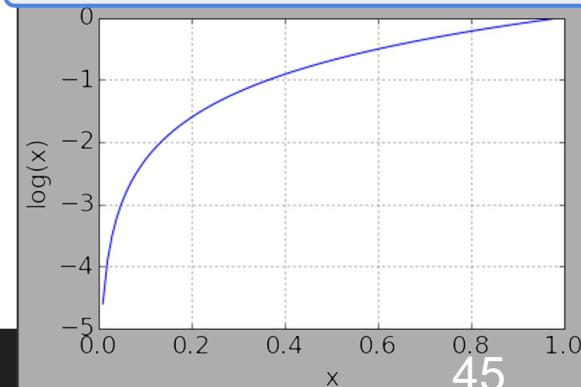
maximize

Real images

**Generator:**

input=random numbers,  
output=synthetic image

Note showing  $\ln(x)$  not  $\log(x)$



**end for**

$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**Discriminator:** (input=generated/real image, output=prediction of real image)

**end for**

Average over  $m$  samples

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by **descending** its stochastic gradient:

$$D(G(z)) = 0.2$$

$$\log(1-0.2) = \log(0.8) = -0.2$$

$$D(G(z)) = 0.8$$

**minimize**

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

$D(G(z))$  scores 0.8 that a generated image is a real image = good :) We fooled  $D(\cdot)$ !

**end for**

Uniform noise vector (random numbers)

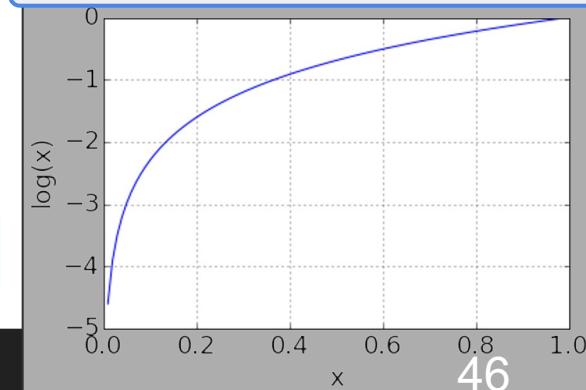
maximize

Real images

**Generator:**

input=random numbers,  
output=synthetic image

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**  
**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**Discriminator:** (input=generated/real image, output=prediction of real image)

**end for**

Average over  $m$  samples

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by **descending** its stochastic gradient:

$$D(G(z)) = 0.2$$

**minimize**

$$\log(1-0.2) = \log(0.8) = -0.2$$

$$D(G(z)) = 0.8$$

$$\log(1-0.8) = \log(0.2) = -1.6$$

**end for**

Uniform noise vector (random numbers)

**maximize**

Real images

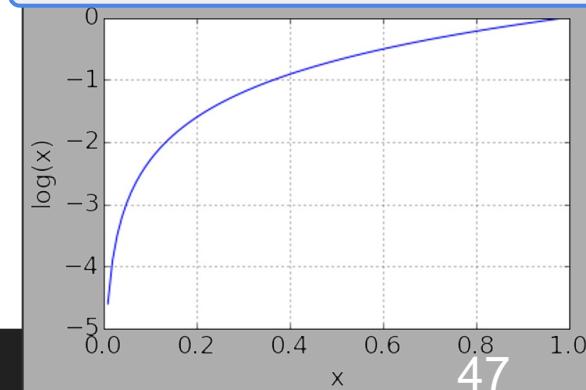
**Generator:**

input=random numbers, output=synthetic image

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

This gives loss of -1.6

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**  
**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**Discriminator:** (input=generated/real image, output=prediction of real image)

**end for**

Average over  $m$  samples

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by **descending** its stochastic gradient:

$$D(G(z)) = 0.2$$

**minimize**

$$\log(1-0.2) = \log(0.8) = -0.2$$

$$D(G(z)) = 0.8$$

$$\log(1-0.8) = \log(0.2) = -1.6$$

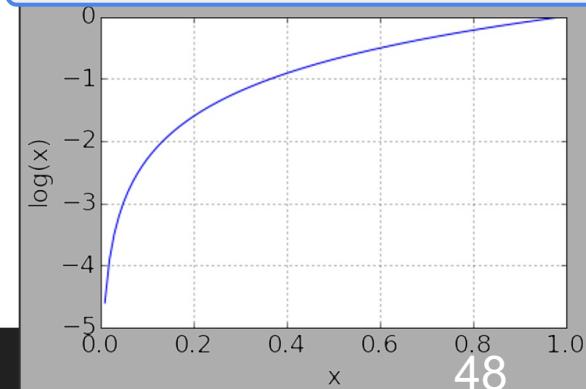
**end for**

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

So minimizing this loss means to generate images that fool  $D(\cdot)$

**Generator:**  
 input=random numbers,  
 output=synthetic image

Note showing  $\ln(x)$  not  $\log(x)$



$D(x)$  = probability that  $x$  is a real image  
 0 = generated image; 1 = real image

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by **maximize** its stochastic gradient:

Gradient w.r.t the parameters of the Discriminator

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

Discriminator: (input=generated/real image, output=prediction of real image)

**end for**

Average over  $m$  samples

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by **descending** its stochastic gradient:

**minimize**

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

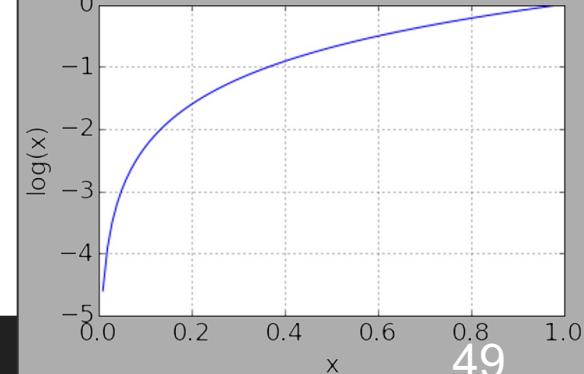
**end for**

Gradient w.r.t the parameters of the Generator

Uniform noise vector (random numbers)

Real images

**Generator:**  
input=random numbers,  
output=synthetic image



# Results from GAN



a)



b)



c)



d)

Training a system for rendering arbitrary images (arbitrary within the bounds of the subject matter that we trained the system on)

# Example 1

Design a GAN to learn Gaussian distributed data with mean = 1.0 and standard deviation = 1.0. The generator receives uniformly distributed 1-dimensional inputs in the range [-1, +1].

Discriminator is a 4-layer DNN:

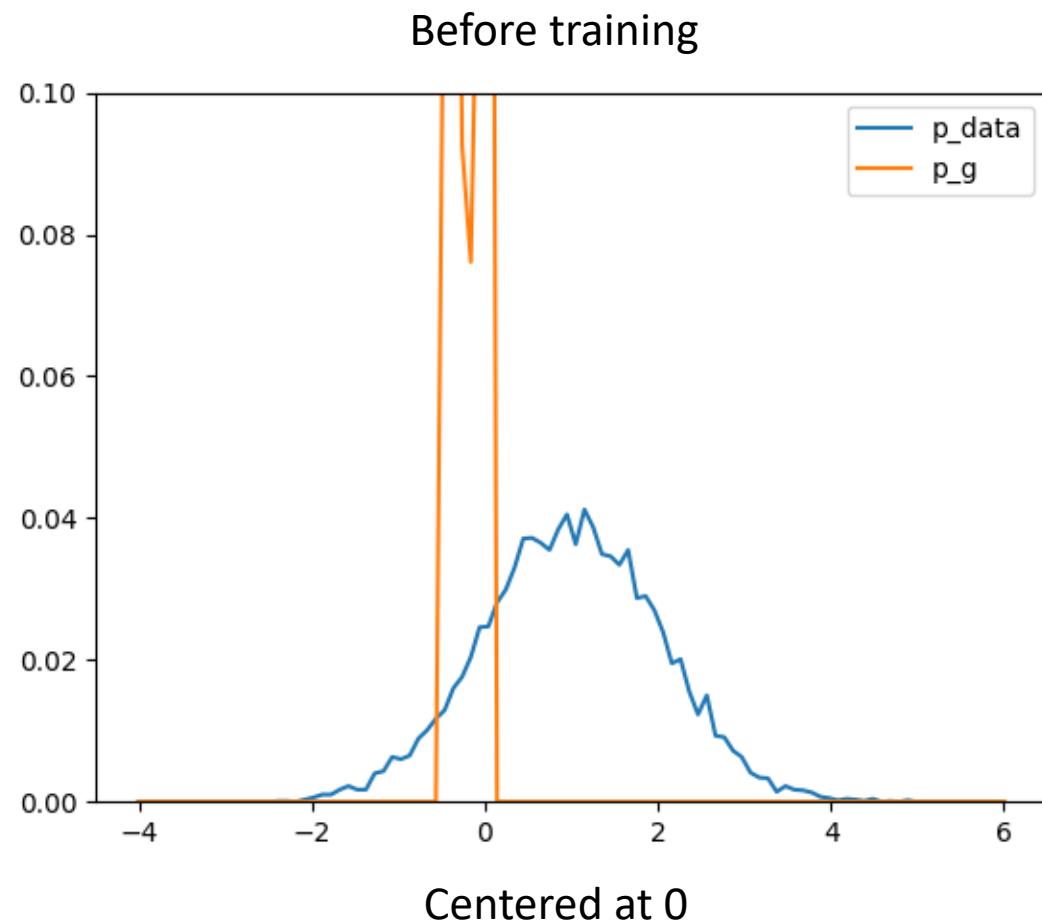
- Input dimension = 1
- Number of hidden neurons in hidden-layer 1 = 10 (activation: Tanh)
- Number of hidden neurons in hidden-layer 2 = 10 (activation: Tanh)
- Output dimension = 1 (binary classification)

Generator is a 4-layer DNN:

- Input dimension = 1
- Number of hidden neurons in hidden-layer 1 = 5 (activation: Tanh)
- Number of hidden neurons in hidden-layer 2 = 5 (activation: Tanh)
- Output dimension = 1 (activation: None)

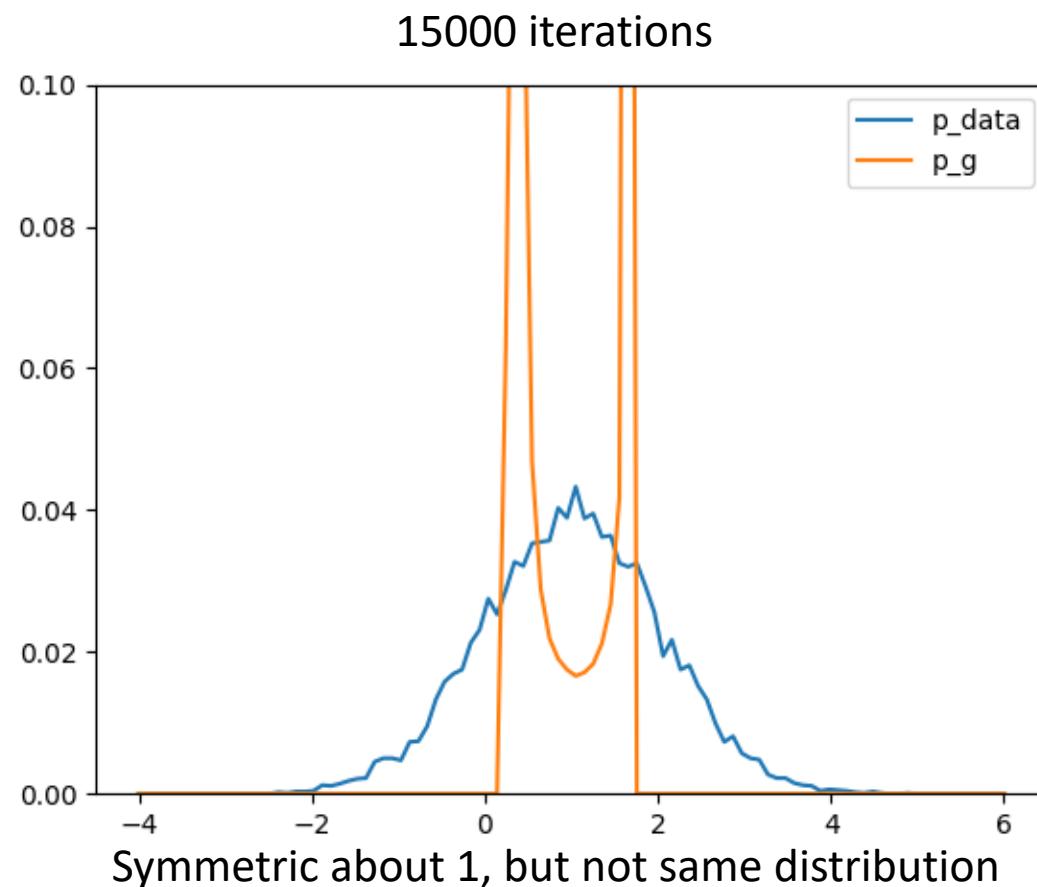
# Example 1

Histogram of 10000 randomly drawn points from  $U[-1, 1]$



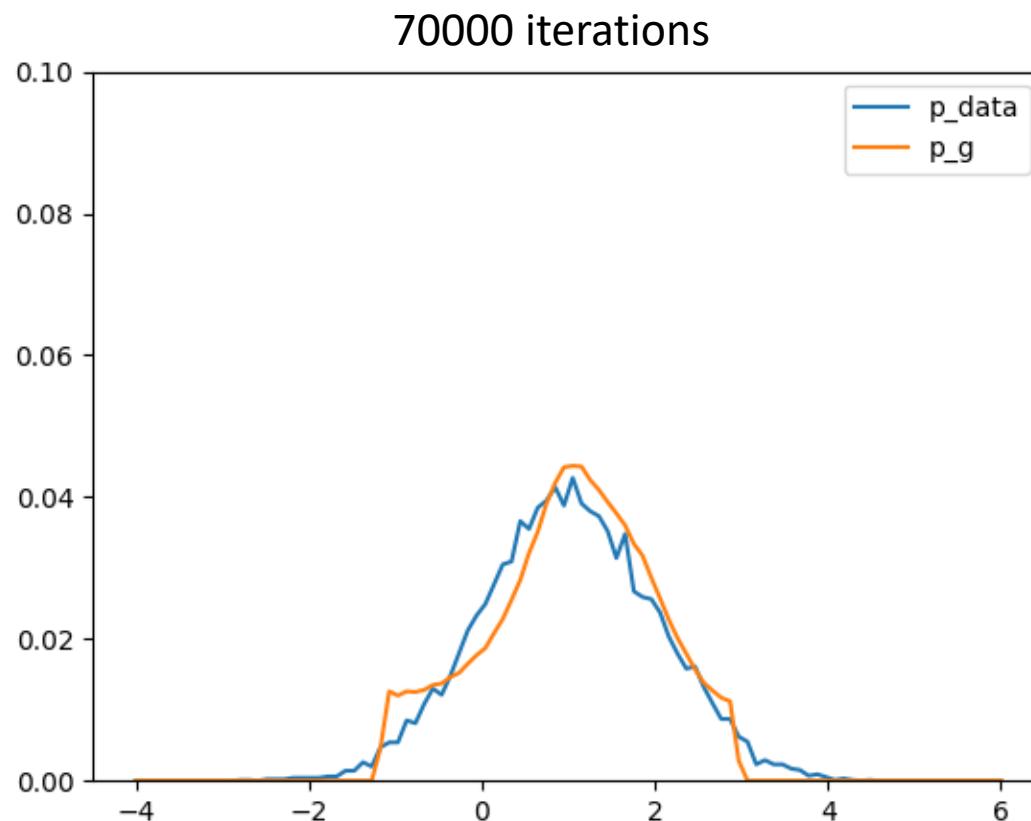
# Example 1

Histogram of 10000 randomly drawn points from  $U[-1, 1]$



# Example 1

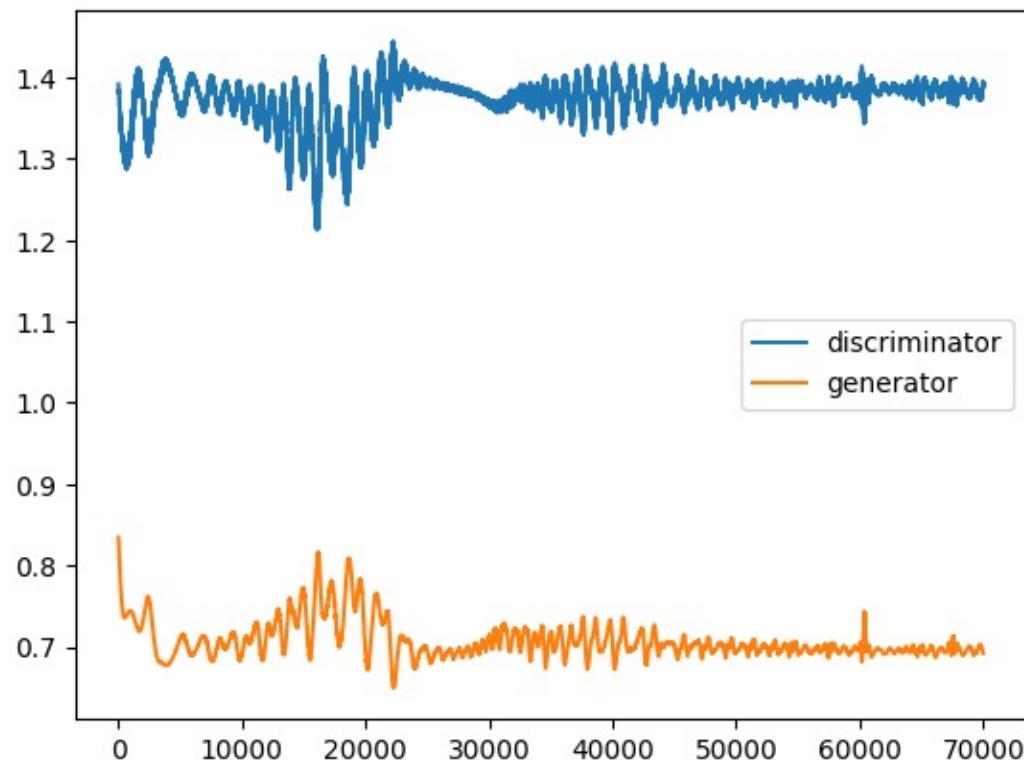
**Histogram of 10000 randomly drawn points from  $U[-1, 1]$**



Not perfect, but somehow similar to the data distribution

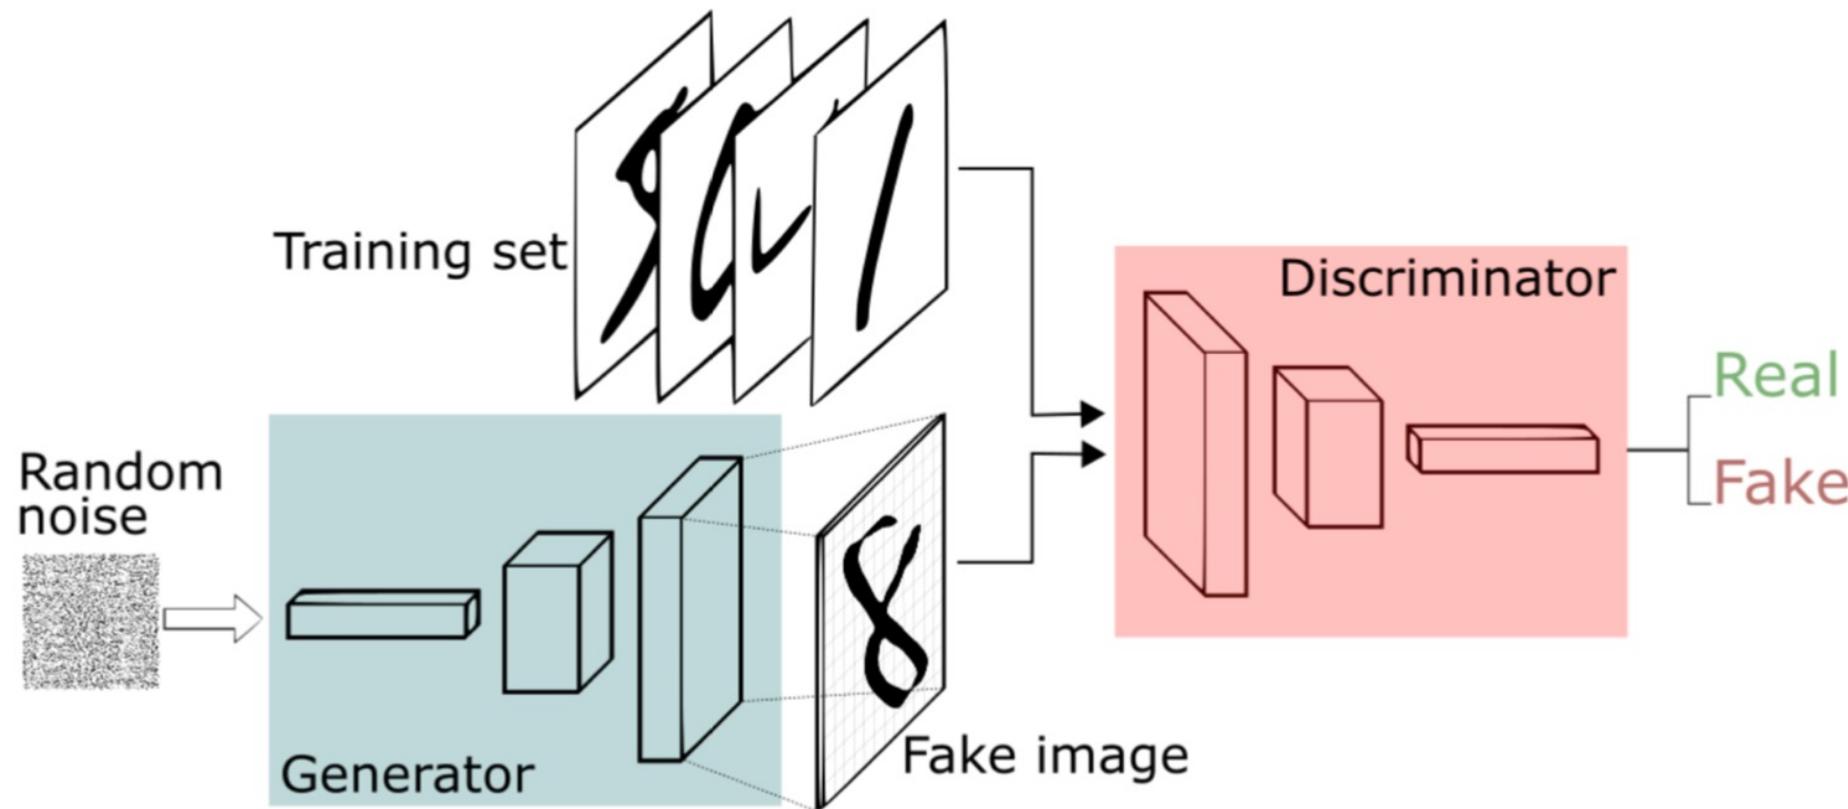
# Example 1

Training Curves



## Example 2

Design a GAN to generate MNIST images from a uniformly distributed noise vector of 100 dimensions.



# Example 2

## Generator:

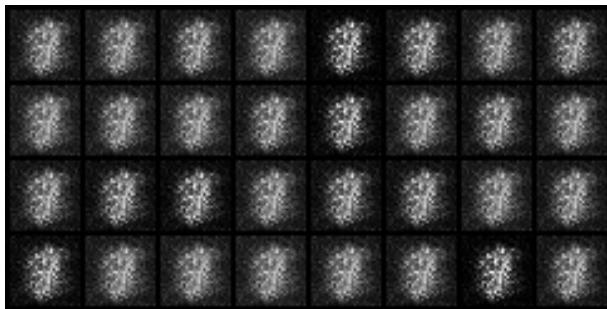
- Input dimension = 100, drawn from a uniform distribution  $U(-1, 1)$
- Hidden layers
  - Number of hidden neurons = 256 (activation: ReLU)
  - Number of hidden neurons = 512 (activation: ReLU)
  - Number of hidden neurons = 1024 (activation: ReLU)
- Output dimension = 784 (activation: Tanh)

## Discriminator:

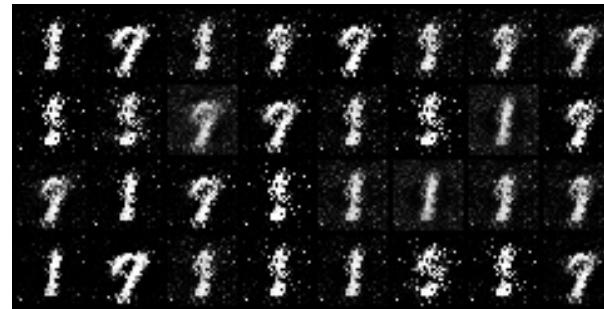
- Input dimension = 784
- Hidden layers
  - Number of hidden neurons = 1024 (activation: ReLU)
  - Number of hidden neurons = 512 (activation: ReLU)
  - Number of hidden neurons = 256 (activation: ReLU)
- Output dimension = 1 (activation: Sigmoid, binary classification)

## Example 2

Start of training



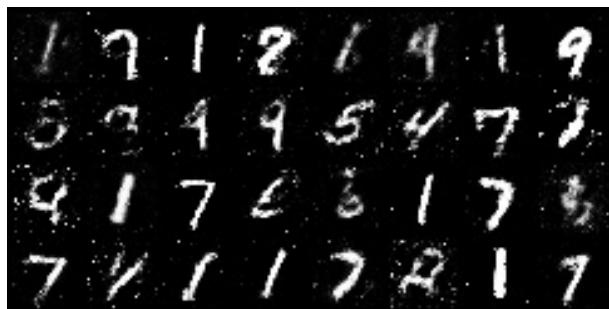
10<sup>th</sup> epoch



20<sup>th</sup> epoch



30<sup>th</sup> epoch



40<sup>th</sup> epoch



50<sup>th</sup> epoch

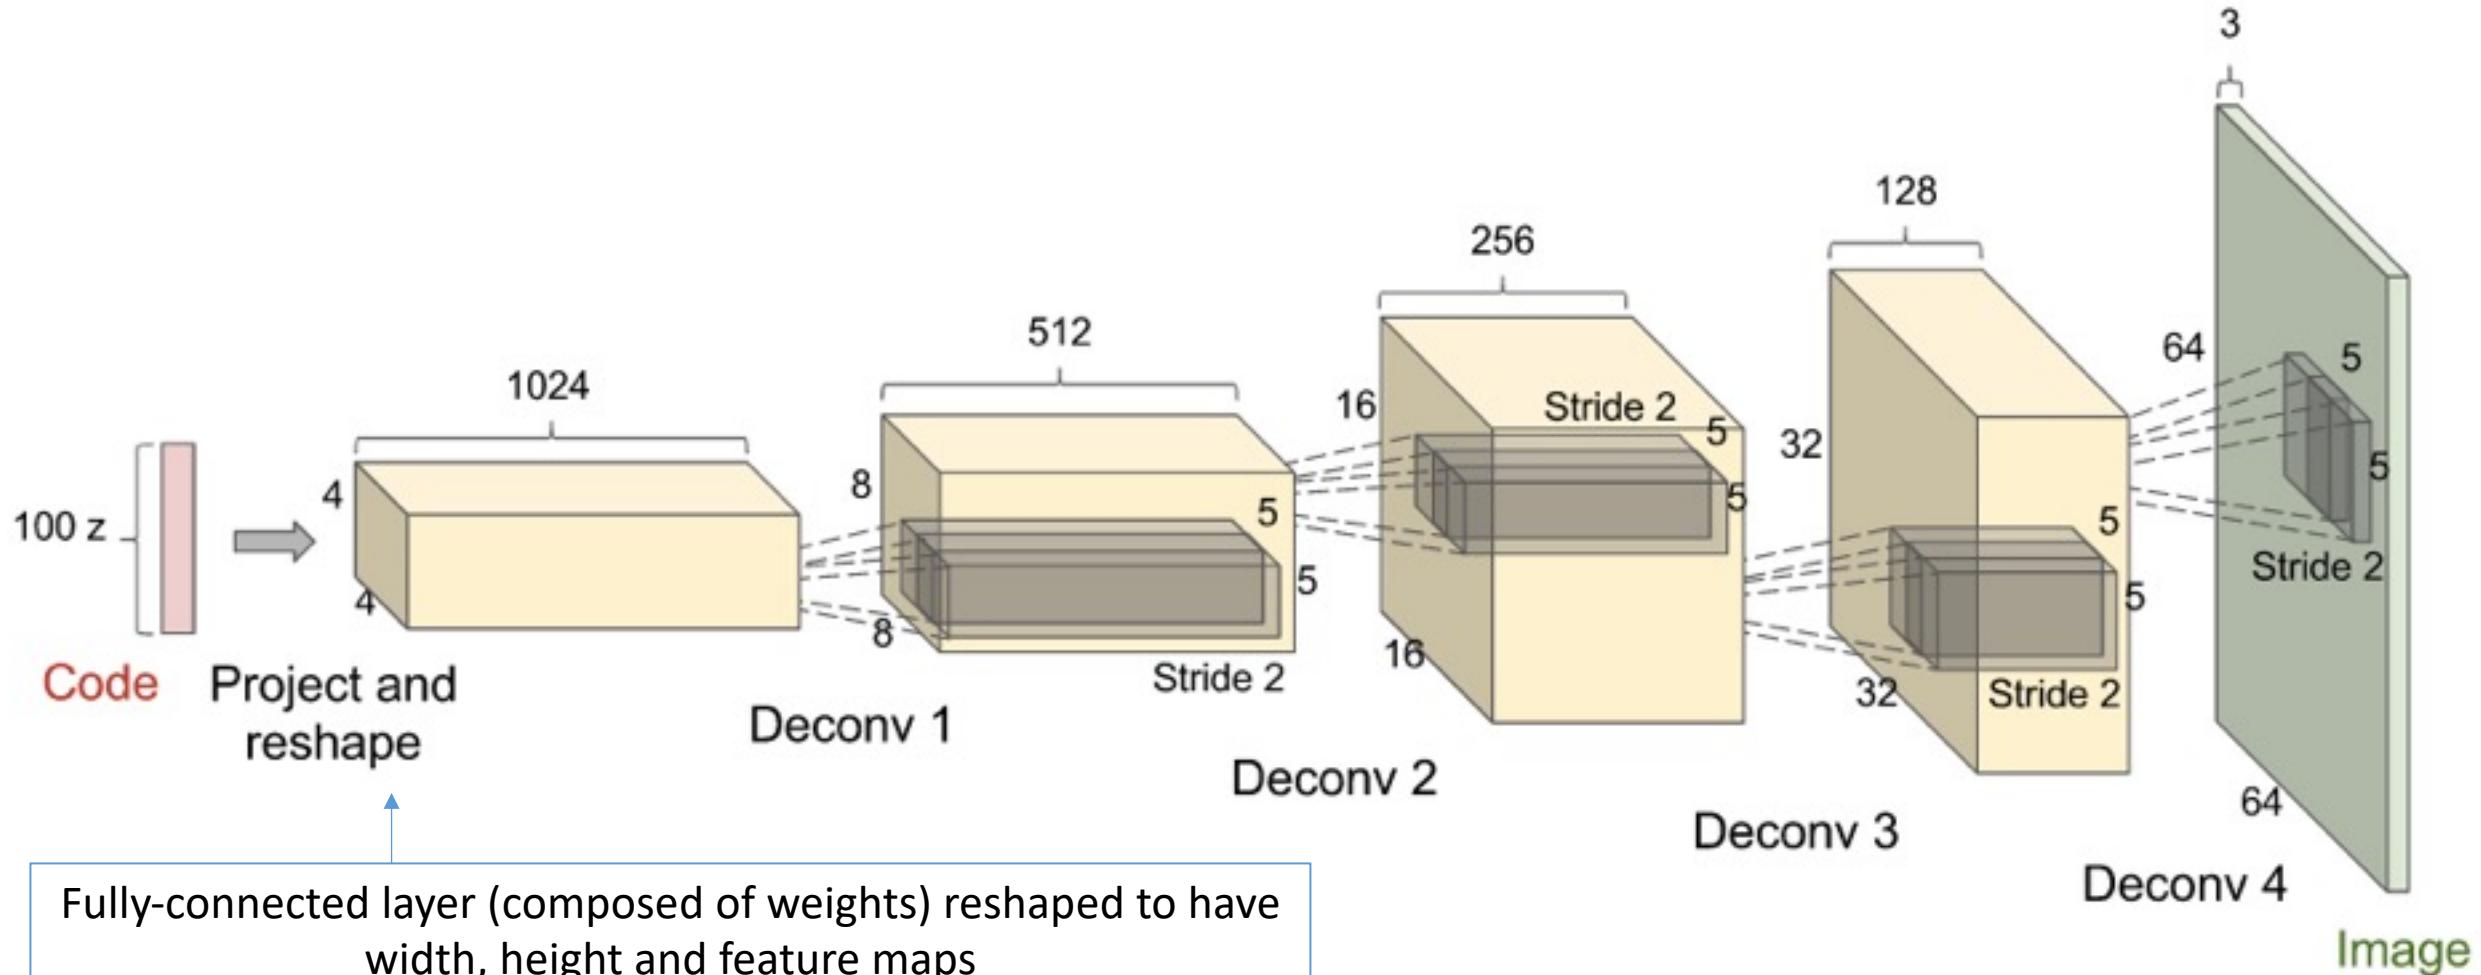


# Outline

- GAN Basics
- GAN Training
- DCGAN
- Mode Collapse

# DCGAN

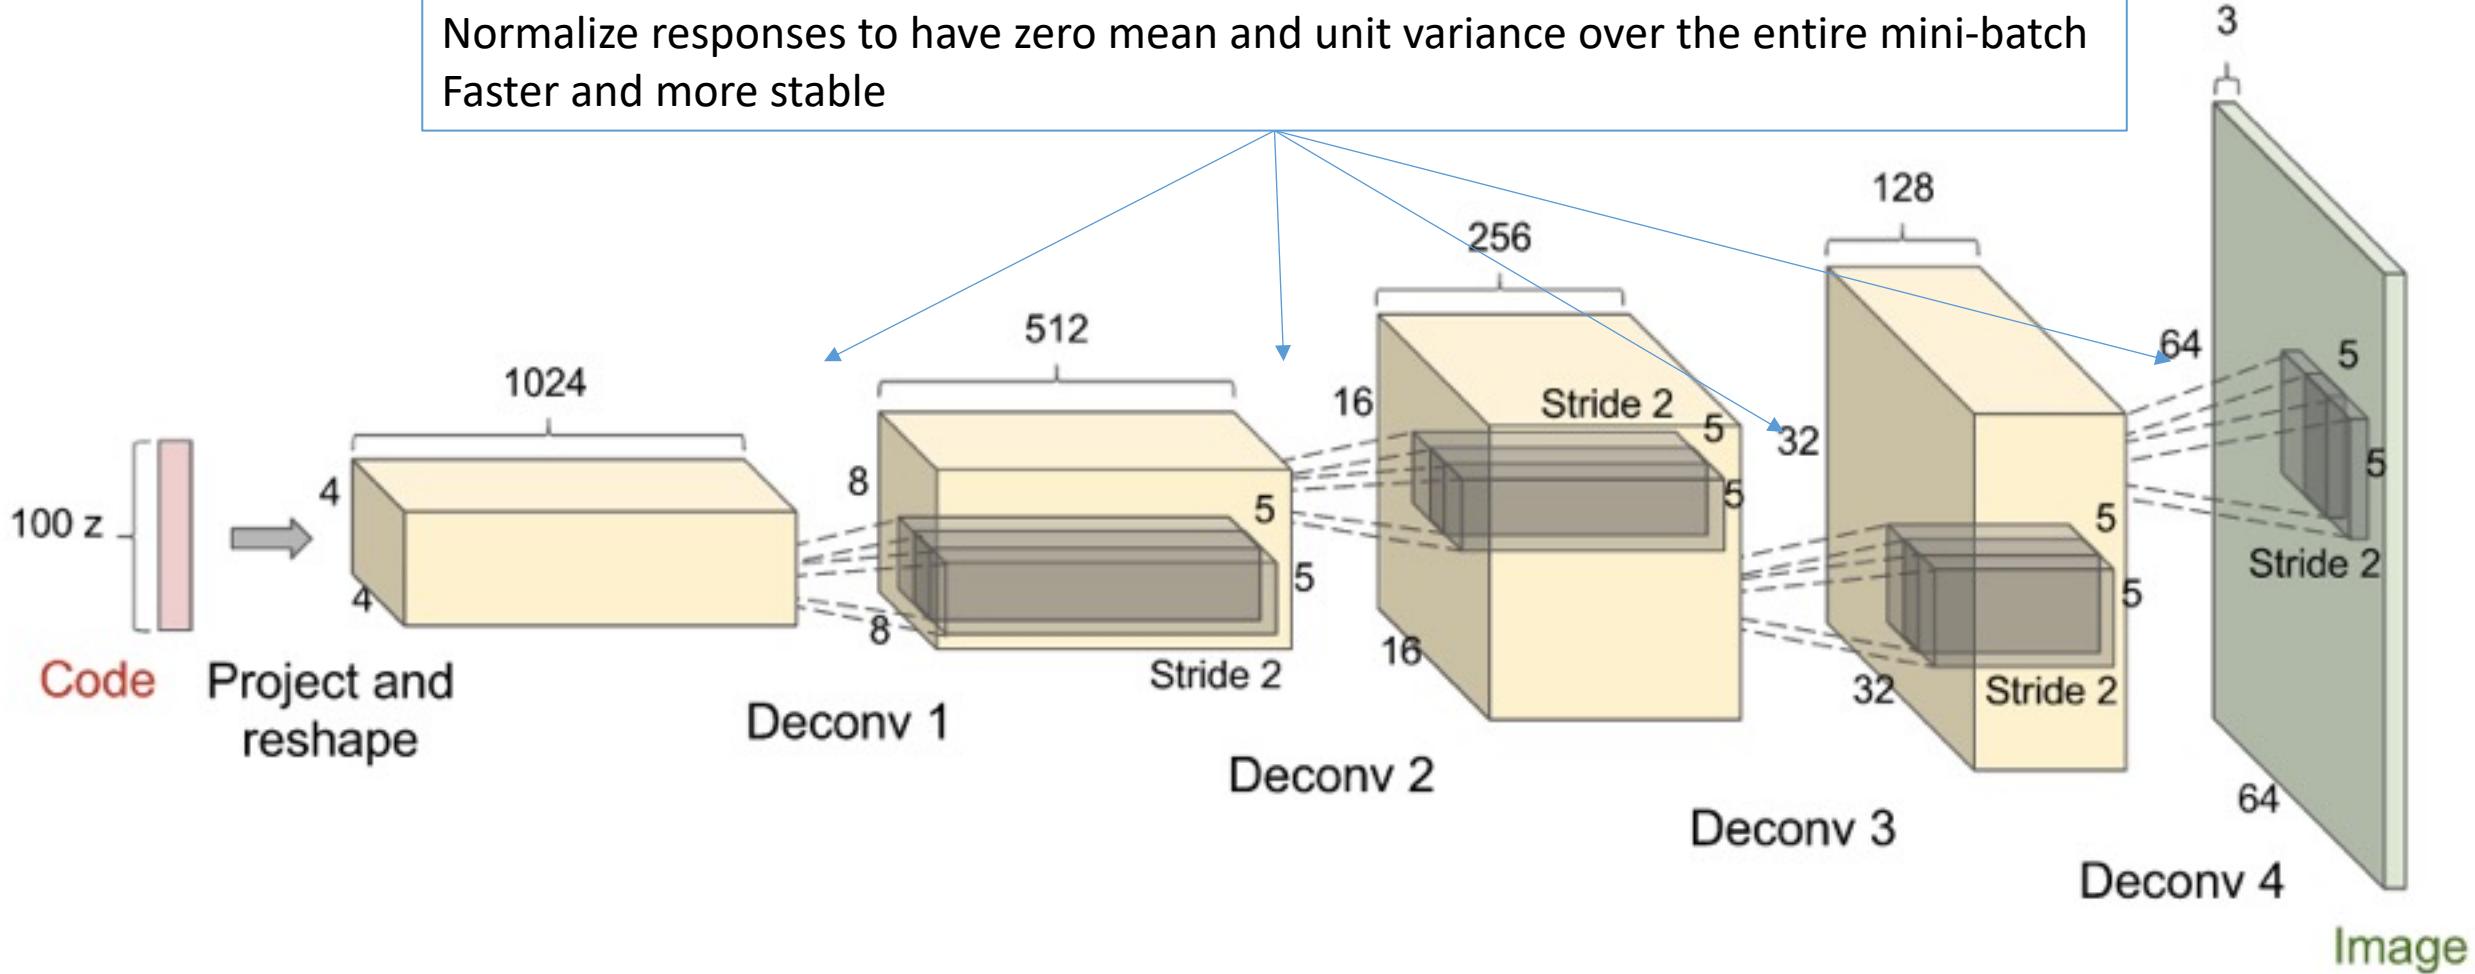
# Deep Convolutional GAN (DCGAN)



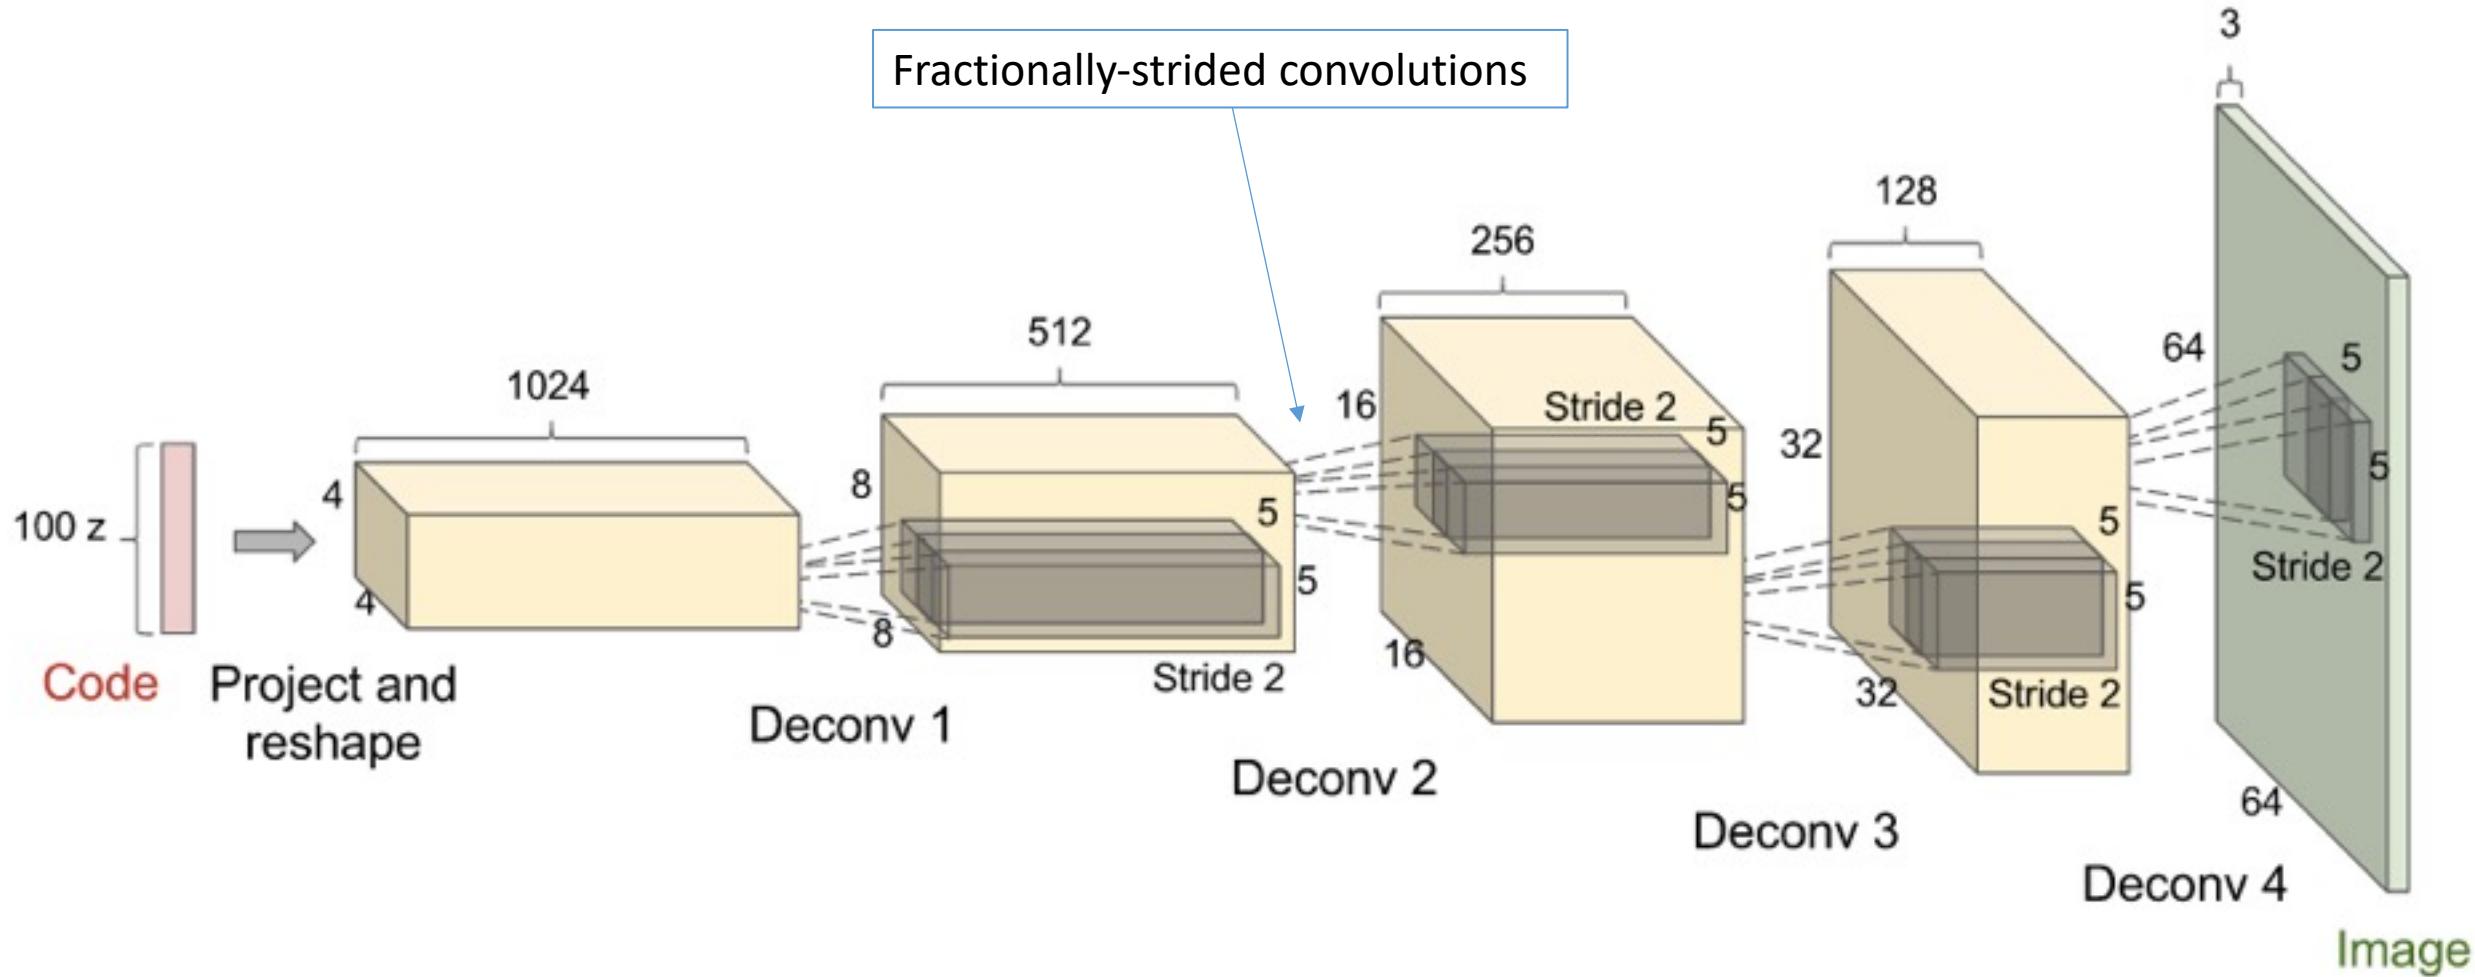
We will go through the code in the tutorial

# Deep Convolutional GAN (DCGAN)

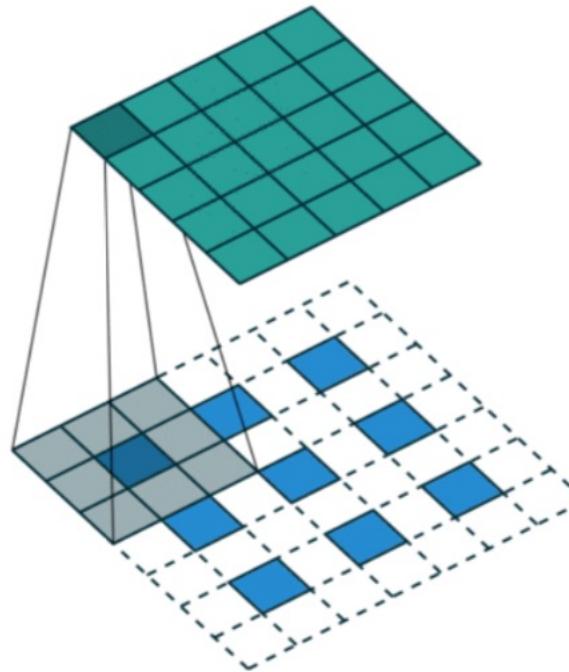
Most “deconv”s are batch normalized:  
Normalize responses to have zero mean and unit variance over the entire mini-batch  
Faster and more stable



# Deep Convolutional GAN (DCGAN)



# Fractionally-strided convolutions



Fractional strides involve: zeros are inserted *between* input units, which makes the kernel move around at a slower pace than with unit strides

<https://pytorch.org/docs/stable/generated/torch.nn.ConvTranspose2d.html>

Up-sampling by fractional striding

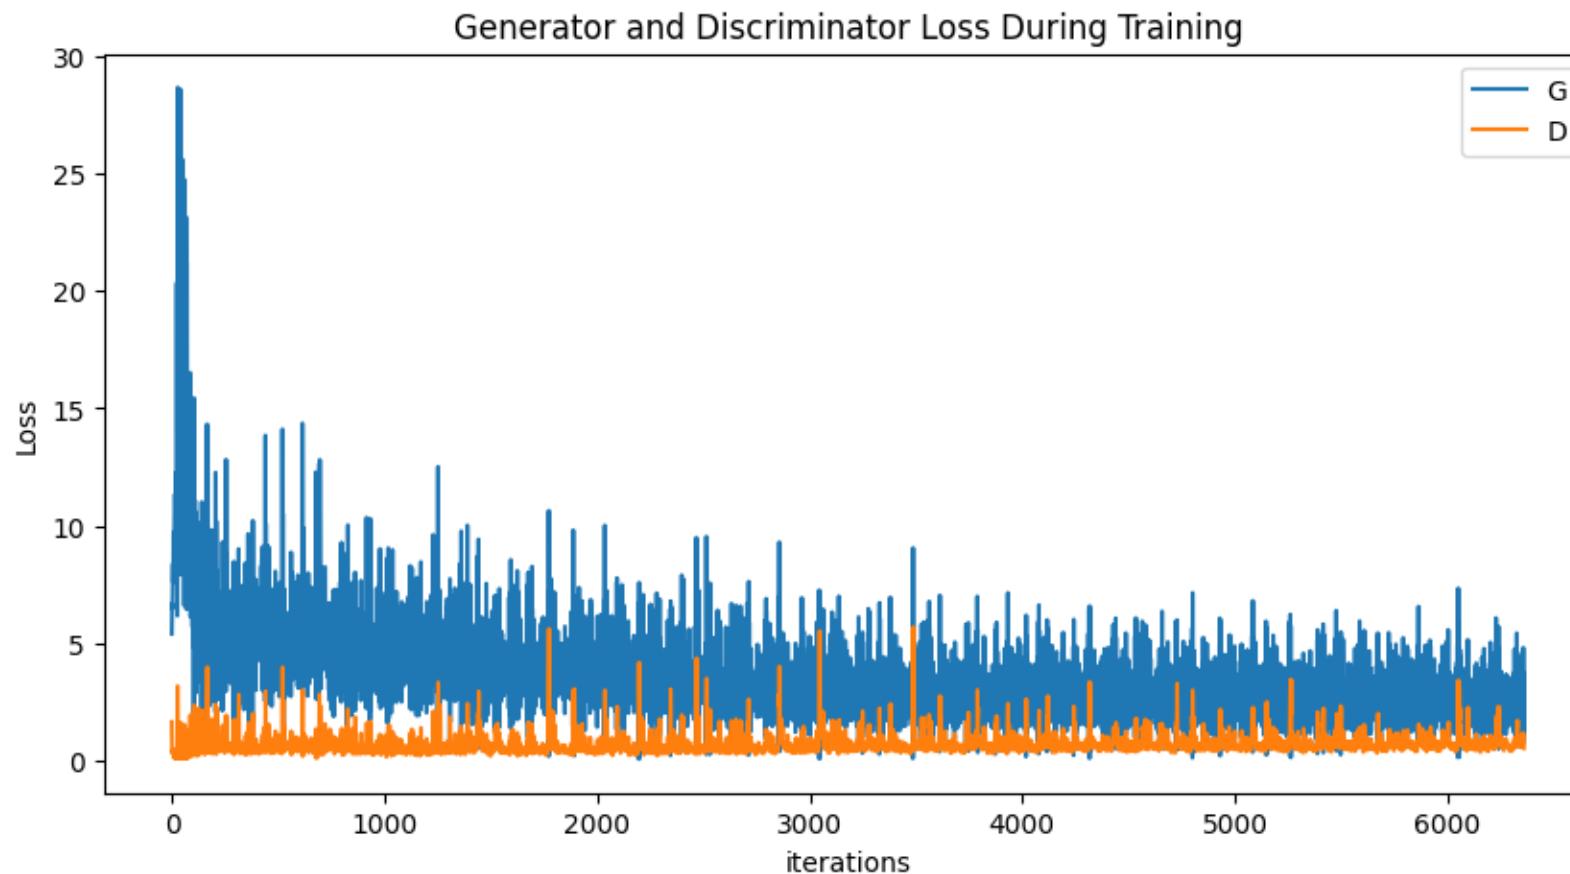
Input 3x3, output = 5x5

stride = 1 convolution window = 3 x 3 (with respect to output)

# Other tricks proposed by DCGAN

- Adam optimizer (adaptive moment estimation) = similar to SGD but with less parameter tuning
- Momentum = 0.5 (usually is 0.9 but training oscillated and was unstable)
- Low learning rate = 0.0002

# DCGAN

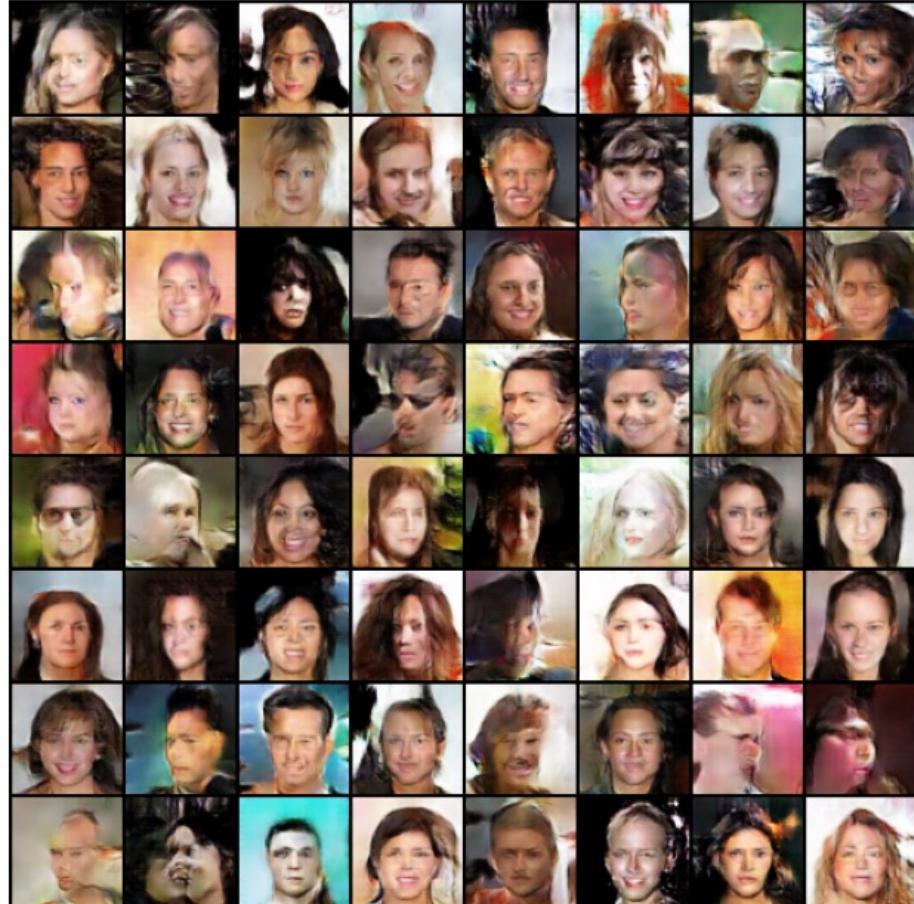


# DCGAN

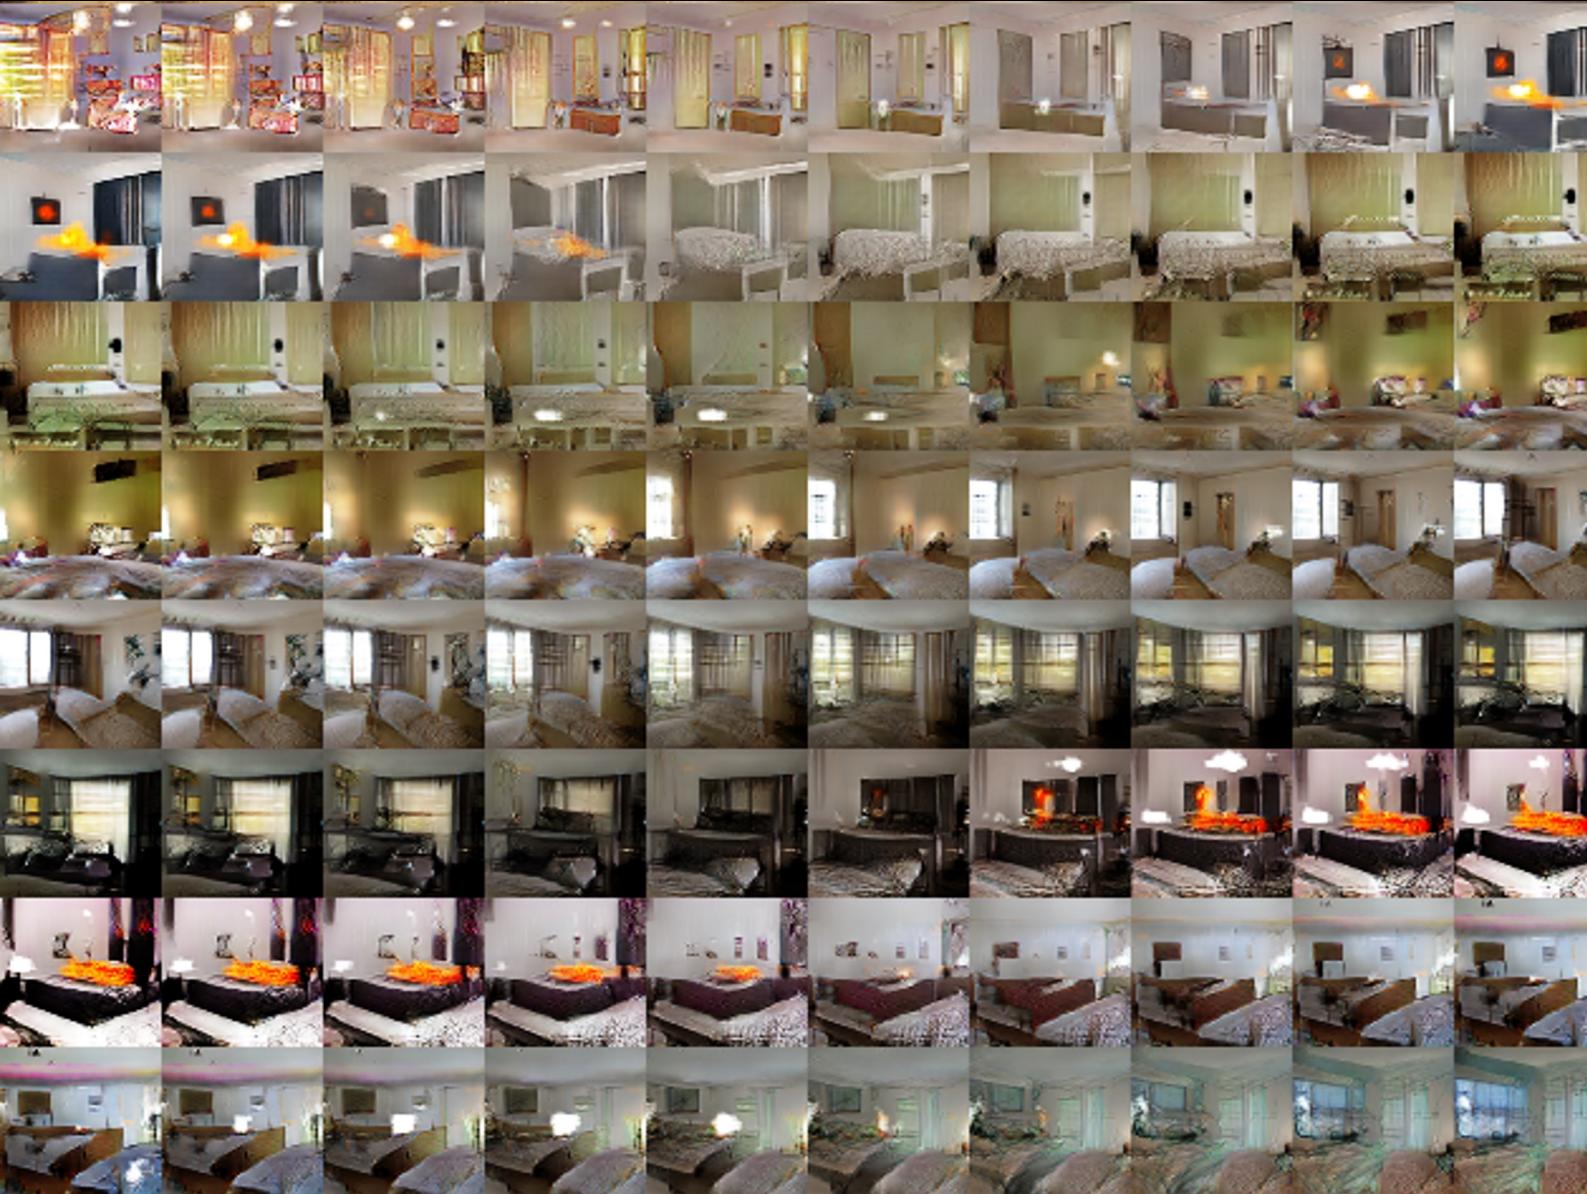
Real Images



Fake Images



# Sanity check by walking on the manifold



Interpolating between a series of random points in  $\mathbf{z}$

( $\mathbf{z}$  = the 100-d random numbers)

e.g.,

$$\mathbf{z}_{new} = m \mathbf{z}_1 + (1 - m) \mathbf{z}_2$$

Note the smooth transition between each scene

←the window slowly appearing

This indicates that the space learned has smooth transitions!  
(and is not simply memorizing)

# Sanity check by vector space arithmetic

$$\text{Man with glasses} - \text{Man} + \text{Woman} = \text{Woman with Glasses}$$


Man with glasses

Man

Woman

=

Woman with Glasses

# Outline

- GAN Basics
- GAN Training
- DCGAN
- Mode Collapse

# Mode Collapse

# Advantages of GANs

- **Plenty of existing work on Deep Generative Models**
  - Boltzmann Machine
  - Deep Belief Nets
  - Variational AutoEncoders (VAE)
- **Why GANs?**
  - Sampling (or generation) is straightforward.
  - Training doesn't involve Maximum Likelihood estimation (i.e. finding the best model parameters that fit the training data the most), believed to cause poorer quality of images (blurry images)
  - Robust to overfitting since the generator never sees the training data.

# Problems with GAN

- **Probability Distribution is Implicit**
  - Not straightforward to compute  $p(x)$
  - Thus **Vanilla GANs** are only good for Sampling/Generation.
- **Training is Hard**
  - Non-Convergence
  - Mode-Collapse

# Non-convergence

- Deep Learning models (in general) involve a single player
  - The player tries to maximize its reward (minimize its loss).
  - Use SGD (with Backpropagation) to find the optimal parameters.
  - SGD has convergence guarantees (under certain conditions).
  - **Problem:** With non-convexity, we might converge to local optima.

$$\min_G L(G)$$

- GANs instead involve two (or more) players
  - Discriminator is trying to maximize its reward.
  - Generator is trying to minimize Discriminator's reward.

$$\min_G \max_D V(D, G)$$

- SGD was not designed to find the Nash equilibrium of a game.
- **Problem:** We might not converge to the Nash equilibrium at all.

# Non-convergence example

$$\min_x \max_y V(x, y)$$

Let  $V(x, y) = xy$

- State 1:

x > 0	y > 0	V > 0
-------	-------	-------

Increase y	Decrease x
------------	------------

- State 2:

x < 0	y > 0	V < 0
-------	-------	-------

Decrease y	Decrease x
------------	------------

- State 3:

x < 0	y < 0	V > 0
-------	-------	-------

Decrease y	Increase x
------------	------------

- State 4 :

x > 0	y < 0	V < 0
-------	-------	-------

Increase y	Increase x
------------	------------

- State 5:

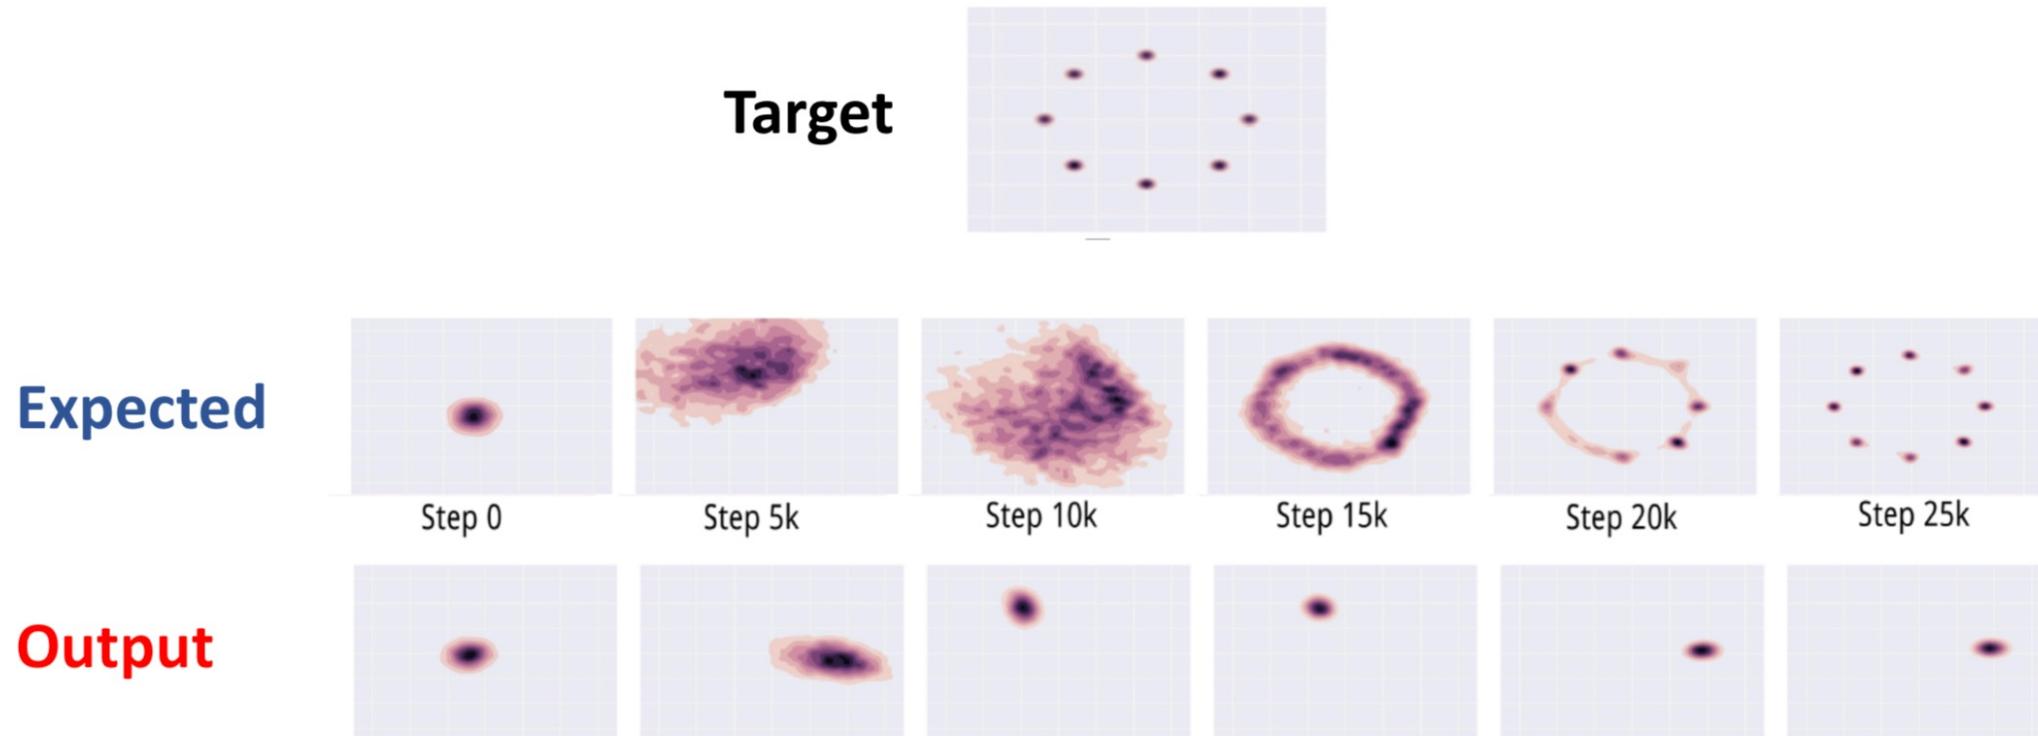
x > 0	y > 0	V > 0
-------	-------	-------

== State 1

Increase y	Decrease x
------------	------------

# Mode Collapse

Generator fails to output diverse samples



# How to reward sample diversity

- **At Mode Collapse,**

- Generator produces good samples, but a very few of them.
- Thus, Discriminator can't tag them as fake.

- **To address this problem,**

- Let the Discriminator know about this edge-case.

- **More formally,**

- Let the Discriminator look at the entire batch instead of single examples
- If there is lack of diversity, it will mark the examples as fake

- **Thus,**

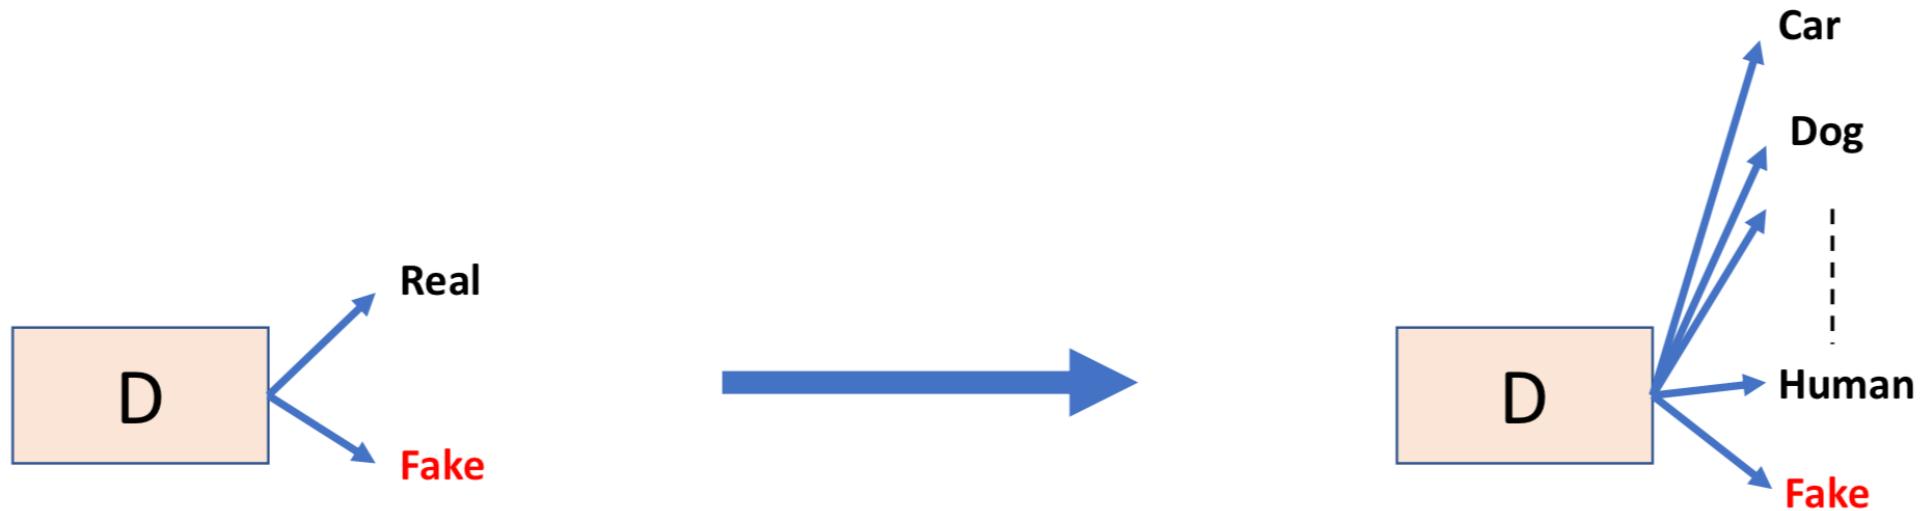
- Generator will be forced to produce diverse samples.

# Mini-Batch GANs

- **Extract features that capture diversity in the mini-batch**
  - For e.g. L2 norm of the difference between all pairs from the batch
  - That is, the errors at an intermediate layer for real and fake samples are minimized.
- **Feed those features to the discriminator along with the image**
- **Feature values will differ b/w diverse and non-diverse batches**
  - Thus, Discriminator will rely on those features for classification
- **This in turn,**
  - Will force the Generator to match those feature values with the real data
  - Will generate diverse batches

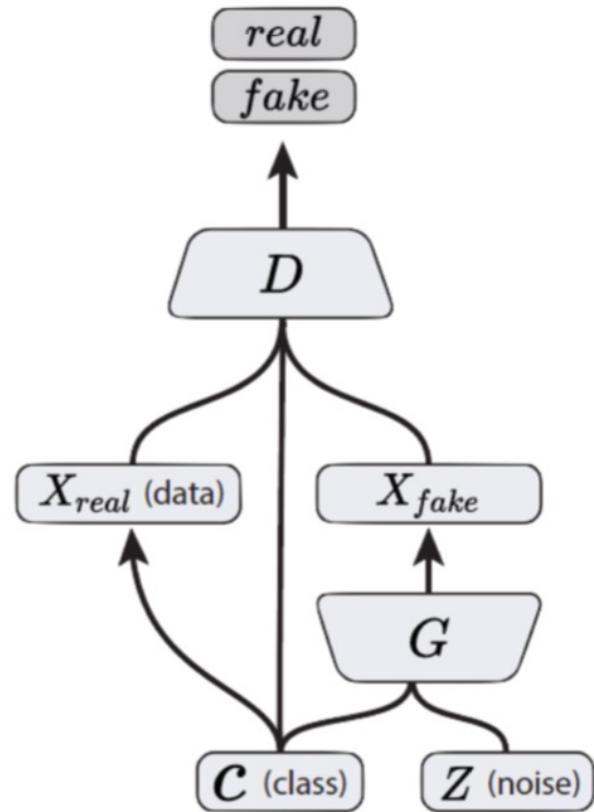
# Supervision with Labels

- Label information of the real data might help



- Empirically generates much better samples

# Conditional GANs



Conditions the inputs to the generator and discriminator with the labels

Conditional GAN  
(Mirza & Osindero, 2014)

# Conditional GANs

MNIST digits generated conditioned on their class label.

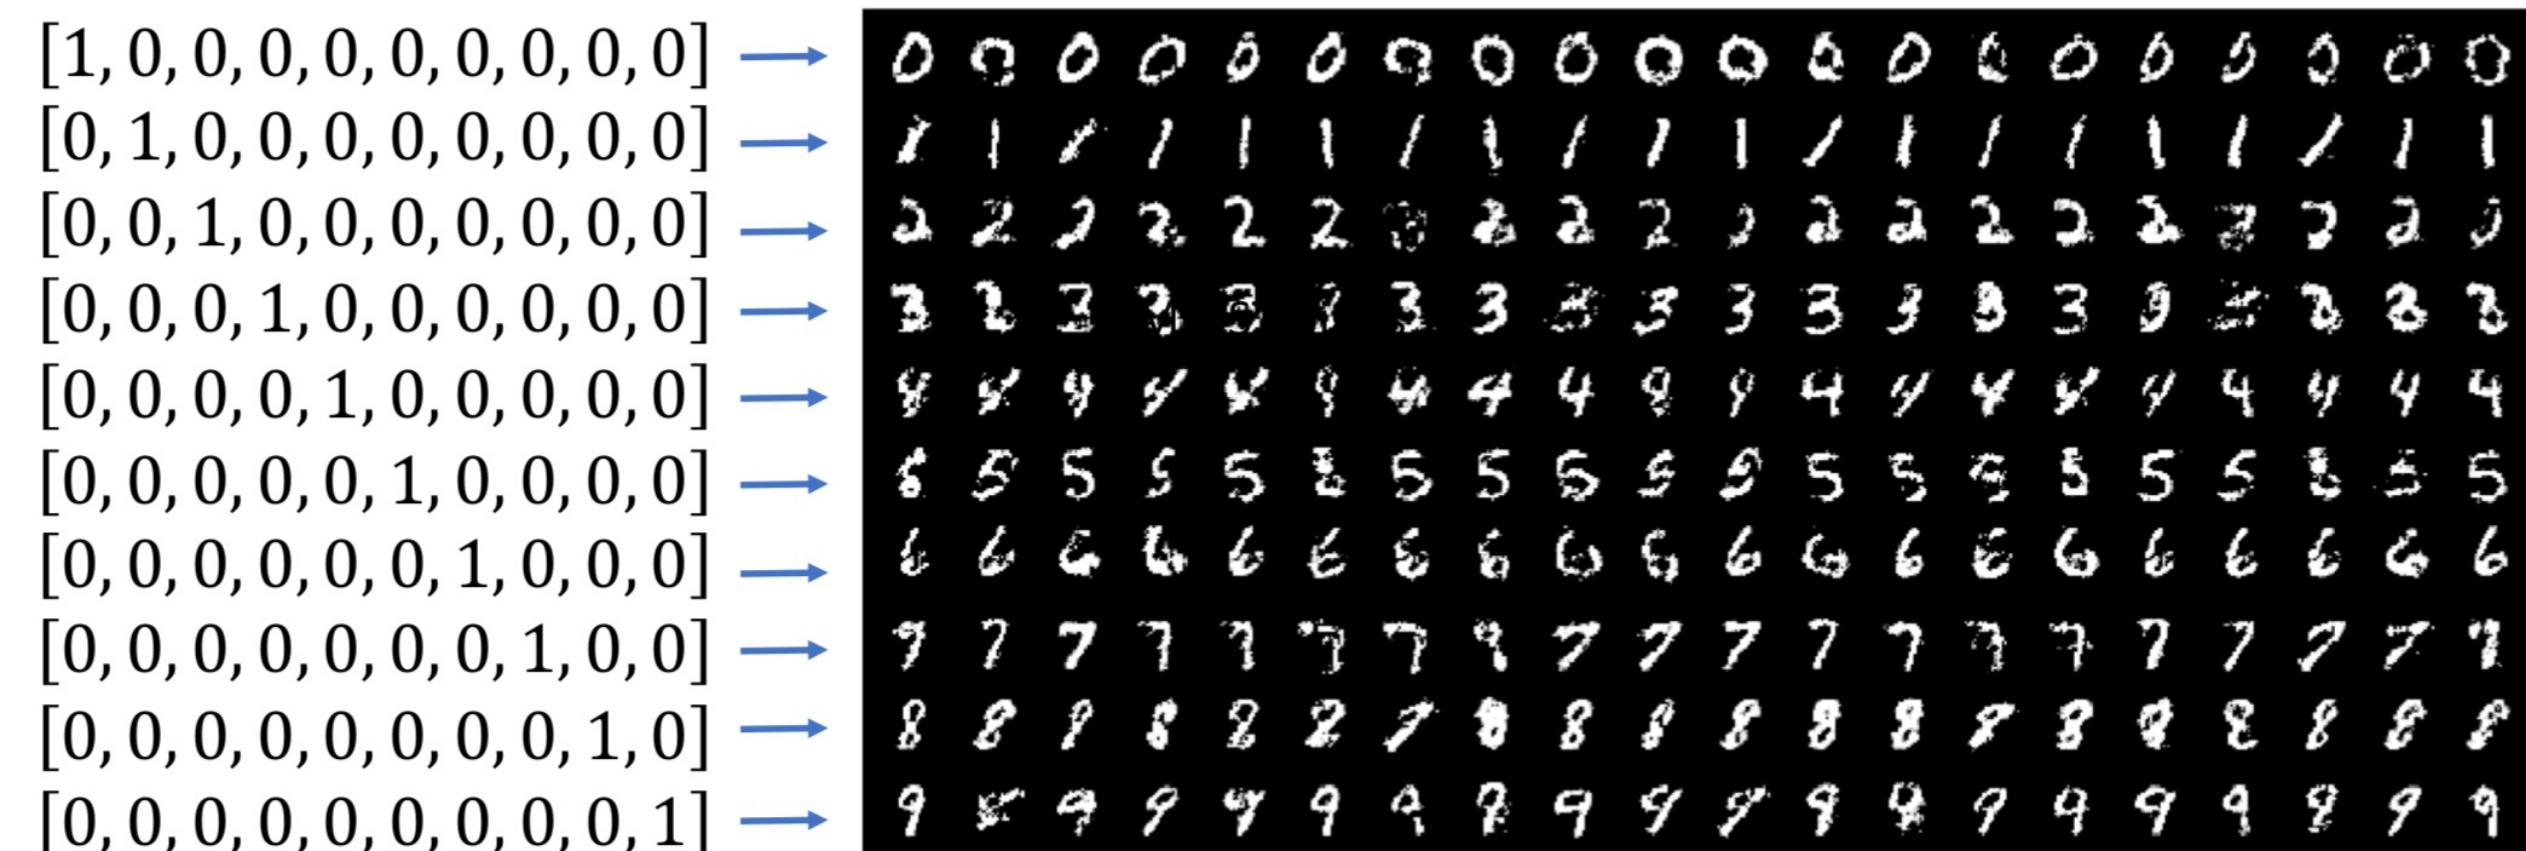
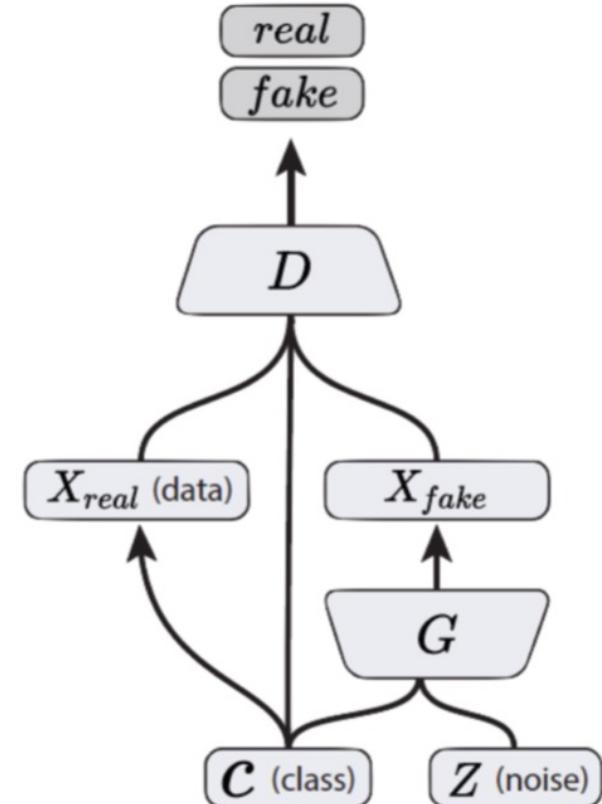


Figure 2 in the original paper.

# Conditional GANs

- Simple modification to the original GAN framework that conditions the model on *additional information* for better multi-modal learning.
- Lends to many practical applications of GANs when we have explicit *supervision* available.



Conditional GAN  
(Mirza & Osindero, 2014)

Image Credit: Figure 2 in Odena, A., Olah, C. and Shlens, J., 2016. Conditional image synthesis with auxiliary classifier GANs. *arXiv preprint arXiv:1610.09585*.