

SC4002 CE4045 CZ4045 **Natural Language Processing**

Regular Expressions

Dr. Sun Aixin



Chatbot

User: I am unhappy.

ELIZA: DO YOU THINK COMING HERE WILL HELP YOU NOT TO BE UNHAPPY

User: I need some help, that much seems certain.

ELIZA: WHAT WOULD IT MEAN TO YOU IF YOU GOT SOME HELP

User: Perhaps I could learn to get along with my mother.

ELIZA: TELL ME MORE ABOUT YOUR FAMILY

User: My mother takes care of me.

ELIZA: WHO ELSE IN YOUR FAMILY TAKES CARE OF YOU

User: My father.

ELIZA: YOUR FATHER

User: You are like my father in some ways.

Weizenbaum (1966)



Regular Expressions (RE)

- RE is a language for specifying text search strings.
 - Used in every computer language, word processor, and text processing tools
 - Formally, a RE is an **algebraic notation** for characterizing a **set** of strings.
- Particularly useful for searching in texts, when we have a **pattern** to search for and a **corpus** of texts to search through.
 - A RE search function will either return **the first match**, or search through the corpus and return **all texts that match the pattern**
 - The corpus can be a **single document** or a **collection**
 - Example: search for **postcode** in **5 documents**

The screenshot shows a web-based Regular Expression testing tool. At the top, there's an 'Expression' input field with the regex `/([A-Z])\w+/'g`. An orange arrow labeled 'Pattern' points to this field. Below it, there's a 'Text' input field with a sample paragraph about RegExr. An orange arrow labeled 'Corpus' points to this text area. The tool shows '29 matches (0.3ms)' for the given pattern and text. Below the text area, there are tabs for 'Tools' including 'Replace', 'List', 'Details', and 'Explain'. A tooltip is visible explaining the 'Capturing group #1' and the 'A-Z Range'.

Basic Regular Expression Patterns

- The simplest RE is a **sequence of simple characters**, like `/test/`
 - Putting characters in sequence is called **concatenation**.
 - `//` is not part of the RE, but the notation to indicate RE.
 - `/test/` will match test in text, but not match Test, or TEST.
 - Regular expressions are **case sensitive**.



RE	Example patterns matched
<code>/woodchucks/</code>	“interesting links to <u>woodchucks</u> and lemurs”
<code>/a/</code>	“M <u>a</u> ry Ann stopped by Mona’s”
<code>/!/</code>	“You’ve left the burglar behind again <u>!</u> ” said Nori

- In our examples, we generally underline the **first exact part** of the match.
 - If the word woodchucks appears to be the first word in a sentence, then *Woodchucks* will not be match.

Disjunction of characters to match

- The string of characters inside the square braces [] specifies a disjunction of characters to match
 - Square brackets [] matches **any single character** from within the class

RE	Match	Example Patterns
<code>/[wW]oodchuck/</code>	Woodchuck or woodchuck	“ <u>W</u> oodchuck”
<code>/[abc]/</code>	‘a’, ‘b’, or ‘c’	“In uomini, in soldati”
<code>/[1234567890]/</code>	any single digit	“plenty of <u>7</u> to 5”

- If there is a well-defined sequence associated with a set of characters, [] can be used with the dash (-) to specify any one character in a range.
 - `/[A-Z]/` matches an upper-case letter
 - `/[a-z]/` matches a lower-case letter
 - `/[0-9]/` matches a single digit

More on square braces in RE

- The square braces can also be used to specify what a single character cannot be matched with caret **^**.
- If the caret **^** is **the first symbol after the open square brace** **[**, the resulting pattern is negated.
 - Pattern **/[^a]/** matches any single character (including special characters) except **a**.
 - This is only true when the caret is the first symbol after the open square brace. If it occurs anywhere else, it usually stands for a caret

RE	Match (single characters)	Example Patterns Matched
/[^A-Z]/	not an upper-case letter	“O y fn pripetchik”
/[^Ss]/	neither ‘S’ nor ‘s’	“ I have no exquisite reason for’t”
/[^.]/	not a period	“ o ur resident Djinn”
/[e^]/	either ‘e’ or ‘^’	“look up ^ now”
/a^b/	the pattern ‘a^b’	“look up a^b now”



Counters in RE, and wildcard expression

- Question mark **?** matches zero or one appearance of the preceding item
- Kleene ***** (generally pronounced “cleany star”) matches zero or more occurrences of the immediately previous character or regular expression
- Kleene **+** matches one or more occurrences of the immediately preceding character or regular expression.
- Examples
 - **/colou?r/** matches color or colour
 - **/ba*!/** matches b!, ba!, baaa!, baaaaa! and more such strings
 - **/ba+!/** matches ba!, baa!, baaaaa! and more such strings
 - **/[ab]+/** means “one or more *a*’s or *b*’s” like aaaa, or ababbb or bbbb
- The period (**/./**) is a wildcard expression that matches any single character (except a carriage return), e.g., **/beg.n/** matches begin, beg’n, begun
 - The wildcard used together with the Kleene star **/.*/** means “any string of characters”.
 - When the period is used inside a negation **/[^.]/**, it means period and not the wildcard expression.



Anchors in RE

- **Anchors** are special characters that anchor regular expressions to particular places in a string
- The caret **^** matches the **start of a line**.
 - The pattern **/^The/** matches the word **The** only at the start of a line.
- The dollar sign **\$** matches the **end of a line**.
 - **/^The dog\.\$/** matches a line that contains only the phrase “**The dog.**”
 - Backslash dot **\.** is escaped character, means **the full stop or period, not the wildcard**.
- Two other anchors: word boundary **\b** and non-word boundary **\B**
 - A “word” for the purposes of a regular expression is defined as any sequence of **digits**, **underscores**, or **letters**, based on the definition of “words” in programming languages.
 - **/\bthe\b/** matches the word **the** but not the word **other**
 - **/\b99\b/** will match “There are **99** bottles” (because **99** follows a space) but not “There are **299** bottles” (since **99** follows a number). But it will match **\$99**, since **99** follows a dollar sign (not part of the word definition above).



Disjunction, Grouping, and Precedence

➤ The disjunction operator, also called the pipe symbol `|`, means “or”.

- The pattern `/cat|dog/` matches **either** the string cat **or** the string dog

➤ Parenthesis operators (and).

- Enclosing a pattern in parentheses makes it **act like a single character** for the purposes of neighboring operators like the pipe `|` and the Kleene*.
- `/gupp(y|ies)/` would specify that we meant the disjunction only to apply to the suffixes **y** and **ies**. The RE matches guppy or guppies
- `/(abb)+/` matches abb, abbabb, abbabbabb

➤ Operator precedence

- Parenthesis ()
- Counters * + ? { }
- Sequences and anchors the ^my end\$
- Disjunction |

REs always match the largest string they can!



More operators

RE	Expansion	Match	First Matches
\d	[0-9]	any digit	Party_of_5
\D	[^0-9]	any non-digit	Blue_moon
\w	[a-zA-Z0-9_]	any alphanumeric/underscore	Daiyu
\W	[^\w]	a non-alphanumeric	!!!!
\s	[\r\t\n\f]	whitespace (space, tab)	
\S	[^\s]	Non-whitespace	in_Concord

RE	Match
*	zero or more occurrences of the previous char or expression
+	one or more occurrences of the previous char or expression
?	exactly zero or one occurrence of the previous char or expression
{n}	<i>n</i> occurrences of the previous char or expression
{n,m}	from <i>n</i> to <i>m</i> occurrences of the previous char or expression
{n,}	at least <i>n</i> occurrences of the previous char or expression
{,m}	up to <i>m</i> occurrences of the previous char or expression

RE	Match	First Patterns Matched
*	an asterisk “*”	“K_A*P*L*A*N”
\.	a period “.”	“Dr. Livingston, I presume”
\?	a question mark	“Why don’t they come and lend a hand?”
\n	a newline	
\t	a tab	



A working example

- Suppose we wanted to write a RE to find the English article *the* in text.

Expression

`/the/g` Missing “The”, and matching “there”

Text Tests

The word the can be a substring of "there" or "the25", or even "the_example".

Expression

`/[tT]he/g` Matching both “the” and “The”, but still matching “there”

Text Tests

The word the can be a substring of "there" or "the25", or even "the_example".



A working example

- Suppose we wanted to write a RE to find the English article *the* in text.

Expression

```
/\b[tT]he\b/g
```

Added word boundary, missing two other cases

Text Tests

The word the can be a substring of "there" or "the25", or even "the_example".

Expression

```
/^[^a-zA-Z][tT]he[^a-zA-Z]/g
```

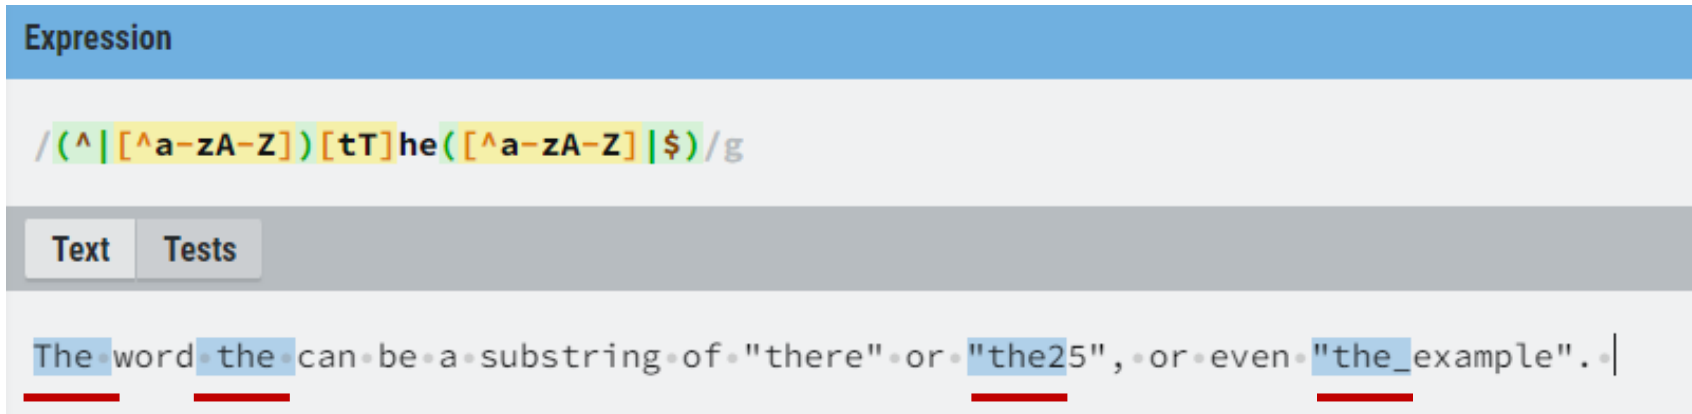
No other alphabetic letters on either side of "the"

Text Tests

The word the can be a substring of "there" or "the25", or even "the_example".

A working example

➤ Suppose we wanted to write a RE to find the English article *the* in text.



- A more complicated (but not perfect) version
- Match “the” or “The” either in text or at the beginning of the line, or at the end of line.
- There is no alphabetic letters on either side of the word “the”.
- The final version can be application dependent.

False Positive and False Negative

- The whole process is to fix two kinds of errors:
 - **False positives:** strings that we incorrectly matched like **other** or **there**,
 - **False negatives:** strings that we incorrectly missed, like **The**.
- Addressing these two kinds of errors is common in implementing language processing systems.
- Reducing the overall error rate for an application thus involves
 - Increasing **precision** (minimizing false positives)
 - Increasing **recall** (minimizing false negatives)



Summary

- RegExr: an online tool to learn, build, & test Regular Expressions
 - <http://regexr.com/>
- Java RegEx API and Tutorial
 - <http://docs.oracle.com/javase/8/docs/api/java/util/regex/package-summary.html>
 - <http://docs.oracle.com/javase/tutorial/essential/regex/>
- Reference: <https://web.stanford.edu/~jurafsky/slp3/>
 - **Chapter 2**, Regular Expressions, Text Normalization, Edit Distance

What can we do?

- Given a document, we are able to search for the matching strings with a query specified in Regular Expression
 - The given document is basically a sequence of characters
 - At this stage, we do not understand words or sentences in the document.
- Next, it would be useful to recognize the words, sentences in the document
 - With the words and sentences, we will be able to understand the structure or meaning of the sentences.

