



# **Natural Language Processing**

## **Tutorial 2: Text Normalization**



Dr. Sun Aixin



# Question I

---

- Consider the following word segmentation algorithm in the lecture notes:
  
- Given a lexicon of Chinese, and a string
  - Start a pointer at the beginning of the string
  - Find the longest word in dictionary that matches the string starting at pointer
  - Move the pointer over the word in string
  - Goto2
  
- Following the algorithm, you perhaps end up with failing to segment a string, if you cannot find a matching.

# Question 1 (cont)

---

## ➤ Example

- String to segment: the tables down there
- lexicon: the table down there bled own.

## ➤ Discuss how to fix the above problem.



# Answer Q1

---

- Start a pointer at the beginning of the string
- Find the longest word in dictionary that matches the string starting at pointer
  - If matched, move the pointer over the word in string
  - If no word is matched, skip to the next letter
- Goto2
  - String to segment: the tables down there
  - lexicon: the table down there bled own.



# Answer 1 (a bit more processing)

---

- Start a pointer at the beginning of the string
- Find the longest word in dictionary that matches the string starting at pointer
  - If matched,
    - Run morphological analysis from the beginning of the word
    - Move the pointer over the morphologically matched part of the string
  - If no word is matched, skip to the next letter
- Goto2
  - String to segment: the tables down there
  - lexicon: the table down there bled own.



## Question 2

---

- Try the tokenization demo on
  - <http://text-processing.com/demo/tokenize/>
  - (you may use other tokenizer APIs).
- Discuss your findings based on the output of different tokenizers.
- <https://textanalysisonline.com/>



# Question 3

---

- Try the stemmer demo on
  - <http://text-processing.com/demo/stem/>
  - (or you may use other stemmer APIs).
- Discuss your findings based on the output of the stemmers.
- <https://textanalysisonline.com/>



# Question 4

---

➤ Write a program to do the following tasks:

- **Download** the Web page of a given link and **extract** the text content of the page
- **Split** the text into sentences and **count** sentences
- **Split** the text into tokens and **count** token types
- **Find** lemmas (or stems) of the tokens and **count** lemma types
- Do **stemming** on the tokens and **count** unique stemmed tokens





# Answer 4

---

## ➤ Example URL:

- [https://en.wikipedia.org/wiki/Natural\\_language\\_processing](https://en.wikipedia.org/wiki/Natural_language_processing)

## ➤ urllib

- <https://docs.python.org/3/library/urllib.html>

## ➤ NLTK

- <https://www.nltk.org/>

## ➤ Beautiful Soup

- <https://www.crummy.com/software/BeautifulSoup/bs4/doc/#porting-code-to-bs4>

### **Alternative libraries:**

StanfordNLP

OpenNLP

spaCy

AllenNLP



# Sample code

```
import urllib.request
import nltk
from bs4 import BeautifulSoup

with urllib.request.urlopen('https://en.wikipedia.org/wiki/Natural_language_processing') as response:
    html=response.read()

text = BeautifulSoup(html, "lxml").get_text()
sentences = nltk.tokenize.sent_tokenize(text)
print ('Number of sentences: '+ str(len(sentences)))

tokens= nltk.tokenize.word_tokenize(text)

print ('Number of tokens: '+ str(len(tokens)))

token_types = list(set(tokens))

print ('Number of token types: '+ str(len(token_types)))

wnl=nltk.stem.WordNetLemmatizer()

stemmer = nltk.stem.porter.PorterStemmer()
lemma_types= set()
stemmed_types= set()

for token_type in token_types:
    lemma_types.add(wnl.lemmatize(token_type))
    stemmed_types.add(stemmer.stem(token_type))

print ('Number of lemma types: '+ str(len(lemma_types)))
print ('Number of stemmed types: '+ str(len(stemmed_types)))
```

**Download** the Web page of a given link and **extract** the text content of the page

**Split** the text into sentences and **count** sentences

**Split** the text into tokens and **count** token types

**Find** lemmas (or stems) of the tokens and **count** lemma types

Do **stemming** on the tokens and **count** unique stemmed tokens



# Sample code

➤ `text = BeautifulSoup(html, "lxml").get_text()`

## Beautiful Soup Documentation

**Beautiful Soup** is a Python library for pulling data out of HTML and XML files. It works with your favorite parser to provide idiomatic ways of navigating, searching, and modifying the parse tree. It commonly saves programmers hours or days of work.

These instructions illustrate all major features of Beautiful Soup 4, with examples. I show you what the library is good for, how it works, how to use it, how to make it do what you want, and what to do when it violates your expectations.

This document covers Beautiful Soup version 4.8.1. The examples in this documentation should work the same way in Python 2.7 and Python 3.2.



<https://beautiful-soup-4.readthedocs.io/en/latest/>



# Sample code

<https://www.nltk.org/api/nltk.tokenize.html>

- `sentences = nltk.tokenize.sent_tokenize(text)`
- `tokens = nltk.tokenize.word_tokenize(text)`

```
nlk.tokenize
nlk.tokenize.api
nlk.tokenize.casual
nlk.tokenize.destructive
nlk.tokenize.legality_principle
nlk.tokenize.mwe
nlk.tokenize.nist
nlk.tokenize.punkt
nlk.tokenize.regexp
nlk.tokenize.repp
nlk.tokenize.sexpr
nlk.tokenize.simple
nlk.tokenize.sonority_sequencing
nlk.tokenize.stanford
nlk.tokenize.stanford_segmenter
nlk.tokenize.texttiling
nlk.tokenize.toktok
nlk.tokenize.treebank
nlk.tokenize.util
```

## nlk.tokenize package

### NLTK Tokenizer Package

Tokenizers divide strings into lists of substrings. For example, tokenizers can be used to find the words and punctuation in a string:

```
>>> from nltk.tokenize import word_tokenize
>>> s = '''Good muffins cost $3.88\nin New York. Please buy me
... two of them.\n\nThanks.'''
>>> word_tokenize(s)
['Good', 'muffins', 'cost', '$', '3.88', 'in', 'New', 'York', '.',
'Please', 'buy', 'me', 'two', 'of', 'them', '.', 'Thanks', '.']
```

This particular tokenizer requires the Punkt sentence tokenization models to be installed. NLTK also provides a simpler, regular-expression based tokenizer, which splits text on whitespace and punctuation:

```
>>> from nltk.tokenize import wordpunct_tokenize
>>> wordpunct_tokenize(s)
['Good', 'muffins', 'cost', '$', '3', '.', '88', 'in', 'New', 'York', '.',
'Please', 'buy', 'me', 'two', 'of', 'them', '.', 'Thanks', '.']
```

We can also operate at the level of sentences, using the sentence tokenizer directly as follows:

```
>>> from nltk.tokenize import sent_tokenize, word_tokenize
>>> sent_tokenize(s)
['Good muffins cost $3.88\nin New York.', 'Please buy me\ntwo of them.', 'Thanks.']
>>> [word_tokenize(t) for t in sent_tokenize(s)]
[['Good', 'muffins', 'cost', '$', '3.88', 'in', 'New', 'York', '.'],
 ['Please', 'buy', 'me', 'two', 'of', 'them', '.'], ['Thanks', '.']]
```



# Sample code

<https://www.nltk.org/api/nltk.stem.html>

<https://www.nltk.org/api/nltk.stem.wordnet.html>

<https://web.stanford.edu/~jurafsky/slp3/>

- `wnl=nltk.stem.WordNetLemmatizer()`
- `stemmer = nltk.stem.porter.PorterStemmer()`

## nltk.stem package

### NLTK Stemmers

Interfaces used to remove morphological affixes from words, leaving only the word stem. Stemming algorithms aim to remove those affixes required for eg. grammatical role, tense, derivational morphology leaving only the stem of the word. This is a difficult problem due to irregular words (eg. common verbs in English), complicated morphological rules, and part-of-speech and sense ambiguities (eg. `ceil-` is not the stem of `ceiling`).

StemmerI defines a standard interface for stemmers.

### Submodules

- `nltk.stem.api` module
- `nltk.stem.arlstem` module
- `nltk.stem.arlstem2` module
- `nltk.stem.cistem` module
- `nltk.stem.isri` module
- `nltk.stem.lancaster` module
- `nltk.stem.porter` module
- `nltk.stem.regexp` module
- `nltk.stem.rslp` module
- `nltk.stem.snowball` module
- `nltk.stem.util` module
- `nltk.stem.wordnet` module

## nltk.stem.wordnet module

`class nltk.stem.wordnet.WordNetLemmatizer`

[\[source\]](#)

Bases: `object`

WordNet Lemmatizer

Lemmatize using WordNet's built-in morphy function. Returns the input word unchanged if it cannot be found in WordNet.

```
>>> from nltk.stem import WordNetLemmatizer
>>> wnl = WordNetLemmatizer()
>>> print(wnl.lemmatize('dogs'))
dog
>>> print(wnl.lemmatize('churches'))
church
>>> print(wnl.lemmatize('aardwolves'))
aardwolf
>>> print(wnl.lemmatize('abaci'))
abacus
>>> print(wnl.lemmatize('hardrock'))
hardrock
```



## Q5: Open-ended Question, for **discussion** only.

---

- In social media (e.g., forums), online users often use informal names or references when mentioning products. Below are example sentences discussing mobile phones, where the words highlighted in bold are the phones being referred to, and the [bracketed text] indicates their official product names.

- 
1. True, **Desire** [HTC Desire] might be better if compared to **X10** [Sony Ericsson Xperia X10] but since I am using **HD2** [HTC HD2], it will be a little boring to use back HTC ...
  2. I just wanna know what problems do users face on the **OneX** [HTC One X]... of course I know that knowing the problems on **one x** [HTC One X] doesn't mean knowing the problems on **s3** [Samsung Galaxy SIII]
  3. Still prefer **ip 5** [Apple iPhone 5] then **note 2** [Samsung Galaxy Note II]...
  4. oh, the mono rich recording at **920** [Nokia Lumia 920] no better than stereo rich recording at **808** [Nokia 808 PureView].
- 



## Q5: Open-ended Question, for **discussion** only.

- The table below shows the number of users who have mentioned a phone using a specific name in a forum.

Name variation	#users	Name variation	#users
1. galaxy s3	553	14. lte s3	46
2. s3 lte	343	15. galaxy s3 lte	45
3. samsung galaxy s3	284	16. s3 non lte	32
4. s iii	242	17. samsung galaxy siiii	32
5. galaxy s iii	225	18. sgs 3	27
6. samsung s3	219	19. samsung galaxy s3 lte	22
7. sgs3	187	20. sg3	21
8. siiii	149	21. gsiiii	16
9. samsung galaxy s iii	145	22. samsung galaxy s3 i9300	15
10. i9300	120	23. samsung i9300 galaxy s iii	13
11. gs3	82	24. s3 4g	11
12. galaxy siiii	61	25. 3g s3	11
13. i9305	52	—	



## Q5: Open-ended Question, for **discussion** only.

- Task: Assume that we have successfully identified the phone mentions (e.g., 's3 lte,' 'sgs3'). How can we **normalize these mentions to their formal names**?

1. True, **Desire** [HTC Desire] might be better if compared to **X10** [Sony Ericsson Xperia X10] but since I am using **HD2** [HTC HD2], it will be a little boring to use back HTC ...
2. I just wanna know what problems do users face on the **OneX** [HTC One X]... of course I know that knowing the problems on **one x** [HTC One X] doesn't mean knowing the problems on **s3** [Samsung Galaxy SIII]
3. Still prefer **ip 5** [Apple iPhone 5] then **note 2** [Samsung Galaxy Note II]...
4. oh, the mono rich recording at **920** [Nokia Lumia 920] no better than stereo rich recording at **808** [Nokia 808 PureView].

Name variation	#users	Name variation	#users
1. galaxy s3	553	14. lte s3	46
2. s3 lte	343	15. galaxy s3 lte	45
3. samsung galaxy s3	284	16. s3 non lte	32
4. s iii	242	17. samsung galaxy siiii	32
5. galaxy s iii	225	18. sgs 3	27
6. samsung s3	219	19. samsung galaxy s3 lte	22
7. sgs3	187	20. sg3	21
8. siiii	149	21. gsiiii	16
9. samsung galaxy s iii	145	22. samsung galaxy s3 i9300	15
10. i9300	120	23. samsung i9300 galaxy s iii	13
11. gs3	82	24. s3 4g	11
12. galaxy siiii	61	25. 3g s3	11
13. i9305	52	–	

