

A complex network diagram with nodes and edges. Nodes are represented by circles in dark blue, red, and grey. Edges are thin lines connecting the nodes, with red lines forming a dense web and grey lines forming a more sparse structure. The background is a light blue-grey gradient.

# **BIG DATA MANAGEMENT**

**CE/CZ4123**

# **OVERVIEW (1<sup>ST</sup> HALF)**


**Siqiang Luo**

**Assistant Professor**

# THIS LECTURE

- ❑ Quickly go through key knowledge we have learnt, including
  - ❑ Big Data 5V's
  - ❑ Data Models
  - ❑ Memory Hierarchy
  - ❑ Column Store

# BIG DATA 5V'S



**Volume:** Large amount of data

**Velocity:** Fast data generation

**Variety:** Various data types/sources

**Veracity:** accurate and truthful

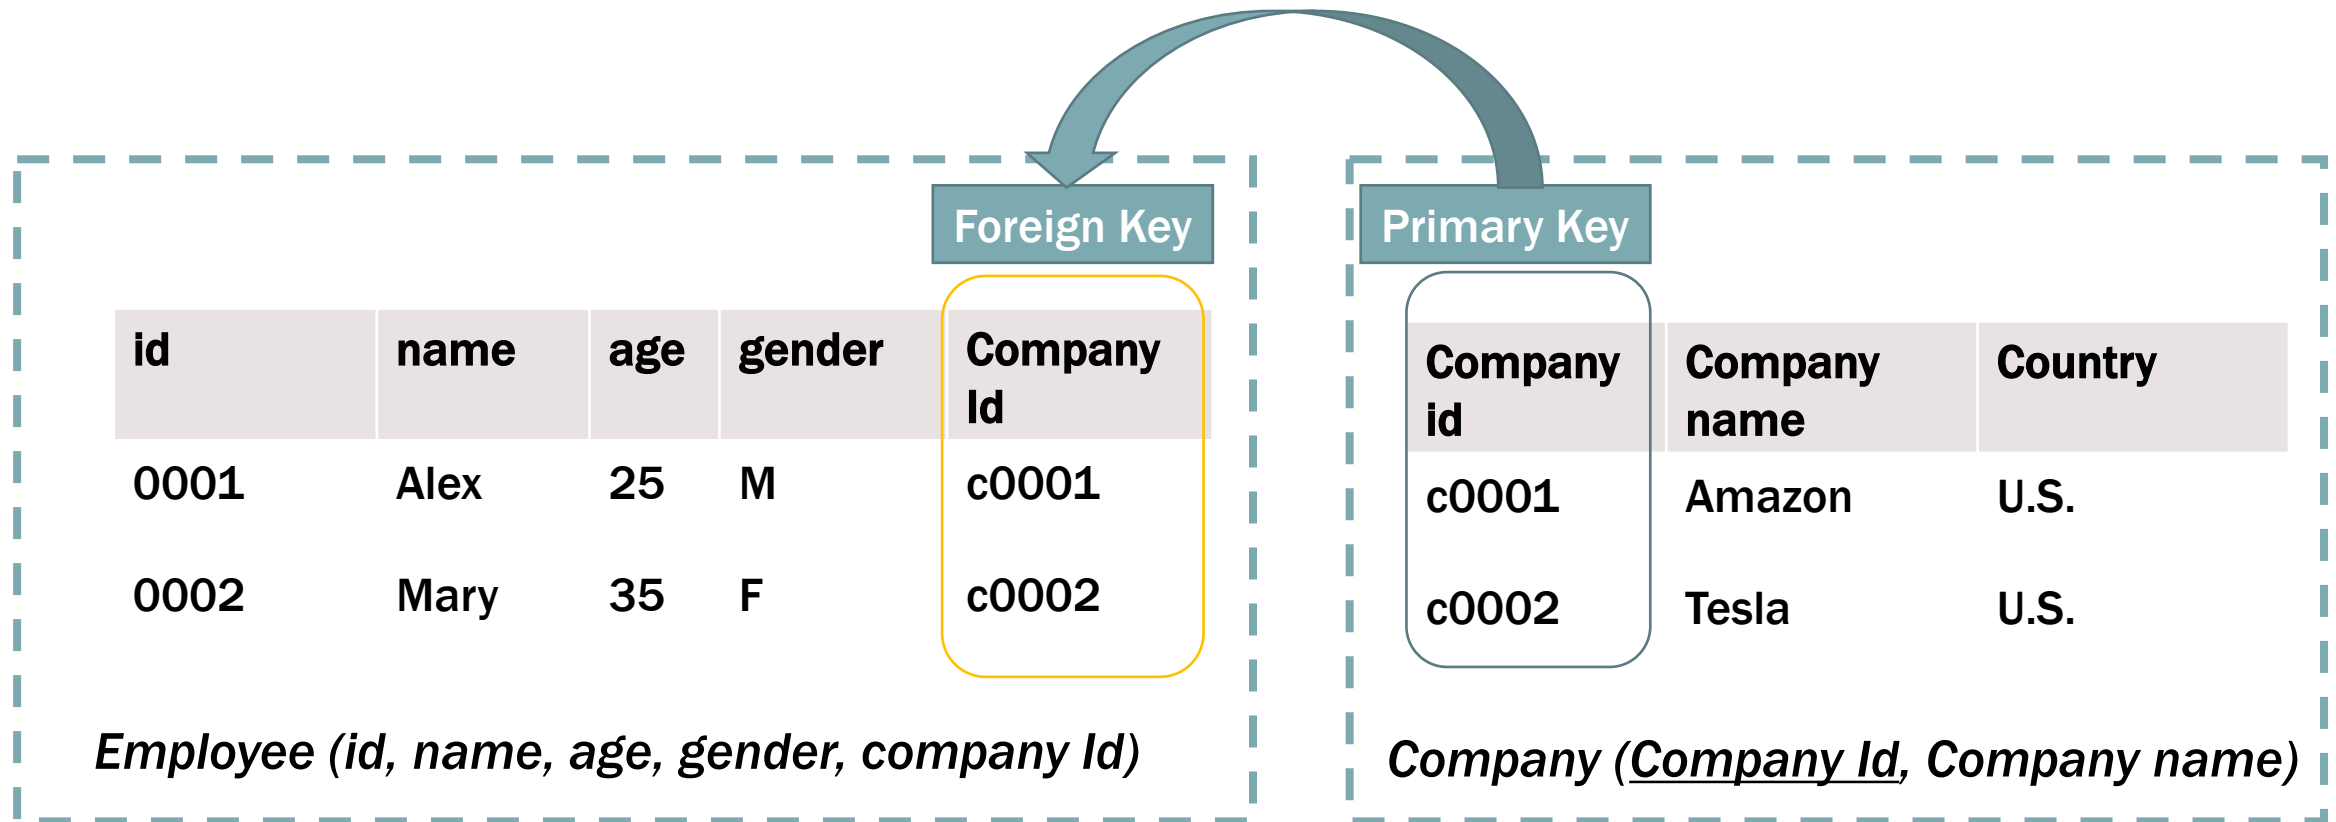
**Value:** Benefits from analyzing the data

# DATA MODELS

- ❑ Relational Data Model
  - ❑ Corresponding to relational database
- ❑ Key-Value Data Model
  - ❑ Corresponding to key-value systems
- ❑ Graph Data Model
  - ❑ Corresponding to graph database

# RELATIONAL DATA MODEL

## Primary key – Foreign key relationship



A foreign key is a set of one or more columns in a table that refer to the primary key in another table.

# KEY-VALUE DATA MODEL

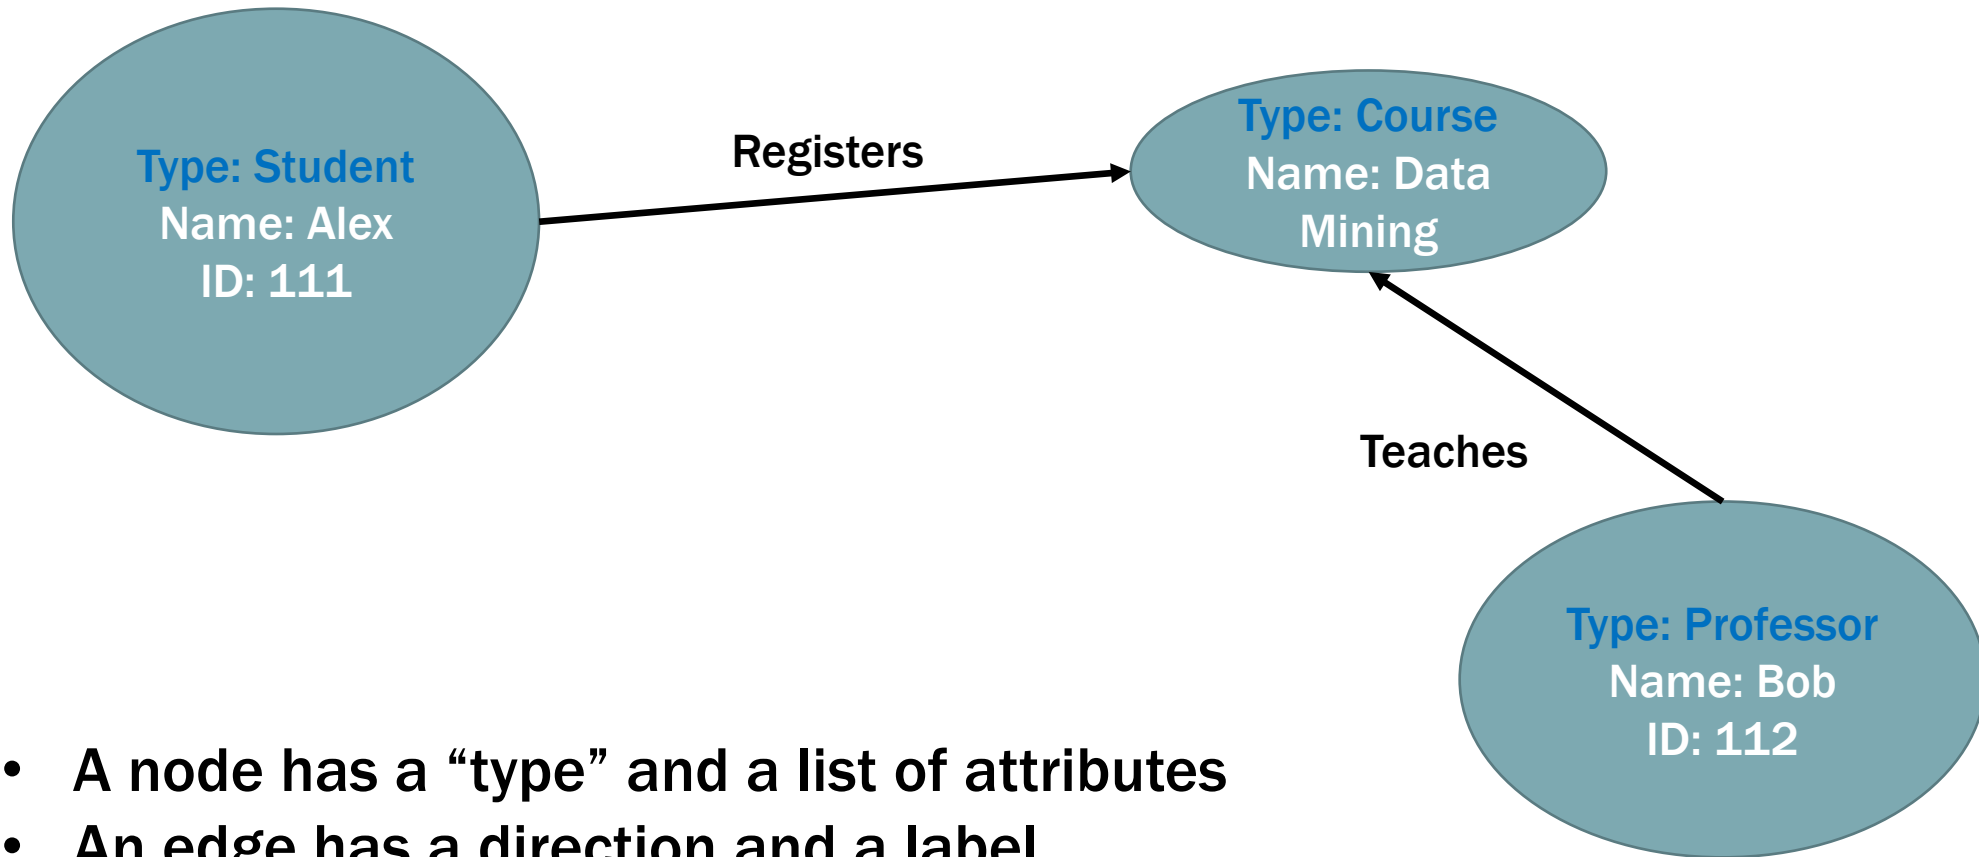
**Key-value Data Model is ubiquitous!**

**For any A that can determine B**

Key

Value

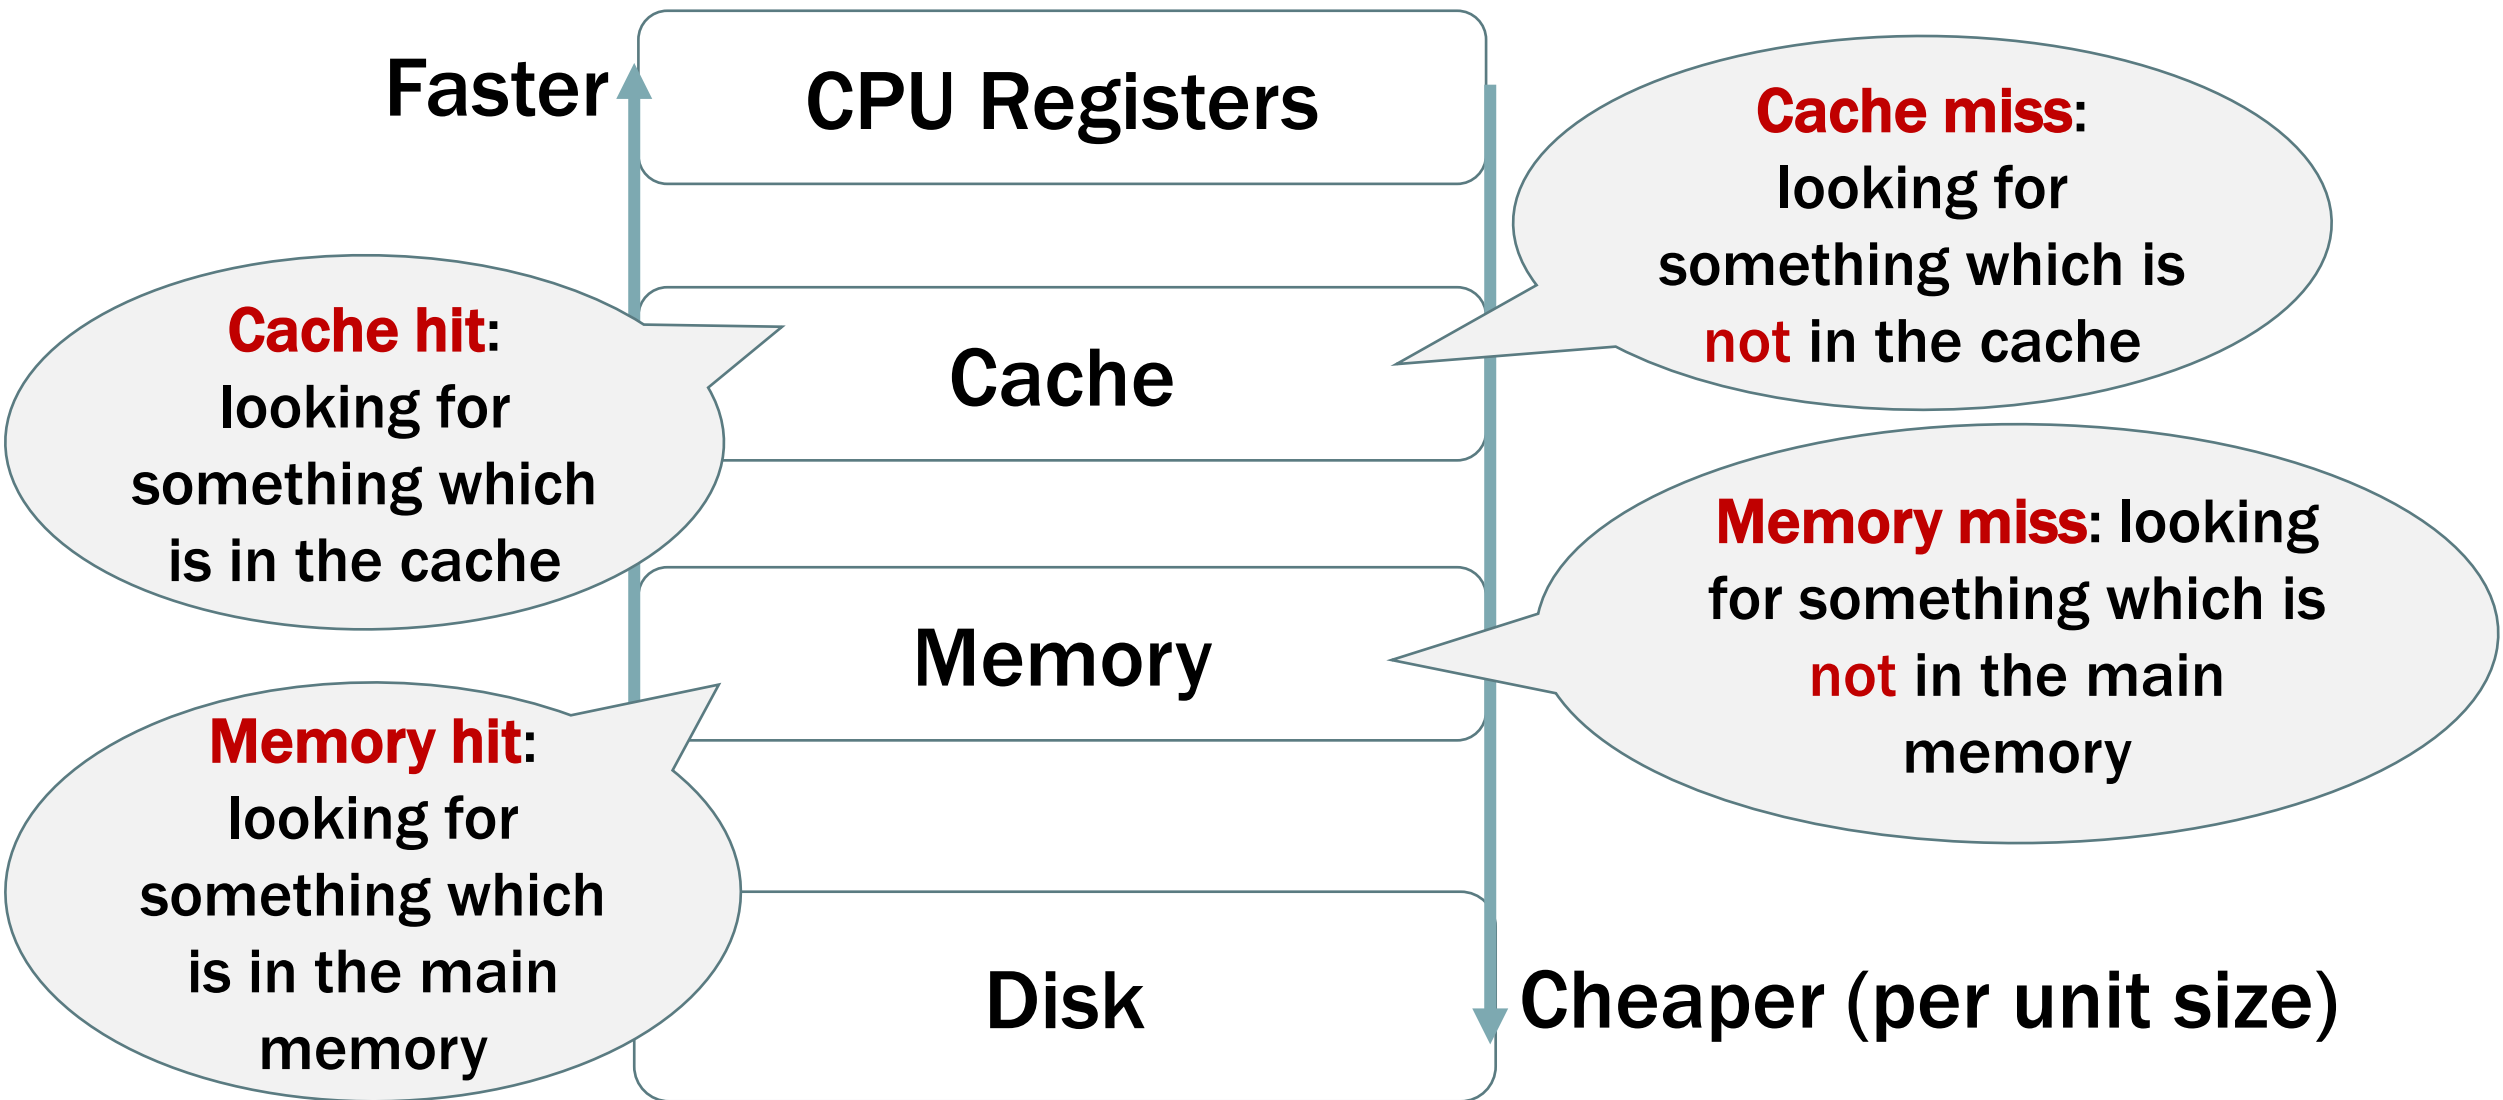
# GRAPHS ARE UBIQUITOUS



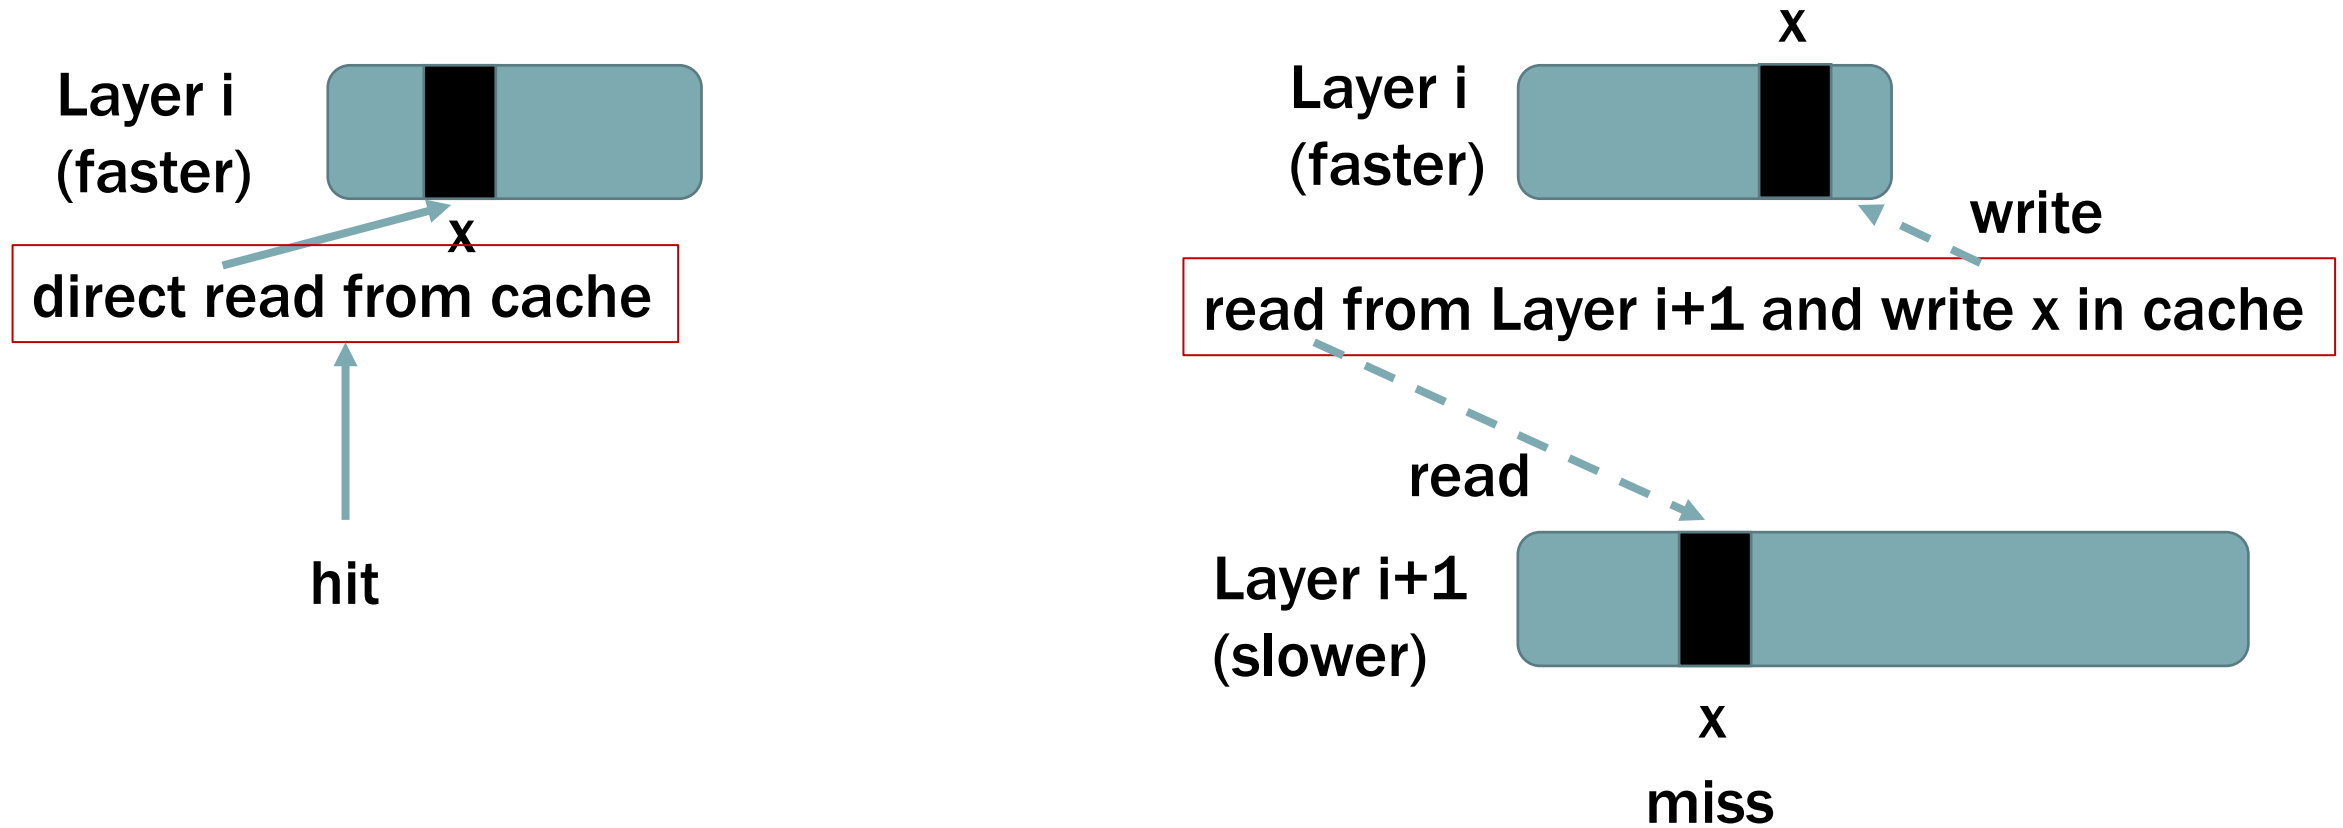
- A node has a “type” and a list of attributes
- An edge has a direction and a label



# MEMORY HIERARCHY



# DATA ACCESS IN MEMORY HIERARCHY



# COST OF A CACHE MISS

$$\square \text{cost\_access\_mem\_overall} = \text{cost\_access\_cache} \times (1 - \text{missrate}) + (\text{cost\_access\_cache} + \text{cost\_access\_mem}) \times \text{missrate}$$

# PAGE-BASED ACCESS EXAMPLE: SCANNING ARRAYS

Query  $x < 4$  from the following data

(size=120 bytes)

Memory Layer i

Memory Layer i+1

5, 10, 7, 4, 12

2, 8, 9, 11, 7

7, 11, 3, 9, 8

Each integer: 8 bytes

Each page: 8 bytes  $\times$  5 = 40 bytes

# PAGE-BASED ACCESS EXAMPLE: SCANNING ARRAYS

Cost: **40 Bytes**

Query  $x < 4$  from the following data

Scan



5, 10, 7, 4, 12

Qualified results



(size=120 bytes)  
Memory Layer i

Memory Layer i+1

5, 10, 7, 4, 12

2, 8, 9, 11, 7

7, 11, 3, 9, 8

Each integer: 8 bytes

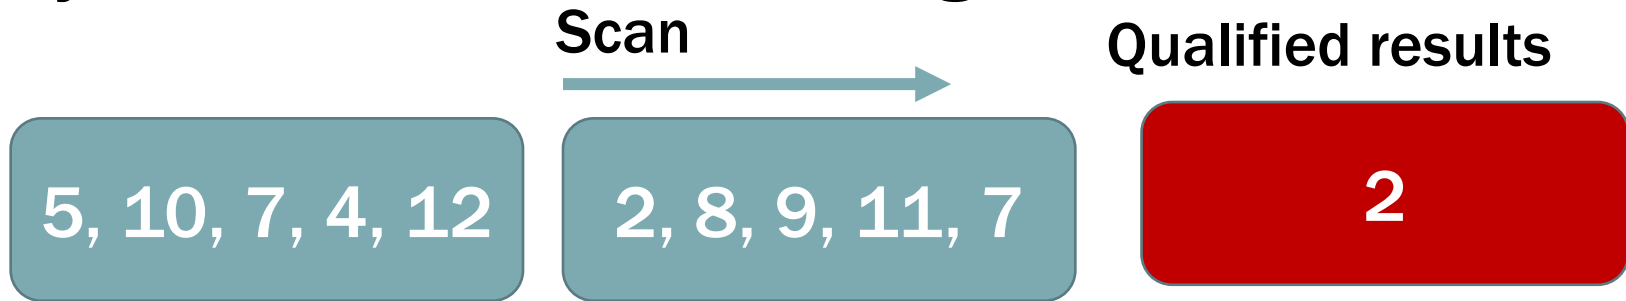
Each page: 8 bytes x 5 = 40 bytes

# PAGE-BASED ACCESS EXAMPLE: SCANNING ARRAYS


Cost: **80 Bytes**

Query  $x < 4$  from the following data

(size=120 bytes)  
Memory Layer i



Memory Layer i+1



Each integer: 8 bytes

Each page: 8 bytes x 5 = 40 bytes

# PAGE-BASED ACCESS EXAMPLE: SCANNING ARRAYS

Cost: **120 Bytes**

Query  $x < 4$  from the following data

Scan



Qualified results

(size=120 bytes)  
Memory Layer i

7, 11, 3, 9, 8

2, 8, 9, 11, 7

2, 3

Memory Layer i+1

5, 10, 7, 4, 12

2, 8, 9, 11, 7

7, 11, 3, 9, 8

Each integer: 8 bytes

Each page: 8 bytes x 5 = 40 bytes

# CACHE CONSCIOUS DESIGNS


- ❑ We show some design principles to make use of memory hierarchy for processing big data. We will discuss following two cases.
  - ❑ Array access patterns
    - ❑ Spatial locality designs
    - ❑ Temporal locality designs
- ❑ Big data sorting



# ACCESS PATTERN 1

**We have a size-10 array stored in main memory,  
cache size = 5, transfer size (cache line)= 5**

Access pattern: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10



# ACCESS PATTERN 1


Access pattern: **1**, 2, 3, 4, 5, 6, 7, 8, 9, 10

Cache



Cache Miss: **1**  
Cache Hit: 0

Memory



# ACCESS PATTERN 1

Access pattern: **1**, **2**, 3, 4, 5, 6, 7, 8, 9, 10


Cache



Cache Miss: 1

Cache Hit: **1**

Memory



# ACCESS PATTERN 1

Access pattern: **1**, **2**, **3**, 4, 5, 6, 7, 8, 9, 10


Cache



Cache Miss: 1

Cache Hit: **2**

Memory



# ACCESS PATTERN 1

Access pattern: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10


Cache



Cache Miss: 1

Cache Hit: 3

Memory



# ACCESS PATTERN 1

Access pattern: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10


Cache



Cache Miss: 1

Cache Hit: 4

Memory



# ACCESS PATTERN 1

Access pattern: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10


Cache



Cache Miss: 2

Cache Hit: 4

Memory



# ACCESS PATTERN 1

Access pattern: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10


Cache



Cache Miss: 2

Cache Hit: 5

Memory





# ACCESS PATTERN 1

Access pattern: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10


Cache



Cache Miss: 2

Cache Hit: 6

Memory



# ACCESS PATTERN 1

Access pattern: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10


Cache



Cache Miss: 2

Cache Hit: 7

Memory



# ACCESS PATTERN 1

Access pattern: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10


Cache



Cache Miss: 2

Cache Hit: 8


Memory



# SORTING

 or  size  $M/6$

$5M/6$  in total for in-memory sorting, remaining  $M/6$  for writing the results



Step 3:

Whenever one of {A,B,C,D,E} is empty, refill it from the source part.

Whenever the write-buffer is full, write it to the disk.

## Details of merge sort


1. 5 iterators scanning each of {A,B,C,D,E} from left to right;
2. Put the smallest-value (pointed by the 5 iterators) in the write-buffer; forward the corresponding iterator;

A	2, 13
B	7, 10
C	6, 11
D	5, 12
E	1, 14


# MAIN DESIGN OF COLUMN STORES

- ❑ Data in column stores are column-oriented
  - ❑ Also called column-oriented database, or columnar database.

**row-store (traditional RDB)**

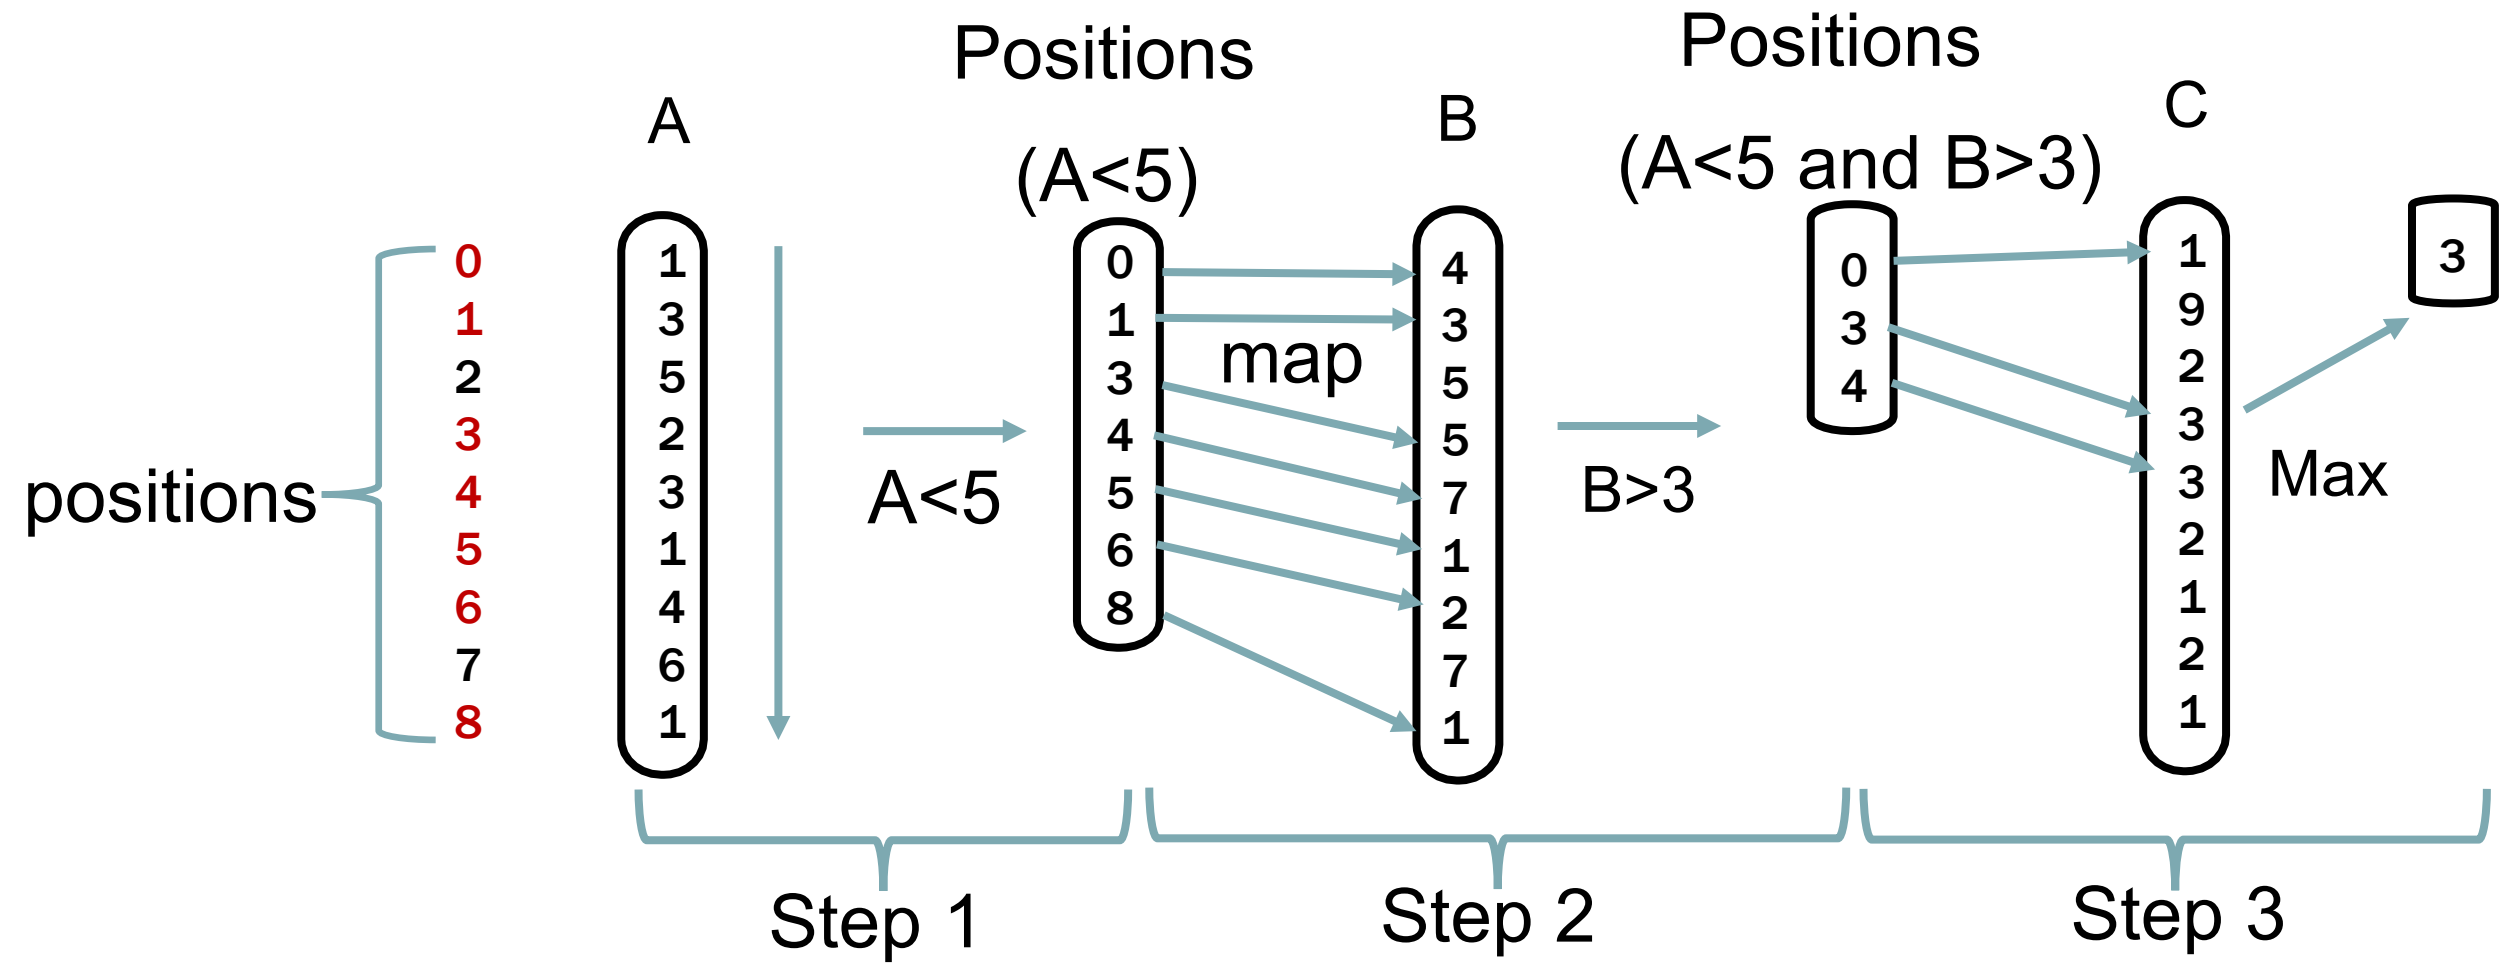


**column-store**



# QUERYING WITH COLUMN STORE (FLOW CHART)

- The whole table is stored in disk
- Have disk cache (memory buffer) whose size is a multiple of page size



# COMPARING QUERYING WITH ROW-STORE AND COLUMN-STORE

Without using an index, handling the query in the row store needs to read through the whole table. Namely, it costs  $Z \cdot w \cdot 3/P$  page accesses.

Cost for column store (number of page access):

$$\begin{aligned} & Zw/P + 2\text{result}(A) \cdot 4/P + \text{result}(A) + 2\text{result}(AB) \cdot 4/P + \text{result}(AB) \\ &= Zw/P + \text{result}(A) \cdot (8/P + 1) + \text{result}(AB) \cdot (8/P + 1) \\ &\approx Zw/P + \text{result}(A) + \text{result}(AB) \end{aligned}$$

Cost for row store (number of page access):

$$3Zw/P$$

Note: the values of  $\text{result}(A)$  and  $\text{result}(AB)$  are typically much smaller than  $Z$ . So we can conclude in this case column store is faster.

# OPTIMIZATIONS

☐ Compression

☐ Shared Scan

☐ Zone Map

☐ Sorting

☐ Indexing




# COMPRESSION

## Original data

8 bytes  
width

value1  
value2  
value3  
value1  
value1  
value4  
value2  
value3  
value5  
...



## Compressed

3 bits  
width

001  
010  
011  
001  
001  
100  
010  
011  
101  
...

## Dictionary

8 bytes  
width

value1  
value2  
value3  
value4  
value5

How many bits?



# SHARED SCANS

## loop fusion

```
for(i=0;i<n;i++)  
    min = a[i]<min ? a[i] : min  
  
for(i=0;i<n;i++)  
    max = a[i]>max ? a[i] : max
```


**Two passes of data**

```
for(i=0;i<n;i++)  
    min = a[i]<min ? a[i] : min  
    max = a[i]>max ? a[i] : max
```

**One pass of data**


# ZONE MAP

- ❑ A common way to help column scan is the zone map.
- ❑ It separates a column into “zones”, each is computed with max and min
- ❑ In filtering, some zones can be skipped.





# ENHANCED WITH SORTING

**SELECT max(C) FROM T WHERE A>10 and A<40 and B>20**



# ENHANCED WITH INDEX

- ❑ Index access is much faster than data access, because index is stored in Main Memory, while data is stored in disk.
- ❑ So it is beneficial to have complicated index access to locate the exact pages where stores the data.



**The End of First-Half.  
Thank you!**