

A complex network diagram with nodes and edges. Nodes are represented by circles of varying sizes in dark blue, red, and grey. Edges are thin lines connecting the nodes, with some being red and others dark blue. The background is a light blue-grey gradient.

BIG DATA MANAGEMENT

CZ/CE4123

Tutorial 7

Distributed Systems and MapReduce



QUESTION 1

Amazon wants to estimate the Top- K best sold products from S purchase records of L products in the form of a list of (User id, Product id) pairs. Suppose there is a distributed system with 1 master machine and M slave machines. Assume that L is a multiple of M . Design a distributed computation procedure to finish the task. Please describe

- (1) how the data is distributed, computed and aggregated?
- (2) how much data is sent across machines?

SOLUTION – MAIN IDEA

Let us roughly think...

First, we let each machine handle L/M products. Distribute the records corresponding to Product 1, 2, ..., L/M to Machine 1, the records corresponding to Product $L/M+1$, ..., $2L/M$ to Machine 2, ...

Second, in each machine, we can simply compute the top- K sold products **locally**.

Third, when there are local Top- K best sold products computed from each machine, aggregate the results to finally form the **global** top- K results.

SOLUTION – DETAILS

Firstly, the master machine distributes the S purchase records to the slave machines, so that the records of Product 1, 2, ..., L/M are sent to Machine 1, the records of Product $L/M+1$, ..., $2L/M$ are sent to Machine 2, ...

Secondly, each machine computes the top- K sold products **locally**. In particular, for each product, we collect the frequencies of it being sold, and sort the frequencies and find the K products *with* top- K frequencies. Only send the top- K pairs of (Product id, Frequency) to master.

SOLUTION – DETAILS

Finally, the master machine aggregates all the MK pairs of (Product id, Frequency). Output the products with the Top- K frequencies.

SOLUTION – DATA SENT

Data Distribution Phase: S records have been sent out.

Aggregation Phase: MK pairs of (Product id, Frequencies) have been sent out.

In total, the data sent out is in the scale of $O(S+MK)$.

QUESTION 2

(1) In a MapReduce job, the output of Map phase is a list of key-value pairs:

(A, 1) (C, 2), (A, 5), (C, 6), (B, 3), (E, 3), (C, 8). Please list the possible input to the Reduce function.

The input to the *Reducer* aggregates the *Map* output keys.

Hence, the possible input to the reducer is (A, {1, 5}); (B, {3}); (C, {2, 6, 8}); (E, {3}).

(2) Based on the answer to Q2(1), write a Reduce function (in pseudocode) so that the MapReduce output is: (2, A), (3, C), (6, A), (7, C), (4, B), (4, E), (9, C).

```
Reduce(int key, iterator values){  
    for(each v in values){  
        Emit(v+1, key);  
    }  
}
```

(3) Consider an employee table containing three columns (EmployeeID, age, monthly-salary) where age and monthly-salary are integers. Use MapReduce to collect the number of employees falling into each of the following two categories:

- Category 1: The age of the employee is between 30 and 40 (including 30 and 40). His/her monthly salary is at most 7000.
- Category 2: The age of the employee is between 40 and 50 (including 40 and 50). His/her monthly salary is more than 7000.

Please use **only one MapReduce Job** to achieve this task and write down the pseudocode of the Map function and Reduce function. The input key and value for Map function are an employee's age and monthly-salary respectively.

```
Map(int age, int salary){  
    if(age>=30 and age<=40 and salary<=7000){  
        Emit-Intermediate(1, 1);  
    }  
    if(age>=40 and age<=50 and salary>7000){  
        Emit-Intermediate(2, 1);  
    }  
}
```

```
Reduce(int category, iterator values){  
    int total_frequency = 0;  
    for(int frequency: values)  
    {  
        total_frequency+=frequency;  
    }  
    Emit(category, total_frequency);  
}
```

QUESTION 3

Design MapReduce algorithms for the multiplication of two matrices A , B of n by n . Elements in the matrices are integers. The input key for *Map* function is *MatrixName*; the input value for *Map* function is in the form of $i;j;v$, indicating that the value of the i -th row and j -th column is v .

(Matrix Multiplication: Given matrix $A[n \times n]$ and matrix $B[n \times n]$, compute matrix C such that $C[i][z] = \sum_{j=0}^{n-1} A[i][j] \times B[j][z]$, for i, z in $[0, n-1]$).


- (1) Use at most two MapReduce jobs to finish the computation.
- (2) Furthermore, can the multiplication be finished using **one MapReduce job**?

SOLUTION – 1ST JOB

```
Map(String MatrixName, String value){  
    int i = get_i(value);// get row  
    int j = get_j(value);// get column  
    int v = get_v(value);// get value  
    if(MatrixName=="A")  
        Emit-Intermediate(j, ToString(MatrixName, i, v));// combine  
    else  
        Emit-Intermediate(i, ToString(MatrixName, j, v));  
}
```

Column as key

Row as key



For each intermediate key j , the values can contain

(1) $A[u][j]$ for any u ;

(2) $B[j][v]$ for any v ;

What needs to be done in `reduce()`?

For any (u, v) , send out $A[u][j] * B[j][v]$ for aggregation of $C[u][v]$

SOLUTION – 1ST JOB

```
Reduce(String key, Iterator<String> values){
```

```
    int A_start[n]; int B_end[n];
```

```
    for(String value : values){
```

```
        String MatrixName=get_first_element(value);
```

```
        if(MatrixName=="A"){
```

```
            int i = get_second_element(value);// get row index in matrix A
```

```
            A_start[i]=get_third_element(value);// get the matrix element
```

```
        }
```

```
    else{
```

```
        int j = get_second_element(value);// get column index in matrix B
```

```
        B_end[j] = get_third_element(value);// get the matrix element
```

```
    }
```

```
}
```

```
}
```

```
for (int u=0;u<n;u++)
```

```
for(int v=0;v<n;v++)
```

```
{
```

```
    Emit(ToString(u,v), A_start[u]*B_end[v]);
```

```
}
```

A[u][j]

B[j][v]



THE 2ND JOB

Need to aggregate all the values for each pair (u,v).

So,

Map() does nothing;

Reduce() conducts the aggregation.

SOLUTION - 2ND JOB

```
Map(String key, String value){  
    Emit-Intermediate(key, value);  
}
```

SOLUTION - 2ND JOB

```
Reduce(String key, Iterator<String> values){  
    int sum=0;  
    for (String value in values){  
        sum+=ToInteger(value);  
    }  
    Emit(key, sum);  
}
```

Can we use a single job only?

Think...

For any $A[i][j]$, which elements of C would it contribute to?

Can we use a single job only?

Think...

For any $A[i][j]$, which elements of C would it contribute to?

$C[i][k]$ for any k

For any $B[i][j]$, which elements of C would it contribute to?

$C[k][j]$ for any k

CAN WE FINISH THESE TASKS WITH ONE JOB ONLY?

```
Map(String MatrixName, String value){  
    int i = get_i(value);  
    int j = get_j(value);  
    int v = get_v(value);  
    if(MatrixName=="A")  
        for(int k=0;k<n;k++)  
            Emit-Intermediate(ToString(i, k), ToString(MatrixName, j, v)); // each A[i][j] may contribute  
                                                                           //to C[i][k] for any k  
    else  
        for(int k=0;k<n;k++)  
            Emit-Intermediate(ToString(k, j), ToString(MatrixName, i, v)); // each B[i][j] may  
                                                                           //contribute to C[k][j] for any k  
}
```

For each intermediate key (pair (i,j)), it aggregates

(1) The values of $A[i][u]$, for any u

(2) The values of $B[u][j]$, for any u

CAN WE FINISH THESE TASKS WITH ONE JOB ONLY?

```
Reduce(String index_pair_for_C, Iterator<String> values){
```

```
    int A_middle[n]; int B_middle[n];
```

```
    for(String value in values){
```

```
        String MatrixName=get_first_element(value);
```

```
        if(MatrixName=="A"){
```

```
            int j = get_second_element(value);
```

```
            A_middle[j]=get_third_element(value);
```

```
        }
```

```
    else{
```

```
        int i = get_second_element(value);
```

```
        B_middle[i] = get_third_element(value);
```

```
    }
```

```
}
```

```
}
```

```
    int sum=0;
```

```
    for (int u=0;u<n;u++)
```

```
        sum+=A_middle[u]*B_middle[u];
```

```
    Emit(index_pair_for_C, sum);
```

EXAMPLE

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \times \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix}$$

EXAMPLE-MAP

$A[0][0] \rightarrow ((0;\textcolor{red}{0}), (A; 0; A[0][0]))$

$\rightarrow ((0;\textcolor{red}{1}), (A; 0; A[0][0]))$

$A[0][1] \rightarrow ((0;\textcolor{red}{0}), (A; 1; A[0][1]))$

$\rightarrow ((0;\textcolor{red}{1}), (A; 1; A[0][1]))$

$A[1][0] \rightarrow ((1;\textcolor{red}{0}), (A; 0; A[1][0]))$

$\rightarrow ((1;\textcolor{red}{1}), (A; 0; A[1][0]))$

$A[1][1] \rightarrow ((1;\textcolor{red}{0}), (A; 1; A[1][1]))$

$\rightarrow ((1;\textcolor{red}{1}), (A; 1; A[1][1]))$

$B[0][0] \rightarrow ((\textcolor{red}{0};0), (B; 0; B[0][0]))$

$\rightarrow ((\textcolor{red}{1};0), (B; 0; B[0][0]))$

$B[0][1] \rightarrow ((\textcolor{red}{0};1), (B; 0; B[0][1]))$

$\rightarrow ((\textcolor{red}{1};1), (B; 0; B[0][1]))$

$B[1][0] \rightarrow ((\textcolor{red}{0};0), (B; 1; B[1][0]))$

$\rightarrow ((\textcolor{red}{1};0), (B; 1; B[1][0]))$

$B[1][1] \rightarrow ((\textcolor{red}{0};1), (B; 1; B[1][1]))$

$\rightarrow ((\textcolor{red}{1};1), (B; 1; B[1][1]))$

EXAMPLE-REDUCE

Take (1,0) for example intermediate key

It aggregates {(A; 0; A[1][0]), (B; 0; B[0][0]), (A; 1; A[1][1]), (B; 1; B[1][0])}

Hence, we can compute C[1][0] by $A[1][0]*B[0][0]+A[1][1]*B[1][0]$