

CE/CZ 4123 Big Data Management Tutorial 10

Key-Value Stores

Question 1:

Consider a leveling LSM-tree with a size ratio 4. The memory buffer (Level 0) can store 5000 key-value pairs. Initially the LSM-tree is empty. After inserting 70000 key-value pairs with distinct keys continuously, how many levels are formed?

Question 2:

Consider a leveling LSM-tree with a size ratio 4. The LSM-tree has 5 levels (excluding the memory buffer level), and it is incorporated with **both** fence pointers and Bloom Filters. Assume that a key-value pair is always entirely stored within a disk page. Consider the procedure of $\text{Get}(K)$ for a key K , which of the followings sequence are possible to be the I/O costs from Level-0 to Level-5? (can select multiple answers)

- (a) 1, 1, 1, 1, 1, 1
- (b) 0, 1, 1, 1, 1, 1
- (c) 0, 0, 1, 0, 0, 1
- (d) 0, 0, 0, 0, 0, 1
- (e) 1, 0, 0, 0, 0, 0

Question 3:

Consider a leveling LSM-tree with a size ratio 4. The LSM-tree has L levels (excluding the memory buffer level), and it is **only** incorporated with fence pointers (without Bloom Filters). Assume that a key-value pair is always entirely stored within a disk page. Consider the procedure of $\text{Get}(K)$ for a key K ,

- (1) What is the possible I/O cost of accessing Level- i (i is in $[1, L]$)?
- (2) If K exists in the LSM-tree, what is the expected I/O cost at Level- i (i is in $[1, L]$)? (hint: divide the cases based on the first-appearing location of the key K)

Question 4:

Consider a leveling LSM-tree with a size ratio 4. The LSM-tree has L levels (excluding the memory buffer level), and it is incorporated with **both** fence pointers and Bloom Filters. Assume that a key-value pair is always entirely stored within a disk page. Consider the procedure of $\text{Get}(K)$ for a key K ,

- (1) What is the possible I/O cost of accessing Level- i (i is in $[1, L]$)?
- (2) If K exists in the LSM-tree and the FPR of the Bloom filter at Level- i (i is in $[1, L]$) is P (P is in $[0, 1]$), what is the expected I/O cost at Level- i ? (hint: divide the cases based on the first-appearing location of the key K)