

A complex network diagram with nodes and edges. Nodes are represented by circles of varying sizes in dark blue, red, and grey. Edges are thin lines connecting the nodes, with red lines forming a dense web and grey lines forming a more sparse structure. The background is a light blue-grey gradient.

BIG DATA MANAGEMENT

CZ/CE4123

Tutorial 10

Key-Value Stores



QUESTION 1

Consider a leveling LSM-tree with a size ratio 4. The memory buffer (Level 0) can store 5000 key-value pairs. Initially the LSM-tree is empty. After inserting 70000 key-value pairs with distinct keys continuously, how many levels are formed?

QUESTION 1-SOLUTION

This question is related to understanding the capacity.

The capacity of a level is defined as “the maximum number of key-value pairs that can be stored in the level”.

Size Ratio = 4

Level 0 capacity = 5000

Level 1 capacity = $5000 \times 4 = 20000$

Level 2 capacity = $5000 \times 4 \times 4 = 80000$

$70000 > \text{Level 0 capacity} + \text{Level 1 capacity} (=25000)$

$70000 < \text{Level 0 capacity} + \text{Level 1 capacity} + \text{Level 2 capacity} (=105000)$

So there are at least 3 levels (Level 0, Level 1, Level 2)

QUESTION 1-SOLUTION

We further verify that Level 3 is not created. (**why?**)

If level 2's actual size is larger than Level 2 capacity – Level 1 capacity, it can already trigger the merge

Before creating Level 3, it must trigger the sort-merge of Level 2. Since the 70000 keys are **distinct**, the actual size of Level 2 must be a multiple of the capacity of Level 1 (i.e., 20000). So only when the actual size of Level 2 reaches 80000 it can trigger a merge, (note that 0, 20000, 40000, 60000 are not larger than {Level 2 capacity – Level 1 capacity}, and hence will not trigger a merge). However, since there are only 70000 keys, it is impossible for Level 2 to reach a size of 80000. Hence, Level 3 will not be created. So finally there are 3 levels (Levels 0, 1, 2).

QUESTION 2

Consider a leveling LSM-tree with a size ratio 4. The LSM-tree has 5 levels (excluding the memory buffer level), and it is incorporated with **both** fence pointers and Bloom Filters. Assume that a key-value pair is always entirely stored within a disk page. Consider the procedure of $\text{Get}(K)$ for a key K , which of the followings sequence are possible to be the I/O costs from Level-0 to Level-5? (can select multiple answers)

- (a) 1, 1, 1, 1, 1, 1
- (b) 0, 1, 1, 1, 1, 1
- (c) 0, 0, 1, 0, 0, 1
- (d) 0, 0, 0, 0, 0, 1
- (e) 1, 0, 0, 0, 0, 0

SOLUTION

(a) and (e) are not possible because Level 0 is memory buffer, and hence no I/O costs.

(b) (c) (d) are possible because if K is not in the LSM-tree, then each level's I/O cost fully depends on the accuracy of the Bloom filter (We consider Fence Pointer will incur 1 I/O when Bloom filter returns true).

- (b) is possible: BFs in Levels 1-5 all generate false-positives
- (c) is possible: BFs in Levels 2 and 5 generate false-positives
- (d) is possible: BF in Level 5 generates a false-positive

QUESTION 3

Consider a leveling LSM-tree with a size ratio 4. The LSM-tree has L levels (excluding the memory buffer level), and it is **only** incorporated with fence pointers (without Bloom Filters). Assume that a key-value pair is always entirely stored within a disk page. Consider the procedure of $\text{Get}(K)$ for a key K ,

- (1) What is the possible I/O cost of accessing Level- i (i is in $[1, L]$)?
- (2) If K exists in the LSM-tree, what is the expected I/O cost at Level- i (i is in $[1, L]$)? (hint: divide the cases based on the first-appearing location of the key K)

SOLUTION FOR Q3(1)

The possible I/O cost at Level- i can be 0 or 1.

0 is possible:

if K first appears in a level smaller than Level- i , then the search is terminated before Level- i . Hence, no I/O cost at Level- i .

1 is possible:

if K first appears in Level- i or larger levels, then fence pointer is used and can incur 1 page read, or 1 I/O. (Note: We assume that using fence pointers always incurs 1 I/O.)

SOLUTION FOR Q3(2)

Let j be the first level that contains key K .

Divided into two cases:

- If $i > j$, then the (expected) I/O cost is 0, because the search ends at Level j .
- If $i \leq j$, then the (expected) I/O cost is 1, because at level i fence pointer is used and incurs 1 I/O.

QUESTION 4

Consider a leveling LSM-tree with a size ratio 4. The LSM-tree has L levels (excluding the memory buffer level), and it is incorporated with **both** fence pointers and Bloom Filters. Assume that a key-value pair is always entirely stored within a disk page. Consider the procedure of $\text{Get}(K)$ for a key K ,

(1) What is the possible I/O cost of accessing Level- i (i is in $[1, L]$)?

(2) If K exists in the LSM-tree and the FPR of the Bloom filter at Level- i is P (P is in $[0, 1]$), what is the expected I/O cost at Level- i (i is in $[1, L]$)? (hint: divide the cases based on the first-appearing location of the key K)

SOLUTION FOR Q4(1)

The possible I/O cost at Level- i can be 0 or 1.

0 is possible: if K is not in the LSM-tree, and the BF in level- i returns FALSE. Then, the search within the disk for this level is skipped.

1 is possible: if K is not in the LSM-tree, and the BF in level- i returns TRUE. Then, fence pointer is used and incurs 1 page read, or 1 I/O. (Note: We assume that when using fence pointers always incurs 1 I/O.)

SOLUTION FOR Q4(2)

Since K exists in the LSM-tree. The cases is divided based on the first-appearing level of key K .

Let Level- j be the level that first contains K .

Case 1: If $i > j$, the search is up to Level- j , and terminates before reaching Level- i , and hence the I/O cost at Level- i is 0;

Case 2: If $j = i$, the I/O cost at level- i must be 1 because the key first appears at Level- i .

Case 3: If $i < j$, the I/O cost depends on the FPR of the Bloom filter:

- Since the key K does not exist in Level- i , then with probability P the Bloom filter returns TRUE, and later the fence pointer incurs 1 I/O.
- Since the key K does not exist in Level- i , then with probability $1-P$ the Bloom filter returns FALSE, and no I/O cost is incurred.
- To summarize, the expected I/O cost is : $P*1+(1-P)*0=P$.